

COMP30030: Introduction to Artificial Intelligence

Neil Hurley

School of Computer Science
University College Dublin
`neil.hurley@ucd.ie`

October 25, 2018

1 Problem Solving by Search

- Uninformed Search
- Informed Search
- Adversarial Search
- Game Playing with Reinforcement Learning

2 Optimisation

- Optimisation Overview
- Combinatorial Optimisation Problems
- Simulated Annealing
- Optimisation Problem Examples
- Convergence of Simulated Annealing
- Genetic Algorithms
- Optimisation in Continuous Spaces

3 Machine Learning

- Supervised Machine Learning

Machine Learning I

- The term **Machine Learning** refers to “a scientific discipline that explores the construction and study of algorithms that can learn from data” (Wikipedia).
- In general, there are two broad ways that we can learn from data:
 - 1 **Supervised** methods learn by generalising from **training** data in the form of sets of inputs with **desired outputs**.
 - 2 **Unsupervised** methods learn from input data, which has no desired outputs associated with it. Instead these methods search for structure in the input data.
- Another important branch of machine learning is **reinforcement learning**, in which a program must interact with a dynamic environment to perform a certain goal or task.

1 Problem Solving by Search

- Uninformed Search
- Informed Search
- Adversarial Search
- Game Playing with Reinforcement Learning

2 Optimisation

- Optimisation Overview
- Combinatorial Optimisation Problems
- Simulated Annealing
- Optimisation Problem Examples
- Convergence of Simulated Annealing
- Genetic Algorithms
- Optimisation in Continuous Spaces

3 Machine Learning

- Supervised Machine Learning

Supervised Machine Learning

- Suppose we are have a problem that we wish to learn from.
- Let's call the set of possible problem *instances*, \mathcal{X} . So an example problem is an element $x \in \mathcal{X}$.
- We can think of the solution to this problem as a map $f : \mathcal{X} \rightarrow \mathcal{Y}$, where \mathcal{Y} is the set of solutions to problems in \mathcal{X} .
- So, given x , the solution to x is $f(x)$.
- We don't know this function $f(.)$ — we would like to **learn** an approximation to it.

Supervised Machine Learning

- A learning algorithm is **supervised**, if, along with a set of instances $x \in \mathcal{X}$, we have their solutions $f(x)$.
- So we have a **training set** $D_{\text{train}} = \{(x, f(x)) | x \in \mathcal{X}\}$.
- If we are lucky, we may have many training examples and supervised Machine Learning algorithms generally work better, the more data that is available.
- The learning algorithm is guided, so *supervised* by the solutions to the training examples that are available.
- The goal is to *generalise* from the training examples, so that solutions to new, unseen instances can be found.

Types of Supervised Learning Task

- While we can imagine many problem types, in fact there are two broad types of machine learning tasks, which apply to many real-world scenarios, which may be distinguished by the type of solution that we seek.
- 1 The solution space is a set of real-valued numbers $\mathcal{Y} = \mathbb{R}^n$ for some dimension n . In this case the problem is referred to as a **regression** problem.
- 2 The solution space is discrete, consisting of a finite set of $k > 0$ labels. $\mathcal{Y} = \{y_1, y_2, \dots, y_k\}$. In this case, the problem is referred to as a **classification** problem.
 - Many approaches have been developed for the simplest case of *binary classification*, where there are just two possible labels.
 - In binary classification, the goal is to determine which, of two classes, each problem instance belongs to.

Types of Supervised Learning Task

- These two problem categories may seem to simple to cover all the many learning tasks we can imagine in the real-world, but actually they are quite general.
- Note that we've (so far), said very little about the **representation** of the problem instances, nor the representation of their solutions.
- All of the above notation simply says that,
 - a problem is a regression problem if we can find a useful representation of its solution as a set of real-valued numbers.
 - a problem is a classification problem if wish to determine class which among a finite set of possibilities a problem instance belongs to.

Real-world Examples

Determining if an email is Spam

- x is a text – a sequence of alphabetic characters, corresponding to an email message.
- $f(x)$ is the binary label indicating, which of the two cases, Spam or Not Spam, applies to the given text.

Real-world Examples

Medical Diagnosis

- x is a set of properties of the patient – symptoms, results of lab tests, previous medical history etc. etc.
- $f(x)$ is a disease. A set of possible diseases has been preselected, e.g. {NO_DISEASE, CANCER, HEART_DISEASE}.
- The problem then is, given x , label x with the correct disease — a classification problem.

Real-world Examples

Making a Computer Read

- x is a text – a sequence of alphabetic characters.
- $f(x)$ is the sound signal corresponding to the utterance of x . $f(y)$ will be represented by an appropriate set of real-valued numbers, corresponding to a representation of the sound signal, or *features* of the sound signal.
- Typically, a database of examples is gathered from human reading of set texts.
- The goal is to train the computer to automatically read texts which are not in the training set.
- In the heart of this problem there is a regression – mapping each x to a real-valued representation of the sound. However, to successfully achieve such a learning task, we need to apply detailed knowledge of speech and linguistics.

Hypothesis Space I

- How is it possible to learn from training examples, $D_{\text{train}} = \{x, f(x)\}$?
 - One way might be to try to create a (close to) exact model for $f(\cdot)$
 - We can write down, for example, models of physical processes, such as fluid dynamics and use these to predict the weather tomorrow.
 - Take our reading example – can we write down the detailed steps by which a text becomes spoken language?
- Instead, we give up on creating such a detailed model of $f(x)$ and instead we choose a set of *candidate* functions, from which a good approximation of $f(x)$ can be selected.

Hypothesis Space II

- We call such a set a **Hypothesis Space**, \mathcal{H}
- The Hypothesis Space should be **simple** enough that it is tractable to select a good function from this space.
- The Hypothesis Space should be **complex** enough that a good candidate function exists in that space that well approximates $f(x)$.

Inductive Learning Hypothesis

- But how will we know that the function we select from \mathcal{H} is really a good approximation to $f(x)$?
 - In practise, we cannot test it on all instances in the problem space.
- The **Inductive Learning Hypothesis**
 - Any $h(.) \in \mathcal{H}$ that approximates $f(x)$ well on the training examples, will also approximate $f(x)$ well on unseen examples.
- Whether or not this holds depends strongly on the training examples (and the learning algorithm applied to them).
- In statistical terms, the training examples need to be a **representative sample** of the full set of problems that occur in \mathcal{X} and this may not always be the case.
- Nevertheless, supervised machine learning algorithms proceed by, as a primary objective, looking for functions in the Hypothesis space that approximate the training examples well.

Some Terminology


- $h(.) \in \mathcal{H}$ is called a **consistent hypothesis** if it agrees with $f(.)$ on all training examples. Given the training data, only some hypothesis in \mathcal{H} are consistent. The set of consistent hypotheses is called the **Version Space**:
$$\{h(.) \in \mathcal{H} : h(x) = f(x) \forall x \in D_{\text{train}}\}$$
- A **consistent learner** always outputs a consistent hypothesis
- The **empirical error** is the fraction of the training examples for which $h(x) \neq f(x)$.
- So, for a consistent learning, the empirical error is 0%.

Example – An ML approach to learning a Boolean Function I

- The following example shows how a machine learning algorithm might proceed to learn a function from examples.
- We assume that there is some unknown Boolean function, over four Boolean variables:
 - $x = (x_1, x_2, x_3, x_4)$ and each variable $x_i \in \{0, 1\}$
 - $y = f(x) \in \{0, 1\}$
- How hard is the problem? Well, how many possible Boolean functions over 4 variables exist?
 - There are $2^4 = 16$ possible inputs to this unknown function. The function is fully defined, once it is specified what the output is for each of these 16 inputs.
 - There are 2 choices for the output $f(x)$. So there are 2 ways to set the output for each possible input. Hence, there are 2^{16} different possible functions.

Example – An ML approach to learning a Boolean Function II

- One of these functions is generating output, but we do not know which one it is – all we have available to us is a set of training instances x and their associated output $f(x)$.
- If there are very few training instances, then we will not be able to learn much.
- However, tractable learning depends not just on the instances we have, but on good selection of a Hypothesis space, that
 - 1 Is likely to contain a good approximation of $f(x)$
 - 2 But which can be searched efficiently for such an approximation.



Training
examples

x_1	x_2	x_3	x_4	y
0	0	1	1	1
1	0	0	1	1
1	0	1	1	1
1	1	0	0	0
1	1	1	0	0

Hypothesis Space of *all* Boolean functions

- We do not restrict the Hypothesis space, and consider all possible Boolean functions among our candidates
- We have $2^{16} = 64k$ possible functions.
- 5 outputs are specified in the training set, so there are $2^{16-5} = 2048$ consistent hypothesis.
- We can do no better than choose randomly between them.
- How likely is it that our unknown function agrees with the chosen one on unseen examples?

x_1	x_2	x_3	x_4	Y
0	0	0	0	?
0	0	0	1	?
0	0	1	0	?
0	0	1	1	?
0	1	0	0	?
0	1	0	1	?
0	1	1	0	?
0	1	1	1	?
1	0	0	0	?
1	0	0	1	?
1	0	1	0	?
1	0	1	1	?
1	1	0	0	?
1	1	0	1	?
1	1	1	0	?
1	1	1	1	?

x_1	x_2	x_3	x_4	Y
0	0	0	0	?
0	0	0	1	?
0	0	1	0	?
0	0	1	1	1
0	1	0	0	?
0	1	0	1	?
0	1	1	0	?
0	1	1	1	?
1	0	0	0	?
1	0	0	1	1
1	0	1	0	?
1	0	1	1	1
1	1	0	0	0
1	1	0	1	?
1	1	1	0	0
1	1	1	1	?

Hypothesis Space : Conjunction of Literals

- A **literal** is a variable x_i or its negation $\neg x_i$
- A **term** is a *conjunction* of literals.
- We define the hypothesis space
$$\mathcal{H} = \{\text{terms over } x_1, x_2, x_3, x_4\}$$
- Now the hypothesis space only contains $3^4 = 18$ possible functions. (Each term can be affirmed or negated or isn't present).
- Learning algorithm
 - 1 Initially $h(.) =$ conjunction of all possible literals
 - 2 Remove literals associated with inconsistent *positive* examples.

Learning Conjunctions

1. $h = \neg x_1 x_1 \neg x_2 x_2 \neg x_3 x_3 \neg x_4 x_4$
2. Observe 1st training example, remove literals x_1 , x_2 , $\neg x_3$, and $\neg x_4$
 $h = \neg x_1 \neg x_2 x_3 x_4$
3. Observe 2nd training example, remove literals $\neg x_1$ and x_3
 $h = \neg x_2 x_4$
4. Observe 3rd training example:
nothing to do ($h = \neg x_2 x_4$)
5. No more positive training examples
Output $h = \neg x_2 x_4$

x_1	x_2	x_3	x_4	y
0	0	1	1	1
1	0	0	1	1
1	0	1	1	1

Learning as Refinement

- Start with a small hypothesis class, such as Boolean conjunctions – we need domain knowledge to choose such a suitable class.
- Use examples to infer the particular function within this class.