

COMP30030: Introduction to Artificial Intelligence

Neil Hurley

School of Computer Science
University College Dublin
`neil.hurley@ucd.ie`

September 25, 2018

1 Problem Solving by Search

- Uninformed Search

- Informed Search

A Quick Review...

- ♦ Search techniques are used in AI to find a sequence of steps that will get us from some initial state to some goal state(s).
- ♦ You can use various algorithms to do the search. We have looked at simple breadth-first and depth-first search.
- ♦ Both are **systematic, exhaustive** search techniques that will eventually try all the nodes in a finite search tree.
- ♦ The appropriate algorithm will depend on the problem you are trying to solve, such as whether you are looking for the shortest path.

A Quick Review...

- ♦ **Are you looking for the shortest path?**
 - If so, breadth-first search may be better as it will find the shortest path first.
- ♦ **Is memory likely to be a problem?**
 - Depth-first search generally requires much less memory.
- ♦ **Do you want to find a solution quickly?**
 - If so, the choice of algorithm gets complex!
 - Depth-first search may be faster if there are many paths that lead to solution states, but all are quite long.
 - Breadth-first search may be faster if there is one short path to the target, but within a large and deep search space.

A Quick Review

- ♦ You may be so uninformed about a search problem that you cannot rule out either a large branching factor or the possibility of following very long useless paths.
- ♦ In such situations you may want to seek a **middle-ground between DFS and BFS**.
- ♦ \Rightarrow **Random queuing** – *nondeterministic search*.

A Quick Review...

- ♦ There are **many variants** of breadth and depth-first that may be useful in certain situations.
- ♦ You can set a **depth-limit** in depth-limited search so it backs up when it gets to nodes further than the specified distance from the initial state.
 - Like DFS, but the search is limited to a predefined depth.
 - The depth of each state is recorded as it is generated. When picking the next state to expand, only those with depth less or equal than the current depth are expanded.
 - Once all the nodes of a given depth are explored, the current depth is incremented.
- ♦ A variant to this is **iterative deepening** which repeats depth-first search with gradually increasing depth limits.

A Quick Review...

- ♦ In spite of the apparent waste in repeating bits of search iterative deepening is a very useful algorithm...

It **requires little memory** (like depth first),

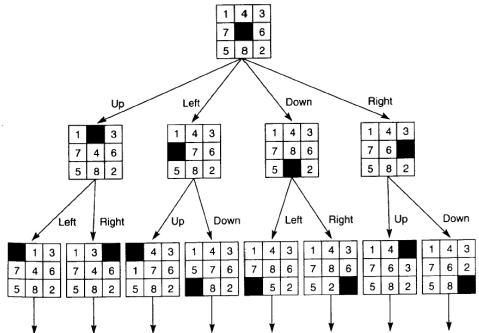
and

it **finds the shortest path first** (like breadth first).

Another Example

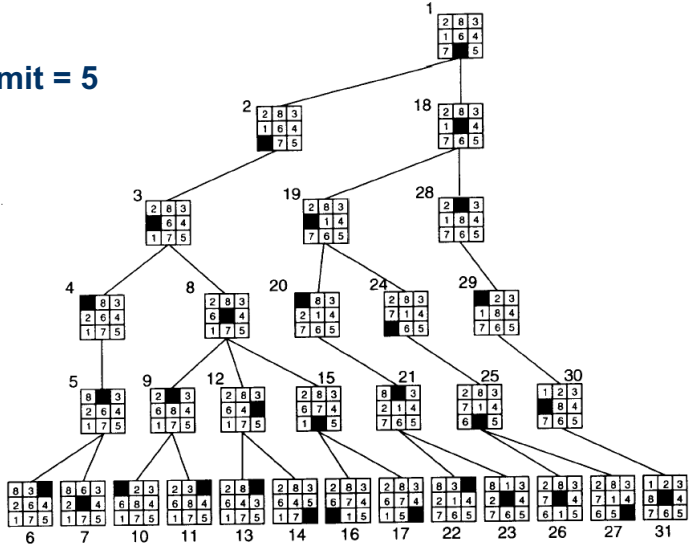
◆ The 8 Sliding Tile Puzzle

- Partial state space representation.
- Generated by the **MoveBlank** operator
 - up, down, left, right




8-Puzzle: Depth First Search

- ◆ Depth-limit = 5



8-Puzzle: Breadth-First Search



Goal


The TSP Example

◆ The Travelling Salesman Problem


- Find the shortest path
- Visit all cities
- Return home

◆ Representing the Problem

- Each arc is marked with the total weight of all paths from the start node (A) to its endpoint.



TSP – Partial State Space



1 Problem Solving by Search

- Uninformed Search

- Informed Search

Heuristic Search

- ◆ **Blind search methods** are simple but they are often **impractical**.
- ◆ They are **uninformed**. That is, they have no information about the state space and so are left to search blindly through it in a systematic fashion.
- ◆ **Heuristic search** methods, on the other hand, improve search efficiency by exploring **the most promising search paths first**.

Heuristic Search

To use heuristic search you need an **evaluation function** that scores a node in a search tree according to how close to the target/goal state it seems to be.

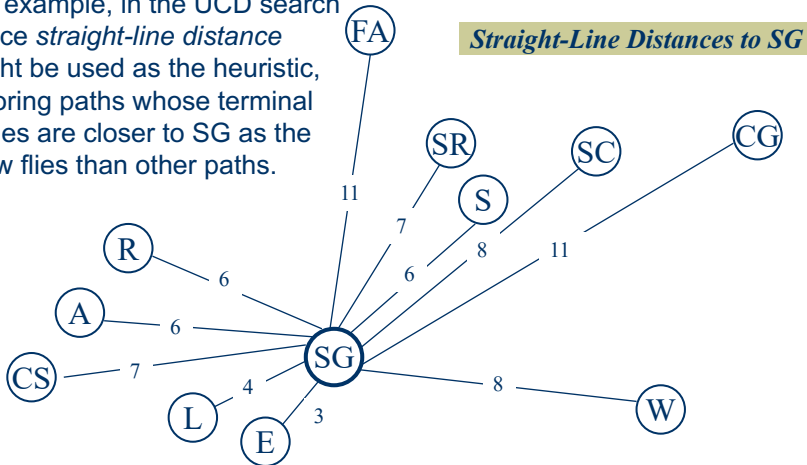
- ◆ NOTE: This will just be a **guess**, but it still should be a useful way of guiding the search.
- ◆ Basic idea...
 - Rather than trying all possible search paths, you try to focus on paths that seem to be getting you nearer your target/goal state.

Heuristic Function Properties

- ♦ They are **problem specific**.
- ♦ They indicate the **promise of states**, not actions.
- ♦ They are **estimates** of the cheapest path to the goal
- ♦ $h(n)=0$ if n is a goal node.

Heuristic Search

- ♦ For example, in the UCD search space *straight-line distance* might be used as the heuristic, favoring paths whose terminal nodes are closer to SG as the crow flies than other paths.




Heuristic Search


- ♦ In short, *heuristics* (heuristic evaluation functions) are simply just *rules of thumb*!
- ♦ They provide us with a means of ordering paths for exploration.
- ♦ They are not guaranteed to find one at all, but *if they are good*, then they help us find an adequate solution **faster!**

The 8-Puzzle – “Tiles-out-of-Place”

State n



Goal



♦ Possible Heuristic

- $h(n)$ = number of tiles out of place in this state relative to the goal state.

The 8-Puzzle – “Manhattan Distance”

State n

2	8	3
1	6	4
7		5

Goal


1	2	7
8		4
3	6	5

♦ **Another Heuristic**


- $h(n)$ = sum of **Manhattan distance** of each tile from its correct location.
= sum over all tiles of *number of rows and number of columns* that tile is from correct position

The 8-Puzzle – “Manhattan Distance

State n




Goal




= 1

The 8-Puzzle – “Manhattan Distance

State n




Goal




$= 1 + 1$

The 8-Puzzle – “Manhattan Distance

State n




Goal




$$= 1 + 1 + 4$$

The 8-Puzzle – “Manhattan Distance

State n




Goal




$$= 1 + 1 + 4 + 0$$

The 8-Puzzle – “Manhattan Distance

State n



Goal



$$= 1 + 1 + 4 + 0 + 0$$

The 8-Puzzle – “Manhattan Distance

State n

2	8	3
1	(1,1)	4
7	(2,1)	5


Goal

1	2	7
8		4
3	6	5


$$= 1 + 1 + 4 + 0 + 0 + 1$$

The 8-Puzzle – “Manhattan Distance

State n




Goal




$$= 1 + 1 + 4 + 0 + 0 + 1 + 4$$

The 8-Puzzle – “Manhattan Distance

State n



Goal



$$= 1 + 1 + 4 + 0 + 0 + 1 + 4 + 2 = 13$$

Admissible Heuristics I

- Notice that the two heuristics for the 8-puzzle are both **optimistic**.
- They *underestimate* the true distance to the goal – we must make *at least* the number of moves indicated by the heuristic.
- Such heuristics are called **admissible**.
- One way to arrive at such heuristics is to consider how to **relax** the original.
- **Relaxation** – removing some of the constraints of the original problem. With fewer constraints, the cost of the optimal solution for the relaxed problem will be at most equal to the cost of the optimal solution to the original problem.

Admissible Heuristics II

- For instance, we can arrive at the “Tiles out of Place” heuristic, by imagining a relaxed version of the 8-puzzle problem where a move consists of *any* two tiles swapping position, rather than just the blank and an adjacent tile.

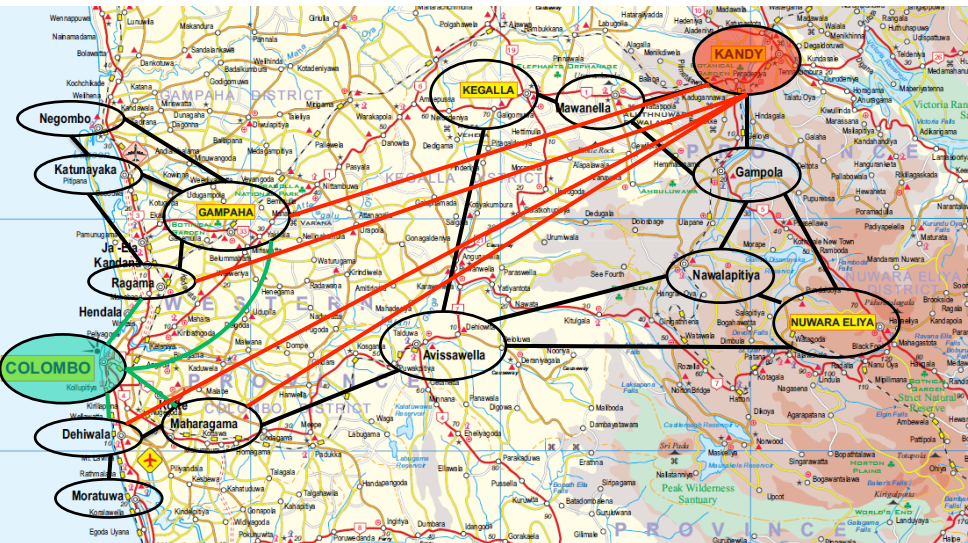
Greedy Best First Search I

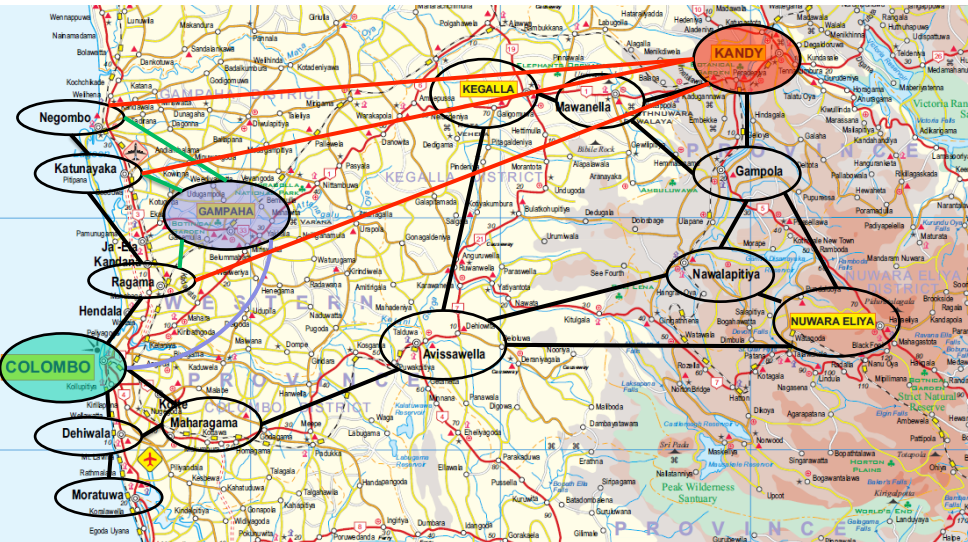
- In **best-first search**, a node, n , is chosen for expansion based on an evaluation function $f(n)$.
- $f(n)$ may depend on the description of the node, the description of the goal, the information gathered by the search up to that point, or any extra knowledge about the problem domain.
- A common approach is to use a **heuristic** that attempts to predict how close the end of a path is to a solution, so that paths which are judged to be closer to a solution are extended first.
- In this case, we will write $h(n)$ for the evaluation function.
- When the evaluation function is a heuristic measure of distance to the goal, the technique is referred to a greedy best first.

Greedy Best First Search II

- In reality we aren't expanding the best node first. We are expanding the one that appears to be the best.
- If we knew the best node, we would have the solution and there would not be any search.
- It resembles depth-first because it plunges in depth and ignores breadth.
- From the point-of-view of the queuing function, we use a **priority queue** for the open list, with the priority given by $h(n)$.

For example, route finding in Sri Lanka . . .





Greedy Best First Search I

- Best-first tells us to go from Colombo to Gampaha
- But there is no route from Gampaha, through its neighbours to Kandy).
- So, what seems the best, may not always be the best – sometimes we might run into dead ends.
- It will back up if it hits a dead end, but only a strategy to avoid repeated states will prevent it choosing Gampaha again.
- So, in general, it is not complete.
- Its performance depends crucially on the heuristic function.
- In our route-finding example, if there's a route from Gampaha to Kegalla, then it will work well...

