

Surprisingly Rational: Probability Theory Plus Noise Explains Biases in Judgment

Fintan Costello
University College Dublin

Paul Watts
National University of Ireland

The systematic biases seen in people's probability judgments are typically taken as evidence that people do not use the rules of probability theory when reasoning about probability but instead use heuristics, which sometimes yield reasonable judgments and sometimes yield systematic biases. This view has had a major impact in economics, law, medicine, and other fields; indeed, the idea that people cannot reason with probabilities has become a truism. We present a simple alternative to this view, where people reason about probability according to probability theory but are subject to random variation or noise in the reasoning process. In this account the effect of noise is canceled for some probabilistic expressions. Analyzing data from 2 experiments, we find that, for these expressions, people's probability judgments are strikingly close to those required by probability theory. For other expressions, this account produces systematic deviations in probability estimates. These deviations explain 4 reliable biases in human probabilistic reasoning (conservatism, subadditivity, conjunction, and disjunction fallacies). These results suggest that people's probability judgments embody the rules of probability theory and that biases in those judgments are due to the effects of random noise.

Keywords: probability, rationality, random variation, heuristics, biases

The capacity to reason with uncertain knowledge (i.e., to reason with probabilities) is central to our ability to survive and prosper in “an ecology that is of essence only partly accessible to foresight” (Brunswik, 1955, p. 236). It is therefore reasonable to expect that humans, having prospered in such an ecology, would be able to reason about probabilities extremely well; any ancestors who could not reason effectively about probabilities would not survive long, and so the biological basis of their reasoning would be driven from the gene pool. Probability theory provides a calculus of chance describing how to make optimal predictions under uncertainty. Taking the argument one step further, it is reasonable to expect that our probabilistic reasoning will follow the rules of probability theory.

The conventional view in current psychology is that this expectation is wrong. Instead, the dominant position is that

in making predictions and judgments under uncertainty, people do not appear to follow the calculus of chance or the statistical theory of prediction. Instead they rely on a limited number of heuristics which sometimes yield reasonable judgments and sometimes lead to severe and systematic errors. (Kahneman & Tversky, 1973, p. 237)

This conclusion is based on a series of systematic and reliable biases in people's judgments of probability, many identified in the 1970s and 1980s by Tversky, Kahneman, and colleagues. This heuristics and biases approach has reached a level of popularity rarely seen in psychology (with Kahneman receiving a Nobel Prize in part for his work in this area). The idea that people do not reason using probability theory but instead follow various heuristics has been presented both in review articles describing current psychological research (Gigerenzer & Gaissmaier, 2011; Shafir & Leboeuf, 2002) and in numerous popular science books summarizing this research for the general public (e.g., Ariely, 2009; Kahneman, 2011). This approach has had a major impact in economics (Camerer, Loewenstein, & Wallsten, 2003; Kahneman, 2003), law (Korobkin & Ulen, 2000; Sunstein, 2000), medicine (Dawson & Arkes, 1987; Eva & Norman, 2005), and other fields (De Bondt & Thaler, 1985; Hicks & Klumper, 2011; Richards, 2012; Williams, 2010). Indeed, the idea that people cannot reason with probabilities has become widespread: For example, the Science Gallery in Dublin recently presented an exhibition on risk, which it described as “enabling visitors to explore our inability to determine the probability of everything from a car crash to a coin toss” (Griffin, 2012).

We have two main aims in this article: to give evidence against the view that people reason about probabilities using heuristics and to give evidence supporting the view that people reason in accordance with probability theory, with bias in people's probability estimates being caused by random variation or noise in the reasoning process. We assume a simple model where people estimate the probability of some event *A* by estimating the proportion of instances of *A* in memory but are subject to random errors in the recall of instances. Although at first glance it may seem that these random errors will result in “nothing more than error variance

Fintan Costello, School of Computer Science and Informatics, University College Dublin; Paul Watts, Department of Mathematical Physics, National University of Ireland.

We would like to thank Mark Keane and John Wixted for their extremely helpful and insightful comments on earlier versions of this article.

Correspondence concerning this article should be addressed to Fintan Costello, School of Computer Science and Informatics, University College Dublin, Belfield, Dublin 4, Ireland. E-mail: fintan.costello@ucd.ie

centered around a normative response” (Shafir & Leboeuf, 2002, p. 501), in fact these random errors cause systematic deviations that push estimates for $P(A)$ away from the correct value in a characteristic way. In our model these systematic deviations explain various biases frequently seen in people’s probabilistic reasoning: conservatism, subadditivity, the conjunction fallacy, and the disjunction fallacy. The general patterns of occurrence of these biases match the predictions of our simple model.

We use this simple model to construct probabilistic expressions that cancel the bias in estimates for one event against the bias in estimates for another. These expressions allow us to test the predictions of the heuristics view of probabilistic reasoning. One such expression involves estimates, for some events A and B , of the individual probabilities $P(A)$ and $P(B)$ and the conjunctive (“and”) and disjunctive (“or”) probabilities $P(A \wedge B)$ and $P(A \vee B)$. People’s estimates for all four of these probabilities are typically subject to various forms of bias. Our account, however, predicts that when combined in the expression

$$X_E(A, B) = P_E(A) + P_E(B) - P_E(A \wedge B) - P_E(A \vee B)$$

(where $P_E(A)$ represents a person’s estimate for $P(A)$, $P_E(B)$ the person’s estimate for $P(B)$, and so on), the various biases on the individual expressions will cancel out on average $X_E(A, B)$ will equal 0 in agreement with probability theory’s addition law, which requires that

$$X(A, B) = P(A) + P(B) - P(A \wedge B) - P(A \vee B) = 0. \quad (1)$$

Notice that the heuristics view assumes that people estimate probabilities using heuristics that in some cases yield reasonable judgments (i.e., judgments in accordance with probability theory) but in other cases lead to systematic biases. To give evidence against the heuristics view it is therefore not enough to show that some of people’s probability judgments agree with probability theory (that is expected in the heuristics view). Instead, our evidence against the heuristics view consists of results showing that, even when people’s probability estimates for a set of events are systematically biased, the results are on average strikingly close to those required by probability theory when those estimates are combined to form expressions like X_E . This cancellation of bias is difficult to explain in the heuristics view: To explain this cancellation, the heuristics view would require some way of ensuring that, when heuristics are applied to estimate the probabilities $P(A)$, $P(B)$, $P(A \wedge B)$, and $P(A \vee B)$ individually, the biases produced in those four probabilities are precisely calibrated to give overall cancellation. Further, to “know” that the bias in these four probabilities should cancel out in this way requires access to the rules of probability theory (as embodied in the addition law in this case). Because the heuristics view by definition does not follow the rules of probability theory, it does not have access to these rules and so has no reason to produce this cancellation.

We also use this model to construct a series of expressions where all but one unit of bias is canceled; our model predicts that the level of bias when people’s responses are combined in these expressions should on average have the same constant value. Our experimental results confirm this prediction, showing the same level of bias across a range of such expressions. Together, these results demonstrate that when noise in recall is canceled, people’s probability estimates follow the rules of probability theory. Thus,

they suggest that biases in those estimates are due to noise. These results are the main contribution of our work.

Note that our evidence against the view that people use heuristics to estimate probabilities is not based on the fact that our model explains the four biases mentioned above (there are many other biases in the literature that our model does not address; see Hilbert, 2012, for a review). Instead, the point is that our experimental results show that the basic idea behind the heuristics view (that people do not follow the rules of probability theory) is contradicted when we use our simple model to cancel the effects of noise.

Bayesian Models of Reasoning

We are not alone in arguing that people reason in accordance with probability theory. Though “the bulk of the literature on adult human reasoning” goes against this view (Cesana-Arlotti, Téglás, & Bonatti, 2012, p. 3), in recent years various groups of researchers have suggested that people follow mathematical models of reasoning based on Bayesian inference, a process for drawing conclusions given observed data in a way that follows probability theory. Bayesian inference applies to conditional probabilities such as the probability of some conclusion H given some evidence E : $P(H|E)$. In Bayesian models these conditional probabilities are computed according to Bayes’ theorem

$$P(H|E) = \frac{P(H) \times P(E|H)}{P(E)},$$

and so the value of the conditional probability $P(H|E)$ depends on the value of the prior $P(H)$ (the probability of H being true independent of the evidence E) and on the value of the likelihood function $P(E|H)$ (the probability of seeing evidence E given that the hypothesis H is true).

The status of these Bayesian models is currently controversial. On one hand, close fits between human responses and Bayesian models have been demonstrated in domains as diverse as categorization, naive physics, word learning, vision, logical inference, motor control, and conditioning (see, e.g., Chater, Tenenbaum, & Yuille, 2006; Oaksford & Chater, 2007; Tenenbaum, Kemp, Griffiths, & Goodman, 2011), leading researchers to conclude that “everyday cognitive judgments follow [the] optimal statistical principles” of probability theory (Griffiths & Tenenbaum, 2006, p. 767). On the other hand, critics have pointed out a range of problems with this Bayesian approach (Bowers & Davis, 2012; Eberhardt & Danks, 2011; Endress, 2013; Jones & Love, 2011; Marcus & Davis, 2013). For example, the estimation of priors and likelihood functions in Bayesian models is problematic: There are “too many arbitrary ways that priors, likelihoods . . . etc. can be altered in a Bayesian theory post hoc. This flexibility allows these models to account for almost any pattern of results” (Bowers & Davis, 2012, p. 394).

It is important to stress that our approach is not connected to this Bayesian view. Our model applies only to the estimation of simple probabilities, such as the probability of some event $P(A)$, and does not involve Bayes’ theorem or conditional probabilities of any form. Neither does our model involve parameter estimation, priors, or likelihood functions. Equally, our results showing that people’s probability estimates follow the requirements of probability theory when noise is canceled do not imply that people follow Bayes’

theorem when estimating conditional probabilities: Bayes' theorem is significantly more complex than the simple probabilities we consider.

Overview

In the first section of the article, we present our model and show how it can explain the observed biases of conservatism, subadditivity, the conjunction fallacy, and the disjunction fallacy. In this section we also discuss other accounts for these biases, some of which are also based on noise (see Costello, 2009a, 2009b; Dougherty, Gettys, & Ogden, 1999; Erev, Wallsten, & Budescu, 1994; Hilbert, 2012; Juslin, Nilsson, & Winman, 2009; Nilsson, Winman, Juslin, & Hansson, 2009). The crucial difference between our account and others is that our account makes specific and testable predictions about the degree of bias in probabilistic expressions and about expressions where that bias will vanish. In the second and third sections we present our model's predictions and describe two experimental studies testing and confirming these predictions. In the final sections we give a general discussion of our work.

Probability Estimation With Noisy Recall

We assume a rational reasoner with a long-term episodic memory that is subject to random variation or error in recall and take $P_E(A)$ to represent a reasoner's estimate of the probability of event A . We assume that long-term memory contains m episodes, where each recorded episode i holds a flag that is set to 1 if i contains event A and set to 0 otherwise, and the reasoner estimates the probability of event A by counting these flags.

We assume a minimal form of transient random noise, in which there is some small probability d that when some flag is read, the value obtained is not the correct value for that flag. We assume that this noise is symmetric, so that the probability of 1 being read as 0 is the same as the probability of 0 being read as 1. We also assume a minimal representation where every type of event, be it a simple event A , a conjunctive event $A \wedge B$, a disjunctive event $A \vee B$, or any other more complex form, is represented by such a flag, and where every flag has the same probability d of being read incorrectly. (We stress here that this type of sampling error is only one of many possible sources of noise. Although we use this simple form of sampling error to motivate and present our model, our intention is to demonstrate the role of noise—from whatever source—in causing systematic biases in probability estimates.)

We take $C(A)$ to be the number of flags marking A that were read as 1 in some particular query of memory and T_A be the number of flags whose correct value is actually 1. Our reasoner computes an estimate $P_E(A)$ by querying episodic memory to count all episodes containing A and dividing by the total number of episodes, giving

$$P_E(A) = \frac{C(A)}{m}. \quad (2)$$

Random error in recall and, hence, in the value of $C(A)$ means that $P_E(A)$ varies randomly. Sampling $P_E(A)$ repeatedly will produce a series of different values, varying due to error in recall. We assume that this estimation process is the same for every form of event: A probability estimate for a simple event A is computed

from the number of flags marking A that were read as 1 in some particular query of memory, a probability estimate for a conjunctive event $A \wedge B$ is computed from number of flags marking $A \wedge B$ that were read as 1 in some particular query of memory, and so on.

We take $P(A)$ to represent the "true" judgment of the probability of A : the estimate that would be given if the reasoner was not subject to random error in recall and produced estimates in a perfect, error-free manner. We take $\overline{P_E}(A)$ to represent the expected value or population mean for $P_E(A)$. This is the value we would expect to obtain if we averaged an infinite number of these randomly varying estimates $P_E(A)$. Finally, we take $\widehat{P_E}(A)$ to represent a sample mean: the average of some finite set of estimates $P_E(A)$. This sample mean $\widehat{P_E}(A)$ will vary randomly around the population mean $\overline{P_E}(A)$, with the degree of random variation in the sample mean decreasing as the size of the sample increases.

For any event A , the expected value of $P_E(A)$ is given by

$$\overline{P_E}(A) = \frac{T_A(1-d) + (m-T_A)d}{m}$$

(as on average $1-d$ of the T_A flags whose value is 1 will be read as 1, and d of the $m-T_A$ flags whose value is 0 will be read as 1). Because by definition

$$P(A) = \frac{T_A}{m},$$

we have

$$\overline{P_E}(A) = P(A) + d - 2dP(A), \quad (3)$$

and the expected value of $P_E(A)$ deviates from $P(A)$ in a way that systematically depends on $P(A)$.

Individual estimates will vary randomly around this expected value, and so for any specific estimate $P_E(A)$, where $C(A)$ flags were read as having a value of 1, we have

$$P_E(A) = \overline{P_E}(A) + e, \quad (4)$$

where

$$e = \frac{C(A) - T_A(1-d) - (m-T_A)d}{m}$$

represents positive or negative random deviation from the expected value across all estimates. Note that this error term e does not introduce an additional source of random error in probability estimates: It simply reflects the difference between the number of flags that were read incorrectly when computing the specific estimate $P_E(A)$ and the number of flags that are read incorrectly on average, across all estimates.

Finally, we can also derive an expression for the expected variance in these randomly varying estimates $P_E(A)$. The expected variance is equal to $d(1-d)/m$ for all events A , $A \wedge B$, $A \vee B$, and so on, and is independent of event probability; see the Appendix for details.

Conservatism

In this section we show how our noisy recall model of probabilistic reasoning explains a reliable pattern of conservatism seen in people's probability estimates.

Probabilities range in value between 0 and 1. A large body of literature demonstrates that people tend to keep away from these extremes in their probability judgments and so are conservative in their probability assessments. These results show that the closer $P(A)$ is to 0, the more likely it is that $P_E(A)$ is greater than $P(A)$, and the closer $P(A)$ is to 1, the more likely it is that $P_E(A)$ is less than $P(A)$. Differences between true and estimated probabilities are low when $P(A)$ is close to 0.5 and increase as $P(A)$ approaches the boundaries of 0 or 1. This pattern was originally seen in research on people's revision of their probability estimates in the light of further data (Edwards, 1968), and it was later found directly in probability estimation tasks. This pattern is sometimes referred to as underconfidence in people's probability estimates (for a review, see Erev et al., 1994; Hilbert, 2012).

Conservatism will occur as a straightforward consequence of random variation in our model. As we saw in Equation 3, the expected value of $P_E(A)$ deviates from $P(A)$ in a way that systematically depends on $P(A)$. If $P(A) = 0.5$, this deviation will be 0. If $P(A) < 0.5$, because d cannot be negative we have $\overline{P_E}(A) > P(A)$, with the difference increasing as $P(A)$ approaches 0. Because estimates $P_E(A)$ are distributed around $\overline{P_E}(A)$, this means that $P_E(A)$ will tend to be greater than $P(A)$, with the tendency increasing as $P(A)$ approaches 0. Similarly, if $P(A) > 0.5$ then $\overline{P_E}(A) < P(A)$ and estimates $P_E(A)$ will tend to be less than $P(A)$, with the tendency increasing as $P(A)$ approaches 1. This deviation thus matches the pattern of conservatism seen in people's probability judgments.

Other accounts. The idea that conservatism can be explained via random noise is not new to our account. It is also found in Erev et al.'s (1994) account based on random error in probability estimates, in the Minerva-DM memory-retrieval model of decision making (Dougherty et al., 1999), and in Hilbert's account based on noise in the information channels used in probabilistic reasoning (Hilbert, 2012). The underlying idea in these accounts is similar to ours. There is, however, a critical difference: Our account predicts no systematic bias for probabilistic expressions with a certain form (see Experiment 1).

Subadditivity

Here we show how our noisy recall model explains various patterns of subadditivity seen in people's probability estimates.

Let $A_1 \dots A_n$ be a set of n mutually exclusive events, and let $A = A_1 \vee \dots \vee A_n$ be the disjunction (the *or*) of those n events. Then, probability theory requires that

$$\sum_{i=1}^n P(A_i) = P(A).$$

Experimental results show that people reliably violate this requirement and do so in a characteristic way. On average, the sum of people's probability estimates for events $A_1 \dots A_n$ is reliably greater than their estimate for the probability of A , with the difference (the degree of subadditivity) increasing reliably as n increases. An additional, more specific pattern is also seen: For pairs of mutually exclusive events A_1 and A_2 whose probabilities sum to 1, the sums of people's estimates for A_1 and A_2 are normally distributed around 1. On average this sum is equal to 1, just as required by probability theory. This pattern is sometimes

referred to as binary complementarity (see Tversky & Koehler, 1994, for a detailed review of these results).

Again, these patterns of subadditivity occur as a straightforward consequence of random variation in our model. From Equation 3 we have

$$\sum_{i=1}^n \overline{P_E}(A_i) = \sum_{i=1}^n [P(A_i) + d - 2dP(A_i)],$$

and using the fact that $P(A_1) + \dots + P(A_n) = P(A)$ this gives

$$\sum_{i=1}^n \overline{P_E}(A_i) = P(A) + nd - 2dP(A).$$

Taking the difference between this expression and that for $\overline{P_E}(A)$ in Equation 3 we get

$$\sum_{i=1}^n \overline{P_E}(A_i) - \overline{P_E}(A) = (n-1)d,$$

and so this difference increases as n increases, producing subadditivity as seen in people's probability judgments. In the case of two mutually exclusive events A_1 and A_2 whose probabilities sum to 1, from Equation 3 we get

$$\overline{P_E}(A_1) + \overline{P_E}(A_2) = P(A_1) + d - 2dP(A_1) + P(A_2) + d - 2dP(A_2) = 1,$$

producing binary complementarity as seen in people's judgments.

Other accounts. The original account for subadditivity given by Tversky and Koehler (1994) explained the general pattern in terms of an unpacking process, which increased the probability of constituent events by drawing attention to their components. This account could not explain the observed pattern of binary complementarity; to account for this observation Tversky and Koehler proposed an additional binary complementarity heuristic, which simply stated that there was no average bias for binary complements.

An alternative explanation for subadditivity is given in the Minerva-DM memory retrieval model of decision making (Bearden & Wallsten, 2004; Dougherty et al., 1999). Minerva-DM is a complex model with a number of components. It provides a two-step process for conditional probability judgments, a parameter controlling the retrieval of items with varying degrees of similarity to the memory probe (the event whose probability is being judged), a complex multivector representation for stored items in memory, a parameter controlling the degree of random error in the initial recording of items in memory, a parameter controlling the degree of random error causing degradation in stored items, and a parameter controlling the degree of detail contained in memory probes. Roughly stated, the Minerva-DM model estimates the probability of some event by counting the number of stored items in memory that are similar enough to that event (whose similarity measure is greater than the similarity criterion parameter). Depending on the value of the similarity criterion, this count will include a number of similar-but-irrelevant items in addition to items correctly matching the target event. Because of these similar-but-irrelevant items, the model will give a probability estimate for the target event that is higher than the true probability, producing a degree of subadditivity that increases with the number of component events in the disjunction, just as required. Note, however, that because this similarity-based account

always increases probability estimates, it cannot explain the observed pattern of binary complementarity in people's probability judgments, which can only be explained if one probability is increased and the other complementary probability is decreased.

More recently, Hilbert (2012) gave an account of subadditivity based on noise in the information channels used for probability computation. Hilbert's model is a very general one, providing for noise at the initial encoding of data, for noisy degradation of stored information, and for noise during the reading of data from memory. The model also specifies three general requirements for the distribution of noise: that the correct probability is most likely, that noise is symmetrical around the correct probability, and that two binary complementary probabilities have the same degree of noise. The last of these requirements allows the model to explain the binary complementarity result observed by Tversky and Koehler (1994). Beyond these requirements, the model leaves the degree and form of noise in the system unspecified. Again, this account is similar to ours but with the crucial difference that our account predicts no systematic bias for certain probabilistic expressions. We give a further comparison between our model, Hilbert's model, and Minerva-DM in the General Discussion.

Conjunction and Disjunction Fallacies

Conservatism and subadditivity both concern averages of people's probability estimates. Here we show how our noisy recall model explains two patterns that involve differences between individual probability estimates: the conjunction and disjunction fallacies.

Let A and B be any two events ordered so that $P(A) \leq P(B)$. Then probability theory's conjunction rule requires that $P(A \wedge B) \leq P(A)$; this follows from the fact that $A \wedge B$ can only occur if A itself occurs. People reliably violate this requirement for some events and commit the conjunction fallacy by giving probability estimates for conjunctions that are greater than the estimates they gave for one or other constituent of that conjunction. Perhaps the best known example of this violation comes from Tversky and Kahneman (1983) and concerns Linda:

Linda is 31 years old, single, outspoken, and very bright. She majored in philosophy. As a student, she was deeply concerned with issues of discrimination and social justice, and also participated in anti-nuclear demonstrations. (p. 297)

Participants in Tversky and Kahneman's study read this description and were asked to rank various statements by their probability. Two of these statements were

Linda is a bank teller. (A)

Linda is a bank teller and active in the feminist movement. ($A \wedge B$)

In Tversky and Kahneman's initial investigation these two statements were presented separately, with one group of participants ranking a set of statements containing A but not $A \wedge B$ and a second group ranking the same set but with A replaced by $A \wedge B$. The results showed that the average ranking given to $A \wedge B$ by the second group was significantly higher than the average ranking given to A by the first group, violating the conjunction rule.

Note that violation of the conjunction rule can occur in averaged data even when very few participants are individually committing

the conjunction fallacy; equally, this violation is not necessarily seen in averaged data even when many participants are individually committing that fallacy. For this reason, Tversky and Kahneman (1983) referred to this comparison of averages as an "indirect" test of the conjunction rule and described violations of that rule in averages as conjunction errors. Surprised by the results of their indirect test, Tversky and Kahneman carried out a series of increasingly direct tests of the conjunction rule. In these direct tests each participant were asked to rank the probability of a set of statements containing both A and $A \wedge B$. Tversky and Kahneman found that in some cases more than 80% of participants ranked $A \wedge B$ as more probable than A , violating the conjunction rule in their individual responses. Tversky and Kahneman used the term *conjunction fallacy* to refer only to these direct violations of the conjunction rule. Most subsequent studies have focused on similar direct tests of the conjunction rule in individual probability estimates (the conjunction fallacy) rather than on indirect tests comparing averages (the conjunction error).

The Linda example is explicitly designed to produce the conjunction fallacy; this fallacy does not occur for all or even most conjunctions. Numerous experimental studies have shown that the occurrence of this fallacy depends on the probabilities of A and B . In particular, the greater the difference between $P(A)$ and $P(B)$, the more frequent the conjunction fallacy, and the greater the conditional probability $P(A|B)$, the more frequent the conjunction fallacy (Costello, 2009a; Fantino, Kulik, Stolarz-Fantino, & Wright, 1997; Gavanski & Roskos-Ewoldsen, 1991).

A similar pattern occurs for people's probability estimates for disjunctions $A \vee B$. Because $A \vee B$ necessarily occurs if B itself occurs, probability theory requires that $P(A \vee B) \geq P(B)$ must always hold. People reliably violate this requirement for some events, giving probability estimates for disjunctions that are less than the estimates they gave for just one of the constituents. Just as for conjunctions, the greater the difference between $P(A)$ and $P(B)$ and the higher the estimated conditional probability $P(A|B)$, the higher the rate of occurrence of the disjunction fallacy (Carlson & Yates, 1989; Costello, 2009b).

The observed patterns of conjunction and disjunction fallacy occurrence arise as a straightforward consequence of random variation in our model. The general idea is that our reasoner's probability estimates $P_E(A)$ and $P_E(A \wedge B)$ will both vary randomly around their expected values $\overline{P}_E(A)$ and $\overline{P}_E(A \wedge B)$. This means that, even though $\overline{P}_E(A \wedge B) \leq \overline{P}_E(A)$ must hold, there is a chance that an estimate for $A \wedge B$ will be greater than an estimate for A , producing a conjunction fallacy. This chance will increase the closer $\overline{P}_E(A \wedge B)$ is to $\overline{P}_E(A)$.

More formally, the reasoner's estimates for probabilities $P_E(A)$ and $P_E(A \wedge B)$ at any given moment are given by

$$P_E(A) = \overline{P}_E(A) + e_A \quad \text{and} \quad P_E(A \wedge B) = \overline{P}_E(A \wedge B) + e_{A \wedge B},$$

where e_A and $e_{A \wedge B}$ represent positive or negative random deviation from the expected estimate at that time (arising due to random errors in reading flag values from memory, as in Equation 4). The conjunction fallacy will occur when $P_E(A) < P_E(A \wedge B)$; that is, when

$$\overline{P}_E(A) + e_A < \overline{P}_E(A \wedge B) + e_{A \wedge B},$$

or, substituting and rearranging, when

$$[P(A) - P(A \wedge B)](1 - 2d) < e_{A \wedge B} - e_A \quad (5)$$

holds. Given that $e_{A \wedge B}$ and e_A vary randomly and can be either positive or negative, this inequality can hold in some cases. The inequality is most likely to hold when $P(A) - P(A \wedge B)$ is low (as in that situation the left side of the inequality is low). Because $P(A \wedge B) = P(A|B)P(B)$, we see that $P(A) - P(A \wedge B)$ is low when $P(A)$ is low and both $P(A|B)$ and $P(B)$ are high (or more strictly, when $P(A)$ is close to 0, $P(B)$ is close to 1 and $P(A|B)$ is close to its maximum possible value of $P(A)/P(B)$). We thus expect the conjunction fallacy to be most frequent when $P(A)$ is low and $P(A|B)$ and $P(B)$ are both high. This is just the pattern seen when the conjunction fallacy occurs in people's probability estimates.

Reasoning in the same way for disjunctions, we see that the disjunction fallacy will occur when

$$\overline{P_E}(B) + e_B > \overline{P_E}(A \vee B) + e_{A \vee B},$$

or, substituting and rearranging as before, when

$$[P(A \vee B) - P(B)](1 - 2d) < e_B - e_{A \vee B}$$

holds. But from probability theory, we have the identity

$$P(A \vee B) - P(B) = P(A) - P(A \wedge B),$$

and substituting we see that the disjunction fallacy will occur when

$$[P(A) - P(A \wedge B)](1 - 2d) < e_B - e_{A \vee B}. \quad (6)$$

Just as with the conjunction fallacy, we expect the disjunction fallacy to be most frequent when $P(A)$ is low and $P(A|B)$ and $P(B)$ are both high. Again, this is the pattern seen when the disjunction fallacy occurs in people's probability estimates.

Note that in our model there is an upper limit on the expected rate of conjunction fallacy occurrence of 50%, which occurs when $\overline{P_E}(A \wedge B) = \overline{P_E}(A)$: Estimates $P_E(A)$ and $P_E(A \wedge B)$ are distributed around the same population means, and so the chance of getting $P_E(A) < P_E(A \wedge B)$ is the same as the chance of getting $P_E(A) > P_E(A \wedge B)$. The same limit holds for the disjunction fallacy and for the same reason. This limit occurs because our simple model assumes (somewhat unrealistically) that the degree of error in recall for examples of A is the same as the degree of error in recall for $A \wedge B$ (extensions of our model that allow for different levels of error in recall for A and $A \wedge B$ would not impose this limit). As we see in the next section, however, experiments that control for various extraneous factors typically give conjunction fallacy rates that are consistent with this limit.

The Reality of the Conjunction Fallacy

A number of researchers have attempted to explain away the conjunction fallacy by pointing to possible flaws in Tversky and Kahneman's (1983) Linda experiment that may have led participants to give incorrect responses. One argument in this line is to propose that the fallacy arises because participants in the experiments understand the word *probability* or the word *and* in a way different from that assumed by the experimenters. A related tactic is to propose that the fallacy occurs because participants, correctly following the pragmatics of communication in their experimental task, interpret the single statement A as meaning $A \wedge \neg B$ (A and

not B). Evidence against these proposals comes from experiments using a betting paradigm, where the word probability is not mentioned and where *and* is demonstrably understood as meaning conjunction, and experiments where participants are asked to choose among three different options A , $A \wedge B$, and $A \wedge \neg B$. Conjunction fallacy rates are typically reduced in these experiments (to between 10% and 50%, as compared to the greater than 80% rate seen in Tversky and Kahneman's, 1983, Linda experiment) but remain reliable (see, e.g., Sides, Osherson, Bonini, & Viale, 2002; Tentori, Bonini, & Osherson, 2004; Wedell & Moro, 2008).

Another approach is to explain away the conjunction fallacy by arguing that Tversky and Kahneman's (1983) probabilistic ranking task is not in a form that is suitable for people's probabilistic reasoning mechanisms, which (in this argument) are based on representations of frequency. The suggestion is that if participants were asked to estimate the frequency with which constituent and conjunctive statements are true, the conjunction fallacy should vanish (Hertwig & Gigerenzer, 1999). In the Linda task, this frequency format could involve giving participants a story about a number of women who fit the description of Linda, then asking them to estimate "how many of these women are bank tellers" and "how many of these women are bank tellers and are active in the feminist movement." However, studies have repeatedly shown that although the occurrence of the conjunction fallacy declines in frequency format tasks (typically going from a rate greater than 80% in Linda tasks to a rate between 20% and 40% in frequency format tasks), the fallacy remains reliable and does not disappear (Fiedler, 1988; Hertwig & Chase, 1998; Mellers, Hertwig, & Kahneman, 2001). Indeed, Tversky and Kahneman's original 1983 paper examined the role of frequency formulations in the conjunction fallacy and found that the fallacy was reduced but not eliminated by that formulation.¹ Taken together, these results suggest that the conjunction fallacy rates of 80% and above found by Tversky and Kahneman are artificially high because of various confounding factors: Studies of the conjunction fallacy that eliminate these factors give fallacy rates that are generally around 50% or lower, in line with our model's expectations.

Other accounts. A large and diverse range of accounts has been proposed for the conjunction and disjunction fallacies. Tversky and Kahneman's original proposal explained these fallacies in terms of a representativeness heuristic, in which probability is assessed in terms of the degree to which an instance is representative of a (single or conjunctive) category. Under Tversky and Kahneman's interpretation, in the Linda example people gave a higher rating to the conjunctive statement because the instance Linda was more representative of (i.e., more similar to members of) the conjunctive category "bank-teller and active-feminist" than the single category "bank teller."

Although the representativeness heuristic remains the routine explanation of the conjunction fallacy in introductory textbooks, a number of experimental results give convincing evidence against this account. Notice that the representativeness heuristic applies

¹ Note, however, that conjunction errors (i.e., violations of the conjunction rule in indirect tests on averages rather than in direct tests on individual responses) can be eliminated by this frequency formulation (Mellers et al., 2001).

only when a question asks about the probability of membership of an instance in a conjunctive category and only when knowledge about representative members of that category is available. Evidence against representativeness comes from results showing that the conjunction fallacy occurs frequently when these requirements do not hold. For example, a series of studies by Osherson, Bonini, and colleagues have shown that the conjunction fallacy occurs frequently when people are asked to bet on the occurrence of unique future events. Such bets are not questions about membership of an instance in a category, and so representativeness cannot explain the occurrence of the fallacy in these cases (Bonini, Tentori, & Osherson, 2004; Sides et al., 2002; Tentori et al., 2004). Gavanski and Roskos-Ewoldsen (1991) found that the conjunction fallacy occurred frequently when people are asked about categories for which no representativeness information is available (questions about imaginary aliens on other planets). Gavanski and Roskos-Ewoldsen (1991) also found that the fallacy occurred frequently when the probability question was not about the membership of an instance in a conjunctive category but about the membership of two separate instances in two separate single categories (rather than asking about the probability of Linda being a bank teller and active feminist, such questions might ask about the probability of Bob being a bank teller and Linda being an active feminist). Again, representativeness cannot explain the occurrence of the conjunction fallacy for such questions (see Nilsson et al., 2009, for a review of research in this area).

The Minerva-DM account gives an alternative explanation for the conjunction fallacy that is based on the role of similarity in retrieval (Dougherty et al., 1999). Minerva-DM estimates the probability of some event by counting the number of stored items in memory whose similarity to the probe event is greater than the similarity criterion parameter. For a conjunction $A \wedge B$, stored items that are members of A alone or members of B alone can be similar enough to the conjunction $A \wedge B$ to be (mistakenly) counted as examples of that conjunction. If there are a large number of such similar-but-irrelevant items, the conjunctive probability estimate $P_E(A \wedge B)$ may be higher than the lower constituent probability $P_E(A)$, producing a conjunction fallacy response. Note, however, that because this similarity-based account always increases probability estimates, it cannot explain the disjunction fallacy (which occurs when a disjunctive probability estimate is lower than one of its constituent probabilities).

Other accounts have been proposed where people compute conjunctive probabilities $P_E(A \wedge B)$ from constituent probabilities $P_E(A)$ and $P_E(B)$ using some equation other than the standard equation of probability theory. In early versions of this approach, the conjunctive probability was taken to be the average of the two constituent probabilities (Carlson & Yates, 1989; Fantino et al., 1997). This averaging approach does not apply to disjunctive probabilities. More recently Nilsson, Juslin, and colleagues (Juslin et al., 2009; Nilsson et al., 2009) have proposed a more sophisticated configural cue model, where conjunctive probabilities are computed by a weighted average of constituent probability values, with a greater weight given to the lower constituent probability, and where disjunctive probabilities are computed by a weighted average, with greater weight given to the higher constituent probability.

Because the average of two numbers is always greater than the minimum of those two numbers and less than the maximum

(except when the numbers are equal), these averaging accounts predict that the conjunction fallacy will occur for almost every conjunction (except when the two constituents have equal probabilities). This is clearly not the case, however: There are many conjunctions for which these fallacies occur rarely if at all. To address this problem, Nilsson et al.'s model also includes a noise component that randomly perturbs conjunctive probability estimates, sometimes moving the conjunctive probability below the lower constituent probability and so eliminating the conjunction fallacy for that estimate. This model thus predicts that fallacy rates should be inversely related to the degree of random variation in people's probability judgments, with fallacy rates being highest when random variation is low and lowest when random variation is high. This contrasts with our account, which predicts that fallacy rates should be high when random variation is high and low when random variation is low. We assess these competing predictions in the Random Variation and Fallacy Rates Across Participants section.

Finally, we should mention an earlier model for the conjunction fallacy proposed by one of the authors (Costello, 2009a). Just as in our current account, this earlier model proposed that people's probability estimates followed probability theory but were subject to random variation. This random variation caused conjunction fallacy responses to occur when constituent and conjunctive probability estimates were close together. Apart from that commonality, the two models are quite different. Unlike our current account, this earlier model was not based on the idea of noise causing random errors in retrieval from memory; instead, that model assumed that estimates $P_E(A)$ for some event A were normally distributed around the correct value $P(A)$, and so the average estimate $\overline{P_E}(A)$ was equal to the true value $P(A)$. That earlier model was therefore unable to account for the patterns of conservatism and subadditivity seen in people's probability estimates. Also unlike our current account, that earlier model assumed that conjunctive and disjunctive probabilities were computed by applying the equations of probability theory to constituent probability estimates, so that $P_E(A \wedge B) = P_E(A) \times P_E(B|A)$. This contrasts with the current model, which computes $P_E(A \wedge B)$ by retrieving episodes of the event $A \wedge B$ from memory.

Experiment 1

Our noisy recall model of probability estimation can explain various patterns of bias in people's probability judgments and can also explain some specific situations in which those biases vanish (when probabilities are close to 0.5, for conservatism; and when two complementary probabilities sum to 1, for subadditivity). We now present a third situation in which our model predicts that bias will disappear.

Consider an experiment where we ask people to estimate, for any pair of events A and B , the probabilities of A , B , $A \wedge B$, and $A \vee B$. For each participant's estimates for each pair of events A and B , we can compute a derived sum

$$X_E(A, B) = P_E(A) + P_E(B) - P_E(A \wedge B) - P_E(A \vee B).$$

We can make a specific prediction about the expected value of $X_E(A, B)$ for all events A and B . This value will be

$$\overline{X_E}(A, B) = \overline{P_E}(A) + \overline{P_E}(B) - \overline{P_E}(A \wedge B) - \overline{P_E}(A \vee B).$$

From Equation 3 we get

$$\begin{aligned}\overline{X_E}(A, B) &= [P(A) + d - 2dP(A)] + [P(B) + d - 2dP(B)] \\ &\quad - [P(A \wedge B) + d - 2dP(A \wedge B)] \\ &\quad - [P(A \vee B) + d - 2dP(A \vee B)] \\ &= (1 - 2d)[P(A) + P(B) - P(A \wedge B) - P(A \vee B)].\end{aligned}$$

However, probability theory requires $P(A) + P(B) - P(A \wedge B) - P(A \vee B) = 0$ for all events A and B , and so $\overline{X_E}(A, B) = 0$. Our prediction, therefore, is that the average value of $X_E(A, B)$ across all pairs of events A and B will be equal to 0. Note that this prediction is invariant: It holds for all pairs of events A and B , irrespective of the degree of co-occurrence or dependency between those events.

What is the distribution of values of $X_E(A, B)$ around this average of 0? Speaking generally, we would expect this distribution to be unimodal and roughly symmetric around the mean of 0 for any pair A, B , because the positive and negative terms in the expression $X_E(A, B)$ are symmetric: $P(A) + P(B) = P(A \wedge B) + P(A \vee B)$.

We examined this expectation in detail via Monte Carlo simulation, by writing a program that simulates the effects of random noise in recall on probability estimations for a given set of probabilities. This program took as input three probabilities, $P(A)$, $P(A \wedge B)$, and $P(\neg A \wedge B)$ (*not A and B*). The program constructed a memory containing 100 items, each item containing flags $A, B, A \wedge B$, and $A \vee B$ indicating whether that item was an example of the given event. The occurrence of those flags in memory exactly matched the probabilities of the given event as specified by the three input probabilities. The occurrence of B , for example, matched the sum of input probabilities $P(A \wedge B)$ and $P(\neg A \wedge B)$. This program also contained a noise parameter d (set to 0.25 in our simulations); when reading flag values from memory to generate some probability estimate $P_E(A)$, the program was designed to have a random chance d of returning the incorrect value.

We carried out this simulation process for a representative set of values for the input probabilities $P(A)$, $P(A \wedge B)$, and $P(\neg A \wedge B)$. These sets consisted of every possible assignment of values from $\{0, 0.1, \dots, 0.9, 1.0\}$ to each input probability, subject to the requirement that both $P(A \wedge B) \leq P(A)$ and $P(\neg A \wedge B) \leq 1.0 - P(A)$ must hold. This requirement ensures that every set of input probabilities was consistent with the rules of probability theory. In total there were 286 sets of input probabilities that were consistent with these requirements. For each such set of input probabilities the program carried out 10,000 runs, on each run generating noisy estimates $P_E(A)$, $P_E(B)$, $P_E(A \wedge B)$, and $P_E(A \vee B)$ and using those estimates to calculate a value for the expression $X_E(A, B)$. These runs give us a picture of the distribution of values of $X_E(A, B)$.

The distribution of X_E values was essentially the same for all these sets of input probabilities: unimodal, approximately symmetric, and centered on 0, just as expected. Figure 1 graphs the frequency distributions of all X_E values across all probability sets. Given that this distribution appears to be essentially independent of the probability values used in our simulations, our prediction is that, in an experiment, the distribution of $X_E(A, B)$ across all pairs of events A and B will be unimodal and approximately symmetric around the mean of 0.




Figure 1. Frequency of different values of $X_E(A, B)$ in the Monte Carlo simulations. This scatterplot shows the raw frequency of occurrence of different values of $X_E(A, B)$ as observed in the Monte Carlo simulations, across runs for a range of different probability values $P(A)$, $P(B)$, $P(A \wedge B)$, and $P(A \vee B)$. (There were 286 sets of probability values, with 10,000 X_E values calculated for each set.) The critical point here is that the distribution of these values is essentially the same across different probability values: unimodal and symmetrically distributed around 0. A -1 to 1 probability scale is used here. Note that Figures 2–4 use the 100-point rating scale from the experiments.

One possible concern with this simulation comes from the common observation that, when estimating probabilities, participants tend to respond in units that are multiples of 0.05 or 0.10 (Budescu, Weinberg, & Wallsten, 1988; Erev et al., 1994; Wallsten, Budescu, Zwick, & Kemp, 1993). To test the impact of this rounding, we modified our simulation program to include a rounding parameter u such that each calculated probability estimate was rounded to the nearest unit of u . We ran the simulation as before but with u set to 0.01, 0.05, and 0.1, and examined the distribution of X_E values for each run. The distribution was essentially the same as that shown in Figure 1, confirming our original simulation results.

Testing the Predictions

We tested these predictions using data from an experiment on conjunction and disjunction fallacies (Experiment 3 in Costello, 2009b). The original aim of this experiment was to examine an attempt by Gigerenzer to explain away the conjunction fallacy as a consequence of people being asked to judge the probability of one-off, unique events (Gigerenzer, 1994). Gigerenzer argued that from a frequentist standpoint the rules of probability theory apply only to repeated events and not to unique events, and so people's deviation from the rules of probability theory for unique events are not, in fact, fallacious. To assess this argument, we examined the occurrence of fallacies in probability judgments for conjunctions and disjunctions of canonical repeatedly occurring events: weather events such as rain, wind, and so on. Contrary to Gigerenzer's argument, participants in these experiments often committed conjunction and disjunction fallacies; these fallacies thus cannot be dismissed as an artifact of researchers using unique events in their studies of conjunctive probability.

This experiment gathered estimates $P_E(A)$, $P_E(B)$, $P_E(A \wedge B)$, and $P_E(A \vee B)$ from 83 participants for 12 pairs A, B of weather

events. Two sets of weather events (the set *cloudy, windy, sunny, thundery* and the set *cold, frosty, sleety*) were used to form these pairs. These sets were selected so that they contained events of high, medium, and low probabilities. Conjunctive and disjunctive weather events were formed by pairing each member of the first set with every member of the second set and placing *and/or* between the elements as required, generating weather events such as *cloudy and cold, cloudy and frosty*, and so on. One group of participants ($N = 42$) was asked questions in terms of probability, of the form

- What is the probability that the weather will be W on a randomly selected day in Ireland?

for some weather event W . This weather event could be a single event such as *cloudy*, a conjunctive event such as *cloudy and cold*, or a disjunctive event such as *cloudy or cold*. The second group ($N = 41$) was asked questions in terms of frequency, of the form

- Imagine a set of 100 different days, selected at random. On how many of those 100 days do you think the weather in Ireland would be W ?

where the weather events were as before. These two question forms were used because of a range of previous work showing that frequency questions can reduce fallacies in people's probability judgments; the aim was to check whether this question form could eliminate fallacy responses for everyday repeated events.

Participants were given questions containing all single events and all conjunctive and disjunctive events, with questions presented in random order on a web browser. Responses were on an integer scale from 0 to 100. There was little difference in fallacy rates between the two forms of question, so we collapse the groups together in our analysis. There were 996 distinct conjunction and disjunction responses in the experiment (83 participants \times 12 conjunctions): A conjunction fallacy was recorded in 49% and a disjunction fallacy was recorded in 51% of those responses.

For every pair of weather events A, B used in the experiment, each participant gave estimates for the two constituents A and B , for the conjunction $A \wedge B$, and for the disjunction $A \vee B$. Each participant gave these estimates for 12 such pairs. For each participant we can thus calculate the value $X_E(A, B)$ for 12 pairs A, B , and so across all 83 participants we have 996 distinct values of $X_E(A, B)$. Our prediction is that the average of these values will equal 0 and that these values will be approximately symmetrically distributed around this average.

Results

Figure 2 graphs the raw frequency of occurrence of values for $X_E(A, B)$ in the experimental data and the average frequency in groups of those values. It is clear from the graph that these values are symmetrically distributed around the mean, just as expected. The mean value of X_E was 0.66 ($SD = 27.1$), within 1 unit of the predicted mean on the 100-point scale used in the experiment and within 0.025 standard deviations of the predicted mean. The predicted mean of 0 lay within the 99% confidence interval of the observed mean. This is in strikingly close agreement with our predictions. Note that the sequence of higher raw frequency values (hollow circles) in Figure 2 falls on units of 5 and represents participants' preference for rounding to the nearest 5 (the nearest




Figure 2. Frequency of different values of $X_E(A, B)$ in Experiment 1. This scatterplot shows the raw frequency of occurrence of different values of $X_E(A, B)$ as observed in the experimental data (hollow circles) and the average frequency across grouped values of $X_E(A, B)$, where each group contained 10 values of $X_E(A, B)$ from $v - 5 \dots v + 5$ for v from -100 to 100 in steps of 10 (filled circles). (The sequence of higher raw frequency values [hollow circles] falls on units of 5 and represents participants rounding to the nearest 5 in their responses.) The critical point here is that these values are symmetrically distributed around 0, as predicted in our model.

unit of 0.05) in their responses; approximately 55% of all responses were rounded in this way.

To examine the relationship between conjunction and disjunction fallacy rates and X_E values, we compared the total number of conjunction and disjunction fallacies produced by each participant with the average value of X_E for that participant. Figure 3 graphs this comparison. There was no significant correlation between the average value of X_E produced by a participant and the number of fallacies produced by that participant ($r = -0.1074$, $p = .34$).

Discussion

The above result is based on a specific expression X_E that cancels out the effect of noise in people's probability judgments. When noise is canceled in this way, we get a mean value for X_E that is almost exactly equal to that predicted by probability theory. This close agreement with probability theory occurs alongside significant conjunction and disjunction fallacy rates in the same data, with values of X_E close to zero even for participants with high conjunction and disjunction fallacy rates (see Figure 3). This cancellation of bias is difficult to explain in the heuristics view. To explain this cancellation, the heuristics view would require some way of ensuring that, when heuristics are used to estimate the four probabilities $P(A)$, $P(B)$, $P(A \wedge B)$, and $P(A \vee B)$ individually, the various biases in those four estimates are calibrated to give overall cancellation. Note that from both the conservatism results and the binary complementarity results described earlier, we know that the bias in estimates $P_E(A)$ and $P_E(B)$ will tend to cancel only when $P_E(A) = 1 - P_E(B)$ (i.e., when A and B are complementary). For the heuristics account to explain cancellation of bias across the four terms in X_E , therefore, it is not enough to say that people




Figure 3. Relationship between conjunction and disjunction fallacies and average $X_E(A, B)$ value in Experiment 1. This scatterplot shows the total number of conjunction and disjunction fallacies produced by each participant versus the average values of $X_E(A, B)$ across all pairs for that participant. The critical point here is that the fallacies occur frequently even when $X_E(A, B) = 0$.

overestimate $P(A \wedge B)$ and underestimate $P(A \vee B)$: It is necessary to calibrate the varying degrees of bias affecting all four probability estimates for $P(A)$, $P(B)$, $P(A \wedge B)$, and $P(A \vee B)$.

Further, to know that the bias in these four probabilities should cancel out in this way requires access to the rules of probability theory (as embodied in the addition law). These results therefore show that people follow probability theory when judging probabilities and that the observed patterns of bias are due to the systematic distorting influence of noise. When distortions due to noise are canceled out, as in expression X_E , no other systematic bias remains.

In the next section we describe a new experiment retesting this result and testing similar predictions for a range of other expressions.

Experiment 2

Our prediction for the derived expression X_E holds because the associated expression X is identically 0 and because there are an equal number of positive and negative terms in the expression (these two requirements are necessary to cancel out the d , or noise terms, in the expression). We now give another expression where these requirements both hold and so for which the same prediction follows.

Consider an experiment where we ask people to estimate, for any pair of events A and B , the probabilities of A , B , $A \wedge B$, $A \vee B$, $A \wedge \neg B$ (A and not B), and $B \wedge \neg A$ (B and not A). One derived sum that involves these probabilities is

$$Y_E(A, B) = P_E(A) + P_E(B \wedge \neg A) - P_E(B) - P_E(A \wedge \neg B),$$

whose expected value will be

$$\overline{Y}_E(A, B) = \overline{P}_E(A) + \overline{P}_E(B \wedge \neg A) - \overline{P}_E(B) - \overline{P}_E(A \wedge \neg B).$$

From Equation 3 we get

$$\begin{aligned} \overline{Y}_E(A, B) &= [P(A) + d - 2dP(A)] + [P(B \wedge \neg A) + d - 2dP(B \wedge \neg A)] \\ &\quad - [P(B) + d - 2dP(B)] - [P(A \wedge \neg B) + d - 2dP(A \wedge \neg B)] \\ &= (1 - 2d)[P(A) + P(B \wedge \neg A) - P(B) - P(A \wedge \neg B)]. \end{aligned}$$

However, probability theory requires $P(A) + P(B \wedge \neg A) = P(B) + P(A \wedge \neg B)$, because each side of the expression is equal to $P(A \vee B)$. We thus have $P(A) + P(B \wedge \neg A) - P(B) - P(A \wedge \neg B) = 0$ for all events A and B , and again we predict that the average value of Y_E across all participants and event pairs will equal 0. Because the positive and negative terms in the expression $Y_E(A, B)$ are symmetric (just as in X_E), we again expect values for $Y_E(A, B)$ to be symmetrically distributed around this mean, just as with X_E (this prediction is supported by Monte Carlo simulations similar to those described earlier). Finally, as X_E and Y_E have the same mean of 0, we predict that the larger combined set of all values of X_E and Y_E across all participants and event pairs will also have an mean of 0 and will be symmetrically distributed around that mean.

We can also consider other derived sums whose values in probability theory are 0, but where there is not an equal number of positive and negative terms in the expression (and so not all d or noise terms will be canceled out). Four such expressions are

$$Z1_E(A, B) = P_E(A) + P_E(B \wedge \neg A) - P_E(A \vee B)$$

$$Z2_E(A, B) = P_E(B) + P_E(A \wedge \neg B) - P_E(A \vee B)$$

$$Z3_E(A, B) = P_E(A \wedge \neg B) + P_E(A \wedge B) - P_E(A)$$

$$Z4_E(A, B) = P_E(B \wedge \neg A) + P_E(A \wedge B) - P_E(B).$$

For the first expression $Z1_E(A, B)$, Equation 3 gives

$$\begin{aligned} \overline{Z1}_E(A, B) &= [P(A) + d - 2dP(A)] + [P(B \wedge \neg A) + d - 2dP(B \wedge \neg A)] \\ &\quad - [P(A \vee B) + d - 2dP(A \vee B)] \\ &= (1 - 2d)[P(A) + P(B \wedge \neg A) - P(A \vee B)] + d \\ &= d, \end{aligned}$$

because from probability theory $P(A) + P(B \wedge \neg A) - P(A \vee B) = 0$ for all A and B . We get the same result for the expressions $Z2$, $Z3$, and $Z4$, and so we have expected values of

$$\overline{Z1}_E(A, B) = \overline{Z2}_E(A, B) = \overline{Z3}_E(A, B) = \overline{Z4}_E(A, B) = d$$

for all pairs A, B . Our prediction, therefore, is that expressions $Z1 \dots Z4$ should all have the same average value in our experiment.

Two other such derived sums are

$$Z5_E(A, B) = P_E(A \wedge \neg B) + P_E(B \wedge \neg A)$$

$$+ P_E(A \wedge B) - P_E(A \vee B)$$

$$Z6_E(A, B) = P_E(A \wedge \neg B) + P_E(B \wedge \neg A)$$

$$+ P_E(A \wedge B)$$

$$+ P_E(A \wedge B) - P_E(A) - P_E(B).$$

Similar computations for these expressions tell us that $Z5$ will have an expected value of $\overline{Z5}_E(A, B) = 2d$ and $Z6$ will have an

expected value of $\overline{Z6_E}(A, B) = 2d$ for all pairs A, B . Our prediction, therefore, is that expressions $Z5$ and $Z6$ should have similar average values in our experiment and that this average should be twice the average for $Z1 \dots Z4$. Note that because these expressions are not symmetric (all have leftover d terms), we do not expect the values of these expressions to be symmetrically distributed.

These last predictions are somewhat similar to the subadditivity results described earlier, in that both involve leftover d terms. The subadditivity results applied only to disjunctions of exclusive events (events that did not co-occur). The current predictions are more general in that they hold for all pairs of events A and B , irrespective of the degree of co-occurrence or dependency between those events.

Method

Participants. Participants were 68 undergraduate students at the School of Computer Science and Informatics, University College Dublin, who volunteered for partial credit.

Stimuli. This experiment gathered people's estimates for $P(A)$, $P(B)$, $P(A \wedge B)$, $P(A \vee B)$, $P(A \wedge \neg B)$, and $P(B \wedge \neg A)$ for 9 different pairs A, B of weather events such as *rainy*, *windy*. We constructed sets of three weather events each (the set *cold*, *rainy*, *icy* and the set *windy*, *cloudy*, *sunny*), selected so that each set contained events of high, medium, and low probabilities. Note that some of these pairs have positive dependencies (it is more likely to be rainy if it is cloudy), some had negative dependencies (it is less likely to be cold if it is sunny), and others were essentially independent. Our predictions apply equally across all cases.

Conjunctive and disjunctive weather events were formed from these sets by pairing each member of the first set with every member of the second set and placing *and/or* between the elements as required, generating weather events such as *cold and windy*, *cold or cloudy*, and so on. Weather events describing conjunctions with negation were constructed by pairing each member of the first set with every member of the second set, taking the two possible orderings of the selected elements, and for each placing *and not* between the elements in each ordering. This approach generated events such as *cold and not windy* and *windy and not cold* for each pair of events.

Procedure. Participants judged the probability of all single events and all conjunctions, disjunctions, and conjunctions with negations. Questions were presented in random order on a web browser. One group of participants ($N = 35$) was asked questions in terms of probability, of the form

- What is the probability that the weather will be W on a randomly selected day in Ireland?

for some weather event W . This weather event could be a single event such as *cloudy*, a conjunctive event such as *cloudy and cold*, a disjunctive event such as *cloudy or cold*, or a conjunction and negation event such as *cloudy and not cold* or *cold and not cloudy*. The second group ($N = 33$) were asked questions in terms of frequency, of the form

- Imagine a set of 100 different days, selected at random. On how many of those 100 days do you think the weather in Ireland would be W ?

where the weather events were as before. Responses were on an

integer scale from 0 to 100. The experiment took around forty minutes to complete.

Results

Two participants were excluded (one because they gave responses of 100 to all but 4 questions and the other because they gave responses of 0 to all but 2 questions), leaving 66 participants in total. There were thus 594 distinct conjunction and disjunction responses for analysis in the experiment (66 participants \times 9 conjunctions). A conjunction fallacy was recorded in 46% and a disjunction fallacy was recorded in 40% of those responses.

For every pair of weather events A, B used in the experiment, each participant gave probability estimates for the two constituents A and B , for $A \wedge B$ and $A \vee B$, and for $A \wedge \neg B$ and $B \wedge \neg A$. Each participant gave these estimates for nine such pairs. For each participant we calculated the values $X_E(A, B)$, $Y_E(A, B)$, and $\overline{Z1_E}(A, B) \dots \overline{Z6_E}(A, B)$ for nine pairs A, B , and so across all 66 participants we have 594 distinct values for each of those expressions.

Expressions X_E and Y_E . The mean value of X_E was -3.90 ($SD = 27.7$), and the mean value of Y_E was 3.82 ($SD = 30.08$). Figure 4 graphs the raw frequency of occurrence of values of X_E and Y_E in the experimental data and the average frequency in groups of those values, as in Figure 2. It is clear that these values are again unimodal and symmetrically distributed around their mean, as predicted. Averaging across all values of X_E and Y_E we get a mean of -0.01 ($SD = 29.2$); the predicted mean of 0 lies with the 99.9% confidence interval of this observed mean. Again, the sequence of higher raw frequency values (hollow circles) in Figure 4 falls on units of 5 and represents participants' preference for rounding to the nearest 5 in their responses; approximately 55% of all responses were rounded in this way.

To examine the relationship between conjunction and disjunction fallacy rates and X_E and Y_E values, we compared the total number of conjunction and disjunction fallacies produced by each participant with the average X_E and Y_E values for that participant. As in the previous experiment, there was no significant correlation between the average values produced by a participant and the




Figure 4. Frequency of different values of $X_E(A, B)$ and $Y_E(A, B)$ in Experiment 2. This scatterplot shows the raw frequency of occurrence of different values of $X_E(A, B)$ and $Y_E(A, B)$, as observed in the experimental data (hollow circles), and the average frequency across groups of 10 values, as in Figure 2. The critical point here is that these values are symmetrically distributed around 0, as predicted in our model.

number of fallacies produced by that participant ($r = -0.073$ and $r = -0.018$, respectively); the results showed values of X_E and Y_E close to zero even for participants with high conjunction and disjunction fallacy rates.

As before, these cancellation of bias results are difficult for the heuristics view to explain. They would require some way of ensuring that, when heuristics are used to estimate the 4 constituent probabilities in X_E and the 4 constituent probabilities in Y_E , the resulting biases are precisely calibrated to give overall cancellation. Further, to know that the bias in these probabilities should cancel requires access to the rules of probability theory, which the heuristics view does not have.

Expressions $Z1_E, \dots, Z4_E$. Recall that in our model estimates for expressions $Z1, \dots, Z4$ should on average all have the same biased value, equal to d (the noise rate). Table 1 gives the average values for these expressions calculated from participants' probability estimates; it is clear that these values are closely clustered (all are less than a 10th of an SD from the mean), just as predicted.

Expressions $Z5_E$ and $Z6_E$. Recall that in our model estimates for expressions $Z5_E$ and $Z6_E$ should on average have the same biased value, equal to $2d$ (twice the noise rate). Our prediction therefore is that values for $Z5_E$ and $Z6_E$ should fall close together and should fall close to twice the overall mean obtained for expressions $Z1_E, \dots, Z4_E$ (as in Table 1). Table 2 gives the average values for these expressions calculated from participants' probability estimates, and compares with twice the overall mean of $Z1_E, \dots, Z4_E$. It is clear that these values are closely clustered around that predicted value (both are less than a 20th of an SD from the predicted value, and their mean is less than 0.001 SD from that predicted value).

According to our model, the mean values of expressions $Z1_E, \dots, Z4_E$ are equal to the average value of d , the rate of random error in recall from memory, and the mean values of expressions $Z5_E$ and $Z6_E$ are equal to twice that value. This raises the interesting possibility of using the values of expressions $Z1_E, \dots, Z6_E$ for a given participant to estimate a value of d for that participant and so estimate the degree of variability due to noise in that participant's probability estimates. We discuss this possibility in the next section.

Random Variation and Fallacy Rates Across Participants

In our model the rate of occurrence of the conjunction fallacy is related to the degree of random variation: If there were no random

Table 1
Average Values of Expressions $Z1_E, \dots, Z4_E$ and Difference From Overall Mean

Expression	M	SD	Difference from overall mean
$Z1_E$	25.37	31.50	0.007
$Z2_E$	22.51	28.25	-0.093
$Z3_E$	26.50	27.95	0.048
$Z4_E$	26.23	29.12	0.037
Overall mean	25.15	29.26	

Note. The difference from overall mean is in units of standard deviation (SD).

Table 2

Average Values of Expressions $Z5_E, Z6_E$ and Difference From Predicted Value of Twice the Overall Mean of $Z1_E, \dots, Z4_E$ or $2 \times 25.15 = 50.3$

Expression	M	SD	Difference from 50.3
$Z5_E$	48.74	42.24	-0.037
$Z6_E$	52.72	49.98	0.048
Overall mean	50.73	46.30	0.001

Note. The difference from 50.3 is in units of standard deviation (SD).

variation in participant estimates the fallacy would never occur, and if there were a high degree of random variation the fallacy would occur frequently. The same prediction applies to the disjunction fallacy.

In this section we test these predictions using the data from our Experiment 2. In this analysis we use each participant's average values for expressions $Z1_E, \dots, Z6_E$ to estimate a value of d , the rate of random variation in recall for that participant. For each participant we can compute 6 estimates for that participant's value of d , by taking that participant's average value for each expression $Z1_E, \dots, Z6_E$ (and dividing the averages for $Z5_E$ and $Z6_E$ by 2). To examine the consistency of these estimates, we computed the pairwise correlation across participants between values of d estimated from each expression. Every pairwise correlation was significant at the $p < .0001$ level, and the average level of correlation was relatively high (mean $r = .79$), indicating that the values for d estimated for each participant from each of these expressions were consistent with each other.

Given this consistency we can produce an average estimate for d for each participant i ,

$$d_i = \frac{Z1_i + Z2_i + Z3_i + Z4_i + Z5_i/2 + Z6_i/2}{6},$$

where $Z1_i$ represents the average value of the derived sum $Z1_E(A, B)$ computed from participant i 's probability estimates for the nine pairs A, B . This gives a reasonable measure of the degree of random variation in recall for that participant.

To test our prediction that conjunctive and disjunctive fallacy rates will rise with the degree of random variation, we measure the correlation between conjunction and disjunction fallacy rates and the d_i random variation measure, across participants. There was a significant positive correlation between conjunction fallacy rates and the random variation measure ($r = .57, p < .0001$) and between disjunction fallacy rates and the random variation measure ($r = .43, p < .0005$), demonstrating that fallacy rates rise with random variation as in our model. This result goes against Nilsson et al.'s model, which predicts that conjunction and disjunction fallacy rates will fall with random variation (Nilsson et al., 2009).

Conjunction Error Rates in Averaged Estimates

In this section we consider the occurrence of the conjunction (and disjunction) fallacy in values that are produced by averaging across a set of probability estimates. Recall that Tversky and Kahneman's initial investigation found that the average probability rankings given to a conjunction $A \wedge B$ by one group of participants

was reliably higher than the average rankings given to a constituent A by another group of participants, producing a conjunction error (Tversky & Kahneman, 1983). More recently Nilsson et al. (2009) carried out a detailed study on the occurrence of conjunction and disjunction fallacies in averaged probability estimates. In Nilsson et al.'s experiments, participants assessed the probability of conjunctions, disjunctions, and constituent events in a test-retest format, with each participant being asked to assess each probability twice, once in Block 1 and once in Block 2. Nilsson et al. calculated the average probability estimate given by each participant for each constituent, conjunction, and disjunction and found that conjunction and disjunction fallacies in the averaged estimates were more frequent than conjunction and disjunction fallacies in the individual probability estimates. In this section we show that our model can account for this pattern of results. We first discuss the factors in our model that cause this pattern and then give simulation results demonstrating their occurrence in the model.

Consider a series of repeated experiments where in each experiment we gather N estimates for A and $A \wedge B$. For each experiment, the sample means $\widehat{P}_E(A)$ and $\widehat{P}_E(A \wedge B)$ represent averages of the N probability estimates obtained in that experiment. Across experiments, these sample means will vary randomly around their population means $\overline{P}_E(A \wedge B)$ and $\overline{P}_E(A)$, with different experiments giving different sample means, just as individual estimates $P_E(A)$ and $P_E(A \wedge B)$ vary randomly around the same population means. The conjunction and disjunction fallacy results described for individual estimates in the Conjunction and Disjunction Fallacies section, which depended on this random variation in individual estimates, thus also apply to sample means; the only difference between the two situations is that the degree of random variation in sample means will decline as the sample size N increases.

The chance of a conjunction fallacy in sample means (i.e., the chance of getting $\widehat{P}_E(A) < \widehat{P}_E(A \wedge B)$ in an experiment) depends on various factors. One factor is the number of individual estimates N being averaged; another is the difference between the probabilities $P(A)$ and $P(A \wedge B)$ being estimated. If $P(A \wedge B)$ and $P(A)$ are far apart, then population means $\overline{P}_E(A \wedge B)$ and $\overline{P}_E(A)$ will also be far apart and a conjunction fallacy can only occur when there is a large degree of variation in the sample means (i.e., when N is low). On the other hand, if $P(A \wedge B)$ and $P(A)$ are close, the population means will be close, and a conjunction fallacy can occur even when the degree of variation in sample means is low (i.e., when N is high). In other words, as sample size N increases, conjunction fallacy rates in sample means will decrease, with the rate of decrease depending on the difference between $P(A)$ and $P(A \wedge B)$.

When $P(A \wedge B) = P(A)$ the situation is different. Recall that in our model the distribution of individual probability estimates for some event depends only on the number of occurrences of that event in memory (see Equation 2). When $P(A \wedge B) = P(A)$ the number of occurrences of $A \wedge B$ is the same as the number of occurrences of A , and so the distribution of estimates for A is identical to the distribution for $A \wedge B$. This means that $\widehat{P}_E(A)$ and $\widehat{P}_E(A \wedge B)$ have the same random distribution around the same population mean. Because population means and distributions are the same, the chance of getting $\widehat{P}_E(A) < \widehat{P}_E(A \wedge B)$ is exactly the

same as the chance of getting $\widehat{P}_E(A) > \widehat{P}_E(A \wedge B)$. Because the first of these two possibilities produces a conjunction fallacy in the averaged data, the chance of getting a conjunction fallacy is

$$\frac{1 - Eq}{2}, \quad (7)$$

where Eq represents the chance of getting exactly the same values for $\widehat{P}_E(A)$ and $\widehat{P}_E(A \wedge B)$. If we assume that $\widehat{P}_E(A)$ and $\widehat{P}_E(A \wedge B)$ are continuous rather than discrete variables, then Eq is negligible, and we see that when $P(A \wedge B) = P(A)$, the chance of getting a conjunction fallacy in averaged data is 0.5 for all sample sizes N .

Consideration of the chance of getting exactly the same values for two sample means $\widehat{P}_E(A)$ and $\widehat{P}_E(A \wedge B)$ brings us to a third factor influencing conjunction error rates: rounding in participant responses. Recall our earlier observation that when estimating probabilities, participants tend to respond in units that are multiples of 0.05 or 0.10. For a small N , this rounding of estimates produces sample means that are not continuous variables but instead take on only a limited range of values; for example, if individual estimates are rounded to units of 0.10, then for $N = 1$ sample means can only have values that are multiples of 0.10, for $N = 2$ sample means can only have values that are multiples of 0.05, for $N = 3$ sample means can only have values that are multiples of 0.03333, and so on. This limitation on the range of possible values for sample means increases the chance of getting exactly the same values for $\widehat{P}_E(A)$ and $\widehat{P}_E(A \wedge B)$; that is, increases the value of Eq . Eq is highest when N is small (when there is only a small range of possible values for the sample means) and declines as N increases. Because a high value for Eq means a low conjunction fallacy rate in sample means (see Equation 7), this rounding effect causes the rate of conjunction fallacies in sample means to increase with increasing sample size N . This rounding effect can thus explain the increase in conjunction fallacy rate when averaging across multiple estimates that was observed by Nilsson et al. (2009). The same reasoning applies to disjunctions.

Simulation of Nilsson et al. (2009), Experiment 2

To test this explanation for Nilsson et al.'s (2009) results we use an extension of the simulation program described earlier (see Experiment 1). This extension simulates Nilsson et al.'s Experiment 2, which directly compared averaged conjunction error rate against individual conjunction fallacy rate.

The stimuli in Nilsson et al.'s (2009) Experiment 2 consisted of 180 components, 90 conjunctions, and 90 disjunctions that were constructed by randomly pairing those components. Nilsson et al. constructed components using a list of 188 countries. Each component consisted of a proposition stating that a given country had a population greater than 6,230,780 (the median population for the list). For example, a component could read "Sweden has a population larger than 6,230,780"; participants in the experiment were asked to indicate whether they thought that statement was true or false and to give their confidence in that judgment on a 5-point scale going from 50% to 100%.

Nilsson et al. (2009) created a unique sample of components for each participant by randomly sampling, with replacement, from

the set of components. Conjunctions and disjunctions were constructed by randomly pairing components (excluding duplication), so one conjunction could read “Sweden has a population larger than 6,230,780 and Spain has a population larger than 6,230,780”; participants were asked to indicate whether they thought that statement was true or false and to give their confidence in that judgment on a 5-point scale going from 50% to 100%.

Nilsson et al. (2009) transformed participants’ responses to a 0% to 100% scale by subtracting from 100 the confidence rating for those items where the participants gave a “false” response. This experiment thus necessarily rounds participants’ responses to units of 10%. The experiment had a test–retest design, where each participant was asked to estimate the probability for every component, conjunction and disjunction twice, in two separate blocks. Nilsson et al.’s primary result was that there was a higher rate of conjunction and disjunction fallacy occurrence when estimates were averaged across the two blocks than there was in the individual blocks alone.

Participant responses in Nilsson et al.’s (2009) experiment represent judgments in the confidence that a given country’s population is above the median (or, for conjunctions, that a pair of countries are above the median). To simulate this experiment we start with a representation of the true confidence that a given population is above the median. This true confidence is then input to our simulation program, which models the effect of random error in causing variation in that confidence. To mirror Nilsson et al.’s experiment as closely as possible, we simulate these true confidence values using a list of 188 highest country populations from Wikipedia and the median population for those countries (“List of Countries,” 2014). To construct simulated confidence judgments analogous to those given by Nilsson et al.’s participants, we took p_i to represent the population of country i and p_m to represent the median population and reasoned that the greater the difference between p_i and p_m , the greater the confidence there should be in judging that country i has a population greater (or less) than the median. For countries with populations greater than the median we therefore took

$$0.5 + \frac{p_i - p_m}{2 \times \max(p_i, p_m)}$$

to represent a simulated measure of confidence in the country’s population being greater than the median. Similarly, for countries with populations less than the median we took

$$0.5 + \frac{p_m - p_i}{2 \times \max(p_i, p_m)}$$

to represent a simulated measure of confidence in the country’s population being less than the median. Note that both these confidence measures run from 0.5 to 1, just as in Nilsson et al.’s experiment. Finally, we transformed these simulated confidence measures onto a 0 to 1 scale by following Nilsson et al.’s procedure and subtracting from 1 the confidence measure for countries with population less than the median. For every country i , this procedure gave a component probability

$$P(i > \text{median}) = 0.5 + \frac{p_i - p_m}{2 \times \max(p_i, p_m)}$$

that corresponds to a simulated measure of confidence that the

population of country i is greater than the median (and where values less than 0.5 represent countries whose population is less than the median). To construct conjunctive and disjunctive probabilities, we simply applied the probability theory equations for conjunction and disjunction to those component probabilities, under the assumption that component probabilities were independent.

On each run our simulation program took as input 180 randomly selected “true” confidence judgments for components (values $P(A)$, $P(B)$) constructed by applying the calculations described above to 180 randomly selected countries, and 90 conjunctive and disjunctive confidence judgments ($P(A \wedge B)$ and $P(A \vee B)$) calculated by applying the equations of probability theory to those components. For each set of components and conjunctive and disjunctive values, the program constructed a memory containing 100 items, each item containing flags A , B , $A \wedge B$, and $A \vee B$ indicating whether that item was an example of the given event. The rate of occurrence of those flags in memory matched the values specified by the input probabilities. The program contained a noise parameter d (set to 0.25 as before) and a rounding parameter set to $u = 0.1$ to match the rounding to units of 10% in Nilsson et al.’s (2009) experiment. To match the test–retest format in Nilsson et al.’s experiment, the program obtained two separate estimates of $P_E(A)$, $P_E(B)$, $P_E(A \wedge B)$, and $P_E(A \vee B)$ for each set of input probabilities. For each run the program returned the proportion of conjunction (and disjunction) fallacy responses in the individual blocks and the proportion of conjunction (and disjunction) fallacy responses in the averages across those two blocks.

Each run of the program thus corresponded to a simulated participant in Nilsson et al.’s experiment. We ran the program 1,000 times and compared the proportion of conjunction and disjunction fallacy responses in the individual blocks against the proportion of conjunction and disjunction fallacy responses in the averages. Just as in Nilsson et al.’s experiment, there was a higher rate of conjunction fallacy responses in averages ($M = 0.19$, $SD = 0.04$) than in individual blocks ($M = 0.16$, $SD = 0.03$), $t(1998) = 18.76$, $p < .0001$, and a higher rate of disjunction fallacy responses in averages ($M = 0.19$, $SD = 0.04$) than in individual blocks ($M = 0.15$, $SD = 0.03$), $t(1998) = 20.13$, $p < .0001$. These simulation results show that our model is consistent with the pattern seen in Nilsson et al.’s experiment.

General Discussion

We can summarize the main results of our experiments as follows: When distortions due to noise are canceled out in people’s probability judgments (as in X_E), those judgments are, on average, just as required by probability theory with no systematic bias. This close agreement with probability theory occurs alongside significant conjunction and disjunction fallacy rates in people’s responses. This cancellation of bias cannot be explained in the heuristics view. To explain this cancellation, the heuristics view would require some way of ensuring that, when heuristics are applied to estimate the various probabilities in expressions like X_E , the biases produced by heuristics are precisely calibrated to give overall cancellation.

Note that cancellation in the expression X_E is required because of probability theory’s addition law (Equation 1), which is equivalent to

$$P(A \vee B) = P(A) + P(B) - P(A \wedge B),$$

the probability theory equation for disjunction. If the heuristics view were able to ensure cancellation for X_E , that would mean the heuristics were embodying this addition law; or in other words, that the heuristics were implementing the probability theory equation for disjunction. However, this undermines the fundamental idea of the heuristics view, which is that people do not reason according to the rules of probability theory. Because the assumption in the heuristics view is that people do not follow probability theory when estimating probabilities, there is no way, in the heuristics view, to know that the terms in X_E should cancel. To put this point another way: If heuristics were selected in some way to ensure cancellation of bias for X_E , they would no longer be heuristics; they would simply be instantiations of probability theory.

Furthermore, our results show that when one noise term is left after cancellation (as in expressions $Z1_E \dots Z4_E$), a constant unit of bias is observed in people's probability judgments, and when two noise terms are left after cancellation (as in expressions $Z5_E$ and $Z6_E$), twice that unit of bias is observed in people's judgments, just as predicted in our simple model. Again, these results cannot be explained by an account where people estimate probabilities using heuristics: Such an account would not predict agreement in the degree of bias across such a range of different expressions. Together, these results demonstrate that people follow probability theory when judging probabilities in our experiments and that the observed conjunction and disjunction fallacy responses are due to the systematic distorting influence of noise and are not systematically influenced by any other factor.

It is worth noting that, although our results demonstrate that people's probability estimates given in our experiments followed probability theory (when bias due to noise is canceled out), we do not think people are consciously aware of the equations of probability theory when estimating probabilities. That is evidently not the case, given the high rates of conjunction and disjunction fallacies in people's judgments. Indeed we doubt whether any of the participants in our experiment were aware of the probability theory's requirement that our various derived sums should equal 0 or would be able to apply that requirement to their estimations. Instead we propose that people's probability judgments are derived from a "black box" module of cognition that estimates the probability of an event A by retrieving (some analogue of) a count of instances of A from memory. Such a mechanism is necessarily subject to the requirements of set theory and therefore implicitly embodies the equations of probability theory.

We expect this probability module to be based on observed event frequencies and to be unconscious, automatic, rapid, parallel, relatively undemanding of cognitive capacity, and evolutionarily "old." Support for this view comes from that fact that people make probability judgments rapidly and easily and typically do not have access to the reasons behind their estimations, from extensive evidence that event frequencies are stored in memory by an automatic and unconscious encoding process (Hasher & Zacks, 1984), and from evidence suggesting that infants have surprisingly sophisticated representations of probability (Cesana-Arlotti et al., 2012). Other support comes from results showing that animals effectively judge probabilities (i.e., the probability of obtaining

food from a given source) and that their judged probabilities are typically close to optimal (Kheifets & Gallistel, 2012).

We also expect this probability module to be subject to occasional intervention by other cognitive systems and particularly by other conscious and symbolic processes that may check the validity of estimates produced. We expect this type of intervention to be both rare and effortful. To quote one participant in an earlier experiment where participants had to bet on a single event or on a conjunction containing that event, "I know that the right answer is always to bet on the single one, but sometimes I'm really drawn to the double one, and it's hard to resist."

Comparing Models of Probabilistic Reasoning

The heuristics and biases approach proposes that people do not follow the rules of probability theory when estimating probabilities. Instead they use various heuristics that sometimes give reasonable judgments but sometimes lead to severe errors in estimation. The results given above directly contradict this proposal, showing that when bias due to noise is canceled, people's probability estimates closely follow the fundamental rules of probability theory. This cancellation of bias cannot be explained in the heuristics view, because to know that the bias in a given probabilistic expression should cancel requires access to the rules of probability theory. It is important to stress that these results are the central point in our argument against the view that people estimate probabilities via heuristics. Our argument against the heuristics view is not that our simple model of noisy retrieval from memory can explain four well-known biases (there are many other biases in the literature that our model does not address: base-rate neglect, the hard-easy effect, confirmation bias, the confidence bias, and so on; see Hilbert, 2012, for a review). Instead, our main point is that our experimental results demonstrate that the fundamental idea behind the heuristics view (that people do not follow the rules of probability theory) is seen to be incorrect when we use our simple model to cancel the effects of noise. Our results also argue against models where people reason about probability using equations different to those of probability theory (Carlson & Yates, 1989; Fantino et al., 1997; Nilsson et al., 2009). Our results give support for models where people follow probability theory in their probabilistic reasoning but are subject to the biasing effects of random noise; such models include Minerva-DM (Dougherty et al., 1999) and Hilbert's noisy channel model (Hilbert, 2012).

We presented our account under the assumption that, for any event A , there is a clear-cut binary criterion for membership in A : a given memory trace is either an instance of A or it is not. Given the complexity of event representation and the graded nature of most natural categories, this assumption is unrealistic: It is more likely that stored instances vary in their degree of membership in the category A and that the process of retrieval from memory would reflect this. Equally, our simple account assumes that there is only one point at which random noise influences probability estimation: the point at which memory is queried for stored instances of event A . Again, this is unrealistic: It is more likely that noise has an influence on initial perceptual processes, on event encoding, on retrieval, on subsequent processing, and on decision making and action. Further, our account applies only to unconditional probabilities, not to conditional probabilities $P(B|A)$; that is, the probability of B given that A has occurred. An important aim

for future work is to see whether any useful predictions could be derived by applying an extended version of our model to conditional probabilities.

A generalized version of our account, taking all of these factors into account, would clearly give more a realistic description of the processes of probability estimation. This realism would come at the cost of increased complexity, however: A more complex generalized model would have various interacting components and parameters that could be tuned in different ways to match behavior. Because of this complexity, it would be difficult to derive clear and testable predictions from such a model. This is the main advantage of our account: Its simplicity allows us to derive clear, specific and verifiable predictions about the impact of random variation on human probabilistic reasoning.

Concluding Remarks

The focus in our work has been on people's estimation of simple, unconditional probabilities. Our results show that patterns of systematic bias in these estimates can be explained via noise in recall and that when this noise is canceled, people's estimates match the requirements of probability theory closely, with no further systematic bias. This result has general implications for research on people's use of heuristics in reasoning. A frequent pattern in that research is to identify a systematic bias in people's responses and to then take that bias as evidence that the correct reasoning process is not being used. We believe that this inference is premature: As we have shown, random noise in reasoning can cause systematic biases in people's responses even when people are using the correct reasoning process. To demonstrate conclusively that people are using heuristics, researchers must show that observed biases cannot be explained as the result of random noise. To put it simply, biases do not imply heuristics, and even a rational reasoning process can produce systematically biased responses solely due to random variation and noise.

References

- Ariely, D. (2009). *Predictably irrational, revised and expanded edition: The hidden forces that shape our decisions*. New York, NY: HarperCollins.
- Bearden, N. J., & Wallsten, T. S. (2004). Minerva-DM and subadditive frequency judgments. *Journal of Behavioral Decision Making*, 17(5), 349–363. doi:10.1002/bdm.477
- Bonini, N., Tentori, K., & Osherson, D. (2004). A different conjunction fallacy. *Mind & Language*, 19(2), 199–210. doi:10.1111/j.1468-0017.2004.00254.x
- Bowers, J. S., & Davis, C. J. (2012). Bayesian just-so stories in psychology and neuroscience. *Psychological Bulletin*, 138(3), 389–414. doi:10.1037/a0026450
- Brunswick, E. (1955). In defense of probabilistic functionalism: A reply. *Psychological Review*, 62, 236–242. doi:10.1037/h0040198
- Budescu, D. V., Weinberg, S., & Wallsten, T. S. (1988). Decisions based on numerically and verbally expressed uncertainties. *Journal of Experimental Psychology: Human Perception and Performance*, 14(2), 281–294. doi:10.1037/0096-1523.14.2.281
- Camerer, C. F., Loewenstein, G., & Rabin, M. (Eds.). (2003). *Advances in behavioral economics*. New York, NY: Princeton University Press.
- Carlson, B., & Yates, J. (1989). Disjunction errors in qualitative likelihood judgment. *Organizational Behavior and Human Decision Processes*, 44(3), 368–379. doi:10.1016/0749-5978(89)90014-9
- Cesana-Arlotti, N., Téglás, E., & Bonatti, L. L. (2012). The probable and the possible at 12 months: Intuitive reasoning about the uncertain future. *Advances in Child Development and Behavior*, 43, 1–25. doi:10.1016/B978-0-12-397919-3.00001-0
- Chater, N., Tenenbaum, J. B., & Yuille, A. (2006). Probabilistic models of cognition: Conceptual foundations. *Trends in Cognitive Sciences*, 10(7), 287–291. doi:10.1016/j.tics.2006.05.007
- Costello, F. (2009a). How probability theory explains the conjunction fallacy. *Journal of Behavioral Decision Making*, 22(3), 213–234. doi:10.1002/bdm.618
- Costello, F. J. (2009b). Fallacies in probability judgments for conjunctions and disjunctions of everyday events. *Journal of Behavioral Decision Making*, 22(3), 235–251. doi:10.1002/bdm.623
- Dawson, N. V., & Arkes, H. R. (1987). Systematic errors in medical decision making. *Journal of General Internal Medicine*, 2, 183–187. doi:10.1007/BF02596149
- De Bondt, W. F. M., & Thaler, R. (1985). Does the stock market overreact? *Journal of Finance*, 40(3), 793–805. doi:10.1111/j.1540-6261.1985.tb05004.x
- Dougherty, M. R. P., Gettys, C. F., & Ogden, E. E. (1999). Minerva-DM: A memory processes model for judgments of likelihood. *Psychological Review*, 106(1), 180–209. doi:10.1037/0033-295X.106.1.180
- Eberhardt, F., & Danks, D. (2011). Confirmation in the cognitive sciences: The problematic case of Bayesian models. *Minds and Machines*, 21(3), 389–410. doi:10.1007/s11023-011-9241-3
- Edwards, W. (1968). Conservatism in human information processing. In B. Kleinmuntz (Ed.), *Formal representation of human judgment* (pp. 17–52). New York, NY: Wiley.
- Endress, A. D. (2013). Bayesian learning and the psychology of rule induction. *Cognition*, 127, 159–176. doi:10.1016/j.cognition.2012.11.014
- Erev, I., Wallsten, T. S., & Budescu, D. V. (1994). Simultaneous over- and underconfidence: The role of error in judgment processes. *Psychological Review*, 101(3), 519–527. doi:10.1037/0033-295X.101.3.519
- Eva, K. W., & Norman, G. R. (2005). Heuristics and biases: A biased perspective on clinical reasoning. *Medical Education*, 39(9), 870–872. doi:10.1111/j.1365-2929.2005.02258.x
- Fantino, E., Kulik, J., Stolarz-Fantino, S., & Wright, W. (1997). The conjunction fallacy: A test of averaging hypotheses. *Psychonomic Bulletin & Review*, 4, 96–101. doi:10.3758/BF03210779
- Fiedler, K. (1988). The dependence of the conjunction fallacy on subtle linguistic factors. *Psychological Research*, 50(2), 123–129. doi:10.1007/BF00309212
- Gavanski, I., & Roskos-Ewoldsen, D. R. (1991). Representativeness and conjoint probability. *Journal of Personality and Social Psychology*, 61(2), 181–194. doi:10.1037/0022-3514.61.2.181
- Gigerenzer, G. (1994). Why the distinction between single-event probabilities and frequencies is relevant for psychology and vice versa. In G. Wright & P. Ayton (Eds.), *Subjective probability* (pp. 129–162). New York, NY: Wiley.
- Gigerenzer, G., & Gaissmaier, W. (2011). Heuristic decision making. *Annual Review of Psychology*, 62, 451–482. doi:10.1146/annurev-psych-120709-145346
- Griffin, D. (2012, October 11). Science gallery bets all on risky show. *Irish Times*, p. 1.
- Griffiths, T. L., & Tenenbaum, J. B. (2006). Optimal predictions in everyday cognition. *Psychological Science*, 17(9), 767–773. doi:10.1111/j.1467-9280.2006.01780.x
- Hasher, L., & Zacks, R. (1984). Automatic processing of fundamental information: The case of frequency of occurrence. *American Psychologist*, 39(12), 1372–1388. doi:10.1037/0003-066X.39.12.1372
- Hertwig, R., & Chase, V. M. (1998). Many reasons or just one: How response mode affects reasoning in the conjunction problem. *Thinking and Reasoning*, 4(4), 319–352. doi:10.1080/135467898394102

- Hertwig, R., & Gigerenzer, G. (1999). The “conjunction fallacy” revisited: How intelligent inferences look like reasoning errors. *Journal of Behavioral Decision Making*, 12, 275–305. doi:10.1002/(SICI)1099-0771(199912)12:4<275::AID-BDM323>3.0.CO;2-M
- Hicks, E. P., & Kluemper, G. T. (2011). Heuristic reasoning and cognitive biases: Are they hindrances to judgments and decision making in orthodontics? *American Journal of Orthodontics and Dentofacial Orthopedics*, 139(3), 297–304. doi:10.1016/j.ajodo.2010.05.018
- Hilbert, M. (2012). Toward a synthesis of cognitive biases: How noisy information processing can bias human decision making. *Psychological Bulletin*, 138, 211–237. doi:10.1037/a0025940
- Jones, M., & Love, B. C. (2011). Bayesian fundamentalism or enlightenment? On the explanatory status and theoretical contributions of Bayesian models of cognition. *Behavioral and Brain Sciences*, 34(04), 169–188. doi:10.1017/S0140525X10003134
- Juslin, P., Nilsson, H., & Winman, A. (2009). Probability theory, not the very guide of life. *Psychological Review*, 116(4), 856–874. doi:10.1037/a0016979
- Kahneman, D. (2003). Maps of bounded rationality: Psychology for behavioral economics. *American Economic Review*, 93(5), 1449–1475. doi:10.1257/000282803322655392
- Kahneman, D. (2011). *Thinking, fast and slow*. New York, NY: Farrar, Straus & Giroux.
- Kahneman, D., & Tversky, A. (1973). On the psychology of prediction. *Psychological Review*, 80, 237–251.
- Kheifets, A., & Gallistel, C. R. (2012). Mice take calculated risks. *Proceedings of the National Academy of Sciences, USA*, 109, 8776–8779. doi:10.1073/pnas.1205131109
- Korobkin, R. B., & Ulen, T. S. (2000). Law and behavioral science: Removing the rationality assumption from law and economics. *California Law Review*, 88, 1051–1144. doi:10.2307/3481255
- List of countries by population. (2014). In *Wikipedia*. Retrieved February 20, 2014, from http://en.wikipedia.org/wiki/List_of_countries_by_population
- Marcus, G. F., & Davis, E. (2013). How robust are probabilistic models of higher-level cognition? *Psychological Science*, 24(12), 2351–2360. doi:10.1177/0956797613495418
- Mellers, B., Hertwig, R., & Kahneman, D. (2001). Do frequency representations eliminate conjunction effects? An exercise in adversarial collaboration. *Psychological Science*, 12(4), 269–275. doi:10.1111/1467-9280.00350
- Nilsson, H., Winman, A., Juslin, P., & Hansson, G. (2009). Linda is not a bearded lady: Configural weighting and adding as the cause of extension errors. *Journal of Experimental Psychology: General*, 138(4), 517–534. doi:10.1037/a0017351
- Oaksford, M., & Chater, N. (2007). *Bayesian rationality: The probabilistic approach to human reasoning*. New York, NY: Oxford University Press.
- Richards, J. (2012). *Psychology of intelligence analysis*. Reston, VA: Pherson Associates.
- Shafir, E., & Leboeuf, R. A. (2002). Rationality. *Annual Review of Psychology*, 53(1), 491–517. doi:10.1146/annurev.psych.53.100901.135213
- Sides, A., Osherson, D., Bonini, N., & Viale, R. (2002). On the reality of the conjunction fallacy. *Memory & Cognition*, 30(2), 191–198. doi:10.3758/BF03195280
- Sunstein, C. (2000). *Behavioral law and economics*. New York, NY: Cambridge University Press.
- Tenenbaum, J. B., Kemp, C., Griffiths, T. L., & Goodman, N. D. (2011, March 11). How to grow a mind: Statistics, structure, and abstraction. *Science*, 331(6022), 1279–1285. doi:10.1126/science.1192788
- Tentori, K., Bonini, N., & Osherson, D. (2004). The conjunction fallacy: A misunderstanding about conjunction? *Cognitive Science*, 28(3), 467–477. doi:10.1207/s15516709cog2803_8
- Tversky, A., & Kahneman, D. (1983). Extensional versus intuitive reasoning: The conjunction fallacy in probability judgment. *Psychological Review*, 90(4), 293–315. doi:10.1037/0033-295X.90.4.293
- Tversky, A., & Koehler, D. J. (1994). Support theory: A nonextensional representation of subjective probability. *Psychological Review*, 101(4), 547–567. doi:10.1037/0033-295X.101.4.547
- Wallsten, T. S., Budescu, D. V., Zwick, R., & Kemp, S. M. (1993). Preferences and reasons for communicating probabilistic information in verbal or numerical terms. *Bulletin of the Psychonomic Society*, 31, 135–138. doi:10.3758/BF03334162
- Wedell, D. H., & Moro, R. (2008). Testing boundary conditions for the conjunction fallacy: Effects of response mode, conceptual focus, and problem type. *Cognition*, 107(1), 105–136. doi:10.1016/j.cognition.2007.08.003
- Williams, B. S. (2010). Heuristics and biases in military decision making. *Military Review*, 9/10, 40–52. Retrieved from militaryreview.army.mil/

(Appendix follows)

Appendix

Variance in Probability Estimates in the Model

In our model, for any event A the count $C(A)$ of the number of flags in memory that are read as 1 is made up of two components: (a) the number of flags whose value is 1 and that are read correctly as 1 and (b) the number of flags whose value is 0 but that are read incorrectly as 1. Because the probability of a flag being read incorrectly is d , the first component is a binomial random variable with distribution $T_A - B(T_A, d)$, where T_A is the number of flags whose true value is 1 and $B(T_A, d)$ represents the binomial distribution of the number of those flags that are incorrectly read. The second component is a binomial random variable with distribution $B(m - T_A, d)$, because there are $m - T_A$ flags whose true value is 0 and $B(m - T_A, d)$ represents the binomial distribution of the number of those flags that are incorrectly read as 1. We thus have

$$C(A) = T_A - B(T_A, d) + B(m - T_A, d).$$

Because the mean of a binomial distribution $B(n, p)$ is pn , this gives

$$\begin{aligned} \overline{P_E(A)} &= \frac{T_A - dT_A + d(m - T_A)}{m} \\ &= (1 - 2d)P(A) + d, \end{aligned}$$

as in Equation 3. Because the expected variance of a binomial $B(n, p)$ is $p(1 - p)n$, the expected variances of these two distributions are $d(1 - d)T_A$ and $d(1 - d)(m - T_A)$, respectively. Each read of a flag is an independent Bernoulli trial with probability d , so these two distributions are independent. For two independent distributions the variance of their difference is equal to the sum of their individual variances, and so the expected variance of $C(A)$ is

$$\begin{aligned} V[C(A)] &= d(1 - d)(T_A + m - T_A) \\ &= d(1 - d)m. \end{aligned}$$

Finally, because variance is defined as the average squared difference from the mean and $P_E(A) = C(A)/m$, we get

$$V[P_E(A)] = \frac{V[C(A)]}{m^2} = d(1 - d)/m.$$

The expected variance in the distribution of $P_E(A)$ is a constant $d(1 - d)/m$ for all values of $P(A)$.

It is useful to check this result via simulation. We did this with a program that simulates the effects of random noise in recall on probability estimations for a given set of probabilities (as in the sections Experiment 1 and Conjunction Error Rates in Averaged Estimates). This program took as input an event probability $P(A)$




Figure A1. Computed variance in estimates $P_E(A)$ in Monte Carlo simulations with $m = 100$, for values of $P(A)$ from 0 to 1, in steps of 0.01, and for $d = 0.4, 0.3, 0.2$, and 0.1 . For each value of $P(A)$ and each value of d , the simulation produced 10,000 estimates $P_E(A)$: Each point in the graph represents the variance of one such set of 10,000 estimates. The critical point here is that this variance is independent of $P(A)$ and, for a given value of d , is equal to the predicted value $d(1 - d)/m$. For $d = 0.4, 0.3, 0.2$, and 0.1 , predicted variances are 0.0024, 0.0021, 0.0016, and 0.0009, respectively. The computed variances in the graph agree almost exactly with those predictions.

and constructed a memory containing $m = 100$ items, each item containing a flag indicating whether that item was an example of the given event. The occurrence of those flags in memory matched the given probability $P(A)$. This program also contained a noise parameter d ; when reading flag values from memory to generate an estimate $P_E(A)$, the program was designed to have a random chance d of returning the incorrect value. For each input probability $P(A)$, ranging from 0 to 1 in steps of 0.01, the program generated 10,000 noisy estimates $P_E(A)$ and used these 10,000 values to estimate $V[P_E(A)]$.

We carried out this simulation process for a range of values of d (0.1, 0.2, 0.3, 0.4). Figure A1 graphs the average variance for these 10,000 noisy estimates $P_E(A)$ for each value of $P(A)$ and for each of those values of d . As is clear from the graph, the calculated variance in the simulation was independent of $P(A)$ and equal to $d(1 - d)/m$, as expected.

Received November 30, 2012

Revision received February 28, 2014

Accepted March 2, 2014 ■