

# Distributed Systems: Logical Clocks

Anca Jurcut

E-mail: [anca.jurcut@ucd.ie](mailto:anca.jurcut@ucd.ie)

School of Computer Science and Informatics  
University College Dublin  
Ireland



# Logical Clocks

- A second approach to dealing with distributed synchronization is based on the concept of “relative time”.
- Based on the notion of ***event ordering***:
  - **Temporal Ordering**: the time of occurrence in real-time
    - induces a ‘happens-before’ relation
  - **Causal Ordering**: cause-effect relationship; implies temporal order
- Note that (with this mechanism) there is no requirement for “relative time” to have any relation to the “real time”.
- Such “clocks” are referred to as ***Logical Clocks***.

# Lamport's Logical Clocks

- Lamport made two observations about the nature of process interaction:
  - **First point:** if two processes do not interact, then their clocks do not need to be synchronized – they can operate concurrently without fear of interfering with each other.
  - **Second (critical) point:** it does not matter that two processes share a common notion of what the “real” current time is. What does matter is that the processes have some agreement on the order in which certain events occur.
- Based upon these observations, he devised the “happens-before” relation (denoted “ $\rightarrow$ ”).

# The “Happens-Before” Relation


- Properties of “ $\rightarrow$ ”:

- If A and B are events in the same process, and A occurs before B, then we can state that **A “happens-before” B** (denoted  $A \rightarrow B$ )
- If A is the event of sending a message, and B is the event of receiving the same message, then  $A \rightarrow B$ .
- If  $A \rightarrow B$  and  $B \rightarrow C$ , then  $A \rightarrow C$  (**transitivity**).
- If A and B are not related by  $\rightarrow$ , then they occurred at the same time (i.e. they are **concurrent** events).

- We can make some deductions about the current clock time based on “ $\rightarrow$ ”:

- If  $C(A)$  is the time when A occurred, then  $C(A) < C(B)$

# Example



happens before:  $p_1 \rightarrow q_2, r_2 \rightarrow r_3, q_4 \rightarrow r_4, r_1 \rightarrow q_9, \dots$

concurrent:  $p_1$  and  $q_1$ ;  $p_3$  and  $q_3, q_4, q_5$ ;  $q_5$  and  $r_3, r_4$ ;  $\dots$

# The “Happens-Before” Relation

- Now, assume three processes are in a Distributed System: A, B and C.
  - All have their own physical clocks (which are running at differing rates due to “clock skew”, etc.).
  - A sends a message to B and includes a “timestamp”.
  - If this sending timestamp is less than the time of arrival at B, things are OK, as the “happens-before” relation still holds (i.e. A “happens-before” B is true).
  - However, if the timestamp is more than the time of arrival at B, things are NOT OK (as A “happens-before” B is not true, and this cannot be as the receipt of a message has to occur after it was sent).

# The “Happens-Before” Relation

- The question to ask is:
  - How can some event that “happens-before” some other event possibly have occurred at a later time??
- The answer is: it can't!
  - So, Lamport's solution is to have the receiving process adjust its clock forward to one more than the sending timestamp value.
  - This allows the “happens-before” relation to hold, and also keeps all the clocks running in a synchronized state.
  - The clocks are all kept in sync relative to each other.

# Example

*P0*

0

*P1*

0

*P2*

0




# Example

<i>P0</i>	<i>P1</i>	<i>P2</i>
0	0	0
6	8	10

- P0 sends message A to P1

# Example




- $P1$  receives message A (everything is OK since  $6 < 16$ )

# Example

<i>P0</i>	<i>P1</i>	<i>P2</i>
0	0	0
6	8	10
12	16	20
18	24	30

- P1 sends message B to P2

# Example



- $P2$  receives message B (everything is OK since  $24 < 40$ )

# Example

*P0*

0
6
12
18
24
30

*P1*

0
8
16
24
32
40

*P2*

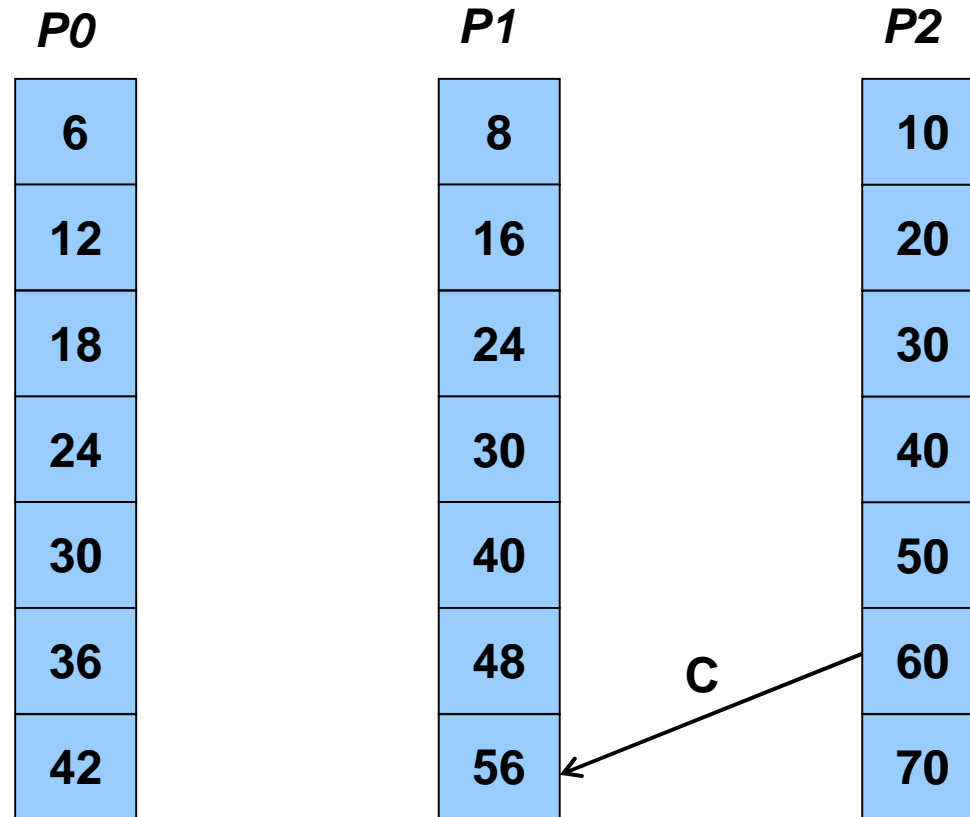
0
10
20
30
40
50

# Example

<i>P0</i>	<i>P1</i>	<i>P2</i>
0	0	0
6	8	10
12	16	20
18	24	30
24	30	40
30	40	50
36	48	60

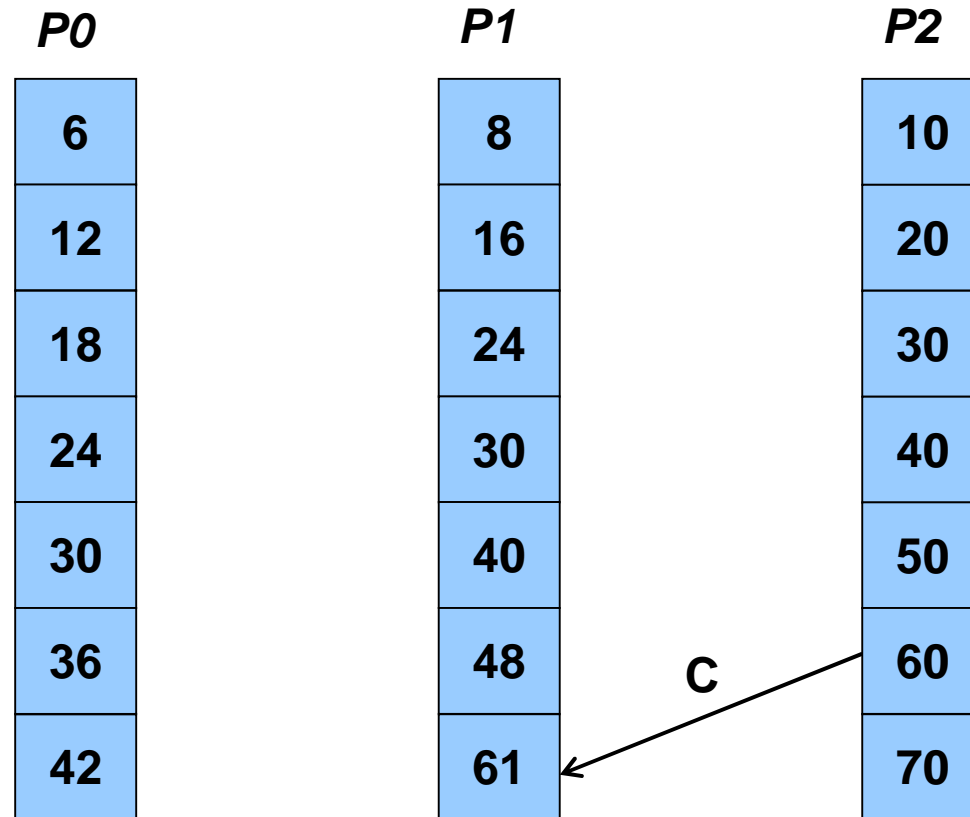
- P2 sends message C to P1

# Example



- $P1$  receives message C (Ouch! The message was sent at time 60 but received at time 56)

# Example



- Logical Time at P1 updated to be 1 greater than the time C was sent at.

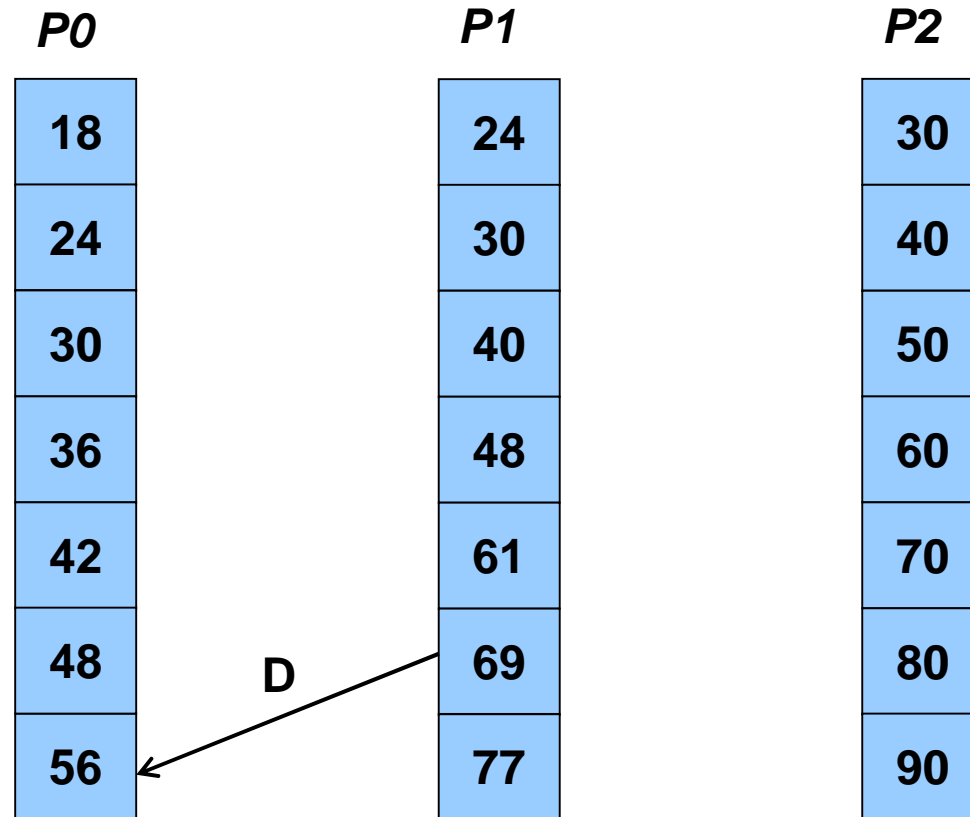


# Example

<i>P0</i>	<i>P1</i>	<i>P2</i>
12	16	20
18	24	30
24	30	40
30	40	50
36	48	60
42	61	70
48	69	80

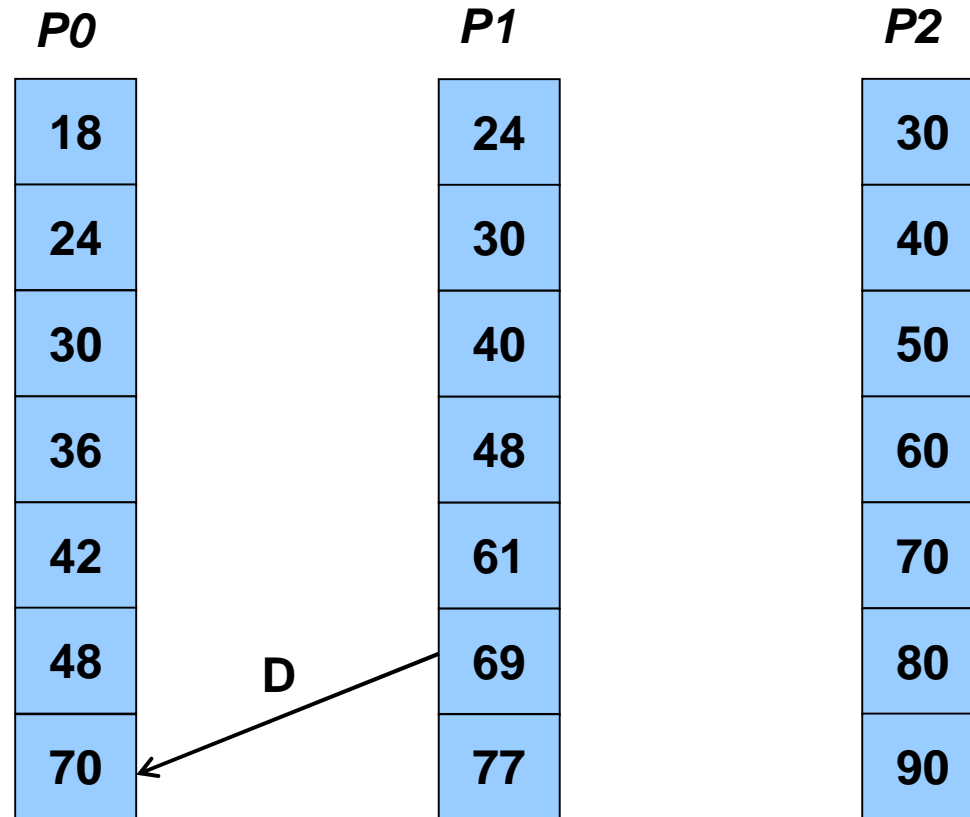
- P1 sends message D to P0

# Example



- $P0$  receives message  $D$  (Ouch!  $56 < 69$ )

# Example



- Logical Time at  $P0$  updated


# Example

<i>P0</i>	<i>P1</i>	<i>P2</i>
24	30	40
30	40	50
36	48	60
42	61	70
48	69	80
70	77	90
76	85	100

● End of Run

# Limitation of Lamport's algorithm


- If  $e_1 \rightarrow e_2$  then  $e_1.TS < e_2.TS$  (TS = "Time Stamp")
- But it is not necessarily the case that if  $e_1.TS < e_2.TS$  then  $e_1 \rightarrow e_2$



$e_1.TS < e_8.TS$ , but is  $e_1 \rightarrow e_8$  true?

Given two events, we cannot say whether they are causally related from their timestamps.

# Vector Time




- Alternative to Lamport's Algorithm in which a global timestamp is used:
- To each event  $e$  we now want to associate timestamps  $V(e)$  such that  

$$e \tilde{E} e' \tilde{O} V(e) < V(e')$$
- We let  $V(e)$  be a vector  $(t_1, \dots, t_N)$ . Intuitively, for each  $1 \leq i \leq N$  the latest event from process  $i$  which could have influenced  $e$  is  $e^{t_i}$
- We define:  

$$(t_1, \dots, t_N) < (t'_1, \dots, t'_N) \tilde{O} t_1 < t'_1, \dots, t_N < t'_N$$

# Vector Time



- The timestamps are formally defined as:

$$\begin{aligned}
 V(e_i^{-1}) &= (\overbrace{0, \dots, 0}^N) \\
 V(e_i^j) &= (\overbrace{0, \dots, 0}^{i-1}, 1, \overbrace{0, \dots, 0}^{N-i}) \\
 &\quad + \begin{cases} V(e_i^{j-1}) & \text{if } e_i^j \neq \text{recv}(m) \\ \max(V(e_i^{j-1}), V(\text{send}(m))) & \text{if } e_i^j = \text{recv}(m) \end{cases}
 \end{aligned}$$

for any  $m, j > 0$  and  $1 \leq i \leq N$ . Here  $\max$  and  $+$  are coordinate wise.

# Distributed Systems: Global State




# Introduction

- Sometimes we need to capture the global state of a distributed system.
  - E.g. Deadlock Detection, Termination Detection, Distributed Debugging, ...
- Capturing the state of the individual processes is relatively easy:
  - We simply write the state of each process to the hard disk.
- How can we synchronise this write activity to capture the state of all the processes simultaneously?


# Introduction

- Naïve Approach: We send a global “capture state message”
- Upon receipt of this message, each process writes its state to disk.
- But, in large systems this may take time:
  - This can lead to “Captured” processes continuing to receive messages from “Uncaptured” processes
- How can we ensure that the state we capture is **consistent**?

# Consistent Global State



(a)



Sender of m2 cannot  
be identified with this cut

(b)

- A consistent cut (a **possible** distributed state)
- An inconsistent cut (it is impossible for the system to ever have been in this state)

# Snapshot algorithm

- Developed by Chandy and Lamport in 1985
  - The paper is provided on moodle— please read
- Assumptions:
  - neither channels nor processes fail and communication is reliable so that every message send is eventually received intact, exactly once
  - channels are unidirectional and provide FIFO-ordered message delivery
  - the graph of processes and channels is strongly connected (there is a path between any two processes)
  - any process may initiate a global snapshot at any time
  - the processes may continue their execution and send and receive normal messages while the snapshot takes place

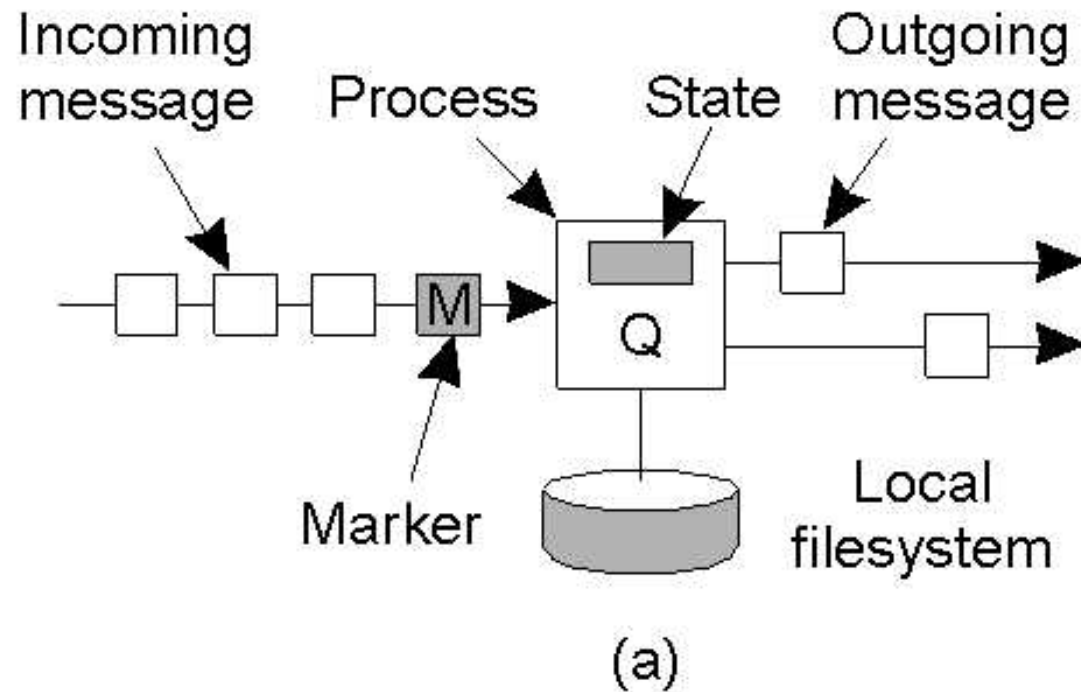
# Snapshot algorithm

- Requires the implementation of two “rules” that together ensure a correct snapshot is taken:
- A “Marker” message is sent around the system to instruct processes to take snapshots, and to get them to record the state of communication channels.
- Marker sending rule for process P
- After P has recorded its state, for each outgoing channel, C:
  - P sends one marker message over C (before it sends any other message)

# Snapshot algorithm


- Requires the implementation of two “rules” that together ensure a correct snapshot is taken:
- Marker-receiving rule for process P
- On P's receipt of a marker message over channel C
  - If P has not yet recorded its state it:
    - Records its process state now
    - Records the state of C as the empty set
    - Turns on recording of messages that arrive over other incoming channels
  - Else
    - P records the state of C as the set of messages it has received over C since it saved its state
  - Endif

# Snapshot algorithm



(a) Process Q receives a marker for the first time and records its local state

# Snapshot algorithm



- b)  $Q$  records its own state and sends a marker down all outgoing communication channels.
- c)  $Q$  records all incoming messages.
- d)  $Q$  receives a second marker on an incoming channel and finishes recording the state of that channel.



# Termination of Snapshot

- We assume that:

- a process that has received a marker message records its state within a finite time, and
- Sends marker messages over each outgoing channel within a finite time.

- If there is a path of communication channels and processes from a process  $p_i$  to a process  $p_j$  ( $j \neq i$ ), then it is clear the  $p_j$  will record its state a finite time after  $p_i$  recorded its state.

- Since we are assuming the graph of processes and channels to be strongly connected:

- It follows that all processes will have recorded their states and the states of incoming channels a finite time after some process initially records its state.

# To Do:

- Read Paper entitled:

“Distributed Snapshots: Determining Global States of Distributed Systems”

K. Mani Chandy and Leslie Lamport

*ACM Transactions on Computer Systems*, Vol. 3, No. 1,  
February 1985, Pages 63-75.

# Distributed Systems

## Coordination: Mutual Exclusion


# Mutual Exclusion

- It is often necessary to protect a shared resource within a Distributed System using “mutual exclusion”
  - For example, it might be necessary to ensure that no other process changes a shared resource while another process is working with it.
- In non-distributed, uniprocessor systems, we can implement “critical regions” using techniques such as semaphores, monitors and similar constructs – thus achieving mutual exclusion.
- These techniques have been adapted to Distributed Systems ...

# Mutual Exclusion: Techniques

- In this course we will consider three techniques:
  - **Centralized**: a single coordinator controls whether a process can enter a critical region.
  - **Distributed**: the group confers to determine whether or not it is safe for a process to enter a critical region.
  - **Token Ring**: processes are organised into a logical loop and use a token to determine when to enter a critical region.
- We will look at each technique over the coming slides...


# Centralized Algorithm




(a)

a) Process 1 asks the coordinator for permission to enter a critical region. Permission is granted by an OK message (assuming it is, of course, OK).

# Centralized Algorithm




(a)




(b)

a) Process 2 then asks permission to enter the same critical region. The coordinator does not reply (but adds 2 to a queue of processes waiting to enter the critical region). No reply is interpreted as a “busy state” for the critical region.


# Centralized Algorithm



(a)



(b)



(c)

a) When process 1 exits the critical region, it tells the coordinator, which then replies to 2 with an OK message.



# The Centralized Algorithm

- Advantages:

- It works.
- It is fair.
- There's no process starvation.
- Easy to implement.

- Disadvantages:

- There's a single point of failure!
- The coordinator is a bottleneck on busy systems.

- **Critical Question:** When there is no reply, does this mean that the coordinator is “dead” or just busy?

# Distributed Mutual Exclusion

- Based on work by Ricart and Agrawala (1981).
- Requirement of their solution: total ordering of all events in the distributed system (which is achievable with Lamport's timestamps).
- Note that messages in their system contain three pieces of information:
  - The critical region ID.
  - The requesting process ID.
  - The current time.


# Distributed Algorithm

- When a process (the “requesting process”) decides to enter a critical region, a message is sent to all processes in the Distributed System (including itself).
- What happens at each process depends on the “state” of the critical region.
  - If not in the critical region (and not waiting to enter it), a process sends back an OK to the requesting process.
  - If in the critical region, a process will queue the request and send back no reply to the requesting process.

# Distributed Algorithm



- If waiting to enter the critical region, a process will:
  - Compare the timestamp of the new message with that in its queue (note that the lowest timestamp wins).
  - If the received timestamp wins, an OK is sent back, otherwise the request is queued (and no reply is sent back).
- When all the processes send OK, the requesting process can safely enter the critical region.
- When the requesting process leaves the critical region, it sends an OK to all the process in its queue, then empties its queue.

# Distributed Algorithm Example






a) Processes 0 and 2 wish to enter the critical region “at the same time”.

# Distributed Algorithm Example



- b) Process 0 wins as it's timestamp is lower than that of process 2.

# Distributed Algorithm Example



- c) When process 0 leaves the critical region, it sends an OK to 2.

# The Distributed Algorithm

- The algorithm works because in the case of a conflict, the lowest timestamp wins as everyone agrees on the total ordering of the events in the distributed system.

- Advantages:

- It works.
- There is no single point of failure

- Disadvantages:

- We now have multiple points of failure!!!
- A “crash” is interpreted as a denial of entry to a critical region.
- (A patch to the algorithm requires all messages to be ACKed).




# The Distributed Algorithm

- Worse still - all processes must maintain a list of the current processes in the group (and this can be tricky)
- Worse still is that one overworked process in the system can become a bottleneck to the entire system – so, everyone slows down.


# Which Just Goes To Show ...

- It isn't always best to implement a distributed algorithm when a reasonably good centralized solution exists.
- Also, what's good in theory (or on paper) may not be so good in practice.
- Finally, think of all the message traffic this distributed algorithm is generating (especially with all those ACKs).
- Remember: every process is involved in the decision to enter the critical region, whether they have an interest in it or not. (Oh dear ... )

# Token-Ring Algorithm



(a)



(b)

- a) An unordered group of processes on a network. Note that each process knows the process that is next in order on the ring after itself.
- b) A logical ring is constructed in software, around which a token can circulate – a critical region can only be entered when the token is held. When the critical region is exited, the token is released.

# Token-Ring Algorithm

- When the ring is initialized, process 0 is given a *token*.
- The token circulates around the ring.
- It is passed from process  $k$  to process  $k+1$ .
- When a process acquires the token from its neighbor it checks to see if it is attempting to enter a critical region.
- If so, the process enters the region, does all the work it needs to and leaves the region.
- Token is passed to the next process in the ring.

# Token-Ring Algorithm

- Advantages:

- It works (as there's only one token, so mutual exclusion is guaranteed).
- It's fair – everyone gets a shot at grabbing the token at some stage.

- Disadvantages:

- Lost token! How is the loss detected (it is in use or is it lost)? How is the token regenerated?
- Process failure can cause problems – a broken ring!
- Every process is required to maintain the current logical ring in memory – not easy.

# Comparison: Mutual Exclusion

Algorithm	Messages per entry/exit	Delay before entry (in message times)	Problems
Centralized	3	2	Coordinator crash
Distributed	$2(n - 1)$	$2(n - 1)$	Crash of any process
Token-Ring	1 to $\infty$	0 to $n - 1$	Lost token, process crash

- *None are perfect – they all have their problems!*

# Comparison: Mutual Exclusion

- The “Centralized” algorithm is simple and efficient, but suffers from a single point-of-failure.
- The “Distributed” algorithm has nothing going for it – it is slow, complicated, inefficient of network bandwidth, and not very robust. It “sucks”!
- The “Token-Ring” algorithm suffers from the fact that it can sometimes take a long time to reenter a critical region having just exited it.
- All perform poorly when a process crashes, and they are all generally poorer technologies than their non-distributed counterparts.
- Only in situations where crashes are very infrequent should any of these techniques be considered.

# Election Algorithms

- Some Distributed Systems requires that one of the processes play a particular role.
  - For example, selecting a process to play the role of “central-server” in a variant of the Centralized Mutual Exclusion algorithm.
- In such situations, we need to employ a mechanism for selecting the “leader” process.
  - This mechanism must allow all relevant processes to participate in the choice.
  - It must also produce a single choice that is accepted by all the processes.
  - Once chosen the “leader” performs the assigned role until either they “retire” or fail.
- We term such a mechanism an **Election Algorithm**.