

# High-Performance Computing

---

COMP 40730

Alexey Lastovetsky

(B2.06, [alexey.lastovetsky@ucd.ie](mailto:alexey.lastovetsky@ucd.ie))

# Course Subject

- Parallel computing technologies
  - Aimed at *acceleration of solving a single problem on available computer hardware*
  - The course focuses on *software tools for developing parallel applications*: optimising compilers, parallel languages, parallel libraries
- Logical view vs. cooking book
  - *Ideas, motivations, models* rather than *technical details*
  - We will follow the evolution of hardware architecture
  - Carefully selected programming systems

# Course Outline

- Vector and Superscalar Processors
  - Programming and performance models
  - Optimising compilers
  - Array-based languages (Fortran 90, C[])
  - Array libraries (**BLAS**)
- Shared-Memory Multiprocessors
  - Programming and performance models
  - Parallel languages (Fortran 95, **OpenMP**)
  - Threads libraries (**Pthreads**)

# Course Outline (ctd)

- Distributed-memory multiprocessors
  - Programming and performance models
  - Parallel languages (High Performance Fortran)
  - Message passing libraries (**MPI**)
- Networks of Computers
  - Hardware and programming issues
  - Parallel computing (HeteroMPI, mpC)
  - High-performance Grid computing (NetSolve/GridSolve)


# Course Output

- You will be able to orient yourselves in parallel computing technologies
  - Mainly theoretical course
  - Some practical assignments (OpenMP, MPI)
    - A cluster of four 2-processor workstations
    - School network of computers

# References

- Reading materials

- A.Lastovetsky. *Parallel Computing on Heterogeneous Networks*.  
John Wiley & Sons, 423 pp, June 2003,  
ISBN: 0-471-22982-2.



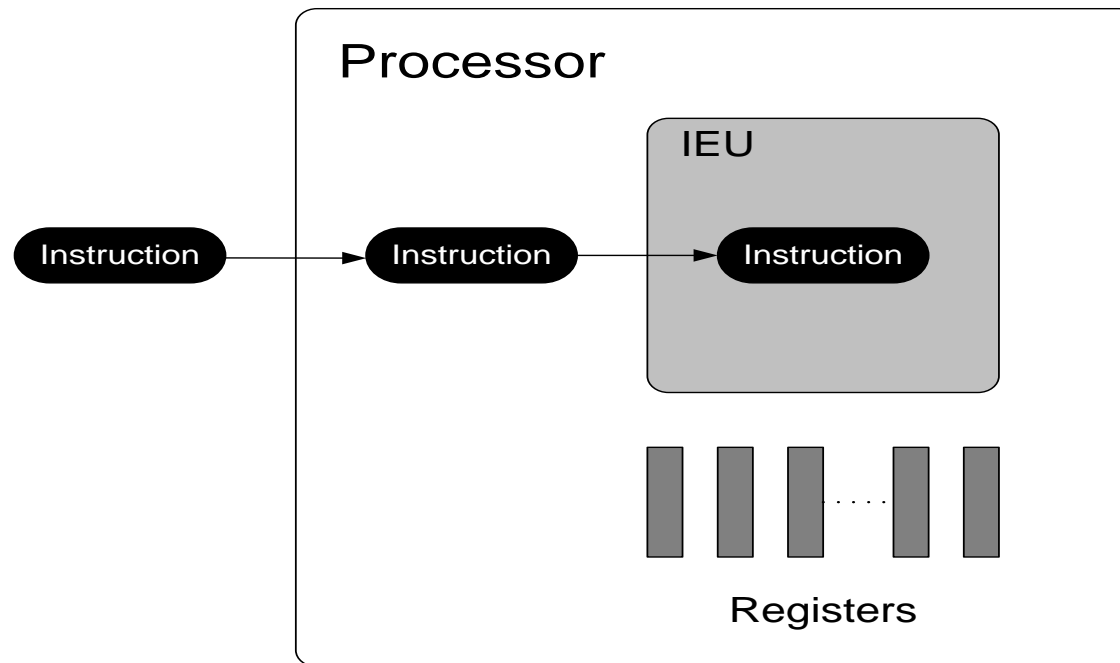
- The course website (lecture notes, assignments, etc.)  
<https://csmoodle.ucd.ie/moodle/course/view.php?id=491>

# Programming Systems for Serial Scalar Processors

---

# Serial Scalar Processor

- Starting-point of evolution of parallel architectures
  - Single control flow with serially executed instructions operating on scalar operands





# Serial Scalar Processor (ctd)

- One instruction execution unit (IEU)
- Next instruction can be only started after the execution of the previous instruction in the flow has been terminated
- A relatively small number of special instructions for data transfer between main memory and registers
- Most of instructions take operands from and put results to scalar registers
- The total time of program execution is equal to the sum of execution times of its instructions
- The performance of that architecture is determined by the clock rate

# Basic Program Properties

- A lot of languages and tools have been designed for programming SSPs
- C and Fortran are the most popular among professionals
- What is so special in these two languages?
  - They support and facilitate the development of software having certain properties considered basic and necessary by most professionals

# Basic Program Properties (ctd)

- Fortran is used mostly for scientific programming.
- C is more general-purpose and widely used for system programming
  - C can be used for programming in the Fortran-like style
  - Fortran 77 can be converted into C (GNU Fortran 77 compiler is implemented as such a convertor).
- The same program properties make Fortran attractive for scientific programming, and C for general-purpose and, especially, for system programming.

# Efficiency

- C allows one to develop highly efficient software.
- C reflects the SSP architecture with completeness resulting in programs of the assembler's efficiency
  - Machine-oriented data types (short, char, unsigned, etc.)
  - Indirect addressing and address arithmetics (arrays, pointers and their correlation)
  - Other machine-level notions (increment/decrement operators, the sizeof operator, cast operators, bit-fields, bitwise operators, compound assignments, registers, etc.)
- C supports *efficient* programming SSPs

# Portability

- C is standardised as ANSI C
  - All good C compilers support ANSI C
  - You can develop a C program running properly on any SSP
- C supports *portable* programming SSPs
- Portability of C applications
  - Portability of source code
  - Portability of libraries
    - Higher level of portability for SSPs running the same OS, or OSs of the same family (Unix)
  - GNU C compiler

# Modularity

- C allows a programmer to develop a program unit that can be separately compiled and correctly used by others without knowledge its source code
- C supports *modular* programming
- Packages and libraries can be only developed with tools supporting modular programming

# Easy to Use

- A clear and *easy-to-use* programming model ensures reliable programming
- Modularity and easy-to-use programming model facilitate the development of really complex and useful applications
- The C language design
  - Provides a balance between efficiency and lucidity
  - Combines lucidity and expressiveness

# Portable efficiency

- Portably efficient C application
  - A portable C application, which runs efficiently on any SSP having a high-quality C compiler and efficiently implemented libraries
- C
  - Reflects all the main features of each SSP affecting program efficiency
  - Hides peculiarities having no analogs in other SSPs (peculiarities of register storage, details of stack implementation, details of instruction sets, etc.)
- C supports *portably efficient* programming SSPs



# Basic Program Properties (ctd)

- There many other properties important for different kinds of software (fault tolerance, testability, etc.)
- 5 primary properties
  - *Efficiency*
  - *Portability*
  - *Modularity*
  - *Easy-to-use programming model*
  - *Portable efficiency*
- We will assess parallel programming systems mainly using the 5 basic program properties

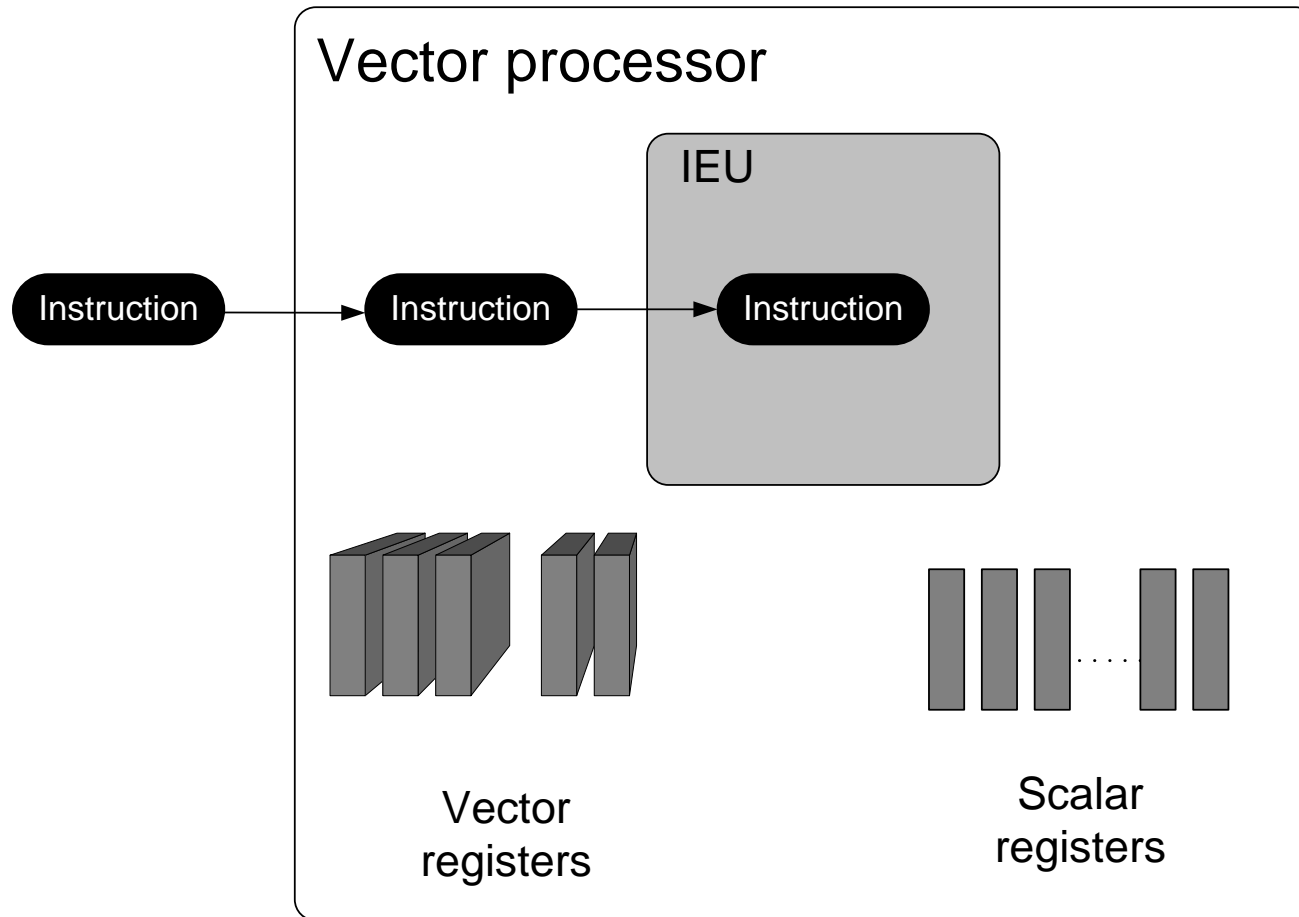
# Vector and Superscalar Processors

---

# Vector Processor

- Vector processor
  - Provides single control flow with serially executed instructions operating on both vector and scalar operands
  - Parallelism of this architecture is at the instruction level
  - Like SSP
    - VP has only one IEU
    - The IEU does not begin executing next instruction until the execution of the current one has completed
  - Unlike SSP
    - Instructions can operate on both scalar and vector operands
    - Vector operand is an ordered set of scalars located on a vector register

# Vector Processor (ctd)



# Vector Processor (ctd)

- A number of different implementations
  - ILLIAC-IV, STAR-100, Cyber-205, Fujitsu VP 200, ATC
- Cray-1 is probably the most elegant vector computer
  - Designed by Seymour Cray in 1976
  - Its processor employs data pipeline to execute vector instructions
  - 80 MFLOPS




# Cray-1 Vector Processor

- Consider the execution of a single vector instruction that
  - performs multiplication of two vector operands
  - takes operands from vector registers **a** and **b** and puts the result on vector register **c** so that  $c_i = a_i \times b_i \ (i=1, \dots, n)$
- This instruction is executed by a pipelined unit able to multiply scalars
  - the multiplication of two scalars is partitioned into  $m$  stages
  - the unit can simultaneously perform different stages for different pairs of scalar elements of the vector operands


# Cray-1 Vector Processor (ctd)

- At the first step, the unit performs stage 1 of the multiplication of elements  $a_1$  and  $b_1$



# Cray-1 Vector Processor (ctd)


- The  $i$ -th step ( $i=2, \dots, m-1$ )






# Cray-1 Vector Processor (ctd)

- The  $m$ -th step




# Cray-1 Vector Processor (ctd)

- Step  $m+j$  ( $j=1, \dots, n-m$ )




# Cray-1 Vector Processor (ctd)

- Step  $n+k-1$  ( $k=2, \dots, m-1$ )



# Cray-1 Vector Processor (ctd)

- Step  $n+m-1$



# Cray-1 Vector Processor (ctd)

- It takes  $n+m-1$  steps to execute this instruction
- The pipeline of the unit is fully loaded only from  $m$ -th till  $n$ -th step of the execution
- Serial execution of  $n$  scalar multiplications with the same unit would take  $nxm$  steps
- Speedup provided by this vector instruction is

$$S = \frac{n \times m}{n + m - 1} = \frac{m}{1 + \frac{m-1}{n}}$$

- If  $n$  is big enough,  $S \approx m$

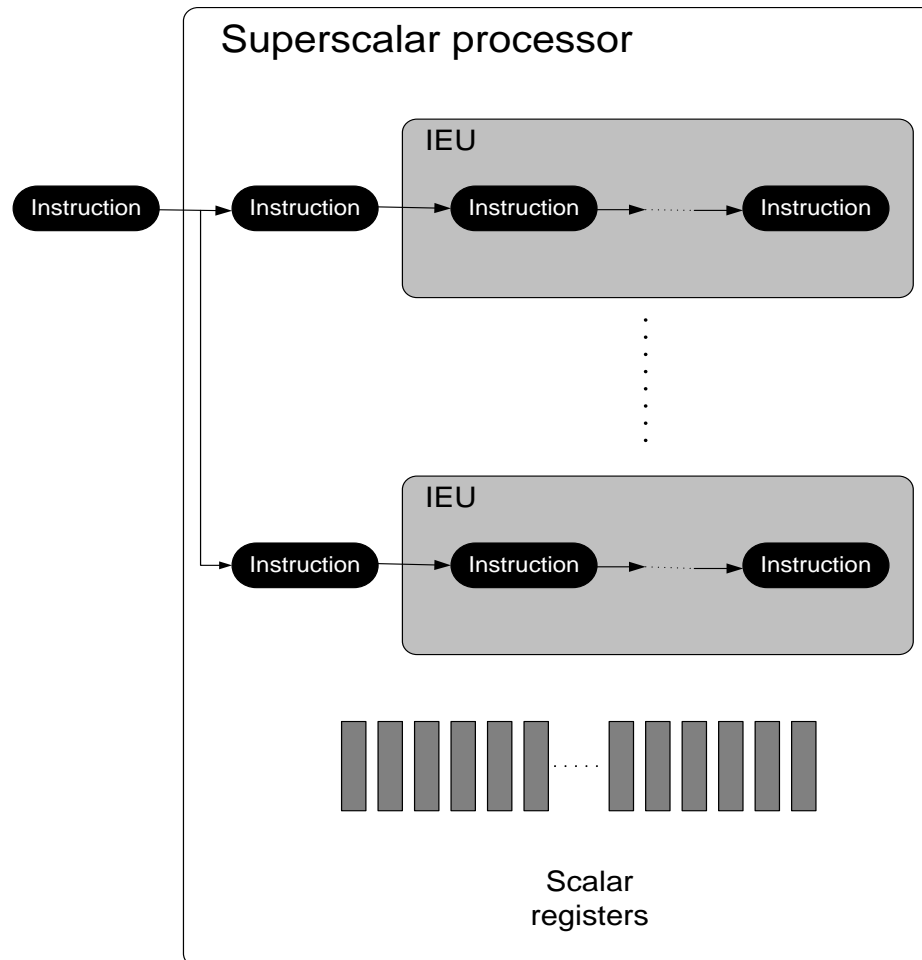
# Vector Processor (ctd)

- VPs are able to speed up applications, whose computations mainly fall into basic element-wise operations on arrays.
- The VP architecture includes the SSP architecture as a particular case ( $n=1$ ,  $m=1$ )

# Superscalar Processor

- Superscalar processor
  - Provides single control flow with instructions operating on scalar operands and being executed in parallel
  - Has several IEUs executing instructions in parallel
  - Instructions operate on scalar operands located on scalar registers
  - Two successive instructions can be executed in parallel by two different IEUs if they do not have conflicting operands
  - Each IEU is characterized by the set of instructions it executes
  - Each IEU can be a pipelined unit

# Superscalar Processor (ctd)






# Instructions Pipeline (ctd)


- A pipelined IEU can execute simultaneously several successive instructions each being on its stage of execution
- Consider the work of a pipelined IEU
  - Let the pipeline of the unit consist of  $m$  stages
  - Let  $n$  successive instructions of the program,  $I_1, \dots, I_n$ , be performed by the unit
  - Instruction  $I_k$  takes operands from registers  $a_k, b_k$  and puts the result on register  $c_k$  ( $k=1, \dots, n$ )
  - Let no two instructions have conflicting operands

# Instructions Pipeline (ctd)

- At the first step, the unit performs stage 1 of instruction  $I_1$




- Step  $i$  ( $i=2, \dots, m-1$ )




# Instructions Pipeline (ctd)

- Step  $m$




- Step  $m+j$  ( $j=1, \dots, n-m$ )




# Instructions Pipeline (ctd)

- Step  $n+k-1$  ( $k=2,\dots,m-1$ )



- Step  $n+m-1$



# Instructions Pipeline (ctd)

- It takes  $n+m-1$  steps to execute  $n$  instructions
- The pipeline of the unit is fully loaded only from  $m$ -th till  $n$ -th step of the execution
- Strictly serial execution by the unit of  $n$  successive instructions takes  $n \times m$  steps
- The maximal speedup provided by this unit is

$$S_{IEU} = \frac{n \times m}{n + m - 1} = \frac{m}{1 + \frac{m-1}{n}}$$

- If  $n$  is big enough,  $S_{IEU} \gg m$

# Superscalar Processor (ctd)

- The maximal speedup provided by the entire superscalar processor having  $K$  parallel IEUs is

$$S_{proc} \approx \sum_{i=1}^K m_i$$

- SPs are obviously able to speed up basic element-wise operations on arrays
  - Successive instructions execute the same operation on successive elements of the arrays

# Superscalar Processor (ctd)

- To efficiently support that type of computation, the processor should have at least

$$R \times \sum_{i=1}^K m_i$$

registers (each instruction uses  $R$  registers)

- CDC 6600 (1964) - several parallel IEUs
- CDC 7600 (1969) - several parallel pipelined IEUs
- Modern microprocessors are superscalar
- The superscalar architecture includes the serial scalar architecture as a particular case ( $K=1$ ,  $m=1$ ).

# Vector and Superscalar Architectures

- Why are vector and superscalar architectures united in a single group?
- The most successful VPs are very close to superscalar architectures
  - Vector pipelined unit can be seen as a specialized clone of the general-purpose superscalar pipelined unit
  - Some advanced superscalar processors (Intel i860) are obviously influenced by the vector-pipelined architecture



# Programming Model

- These architectures share the same programming model
  - A good program for vector processors widely uses basic operations on arrays
  - A program intensively using basic operations on arrays is perfectly suitable for superscalar processors
  - More sophisticated mixtures of operations able to efficiently load pipelined units of SPs are normally
    - Not portable
    - Quite exotic in real-life applications
    - Too difficult to write or generate

# Programming Systems

- Vector and superscalar processors are an evolution of the serial scalar processor → no wonder that programming tools for the architectures are mainly based on C and Fortran
- The programming tools are
  - Optimising C and Fortran 77 compilers
  - Array-based libraries
  - High-level parallel extensions of Fortran 77 and C