

Introduction to Software Defined Networking

Sofiane Zemouri

John Murphy

Outline

- Introduction & Traditional Networking
- What is SDN?
- OpenFlow basics
- SDN Challenges

INTRODUCTION


The Internet: A Remarkable Story

- Tremendous success
 - From research experiment to global infrastructure
- Brilliance of under-specifying
 - Network: best-effort packet delivery
 - Hosts: arbitrary applications
- Enables innovation in applications
 - Web, P2P, VoIP, social networks, virtual worlds
- But, change is easy only at the edge...



Inside the 'Net: A Different Story...


- Closed equipment
 - Software bundled with hardware
 - Vendor-specific interfaces
- Over specified
 - Slow protocol standardization
- Few people can innovate
 - Equipment vendors write the code
 - Long delays to introduce new features



Impacts performance, security, reliability, cost...

Networks are Hard to Manage

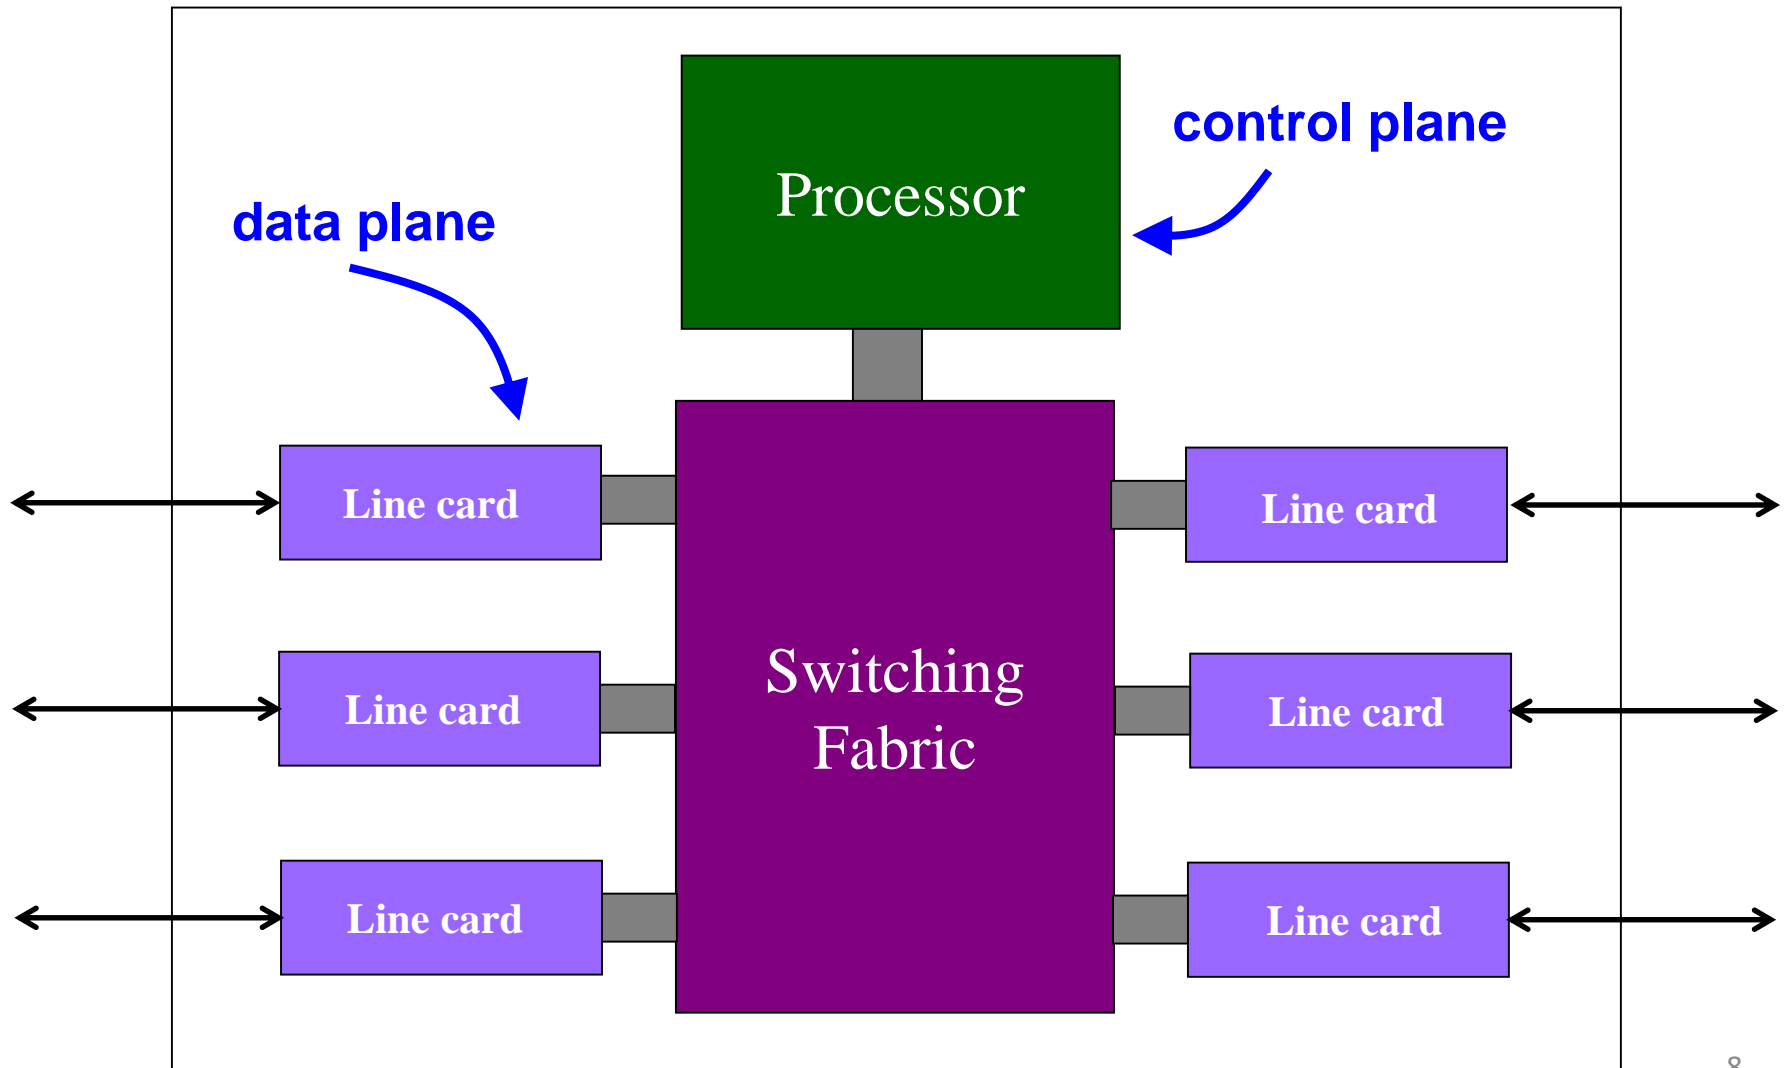
- Operating a network is expensive
 - More than half the cost of a network
 - Yet, operator error causes most outages
- Buggy software in the equipment
 - Routers with 20+ million lines of code
 - Cascading failures, vulnerabilities, etc.
- The network is “in the way”
 - Especially a problem in data centers
 - ... and home networks



The Networking “Planes”

- **Data plane:** processing and delivery of packets with local forwarding state
 - Forwarding state + packet header → forwarding decision
 - Filtering, buffering, scheduling
- **Control plane:** computing the forwarding state in routers
 - Determines how and where packets are forwarded
 - Routing, traffic engineering, failure detection/recovery, ...
- **Management plane:** configuring and tuning the network
 - Traffic engineering, ACL config, device provisioning, ...


Data and Control Planes



TRADITIONAL COMPUTER NETWORKS

Traditional Computer Networks


Data plane:
Packet
streaming



Forward, filter, buffer, mark,
rate-limit, and measure packets

Traditional Computer Networks


Control plane:
Distributed algorithms



Track topology changes, compute routes,
install forwarding rules

Traditional Computer Networks

Management plane: Human time scale




Collect measurements and configure the equipment

Timescales

	Data	Control	Management
Time-scale	Packet (nsec)	Event (10 msec to sec)	Human (min to hours)
Location	Linecard hardware	Router software	Humans or scripts


Data Plane

- Streaming algorithms on packets
 - Matching on some header bits
 - Perform some actions
- Example: **IP Forwarding**




Control Plane

- Compute paths the packets will follow
 - Populate forwarding tables
 - Traditionally, a distributed protocol
- Example: **Link-state routing (OSPF, IS-IS)**
 - Flood the entire topology to all nodes
 - Each node computes shortest paths
 - Dijkstra's algorithm




1. Figure out which routers and links are present.
2. Run Dijkstra's algorithm to find shortest paths.

"If a packet is going to B,
then send it to output 3"



Management Plane

- **Traffic Engineering:** setting the weights
 - Inversely proportional to link capacity?
 - Proportional to propagation delay?
 - Network-wide optimization based on traffic?



Traditional Networking Issues

(Too) many task-specific control mechanisms

- No modularity, limited functionality

Indirect control

- Must invert protocol behavior, “coax” it to do what you want
- Ex. Changing weights instead of paths for TE

Uncoordinated control

- Cannot control which router updates first

Interacting protocols and mechanisms

- Routing, addressing, access control, QoS

Traditional Networking Issues

The network is

- Hard to reason about
- Hard to evolve
- Expensive

WHAT IS SDN?

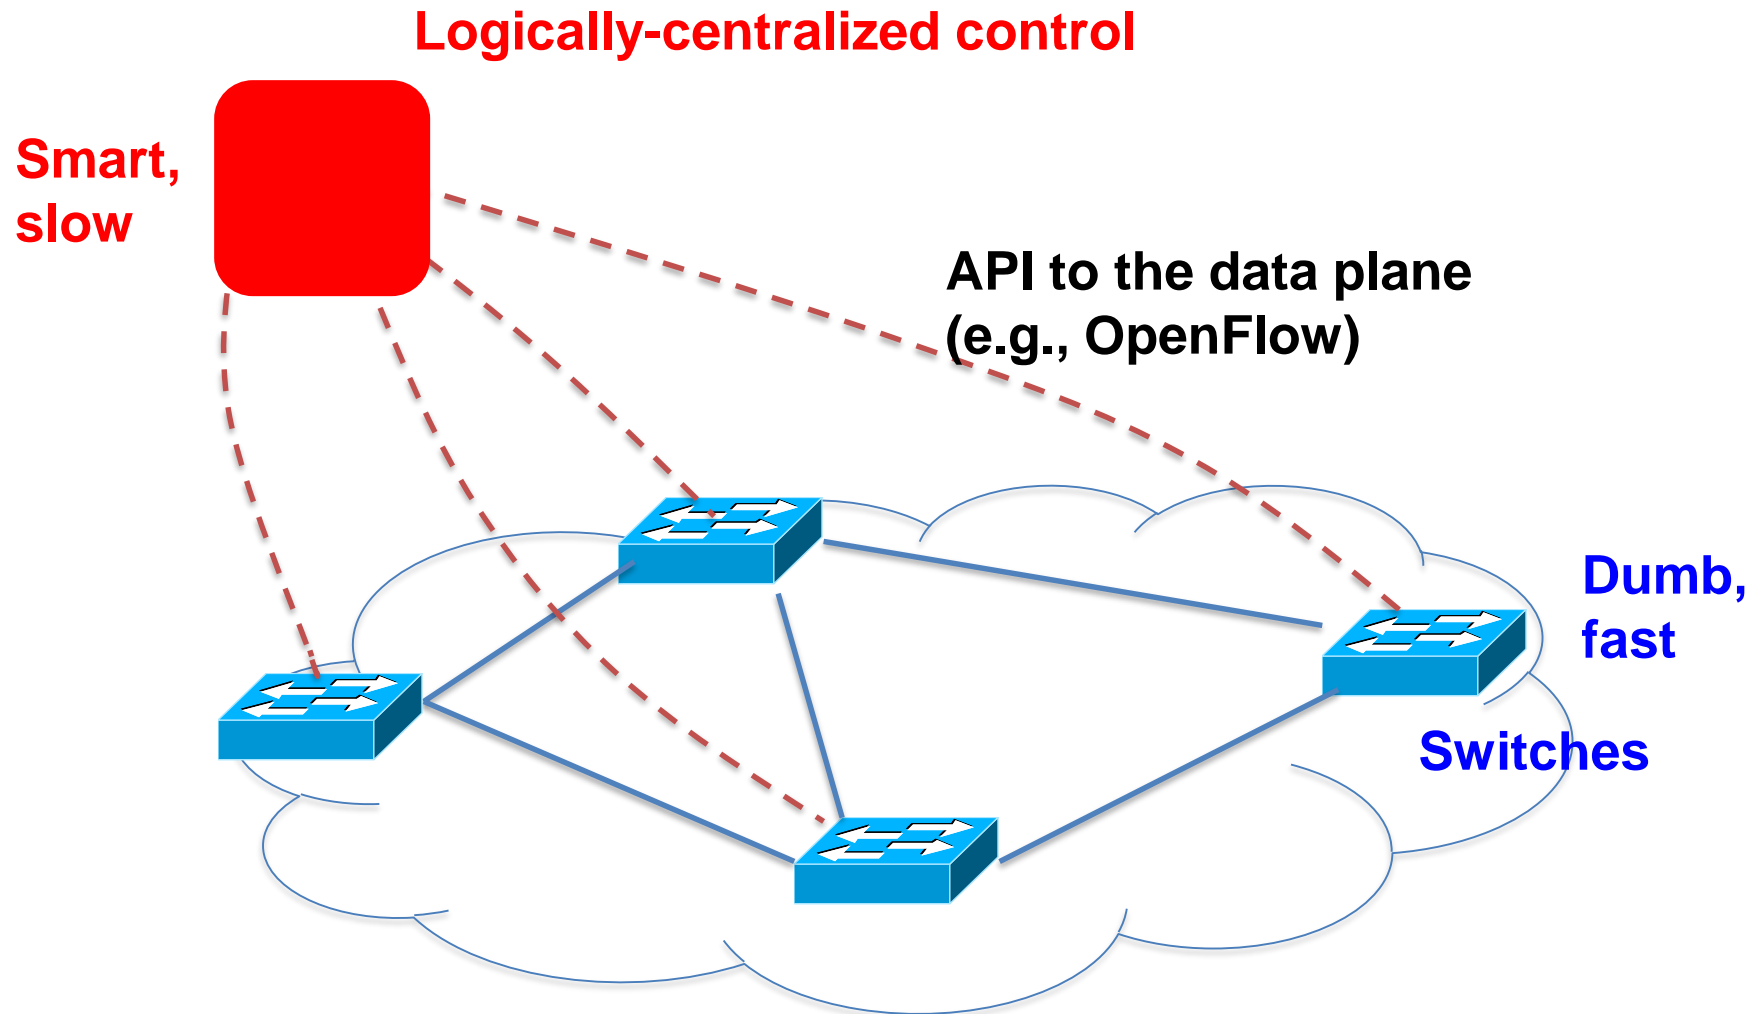
Software Defined Networks

A network in which the control plane is physically separate from the data plane.

and


A single (logically centralized) entity controls several forwarding devices.

Software Defined Networks



Centralization of the Control Plane

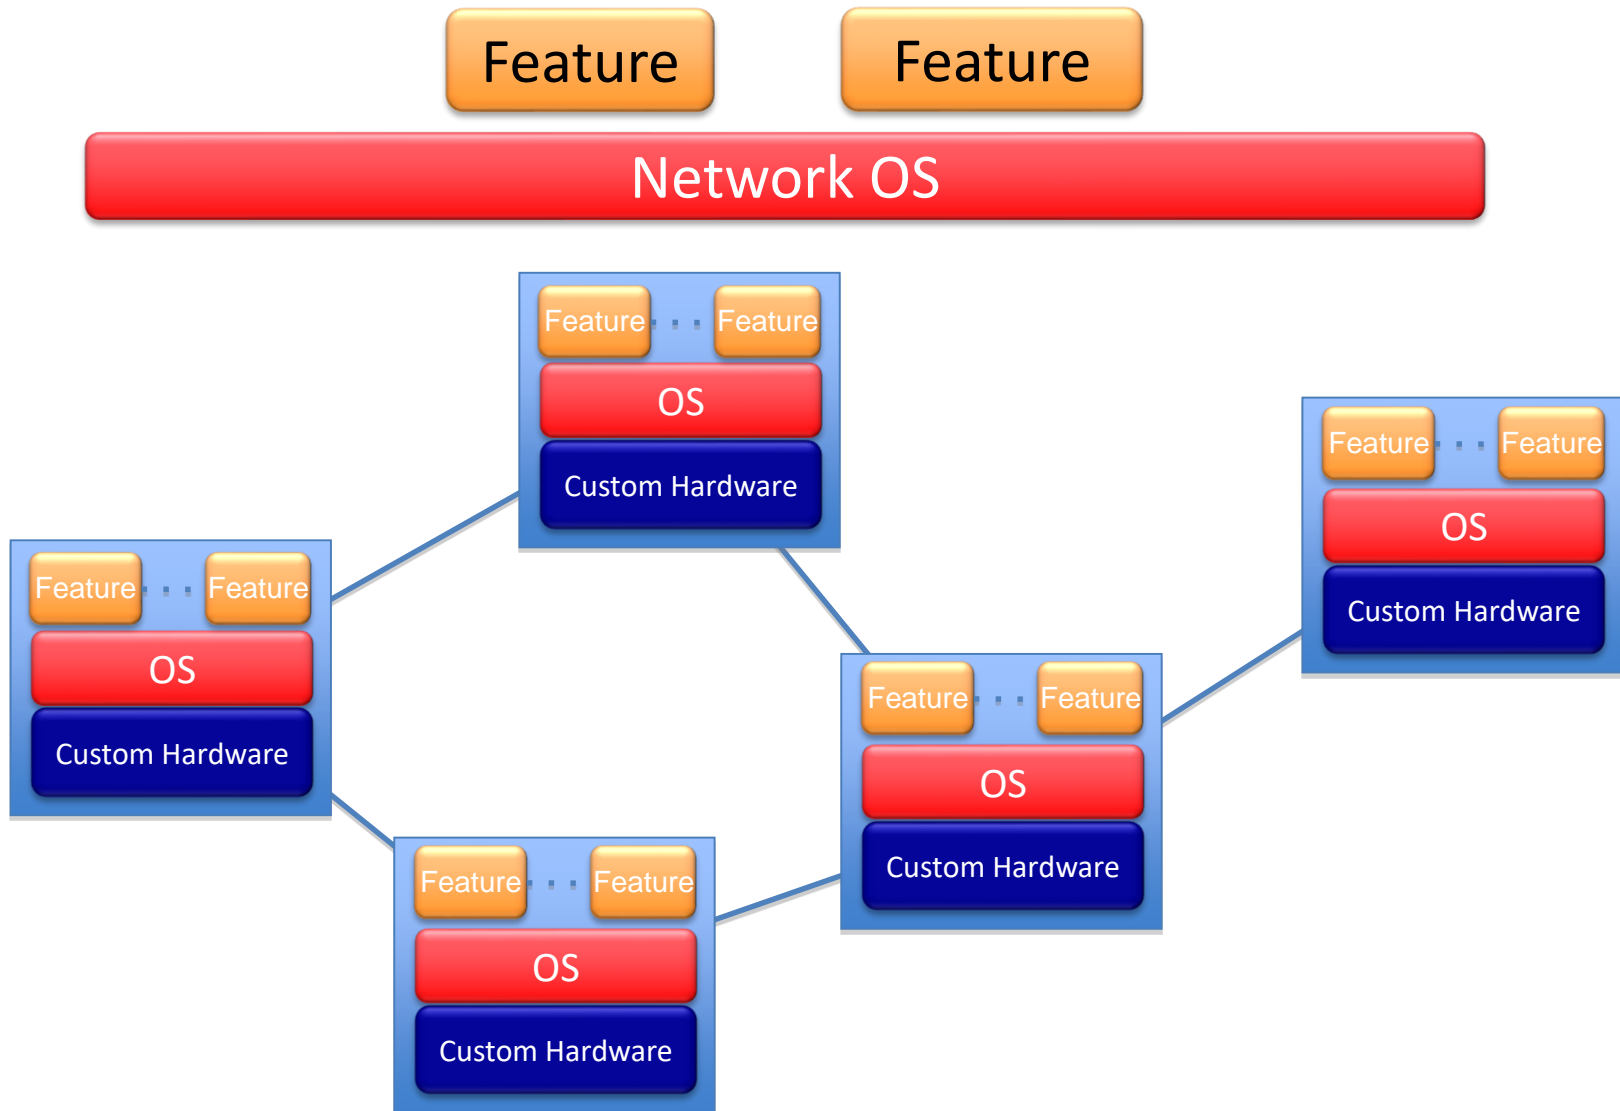
- **Simpler management**
 - No need to “invert” control-plane operations
- **Faster pace of innovation**
 - Less dependence on vendors and standards
- **Easier interoperability**
 - Compatibility only in “wire” protocols
- **Simpler, cheaper equipment**
 - Minimal software



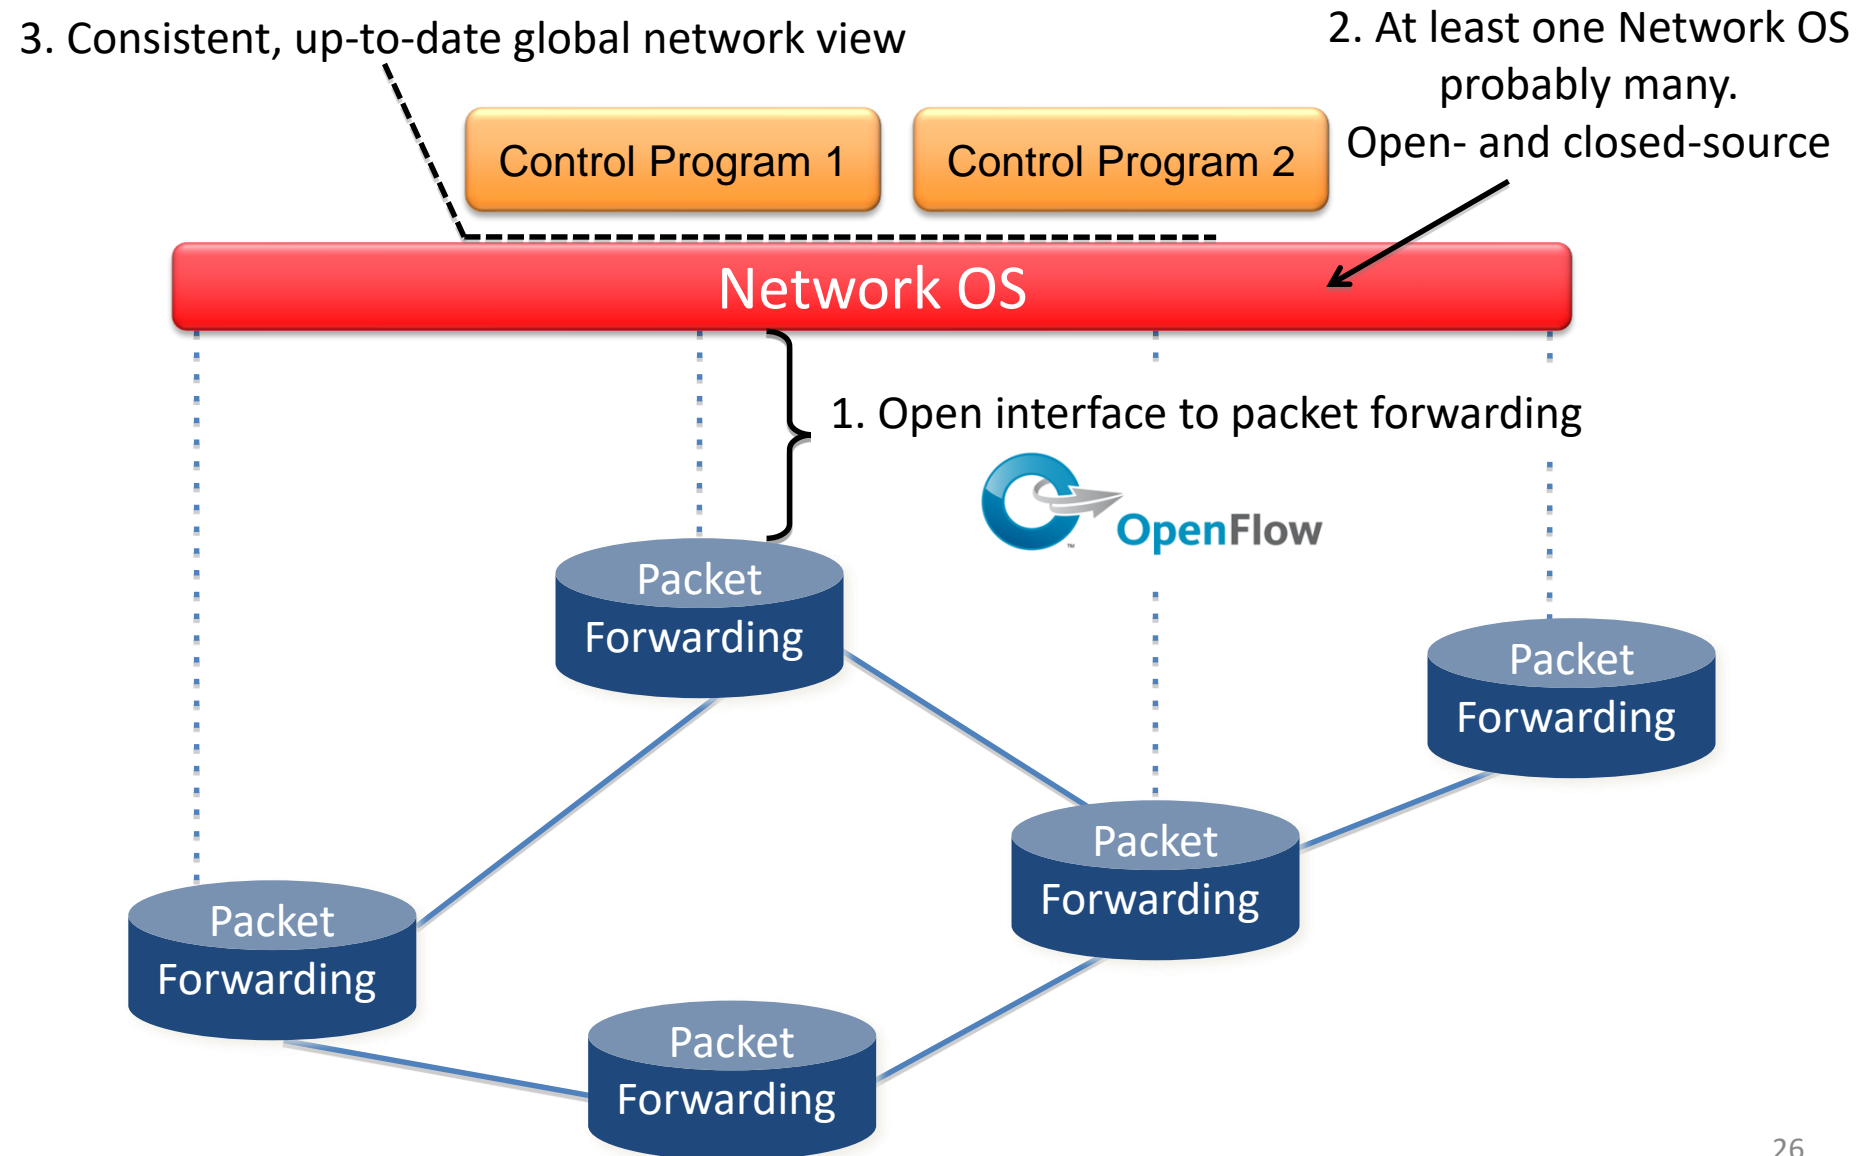
SDN advantages

- **Flexibility:** topologies, routing and forwarding architecture, independent configuration
- **Manageability:** separate policy and mechanism
- **Scalability:** maximize number of co-existing virtual networks
- **Security and Isolation:** isolate both the logical networks and the resources
- **Programmability:** programmable routers, etc.
- **Heterogeneity:** support for different technologies

How SDN changes the network



How SDN changes the network

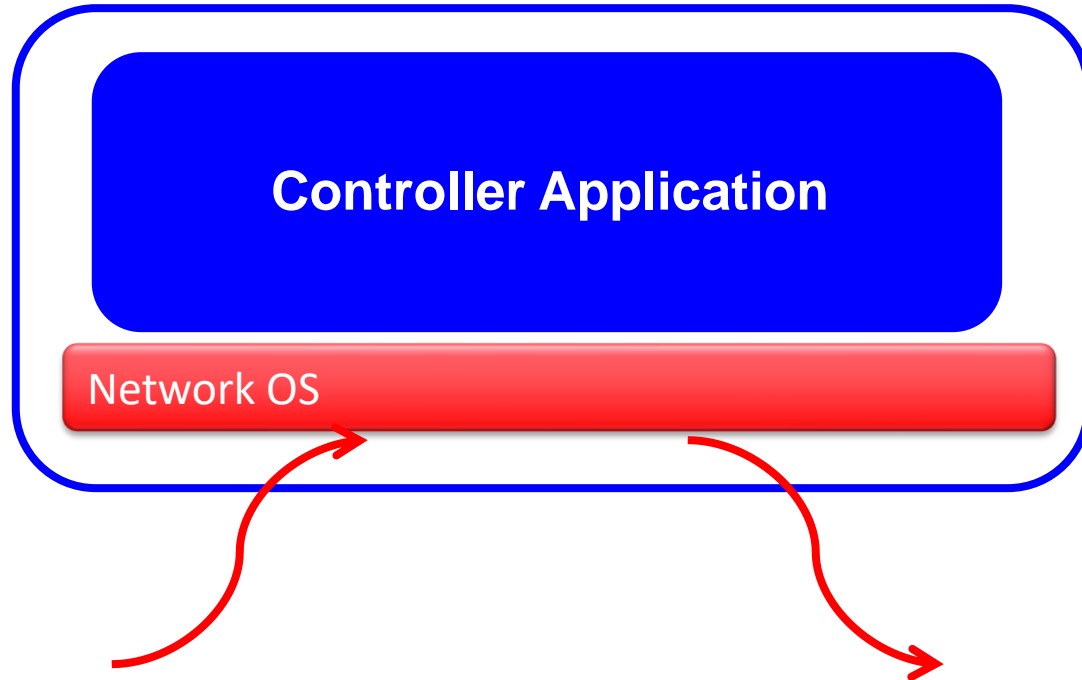


Control Program (controller app)

Control program operates on view of network

- **Input:** global network view (graph/database)
- **Output:** configuration of each network device

Controller: Programmability



Events from switches

Topology changes,
Traffic statistics,
Arriving packets

Commands to switches

(Un)install rules,
Query statistics,
Send packets

Control Program (controller app)

Control program is not a distributed system

- Abstraction hides details of distributed state


Abstraction of the physical network

- Support for multiple logical networks running on a common shared physical substrate


Aspects of the network that can be virtualized

- **Nodes:** Virtual machines
- **Links:** Tunnels (e.g., Ethernet GRE)
- Storage

Network virtualization



Virtualization simplifies control program



Does SDN Simplify the Network?

Abstraction doesn't eliminate complexity

- Network OS and Hypervisor are still complicated pieces of code

SDN main achievements

- Simplifies interface for control program (user-specific)
- Pushes complexity into reusable code (SDN platform)

Just like compilers....

Forwarding Abstraction

Purpose: Standard way of defining forwarding state

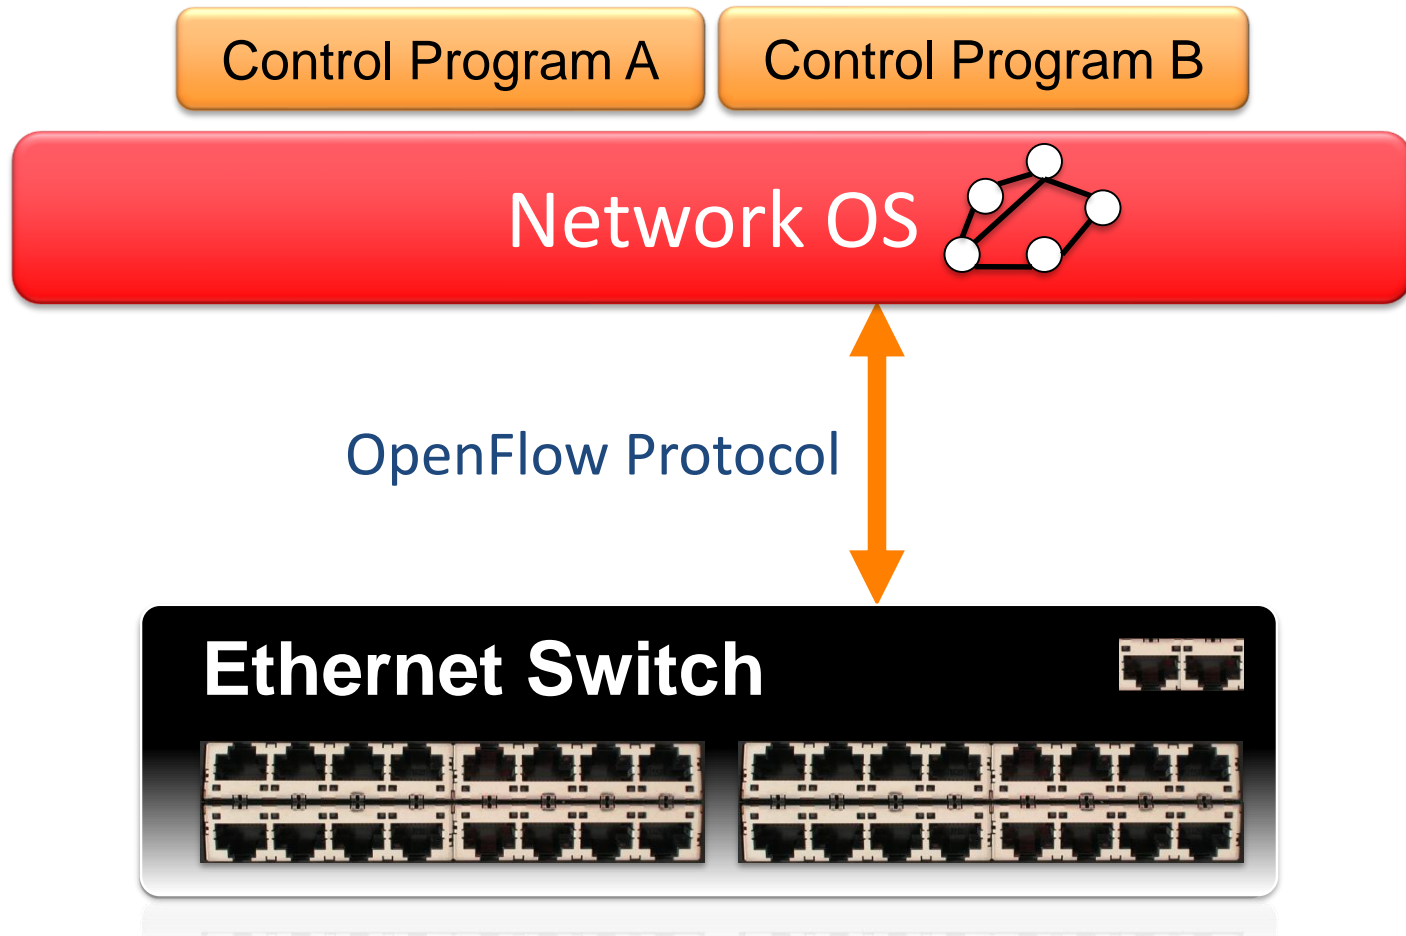
- Flexible
 - Behavior specified by control plane
 - Built from basic set of forwarding primitives
- Minimal
 - Streamlined for speed and low-power
 - Control program not vendor-specific
- OpenFlow is an example of such an abstraction

OPENFLOW BASICS

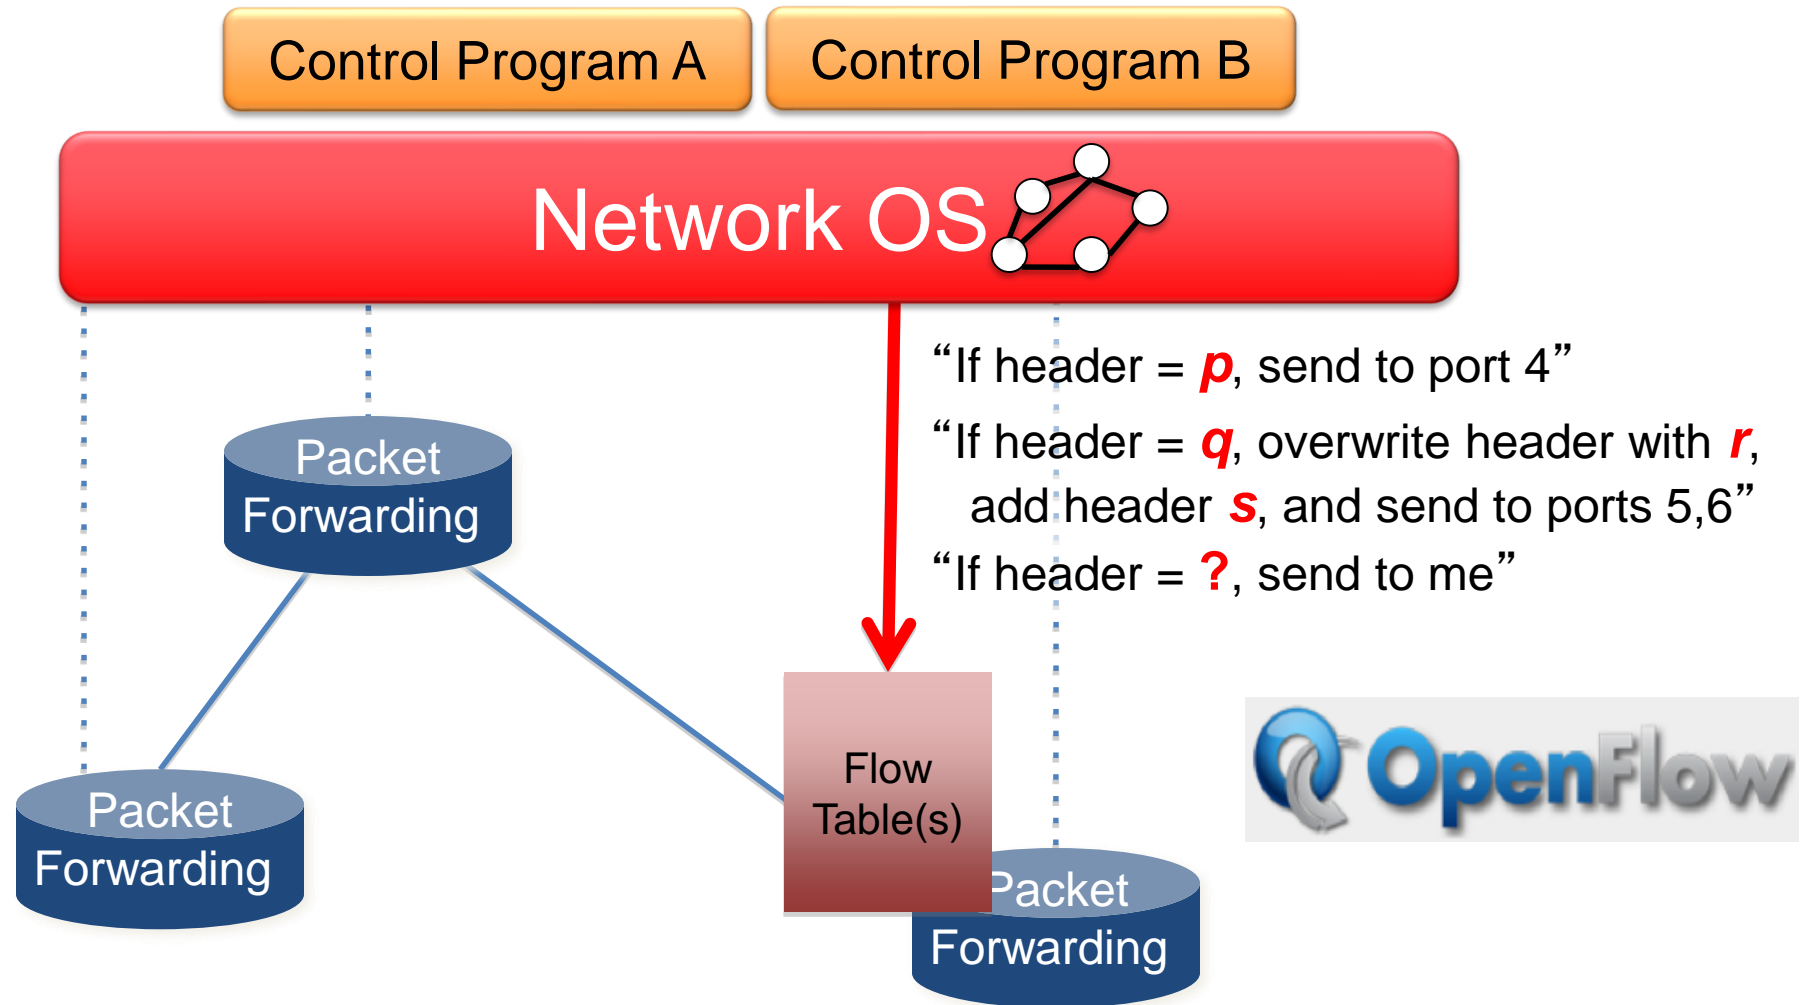
OpenFlow in the Wild

- **Open Networking Foundation**
 - Google, Facebook, Microsoft, Yahoo, Verizon, Deutsche Telekom, and many other companies
- **Commercial OpenFlow switches**
 - HP, NEC, Quanta, Dell, IBM, Juniper, ...
- **Network operating systems**
 - NOX, Beacon, ONIX, Floodlight, Ryu, OpenDaylight
- **Network deployments**
 - Campuses, and research backbone networks
 - Commercial deployments (e.g., Google backbone)

OpenFlow Basics

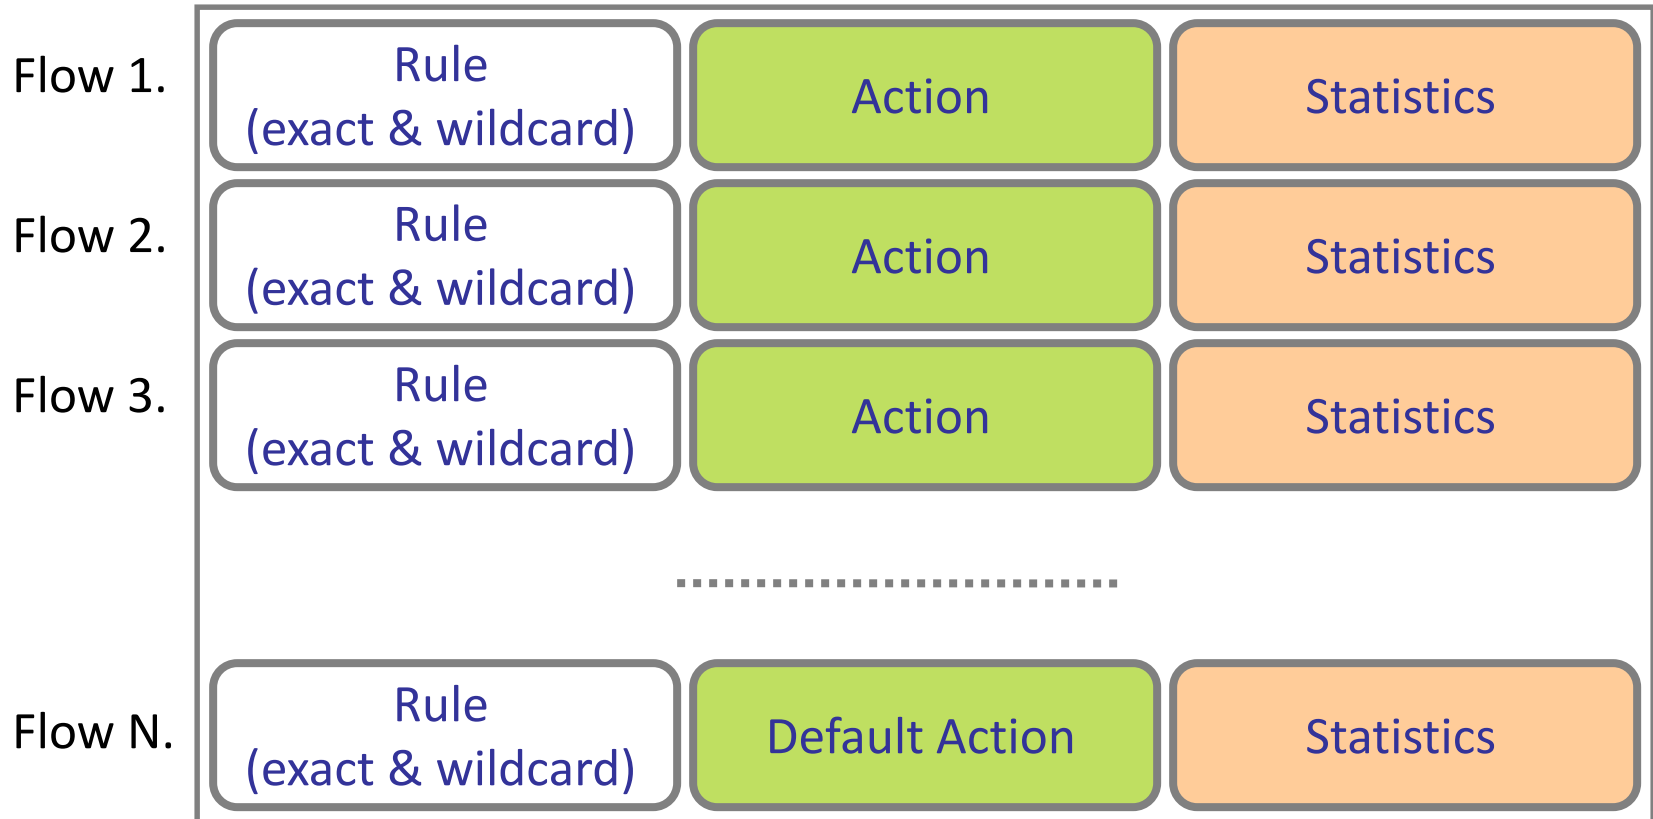


OpenFlow Basics




OpenFlow Rules

Exploit the flow table in switches, routers, and chipsets



Primitives <match, action>

Match arbitrary bits in headers:



Match: 1000x01xx0101001x

- Match on any header, or new header
- Allows any flow granularity

Action

- Forward to port(s), drop, send to controller
- Overwrite header with mask, push or pop
- Forward at specific bit-rate

Data-Plane: Simple Packet Handling



- Simple packet-handling rules
 - Pattern: match packet header bits
 - Actions: drop, forward, modify, send to controller
 - Priority: fixes overlapping patterns
 - Counters: #bytes and #packets



1. `src=1.2.*.*`, `dest=3.4.5.*` → drop
2. `src = *.*.*.*`, `dest=3.4.*.*` → forward(2)
3. `src=10.1.2.3`, `dest=*.*.*.*` → send to controller

Unifies Different Kinds of Boxes

- Router
 - Match: destination IP prefix
 - Action: forward out a link
- Switch
 - Match: destination MAC address
 - Action: forward or flood
- Firewall
 - Match: IP addresses and TCP/UDP port numbers
 - Action: permit or deny
- NAT
 - Match: IP address and port
 - Action: rewrite address and port

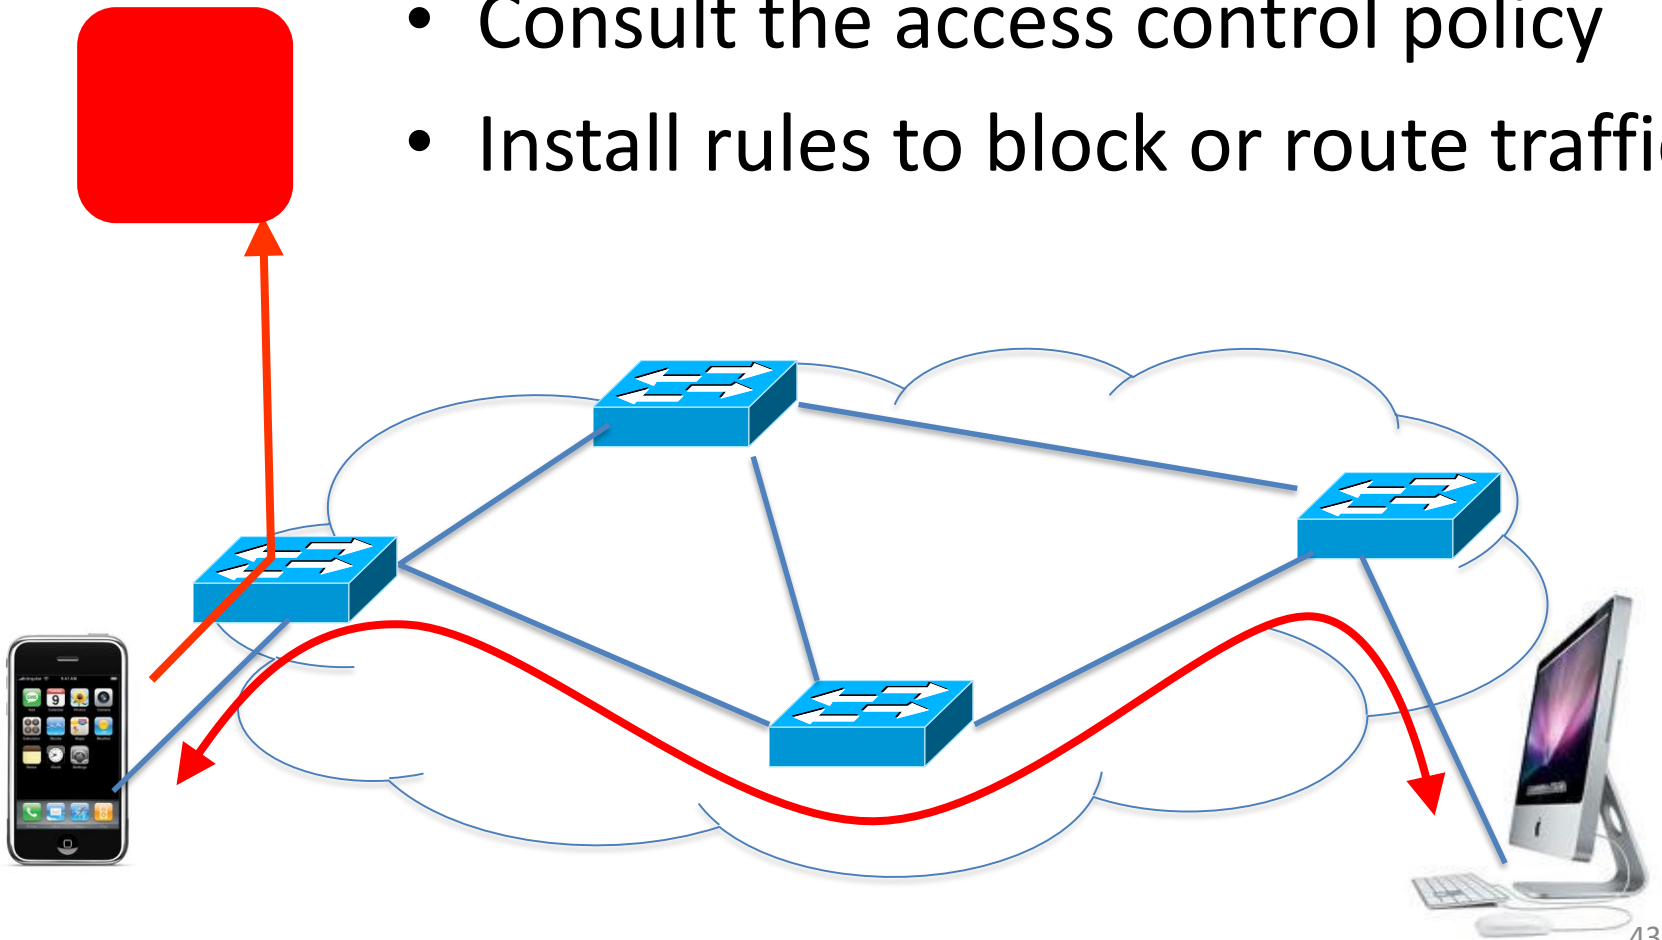
Example OpenFlow Applications

- **Dynamic access control**
- **Seamless mobility/migration**
- **Server load balancing**
- **Network virtualization**
- Using multiple wireless access points
- Energy-efficient networking
- Adaptive traffic monitoring
- Denial-of-Service attack detection


See <http://www.openflow.org/videos/>

E.g.: Dynamic Access Control


- Inspect first packet of a connection
- Consult the access control policy
- Install rules to block or route traffic



E.g.: Seamless Mobility/Migration




- See host send traffic at new location
- Modify rules to reroute the traffic



E.g.: Server Load Balancing

- Pre-install load-balancing policy
- Split traffic based on source IP




E.g.: Network Virtualization

Controller #1

Controller #2

Controller #3


Partition the space of packet headers



SDN CHALLENGES


Heterogeneous Switches

- Number of packet-handling rules
- Range of matches and actions
- Multi-stage pipeline of packet processing
- Offload some control-plane functionality (?)


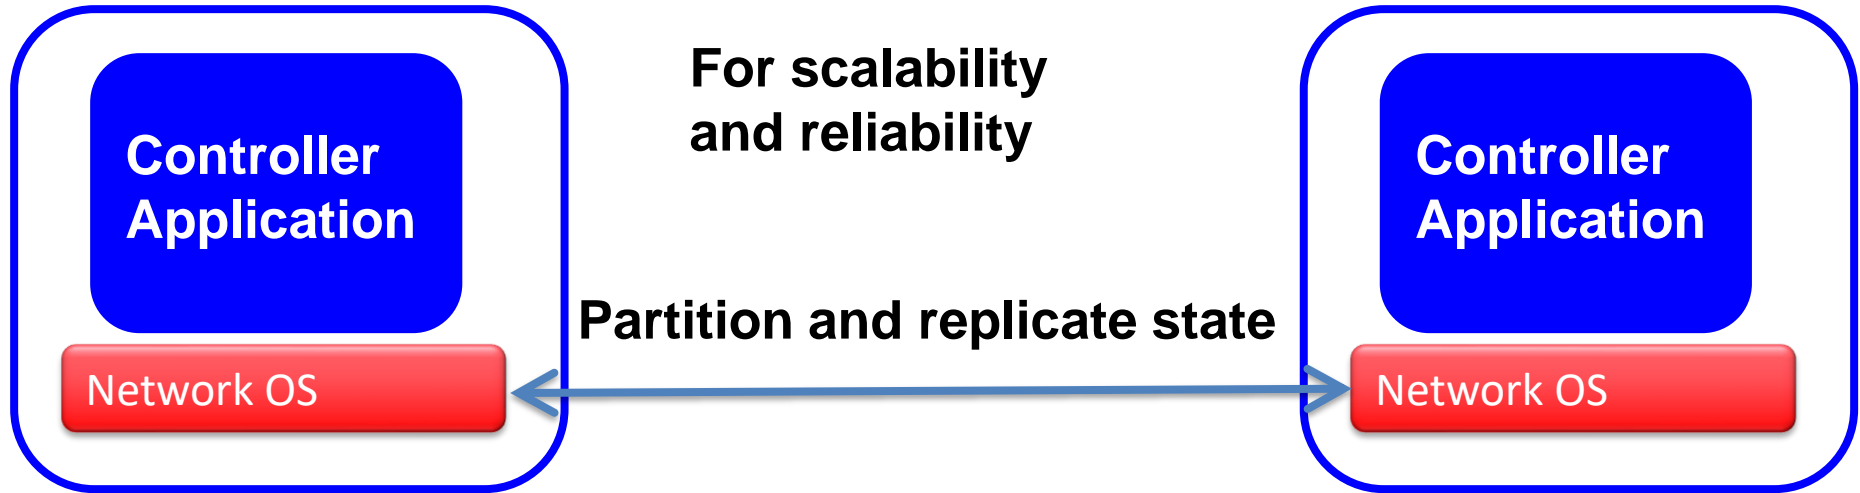


Controller Delay and Overhead

- Controller is much slower than the switch
- Processing packets leads to delay and overhead
- Need to keep most packets in the “fast path”



Distributed Controller




Testing and Debugging

- OpenFlow makes programming the network possible
 - Network-wide view at controller
 - Direct control over data plane
- Plenty of room for bugs
 - Still a complex, distributed system
- Need for testing techniques
 - Controller applications
 - Controller and switches
 - Rules installed in the switches

Programming Abstractions

- **Controller APIs are low-level**
 - Thin veneer on the underlying hardware

- **Need better languages**
 - Composition of modules
 - Managing concurrency
 - Querying network state
 - Network-wide abstractions



Conclusion

- Rethinking networking
 - Open interfaces to the data plane
 - Separation of control and data
 - Leveraging techniques from distributed systems
- Significant momentum
 - In both research and industry