

# Process Management I

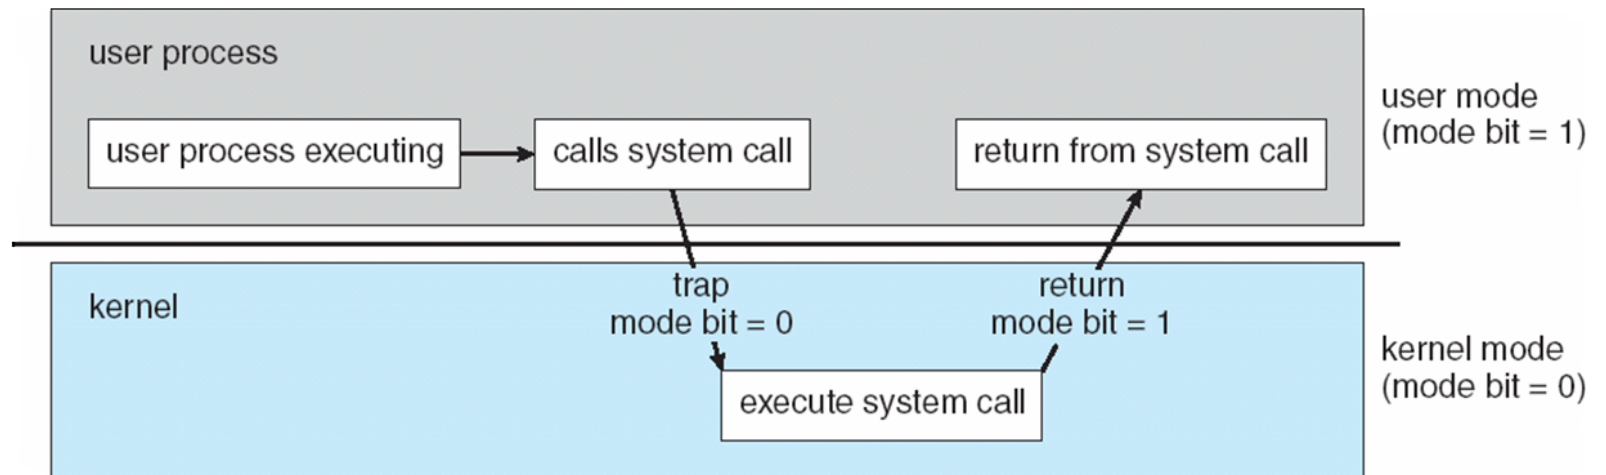
## Processes



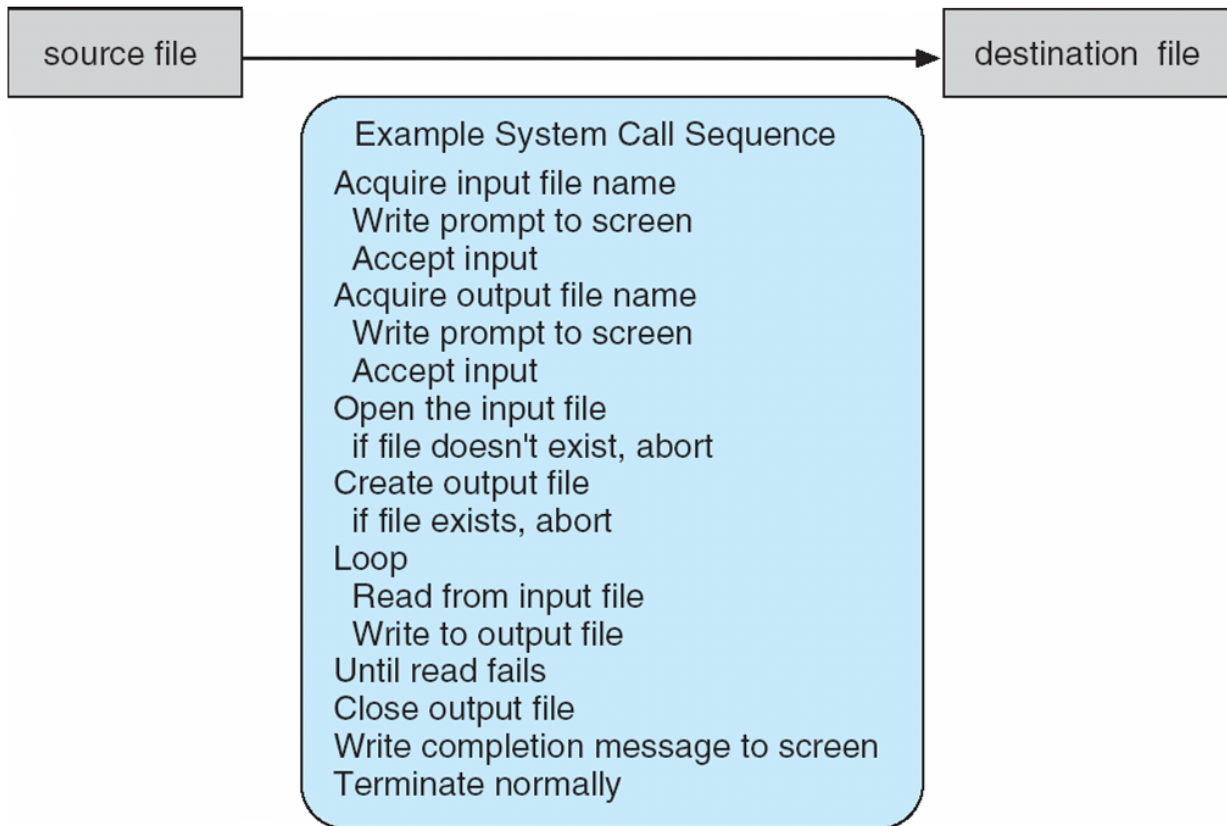
**School of Computer Science,  
UCD**

**Scoil na Ríomheolaíochta,  
UCD**

# Last Week...

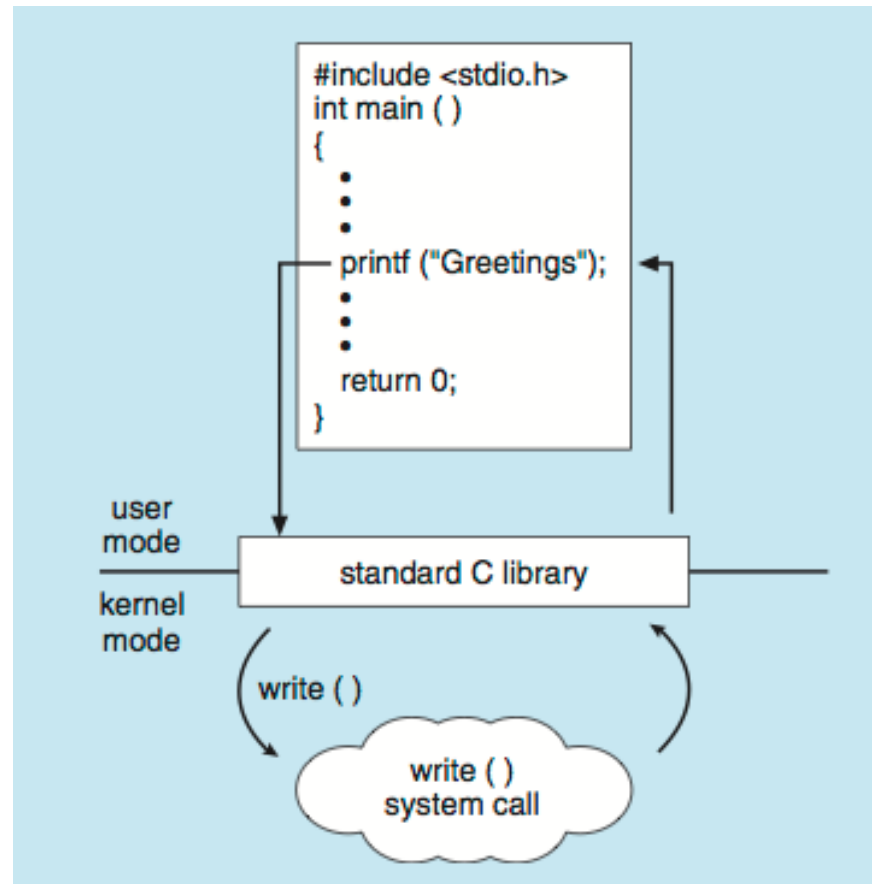


# Last Week...



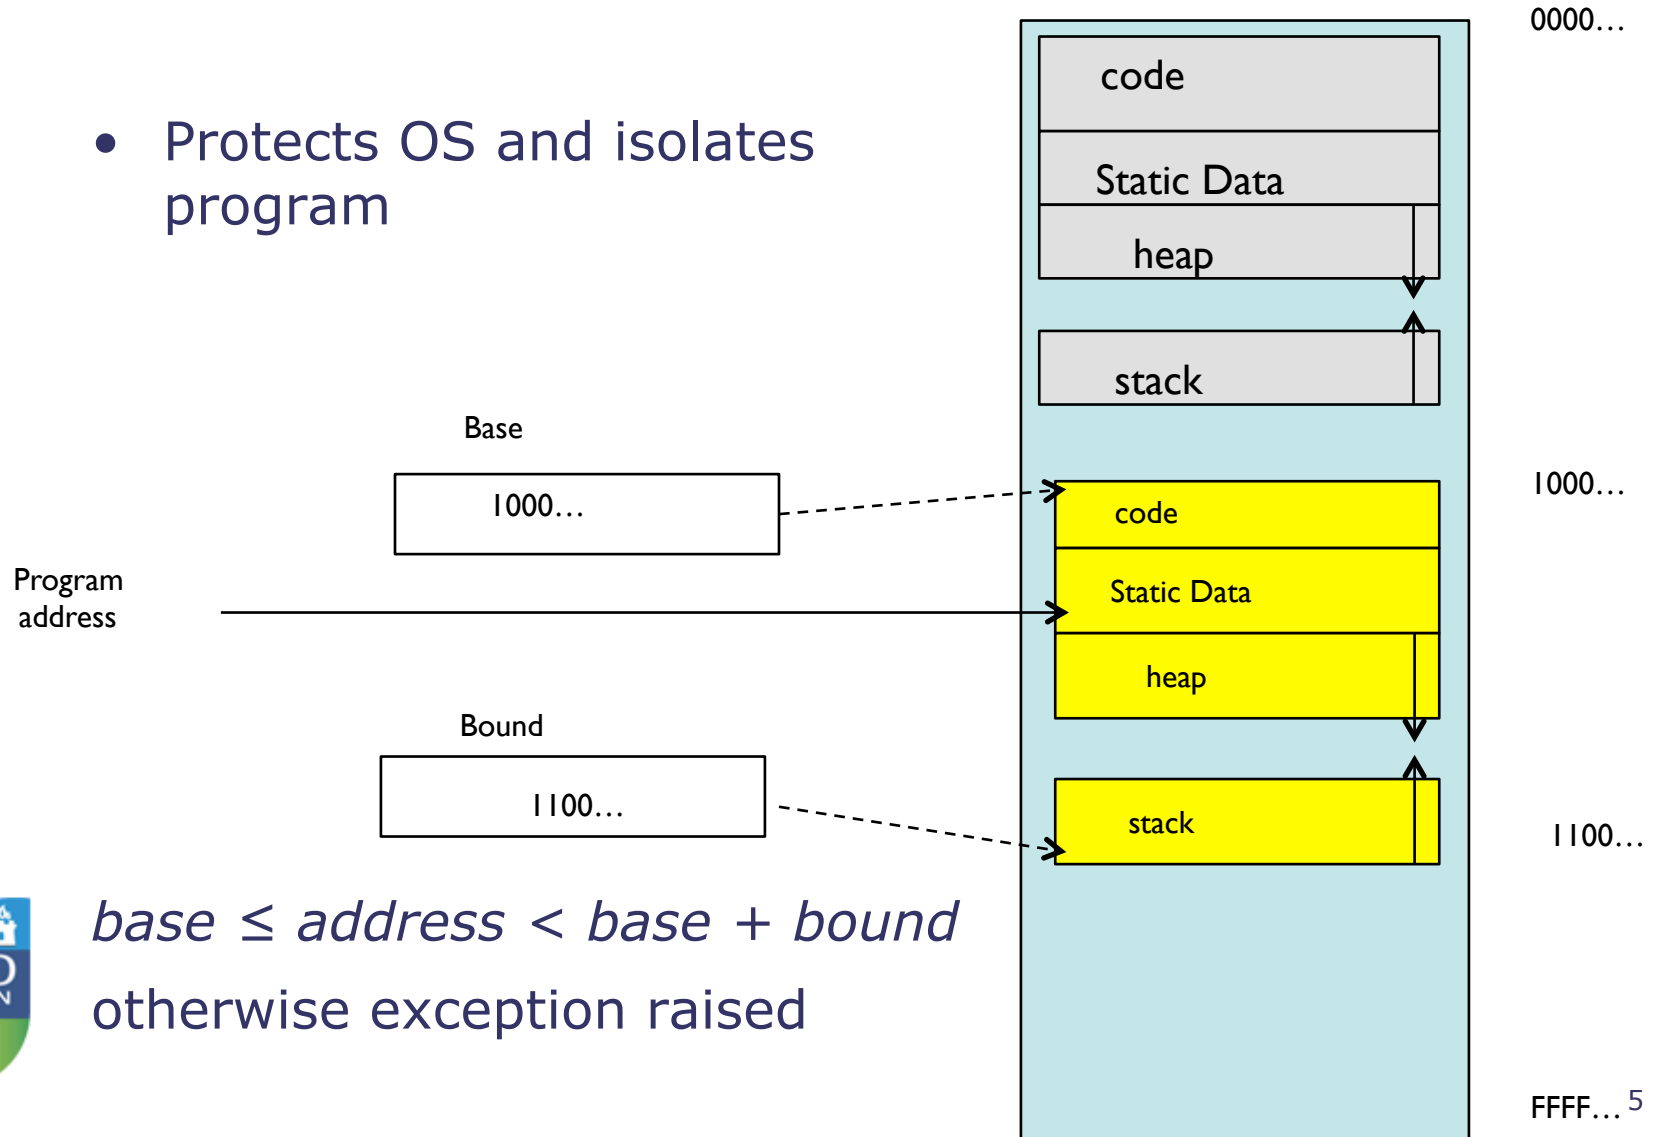
# Last Week...

- C program invoking printf() library call, which calls write() system call



# Last Week...

- Protects OS and isolates program



## Last Week...

- Apart from protecting memory and I/O, we must ensure that the OS always maintains control
  - A user program might get stuck into an infinite loop and never return control to the OS
- **Timer:** Generates an interrupt after a fixed or variable amount of execution time
  - OS may choose to treat the interrupt as a fatal error (and stop program execution) or allocate more execution time
  - note: with time-sharing a timer interrupt is periodically generated in order to schedule a new process



# Last Week...

- A script bash is a bash “text” in a text file
  - First line always contains the following:  
***#!/bin/bash***
  - a bash program needs to be made executable:  
***chmod u+x my\_script.sh***
  - the .sh extension is just a convention
  - to execute the script:  
***./my\_script.sh [arg1 arg2 ...]***



# Outline

- What is a Process (Process/Program)
- Program Control Block
- Process features:
  - Process Switching
  - Process Creation and Termination
  - Communication between Processes

Take home message:

*A process is a program in execution, which forms the basis of all computation*





# OS Activities

- Early computer systems executed one program at a time
- All modern OS execute many kinds of activities concurrently (basis of multiprogramming)
  - user programs
  - batch jobs and command scripts
  - system programs
- Each “execution entity” is encapsulated in a process
  - in different OSs also called: task, job, actor, . . .
- The OS takes care of most aspects concerning processes:
  - it creates, deletes, suspends and resumes processes
  - it schedules & manages processes (resources, IPC, etc.)



# Program vs. Process(s)

- Program is ***passive*** entity stored on disk (***executable file***), process is ***active***
  - Program becomes process when executable file loaded into memory
- ***A process is one instance of a program in execution***
  - A process requires a section of main memory to run
  - At any time, there may be more than one process running a different instance of the same program
    - e.g. several processes running the same editor
  - One program can spawn several processes




# Process Abstraction

- **Process**: an instance of a program, running with limited rights
  - **Thread**: a sequence of instructions within a process
    - Potentially many threads per process (for now 1:1)
  - **Address space**: set of rights of a process
    - Memory that the process can access
    - Other permissions the process has (e.g., which system calls it can make, what files it can access)



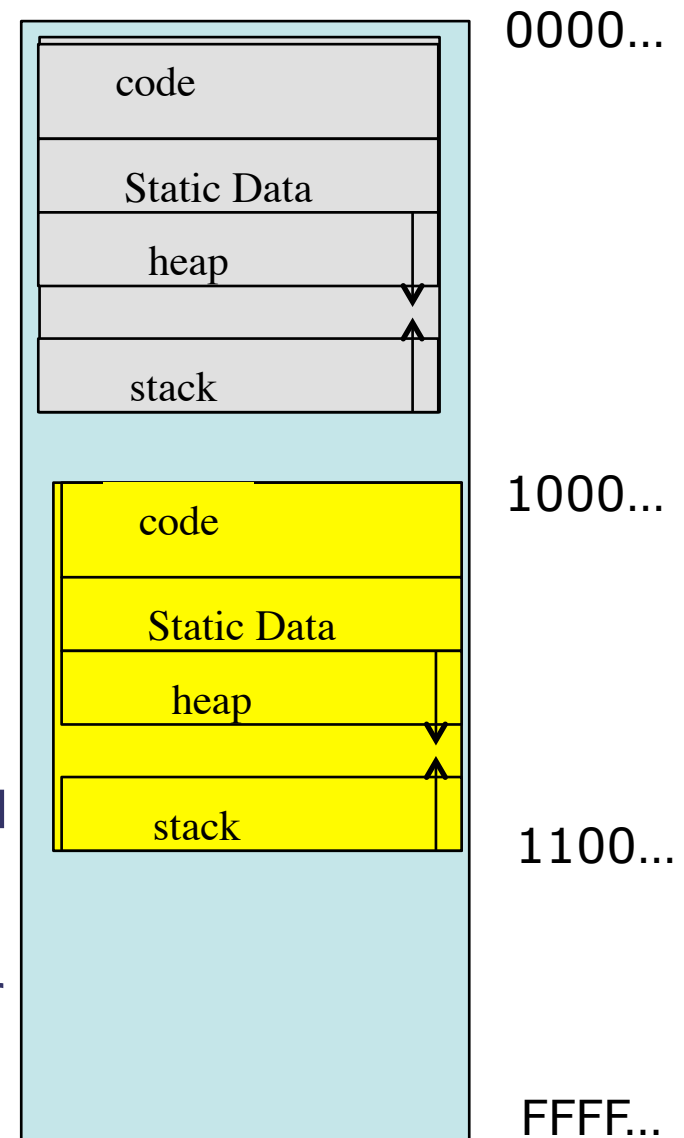
# Process Concept

- A process includes all information needed to run its program. Process address space (main memory section), divided into:
  - **Text region** (executable code or program)
  - **Data region** (global variables)
  - **Heap** (containing memory dynamically allocated during run time)
  - **Stack region** (return address and local variables for active procedure calls)



# Process Concept

- ***Each process runs in its own address space***
  - The address space of two different processes translates into two different ***physical*** addresses in main memory
  - A process cannot directly read or write memory addresses belonging to another process
  - Example: with base and bound registers, a process' address space goes from ***base to bound*** (which are different for all processes)



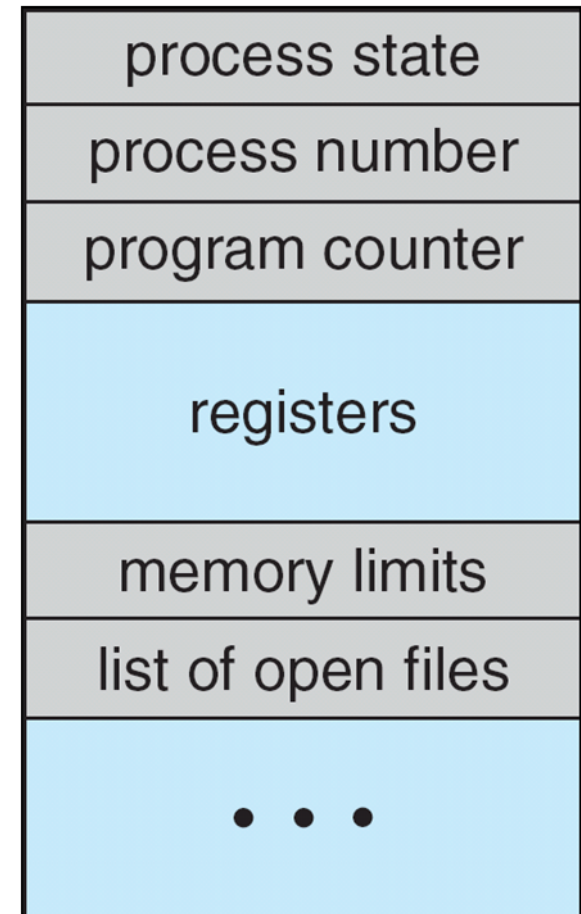
# Process States and PCB

- The states a process may be in are either active or suspended
- **Active** states of a process:
  - **Running**: using a processor to execute instructions
  - **Ready**: executable, but all system processors (CPUs) are currently in use; usually, many more processes than processors in a system
  - **Blocked**: waiting for an event to occur
- The process state is represented by a dynamic data structure called **process control block (PCB)**, or process descriptor
  - PCB describes the status of all resources used by a process
  - PCB contains critical info. and must be kept in memory protected from user access
- Kernel Scheduler maintains data structure containing PCBs (**Process Table**)
  - Scheduling algorithm selects the next one to run
  - each entry contains a PID and a pointer giving the address of that process's PCB in memory



# Process Control Block (PCB)

- Information associated with each process (also called **task control block**)
- **Process state** – running, waiting, etc.
- **Program counter** – location of instruction to execute next
- **CPU registers** – contents of all process-centric registers
- **CPU scheduling information**- priorities, scheduling queue pointers
- **Memory-management information** – memory allocated to the process
- **Accounting information** – CPU used, clock time elapsed since start, time limits
- **I/O status information** – I/O devices allocated to process, list of open files



# Process Operations

- A number of operations can be performed to create/ delete a PCB (and thus a process), or to change process information in the PCB
  - **create**: creates and initialises a new PCB
  - **delete (destroy, kill)**: removes a process from the system
  - **wait (block)**: execution ceases until a specific event occurs
  - **signal**: indicates that a specific event has occurred
    - event: a relevant change in the status of an entity that a process requires in order to execute
    - two examples of I/O events:
      - event triggering a wait operation: “write file”
      - event completion, triggering a signal operation: “write file complete” (hardware interrupt)



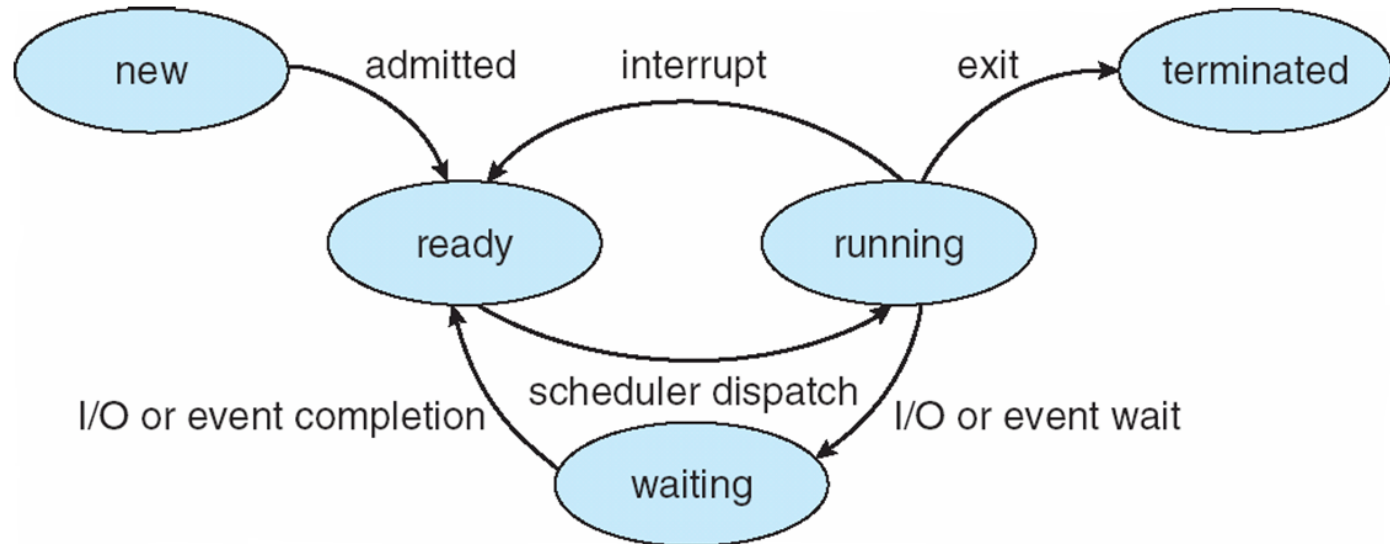


# Process Operations (ii)

- Operations in the PCB (con'td)
  - ***schedule (dispatch)***: assigns ready process to CPU
  - ***change priority***: alters the scheduling priority of a process
    - schedule is usually triggered by some system change altering the decision on which process should be currently running
    - it depends on the scheduling algorithm used by the OS
    - a time-out can trigger a schedule operation
- Operations involving suspended states:
  - ***suspend (sleep)***: suspends an active process
  - ***resume (wake up)***: places a suspended process into blocked or ready state
    - on being suspended, a process is removed from contention for time on a processor (i.e., contention for being in running state), without being actually deleted



# Process State Diagram (Life Cycle of a Process)



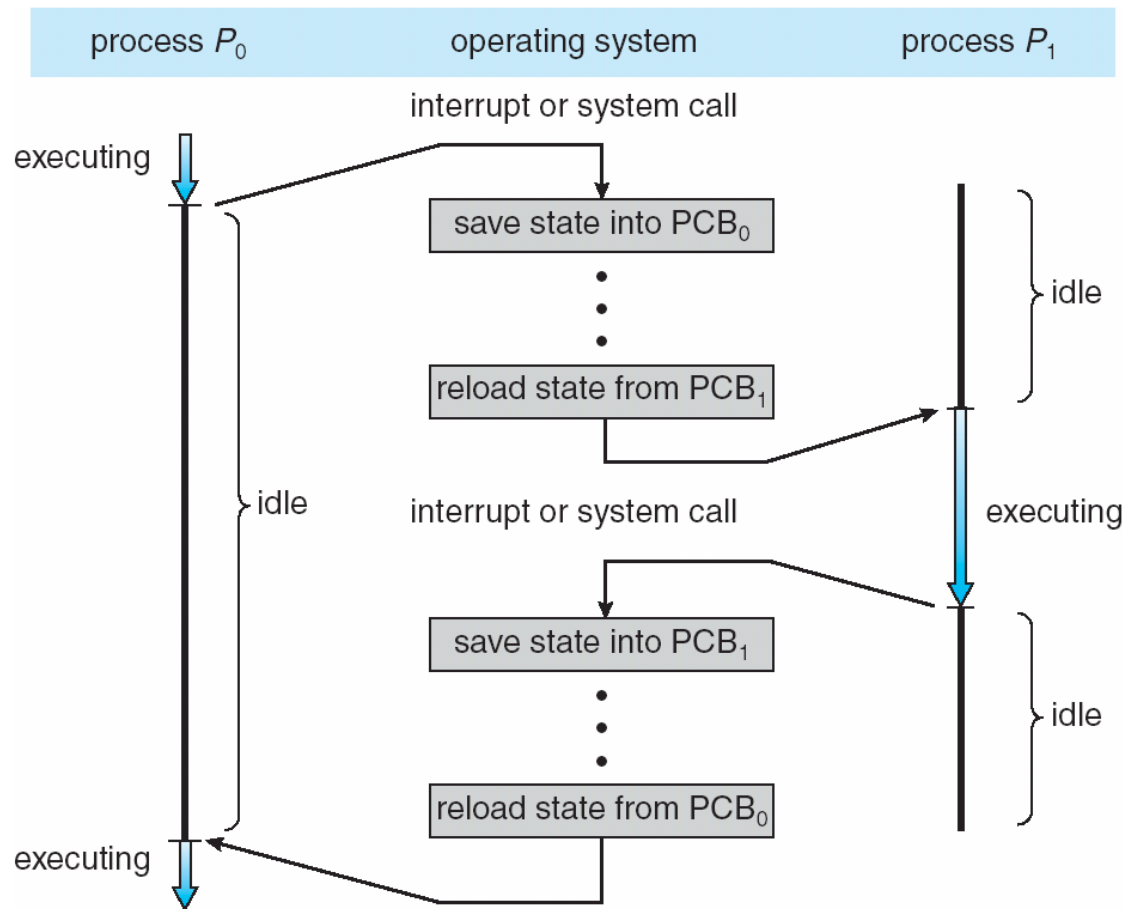
- As a process executes, it changes state
  - **New**: The process is being created
  - **Ready**: The process is waiting to run
  - **Running**: Instructions are being executed
  - **Waiting**: Process waiting for some event to occur
  - **Terminated**: The process has finished execution

# Context Switch

- A **context switch** happens when a running process stops execution, and another process (initially in the ready state) starts execution
  - CPU Switch From Process to Process
- Dispatching the ready process requires that the system:
  - Saves the state of the running process in its PCB
  - Loads the saved state from the ready process' PCB
- Context-switch time is **overhead**. The system does not do useful work during a context switch; therefore:
  - Context switches should be quick (**hardware dependent**)
  - Number of context switches per unit time should be minimised (**software dependent**); however, consider interactivity
  - Overhead sets minimum practical switching time



# Context Switch



# Process Creation and Termination

- Processes are **created** by two main events:
  - System boot
  - Execution of process creation system call by another process
    - examples: user request to create another process, initiation of a batch job, . . .
- Processes are **terminated** due to
  - Voluntary conditions: normal exit, error exit
  - Involuntary conditions: fatal error, killed by another process
- Creation and termination **states**:
  - A created process goes to the ready or suspended-ready states (the very first process can go to the running state directly)
  - Voluntary termination is from the running state
  - Involuntary termination may happen in any state

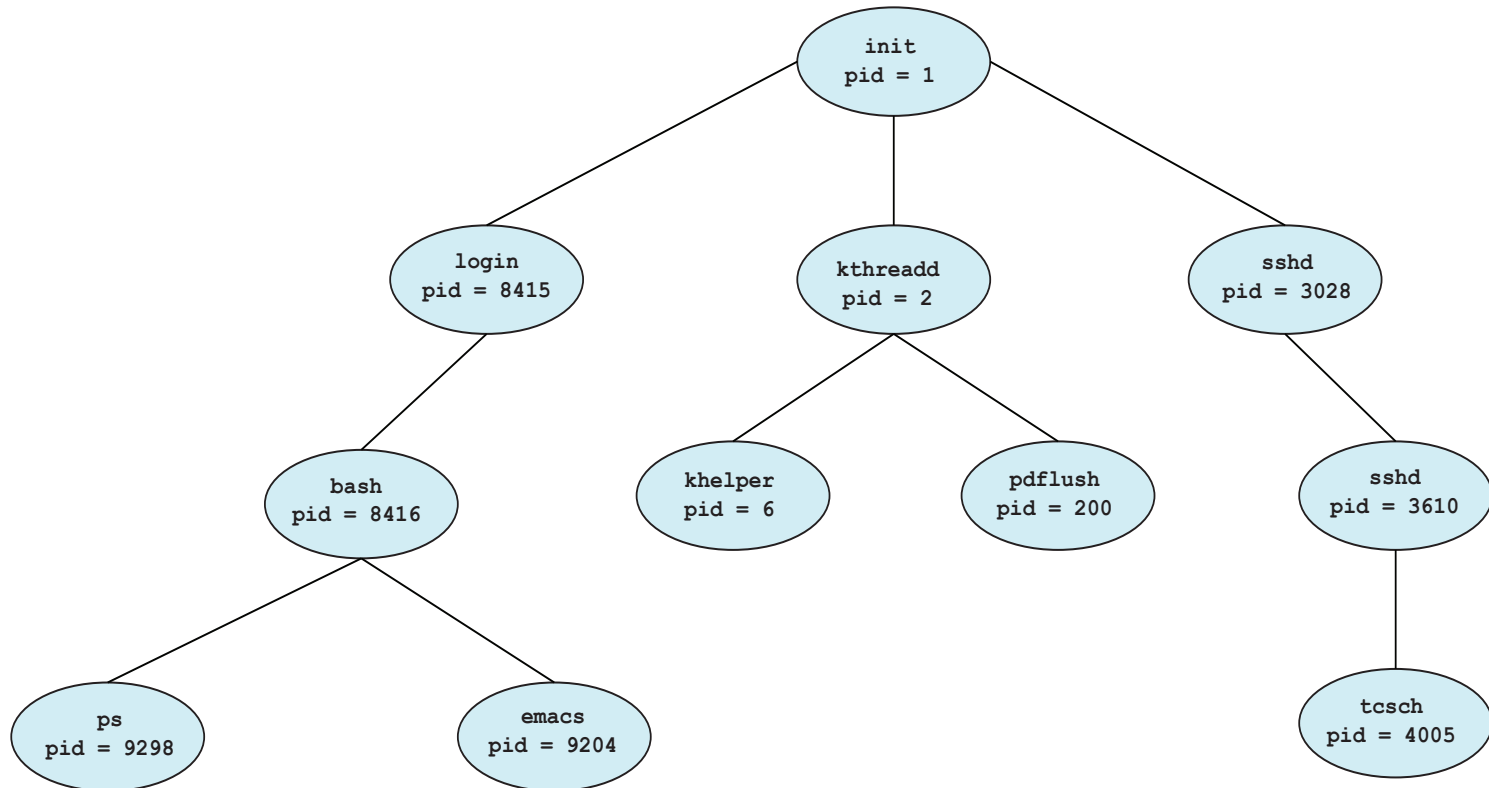


# Process Creation

- **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- Generally, process identified and managed via a **process identifier (pid)**
- Resource sharing options
  - Parent and children share all resources
  - Children share subset of parent's resources
  - Parent and child share no resources
- Execution options
  - Parent and children execute concurrently
  - Parent waits until children terminate



# A Tree of Processes in Linux



# pstree


```
Documents — anthony@hibernia: ~ — ssh hibernia.ucd.ie — 64x22
anthony@hibernia:~$ pstree
init--NetworkManager--2*[{NetworkManager}]
    |--acpid
    |--apache2--10*[apache2]
    |--at-spi-bus-laun--dbus-daemon
    |                       |--3*[{at-spi-bus-laun}]
    |--at-spi2-registr--{at-spi2-registr}
    |--atd
    |--avahi-daemon--avahi-daemon
    |--bluetoothd
    |--console-kit-dae--64*[{console-kit-dae}]
    |--cron
    |--cupsd
    |--2*[dbus-daemon]
    |--dconf-service--2*[{dconf-service}]
    |--dhclient
    |--dnsmasq
    |--fail2ban-server--2*[{fail2ban-server}]
    |--gconfd-2
    |--6*[getty]
    |--gnome-screensav--2*[{gnome-screensav}]
    |--gvfsd--{gvfsd}
```





# Process Creation (cont'd)

- **Address space**
  - Child duplicate of parent
  - Child has a program loaded into it
- UNIX examples
  - **fork()** system call creates new process
  - **exec()** system call used after a fork() to replace the process' memory space with a new program



# Process in an OS

The OS manages for the process:

- **Resources**
  - I/O operations: open, close, read, write, . . .
  - memory operations: sbrk (used by malloc and free)
- **IPC**
  - global data: shared memory segments (shmget, shmop, . . . )
  - message-based communications: pipes, sockets, streams
  - synchronisation: semaphores, etc. (semget, semop)
  - (note: examples above are Unix system calls)
- Process is the elementary resource sharing agent in an OS



# Interprocess Communication (IPC)

Two basic IPC models:

- ***Shared memory***

- A shared memory region usually resides in the space address of the process creating a shared segment
- Fast (only creating and attaching a shared segment are system calls)
- but: data format and access is controlled by the processes sharing the segment, not the OS (thus conflicts are possible)

- ***Message passing***

- Messages are sent by one process and received by another via a buffer/mailbox/port in kernel space
- Slower (send/receive operations are system calls)
- but: processes never share memory and thus do not need to trust each other



# Conclusion

- An instance of an executing program is a process consisting of an address space and one or more threads of control
- As a process executes it changes state (***new, ready, running, waiting, terminated***) and transitions between states with operations (***create, delete (destroy, kill), wait (block), signal, schedule (dispatch), change priority, suspend (sleep), resume (wake up)***)
- ***Parent*** process create children processes, which, in turn create other processes, forming a ***tree*** of processes
- Interprocess communication can take place using shared memory and message passing

