

Process Management II

Threads



School of Computer Science,
UCD

Scoil na Ríomheolaíochta,
UCD

Outline

- Explain the challenges of process cooperation
- Understand the concept of a thread, its motivation, lifecycle, states and role in concurrency
- Discuss single vs. multithreading
- Explore concurrency accomplished by multiplexing CPU Time

Take home message:

Program -> Process -> Thread



Independent Processes

- **Definition:** Independent processes are those that can neither affect nor be affected by the rest of the system
 - Two independent processes cannot share system state or data
 - Processes running on different non-networked computers
- Properties:
 - **Deterministic behaviour:** only the input state determines the results → **reproducible**
 - Can be stopped and restarted with no detrimental effects



Cooperative Processes

- **Definition:** Cooperative processes are those that share something (not necessarily for a purpose)
 - Two processes are cooperative if the execution of one of them may affect the execution of the other
 - Processes that share a single file system
- Properties:
 - **Nondeterministic behaviour** (from the point of view of one of the processes) → it may be **difficult to reproduce**
 - hence: testing & debugging may be troublesome
 - Subject to **race conditions**: the outcome of the process may depend on the sequence or timing of events in other processes



Why Allow Processes to Cooperate?

- Resources and information sharing
 - One computer, many users
- System speed-up, by introducing concurrency into program execution
 - Overlap I/O with computations
- Convenience
 - Editing a file at the same time it is being printed
- Modularity, simplicity, divide & conquer
 - structured programming, object programming, etc.



Process and Collaboration

- Multiple concurrent cooperative activities necessarily happen within any OS
- Why not define each and every one of these concurrent activities as a different process?
- ***Processes are not ideal for cooperation:***
- Processes are not very efficient:
 - Creation of a new process is costly
 - All the process structures must be allocated upon creation
- Processes don't (directly) share memory
 - Each process runs in its own address space
 - Parallel and concurrent processes often want to manipulate same data
 - Most communications go through the OS: slow



Motivation for Threads

- Example: consider a file server process that occasionally has to block waiting for the disk to respond (I/O wait)
 - Assume that it keeps a cache of recently used files in its memory space, in order to speed up future operations
 - We might think of running a second concurrent file server to serve files while the first file server waits for disk I/O (principle of multiprogramming)
- However concurrency is not efficient using processes
 - They would have to run in the same address space to efficiently share a common cache (but they do not)
 - We could use shared memory, but slow system calls would be required
- ***Solution: threads***
 - More than one thread of control within a single process
 - i.e.: more than one active entity within a single process



Threads

- Modern OSs support both entities (process & thread):
multi-threaded OS
 - ***Process***: defines the address space and general process attributes
 - ***Thread***: defines a single sequential execution stream within a process
 - ***Multithreading***: a single program made up of a number of different concurrent activities
- Concurrency in some existing OS:
 - MS-DOS: one address space, one thread
 - Unix (originally): multiple address spaces, one thread per address space
 - Mach, Chorus, Solaris, NT: multiple address spaces, multiple threads per address space (multi-threading)



More on Threads

Threads' features:

- ***Cheap to create*** (no need to allocate PCB, new address space); they can be created statically or dynamically (by a process or by another thread)
- Threads ***do not exist on their own***, they belong to processes: number of threads in a process ≥ 1
- Threads can ***communicate with each other efficiently*** through the process global variables or through common memory, using simple primitives
- Threads ***facilitate concurrency***, and therefore are useful even on single processor systems
- If a thread needs a service provided by the OS (system call) it ***acts on behalf of the process*** it belongs to

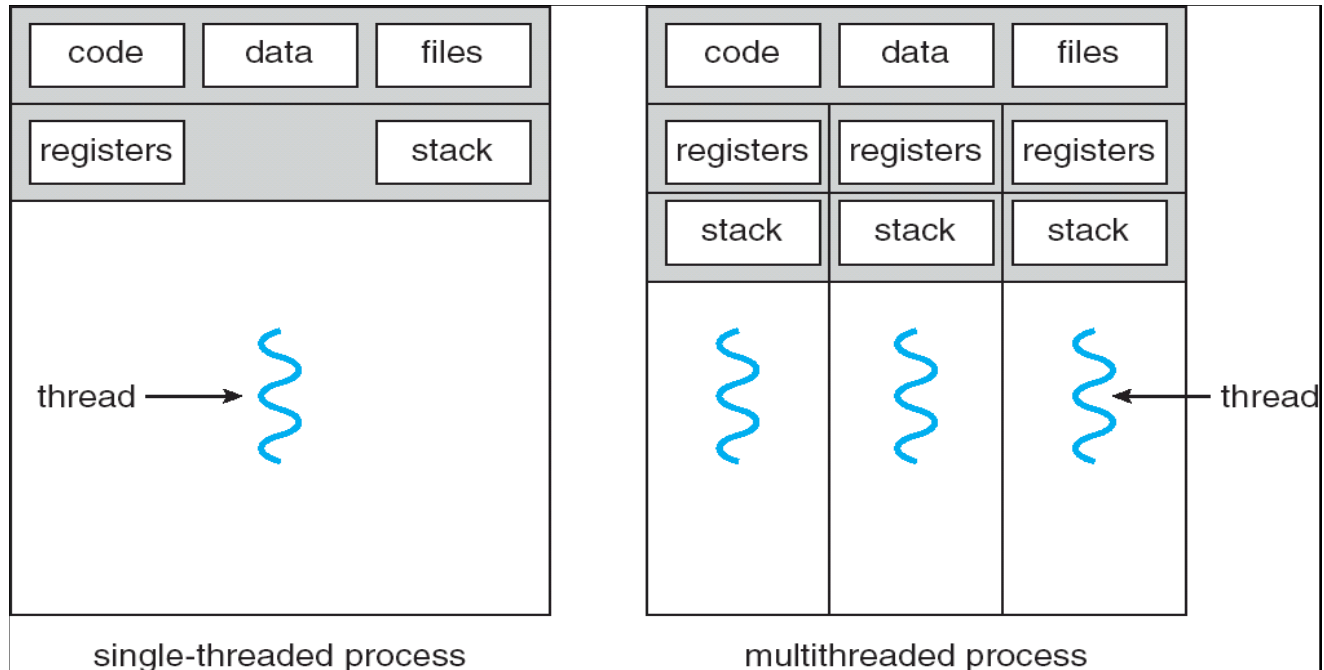


Examples of Multithreaded Programs

- Embedded systems
 - Elevators, Planes, Medical systems, Wristwatches
 - Single Program, concurrent operations
- Most modern OS kernels
 - Internally concurrent because have to deal with concurrent requests by multiple users
 - But no protection needed within kernel
- Database Servers
 - Access to shared data by many concurrent users
 - Also background utility processing must be done



Single and Multithreaded Processes



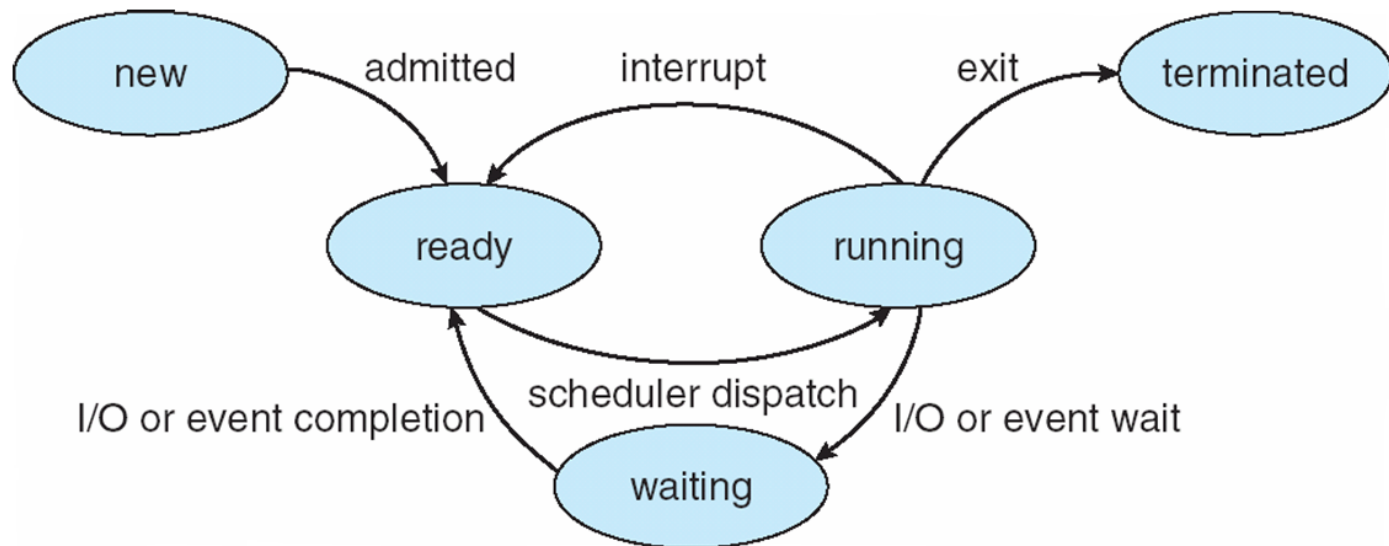
- Threads encapsulate concurrency: “Active” component
- Address spaces encapsulate protection: “Passive” part
 - Keeps buggy program from trashing the system
- Why have multiple threads per address space?

Thread State

- **State shared** by all threads in process/address space
 - Address space (code and data)
 - Content of memory (global variables, heap)
 - I/O state (file descriptors, open files, network connections, etc)
 - Privileges
 - Timers, Signals, Semaphores
 - Accounting information
- **State "private"** to each thread
 - Kept **in TCB = Thread Control Block**
 - CPU register set, in particular:
 - program counter (PC)
 - stack pointer (SP)
 - interrupt vectors
 - Stack (Parameters, temporary variables)
 - State
 - Child threads (threads can spawn new threads)





Lifecycle of a Thread (or Process)



- As a thread executes, it changes state:
 - **New**: The thread is being created
 - **Ready**: The thread is waiting to run
 - **Running**: Instructions are being executed
 - **Waiting**: Thread waiting for some event to occur
 - **Terminated**: The thread has finished execution
- “Active” threads are represented by their TCBs
 - TCBs organized into queues based on their state


Thread Abstraction

- Infinite number of processors
- Threads execute with variable speed
 - Programs must be designed to work with any schedule



Multiprocessing vs. Multiprogramming

- Definitions:
 - Multiprocessing \equiv Multiple CPUs
 - Multiprogramming \equiv Multiple Jobs or Processes
 - Multithreading \equiv Multiple threads per Process
- What does it mean to run two threads “concurrently”?
 - Scheduler is free to run threads in any order and interleaving: FIFO, Random, ...
 - Dispatcher can choose to run each thread to completion or time-slice in big chunks or small chunks



Programmer vs. Processor View

Programmer's View	Possible Execution #1	Possible Execution #2	Possible Execution #3
.	.	.	.
.	.	.	.
.	.	.	.
$x = x + 1;$	$x = x + 1;$	$x = x + 1$	$x = x + 1$
$y = y + x;$	$y = y + x;$	$y = y + x$
$z = x + 5y;$	$z = x + 5y;$	thread is suspended
.	.	other thread(s) run	thread is suspended
.	.	thread is resumed	other thread(s) run
.	thread is resumed
		$y = y + x$
		$z = x + 5y$	$z = x + 5y$



Conclusion

- It is difficult for processes to cooperate
- **Thread**: defines a single sequential execution stream within a process
- Multithreading: a single program made up of a number of different concurrent activities
- Threads encapsulate concurrency: the “Active” component of a process

