

Process Management IV

Synchronisation (2)



School of Computer Science,
UCD

Scoil na Ríomheolaíochta,
UCD

Announcements

- ***In-class test:***
 - 25th or 26th October(room not known yet)
 - 50 minutes
 - 20%



Outline

- Understand the Producer Consumer synchronisation problem
- Understand and be able to use higher-level synchronisation techniques:
 - monitors
 - messages

Take home message:

Semaphores can lead to deadlock and monitors can solve this.




Producer/Consumer Problem

- A common synchronisation problem
- Consider a **producer process** supplying a resource (also called message) to a **consumer process**
 - **Producer**: creates instances of a resource
 - **Consumer**: uses up (destroys) instances of a resource
- Producer and consumer share a **buffer** into which resources are placed by producer and from which resources are removed by consumer
 - Buffer has finite size (fixed number of slots)
- Example: use of a printer
 - Printing process: producer
 - Spooler: buffer
 - Printer: consumer




Producer/Consumer Problem Constraints

- Consumer must wait for producer if the buffer is empty



- Producer must wait for consumer if the buffer is full



- Synchronisation:** keeping producer and consumer in step
 - Buffer constitutes a CS of the Producer/Consumer problem
 - Usually need mutually exclusive access to the buffer
- Note: here we have assumed just one producer and one consumer; in general there can be more than one producer and more than one consumer sharing a buffer

Producer/Consumer Problem: Naive Solution Attempt with Semaphores

```
semaphore S(1,NULL);
```

Producer

```
while(true) {  
    msg=produce_message();  
    P(S);  
    if(buffer_full)  
        wait();  
    put_message(msg);  
    V(S);  
}
```

Consumer

```
while(true) {  
    P(S);  
    if(buffer_empty)  
        wait();  
    msg=get_message();  
    V(S);  
    consume_message(msg);  
}
```

- Put message() and get message(): put a message in/ get a message from the buffer, respectively
- If buffer is full or empty “solution” above does not work



Producer/Consumer Algorithm with Semaphores

- Three semaphores are needed for a true solution:

<pre>int N=buffer_size; semaphore S(1,NULL); semaphore full_s(0,NULL); semaphore empty_s(N,NULL);</pre>	<p>(number of slots in the buffer) (anybody accessing the buffer?) (are there full slots?) (are there empty slots?)</p>
---	---

Producer

```
while(true) {  
    msg=produce_message();  
    P(empty_s);  
    P(S);  
    put_message(msg);  
    V(S);  
    V(full_s);  
}
```

Consumer

```
while(true) {  
    P(full_s);  
    P(S);  
    msg=get_message();  
    V(S);  
    V(empty_s);  
    consume_message(msg);  
}
```

- full_s and empty_s are **counting semaphores**



Producer/Consumer Algorithm with Semaphores (Optimisation)

- It is also possible to allow simultaneous access to the buffer to one producer and one consumer: four semaphores are required

```
int N=buffer_size;  
semaphore Sp(1,NULL);  
semaphore Sc(1,NULL);  
semaphore full_s(0,NULL);  
semaphore empty_s(N,NULL);
```

(a producer accessing the buffer?)
(a consumer accessing the buffer?)

Producer

```
while(true) {  
    msg=produce_message();  
    P(empty_s);  
    P(Sp);  
    put_message(msg);  
    V(Sp);  
    V(full_s);  
}
```

Consumer

```
while(true) {  
    P(full_s);  
    P(Sc);  
    msg=get_message();  
    V(Sc);  
    V(empty_s);  
    consume_message(msg);  
}
```



Monitors

- Semaphores are nice but they have drawbacks:
 - Joint use of more than one semaphore can lead to **deadlocks**, if the order of Ps is not correctly set
 - → difficult to program with semaphores
- Monitors are higher-level sync primitives (i.e., tied to a higher-level programming language) proposed by Hoare (1974)
 - A **monitor** is a collection of procedures, variables and data grouped together in a special kind of structure
 - Goal is to avoid the “catches” that can arise when protecting critical sections (CS) with semaphores



Monitor Structure

- Example: Monitor in pseudo-code

```
monitor example {  
    int i;                (example internal data)  
    condition c;          (example condition variable)  
  
    void p() {             (example monitor procedure)  
        ...  
    }  
}
```

- Rules:

- Processes may call monitor procedures whenever they wish
- Processes can't access internal monitor data (variables, etc) using external procedures
- At most, only one process at any time can be active inside a monitor
 - This means ME inside the monitor, and it is guaranteed by the compiler



Blocking in Monitors

- ME is implemented relying on internal **condition** variables, which have two special operations associated:
 - 1. *wait(condition)***: executed when the monitor discovers that a process cannot proceed
 - The monitor blocks a process calling **wait()** and makes it wait on **condition** (blocked, hence out of the monitor)
 - Another process can be then allowed into the monitor
 - A process that invokes **wait()** is **always blocked** (note: this is unlike invoking P in semaphores)
 - 2. *signal(condition)***: executed to wake up a process waiting on condition
 - After executing **signal()** the calling process must exit the monitor immediately (in order to guarantee ME)



Producer/Consumer Algorithm with Monitors

```
monitor pr_co {  
    int count;  
    condition full_s, empty_s;  
  
    void put(msg) {  
        if(count==N) wait(empty_s);  
        put_message(msg);  
        count++;  
        if(count==1) signal(full_s);  
    }  
  
    msg get() {  
        if(count==0) wait(full_s);  
        msg=get_message();  
        count--;  
        if(count==N-1) signal(empty_s);  
    }  
}
```

Producer

```
while(true) {  
    msg=produce_message();  
    pr_co.put(msg);  
}
```

Consumer

```
while(true) {  
    msg=pr_co.get();  
    consume_message(msg);  
}
```

- Simplest possible consumer and producer code
- Monitor takes care of any issues, not producer or consumer




Message System

- Synchronisation with both semaphores and monitors relies on ***shared memory***
- The message system is a mechanism for inter-process communication (IPC), that can also be relied upon for synchronisation ***without having to share memory***
- Elements in the message system
 - ***Message:*** Information that can be exchanged between two (or more) processes or threads
 - ***Mailbox:*** A place where messages are stored between the time they are sent and the time they are received
 - It usually resides in kernel space




Message System Scheme

- Process A sends messages to Process B using a mailbox:



- Example: a **pipe** can be implemented with messages



- Remarks:
 - One process or another owns the data; never two at the same time (as it happens with data in shared memory)
 - Duplex communication (streams) would require two mailboxes
 - The message system is an example of a Producer /Consumer setting

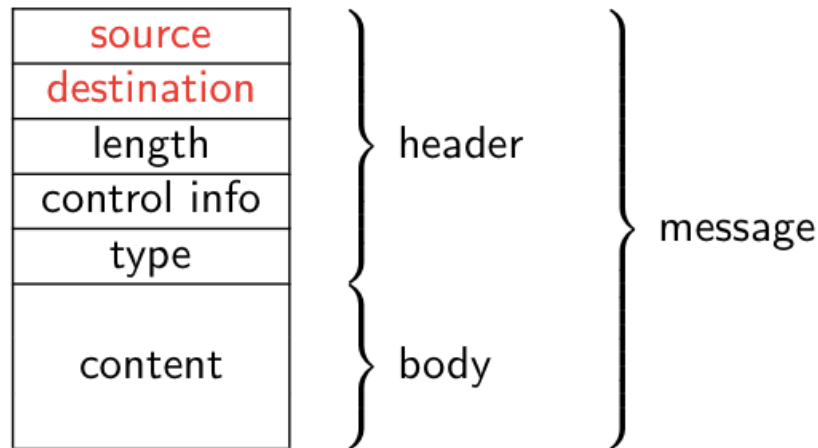
Message System: Operations & Addressing

- Two basic operations (aka primitives):
 - **send**(mailbox, message): place message in mailbox
 - **receive**(mailbox, message): remove message from mailbox
- Addressing
 - Message identifies source and destination processes (for instance through their PIDs)
 - May be multiple destinations, or destination may be unspecified (broadcast)
- In practice we also need the system calls **create**(mailbox), **delete**(mailbox)



Message Format

- Depends on the objectives of the message facility
 - Fixed-length: minimise processing and storage overhead
 - Variable-length: more flexible approach for Oss



- No constraints with respect to content (it can be empty)



Why Use Messages?

- Many kinds of applications fit into the model of processing a sequential flow of information (e.g., Unix filters)
- Communicating processes sometimes need to be separate:
 - They do not trust each other (e.g., OS vs. user)
 - The programs were written at different times by different programmers who knew nothing about each other
 - They run on different processors (locally or on a network)
 - Less error-prone without shared memory
- Apart from sharing memory semaphores and monitors have other issues:
 - Semaphores are too low level (system calls very specific to OS)
 - Monitors need same structures/programming language for all processes



Mutual Exclusion with Messages (Example)

- Example: ME of **L processes** $P(\text{int } I)$, sharing a resource (CS)

parent process

```
create_mailbox(mbox);  
send(mbox, NULL);  
parallel {  
    P(1);  
    P(2);  
    ...  
    P(L);  
}
```

process $P(\text{int } I)$

```
while(true) {  
    receive(mbox, msg);  
    CS  
    send(mbox, NULL);  
}
```

- An empty message (NULL) in the mailbox is the token required to access the CS



Producer/Consumer Algorithm with Messages

- Not only do we need to ensure ME but also to keep a count:
 - Remember: Producer/Consumer buffer has finite size N
- Two mailboxes are needed
 - One mailbox for consumers: data produced and consumed
 - One mailbox for producers: accountancy purposes
- Producer and Consumer synchronise using these mailboxes as follows:
 - Producers generate data and send it to the consumer's mailbox when the producer's mailbox indicates there is at least one free slot in the buffer
 - Consumers take messages from the consumer's mailbox when available, and then send **empty** messages to the producer's mailbox (to signal new free slots)



Producer/Consumer Algorithm with Messages (II)

Parent Process

```
create_mailbox(cons_mbox);
create_mailbox(prod_mbox);
for(l=0;l<N;l++)
    send(prod_mbox,NULL);
parallel {
    producers();
    consumers();
}
```

producer()

```
while(true) {
    msg=produce_message();
    receive(prod_mbox,token);
    send(cons_mbox,msg);
}
```

consumer()

```
while(true) {
    receive(cons_mbox,msg);
    consume_message(msg);
    send(prod_mbox,NULL);
}
```



Conclusion

- Explored the resource Producer Consumer synchronisation problem and approaches to solve it
- Semaphores have drawbacks that can lead to deadlocks, if the order of Ps is not correctly set.
- Monitors are higher-level synchronisation primitives that avoid the “catches” that can arise when protecting critical sections (CS) with semaphores
- Messages are a mechanism for inter-process communication (IPC), that can also be relied upon for synchronisation without having to share memory

