

Deadlock and Starvation (I)



**School of Computer Science,
UCD**

**Scoil na Ríomheolaíochta,
UCD**

Announcements

- ***Extra lab session:***
 - Friday 2nd November: 10 a.m. – 11 a.m.
 - H1.51 Science Hub



Last week...

Conditions for True Solution to CS Problem (Dijkstra)

1. **Mutual exclusion**

- One process at most inside the CS at any time

2. **Progress**

- A process in execution out of a CS cannot prevent other processes from entering it
- If several processes are attempting to enter a CS simultaneously the decision on which one goes in cannot be indefinitely postponed
- A process may not remain in its CS indefinitely (neither terminate inside it)

3. **Bounded waiting** (no starvation)

- A process attempting to enter its CS will eventually do so

Notes:

- These are **necessary and sufficient** conditions, provided that basic operations are atomic
- No assumptions are made about: number of processes, relative speed of processes, or underlying hardware



Last Week...

Mechanisms for Implementing ME in a CS

Three basic mechanisms:

1. Semaphores

- Simple, but hard to program with (low level)

2. Monitors

- More abstract, higher level mechanism (language support)

3. Messages


- Very flexible and simple method of interprocess communication (IPC) & synchronisation



Last Week...

Semaphores

- A CS may be protected by a semaphore → we may implement ME by means of semaphores; example:
 - Initialise $S = 1$
 - To enter the CS, execute **P** on its semaphore
 - When leaving the CS, execute **V** on its semaphore
 - Therefore, two or more processes sharing a CS and this semaphore achieve ME by executing



Last Week...

Monitors

```
monitor pr_co {  
    int count;  
    condition full_s, empty_s;  
  
    void put(msg) {  
        if(count==N) wait(empty_s);  
        put_message(msg);  
        count++;  
        if(count==1) signal(full_s);  
    }  
  
    msg get() {  
        if(count==0) wait(full_s);  
        msg=get_message();  
        count--;  
        if(count==N-1) signal(empty_s);  
    }  
}
```

Producer

```
while(true) {  
    msg=produce_message();  
    pr_co.put(msg);  
}
```

Consumer

```
while(true) {  
    msg=pr_co.get();  
    consume_message(msg);  
}
```

- Simplest possible consumer and producer code
- Monitor takes care of any issues, not producer or consumer



Outline

- Deadlock and Starvation
- Four condition for deadlocks
- Resource-allocation Graphs (RAGs)

Take home message:

Deadlock (2+ processes waiting indefinitely for an event that can be caused only by a waiting process) can occur only if four necessary conditions hold simultaneously: mutual exclusion, hold and wait, no preemption and circular wait.




Deadlock and Starvation

- Mutual exclusion (ME) mechanisms for synchronisation guarantee that processes do not clash when using ***shared resources***
 - However, we saw that severe problems may still arise
- Definitions:
 - ***Deadlock***: a set of processes is in a deadlock state when every process in the set is blocked forever, waiting for the availability of resources held by other processes in the set
 - ***Starvation (indefinite postponement)*** occurs when a process waits for resources that periodically become available, but are never allocated to that process due to some ***scheduling policy***



Starvation vs. Deadlock

- Starvation: process/thread waits indefinitely
 - Example, low-priority process/thread waiting for resources constantly in use by high-priority process/threads
- Deadlock: circular waiting for resources
 - Process A owns Res 1 and is waiting for Res 2
 - Process B owns Res 2 and is waiting for Res 1



- Deadlock \Rightarrow Starvation but not vice versa
 - Starvation can end (but doesn't have to)
 - Deadlock can't end without external intervention

Resources

Processes request and are assigned resources; two types:

Reusable resources

- Those used by only one process at a time, and not depleted by that use
- After their use, they are released for reuse by other processes
- *Examples:* processors, main and secondary memory, devices, data structures such as databases and semaphores

Consumable resources

- Those created (produced) by one process and destroyed (consumed) by another
- Unbounded number of instances
- No need to release them
- *Examples:* interrupts, signals, messages



Deadlock and starvation are possible with both types of resources ¹⁰

Example 1 (Reusable Resource)

- Two processes P1 and P2 requesting memory allocation
- Assume that the available memory space is 200MB, and that the processes execute the following pseudocode:

P_1

```
...  
Request 80MB  
...  
Request 60MB
```


P_2

```
...  
Request 70MB  
...  
Request 80MB
```

- The processes are correctly designed, since neither requests more than the total space in the system
- However: deadlock occurs if both processes have gone through their first requests and progress to their second requests (without having released the memory initially allocated)

Example 2 (Consumable Resource)


- Consumable resource: message (we may send and receive as many as we want: they are produced and consumed)
- Assume that the receive operation is blocking (i.e. not asynchronous), and that two processes execute the pseudocode below:



- Deadlock occurs, because each process requests a resource held by the other process, while at the same time holding a resource requested by the other process

Dining Philosophers

- Classic example proposed by Dijkstra (1971) to illustrate deadlock & starvation
- 5 philosophers living together
- Their life consists of two states: thinking or eating
- They have a position assigned at a round table on which there are 5 plates of noodles and 5 chopsticks
- A philosopher wishing to eat takes **both chopsticks** on the sides of their plate, and eats noodles
- No two philosophers can share a chopstick simultaneously



Philosopher = Process

Chopstick = Shared resource

Dining Philosophers: Getting Around Deadlock

- *Idea:* Make the philosophers release their left chopstick if after having grabbed it they detect that the right chopstick is in use
 - After waiting for some ***fixed time*** the philosopher would try again to grab both chopsticks
- *Issue:* No deadlock, but ***starvation*** is possible
 - If all philosophers start the algorithm simultaneously, no philosopher ever grabs both chopsticks
 - The philosophers are caught in an endless cycle
- *Solution:* Make waiting time ***random***



Dining Philosophers: Other Possible Solutions


Many ad-hoc solutions to deadlock and starvation are possible:

1. *Limit the number of philosophers* at the table
 - Allow at most four philosophers at a time at the table
 - Pass a token around the table so that only the philosopher holding the token can eat (i.e. enforce a sequence)
2. *Asymmetric solution:*
 - Even philosophers try left chopstick first
 - Odd philosophers try right chopstick first
3. Use five *counting semaphores*, each counting the number of available chopsticks per philosopher (careful initialisation needed)
4. etc
 - Formalisation is needed to handle the complexity of the problem




Resource-allocation Graph

- Deadlocks can be described more formally in terms of a directed graph called Resource-Allocation Graph (RAG)
- A RAG completely describes state of a system in terms of
 1. What resources are allocated to what processes
 2. What processes are waiting for what resources
- A usual directed graph is just a set of vertices (i.e., nodes) plus set of directed edges (i.e., arrows) connecting the vertices; example:




Resource-allocation Graph

- A RAG is a **special directed graph**, since both vertices and directed edges are partitioned into **two sets**
- Vertices:
 1. $P = \{P_1, \dots, P_n\}$, all active processes in the system:
 2. $R = \{R_1, \dots, R_m\}$, all resource types in the system:
 - also: number of • inside \square is the number of instances of the resource (two displays, three printers. . .)
- Directed edges:
 1. $P_i \rightarrow R_j$: request edge
 2. $R_j \rightarrow P_i$: assignment edge



Request



Assignment

Example

- Two processes P_1 and P_2 , two shared resources R_1 and R_2 (each resource has one instance only)
- Assume the following scenario:

P_1

```
0: ...  
while(true) {  
  1: request( $R_1$ );  
  2: ...  
  3: request( $R_2$ );  
  4: ...  
    release( $R_2$ );  
  5: ...  
    release( $R_1$ );  
  0: ...  
}
```

P_2

```
0: ...  
while(true) {  
  1: request( $R_2$ );  
  2: ...  
  3: request( $R_1$ );  
  4: ...  
    release( $R_1$ );  
  5: ...  
    release( $R_2$ );  
  0: ...  
}
```



Example: System States

State	Situation of P_1
0	Holds no resources
1	Holds none, requests R_1
2	Holds R_1
3	Holds R_1 , requests R_2
4	Holds R_1 and R_2
5	Holds R_1 , R_2 released

State	Situation of P_2
0	Holds no resources
1	Holds none, requests R_2
2	Holds R_2
3	Holds R_2 , requests R_1
4	Holds R_2 and R_1
5	Holds R_2 , R_1 released


- Each pair of possible states corresponds to one RAG
- Note: some states such as (4,4) are impossible in this example



Example: RAG


(State P_1 , State P_2)

- (1,0): P_1 asks for R_1 which is free




Example: RAG (2)

- (State P_1 , State P_2)
- $(4,1)$: P_2 blocked (waiting for P_1 to release R_2)




Example: RAG (3)

- (State P_1 , State P_2)
- (2,3): Technically not a deadlock, but next instruction in P_1 (in the example) leads inevitably to a deadlock



Example: RAG (4)

- (State P_1 , State P_2)
- $(3,3)$: deadlock



Four Necessary Conditions for Deadlock

1. **Mutual exclusion** (limited access)
 - At least one resource may be acquired exclusively by only one process at a time
2. **Hold-and-wait** (wait-for)
 - Processes may ask for resources while holding other resources
3. **No preemption**
 - Once allocated, resources are released only voluntarily by the process holding the resource, after process is finished with it
4. **Circular chain of request** (circular-wait)
 - Two or more processes locked in a circular chain in which each process is waiting for one or more resources that the next process in the chain is holding
 - Equivalent to a cycle in the RAG



Note: ***These are not sufficient conditions***

Cycles in RAG and Deadlock

Possibilities:


1. There are no cycles in the RAG: then there is no deadlock
 - A cycle is **a necessary condition** for a deadlock
 2. There is a cycle in the RAG and
 - A. There is **a single instance** of each resource in the cycle: then there is a deadlock
 - B. There are **several instances** of at least one resource in the cycle: then there may or there may not be a deadlock
- In general, a cycle is **not a sufficient condition** for a deadlock
 - An exception is case A) above



Examples: RAG Cycles and Multiples Resource Instances


- **Note:** There are two instances of R_2
- Cycle but no deadlock

$$P_1 \rightarrow R_1 \rightarrow P_2 \rightarrow R_2 \rightarrow P_1$$



Examples: RAG Cycles and Multiple Resource Instances (2)

- **Note:** There are two instances of R_2
- Cycle broken if P_3 releases R_2




Examples: RAG Cycles and multiple Resource Instances (3)

Cycle(s) & deadlock:

$P_1 \rightarrow R_1 \rightarrow P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_1$

$P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_2$



Conclusion

- Starvation vs. Deadlock
 - Starvation: thread/process waits indefinitely
 - Deadlock: circular waiting for resources
- Four conditions for deadlocks
 - **Mutual exclusion**
 - Only one thread/process at a time can use a resource
 - **Hold and wait**
 - Thread/process holding at least one resource is waiting to acquire additional resources held by other threads/processes
 - **No preemption**
 - Resources are released only voluntarily by the threads/processes
 - **Circular wait**
 - \exists set $\{T_1, \dots, T_n\}$ of threads/processes with a cyclic waiting pattern

