

Anthony Ventresque

anthony.ventresque@ucd.ie

Operating Systems

COMP30640

Operating System Components and Structures



School of Computer Science,
UCD

Scoil na Ríomheolaíochta,
UCD

Announcements

- Everybody should have access to a command line interface now
 - at least on a remote server (ssh/PuTTY)
 - and/or on your own machine (having a network connection will be important too)
 - if you don't, make sure to catch up at this Friday's lab session
- First quiz (**1%** of the final mark) will be open from this Friday until before the lecture next Tuesday
 - no time limit
 - no limit in the number of attempts
- **No lecture** next Tuesday (26th September)
- **Extra session** on Friday 29th September (time TBC; Bash commands and scripts)



Last Week...

- An OS is a ***program*** that acts as an ***intermediary*** between a user of a computer and the computer hardware
- Goals: Execute user programs, make the computing system easy to use, utilise hardware efficiently
- OS is:
 - ***Resource allocator***: decides between conflicting requests for efficient and fair resource use
 - ***Control program***: controls execution of programs to prevent errors and improper use of computer

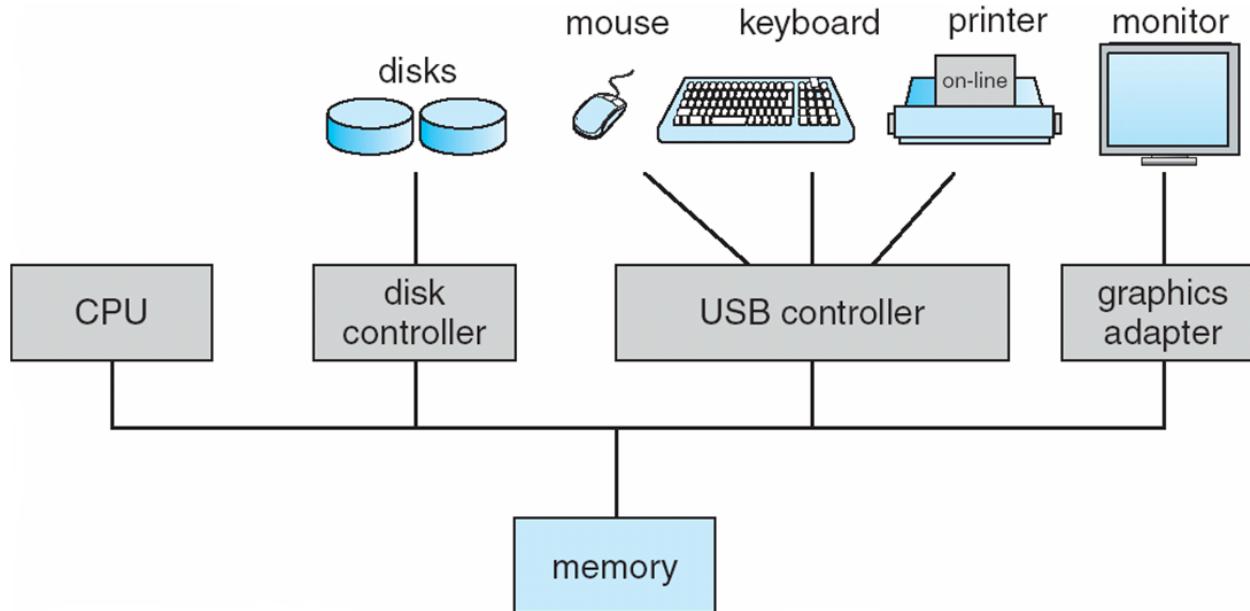


Last Week...

- ***Kernel***: the one program running at all times on the computer
- ***Bootstrap program***: loaded at power-up or reboot



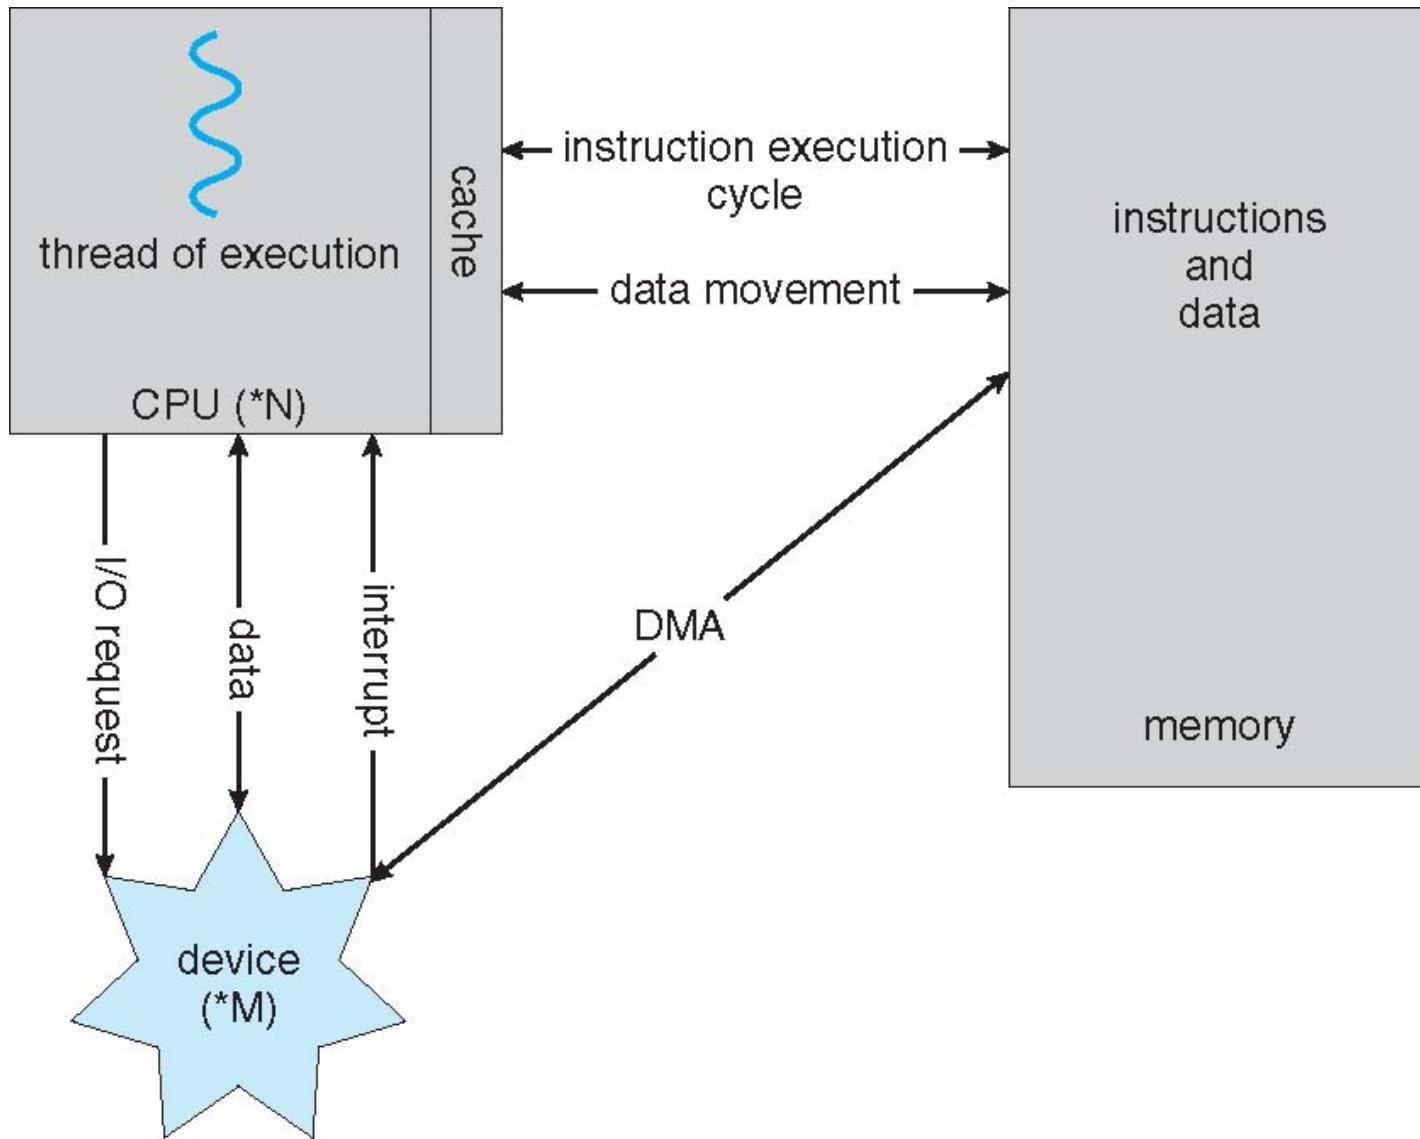
Last Week...



- Device controllers inform CPU that it is finished by causing an ***interrupt***
 - ***Trap*** is a software generated interrupt caused by error or user request
- ***System call***: request to the operating system to allow user to wait for I/O completion



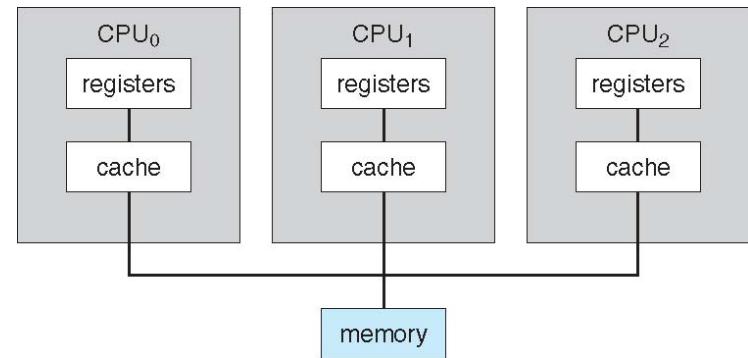
Last Week...



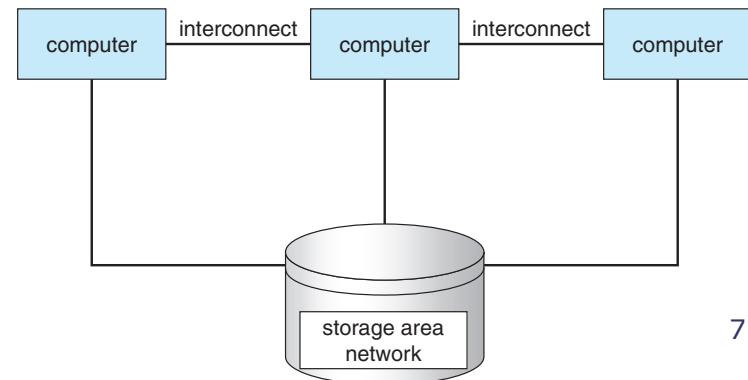
Last Week...

- **Multiprocessor** Systems:
Increased throughput,
economy of scale, increased
reliability

- Can be asymmetric (1 processor=1 task) or symmetric (each processor performs all tasks)

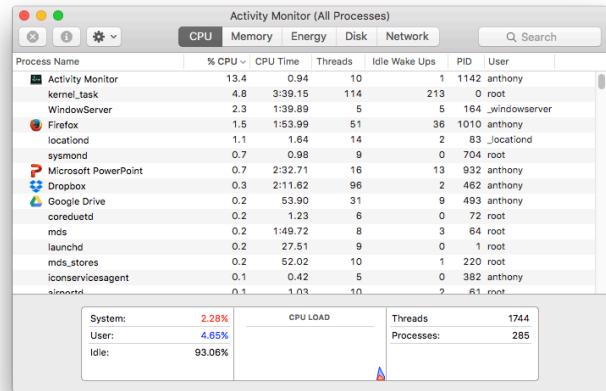


- Clustered systems –
Linked multiprocessor
systems



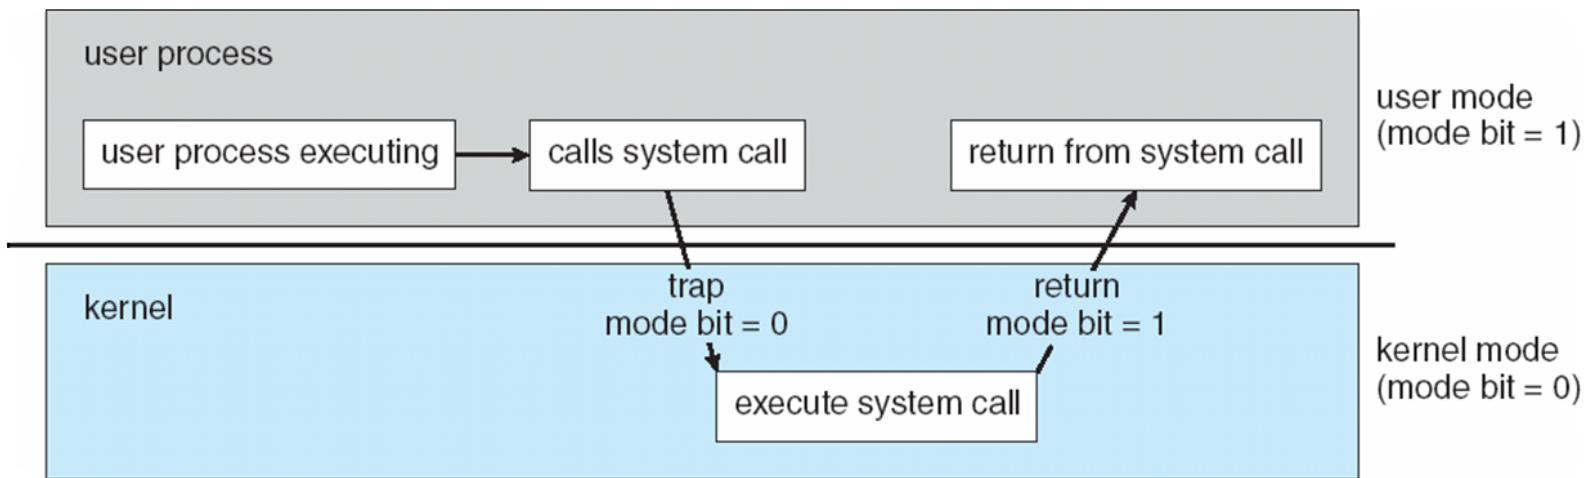
Last Week...

- **Multiprogramming**
(multitasking) – Provides efficiency via job scheduling
 - When OS has to wait (ex: for I/O), switches to another job
- **Timesharing** – CPU switches jobs so frequently that each user can interact with each job while it is running (interactive computing)



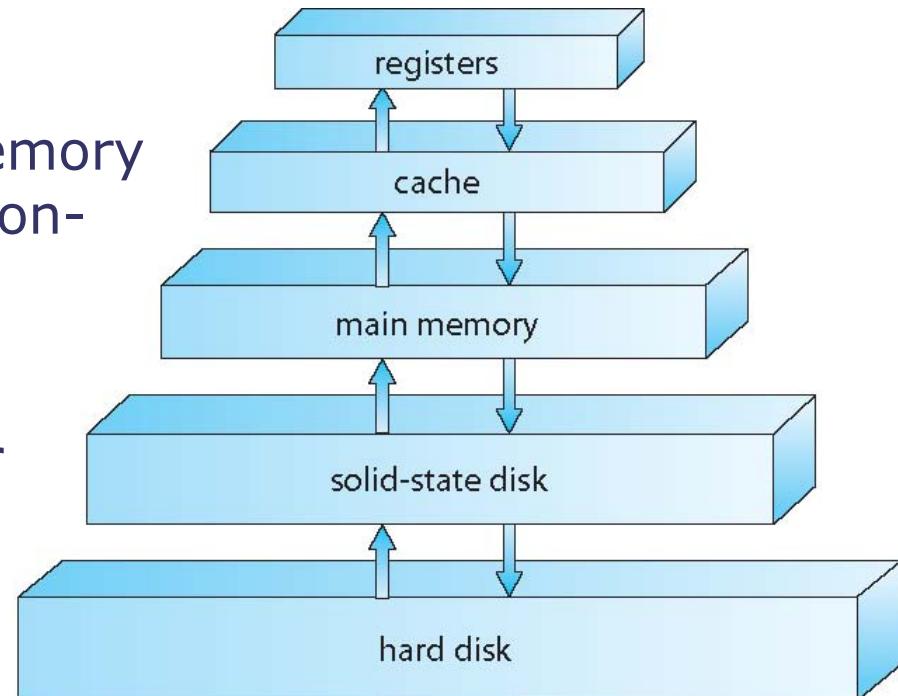
Last Week...

- **Dual-mode operation** allows OS to protect itself and other system components – User mode and kernel mode
 - Some instructions are only executable in kernel mode, these are privileged
 - Nowadays most CPUs can support multi-mode operations (e.g., VM managers)



Last Week...

- Storage structure:
 - Main memory – random access, volatile
 - Secondary storage – extension of main memory That provides large non-volatile storage
- **Caching** – copying information into faster storage system



...



Outline

- To describe the services an OS provides to users, processes and other systems
- To discuss the various ways of structuring an OS
- Understanding the different principles of OS organisation including abstractions and layering

Take home message:

An OS provides a number of components to users, processes, and other systems

Operating systems can be structured using different principles of organisation (i.e., Simple, Monolithic, Layered, Microkernel, Modular)



OS Structure

- How to organise an OS?
 - What are the essential components involved?
 - Where do they exist?
 - How do they cooperate?
- The task is not trivial if all system goals must be met:
 - Efficiency (high throughput)
 - Interactivity
 - Robustness (fault tolerance & reliability), security
 - Scalability, extensibility, portability



OS Services

- Operating systems provide an environment for execution of programs and services to programs and users
- One set of OS services provides functions that are helpful to the user:
 - **User interface** - Almost all operating systems have a user interface (UI).
 - Varies between **Command-Line** (CLI), **Graphics User Interface** (GUI), **Batch**
 - **Program execution** - The system must be able to load a program into memory and to run that program, end execution, either normally or abnormally (indicating error)
 - **I/O operations** - A running program may require I/O, which may involve a file or an I/O device



OS Services

- One set of operating-system services provides functions that are helpful to the user (Cont.):
 - **File-system manipulation** - Programs need to read and write files and directories, create and delete them, search them, list file Information, permission management.
 - **Communications** – Processes may exchange information, on the same computer or between computers over a network
 - Communications may be via shared memory or through message passing (packets moved by the OS)
 - **Error detection** – OS needs to be constantly aware of possible errors
 - May occur in the CPU and memory hardware, in I/O devices, in user program
 - For each type of error, OS should take the appropriate action to ensure correct and consistent computing
 - Debugging facilities can greatly enhance the user's and programmer's abilities to efficiently use the system

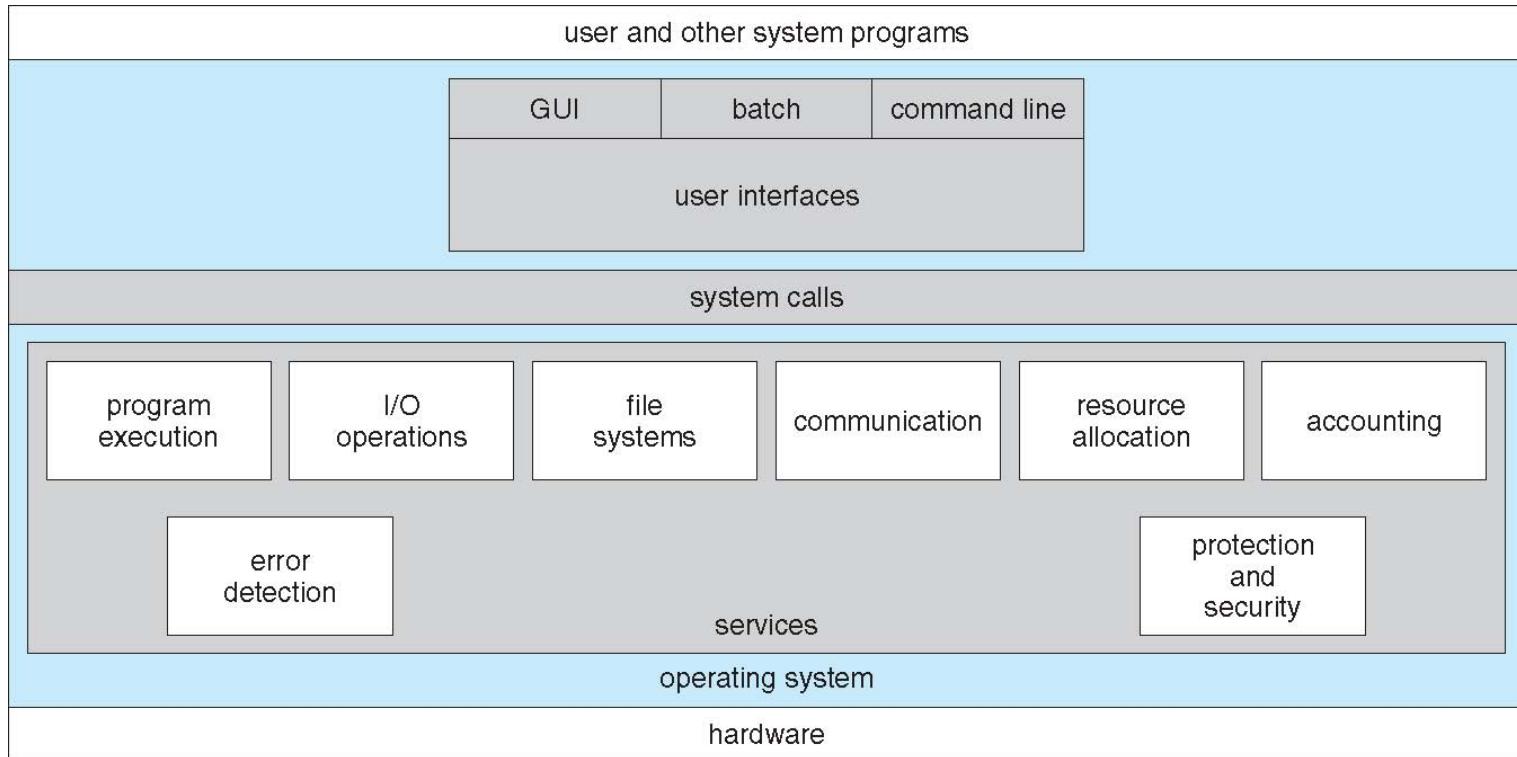


OS Services

- Another set of OS functions exists for ensuring the efficient operation of the system itself via resource sharing
 - **Resource allocation** - When multiple users or multiple jobs running concurrently, resources must be allocated to each of them
 - Many types of resources - CPU cycles, main memory, file storage, I/O devices.
 - **Accounting** - To keep track of which users use how much and what kinds of computer resources
 - **Protection and security** - The owners of information stored in a multiuser or networked computer system may want to control use of that information, concurrent processes should not interfere with each other
 - Protection involves ensuring that all access to system resources is controlled
 - Security of the system from outsiders requires user authentication, extends to defending external I/O devices from invalid access attempts



A View of OS Services



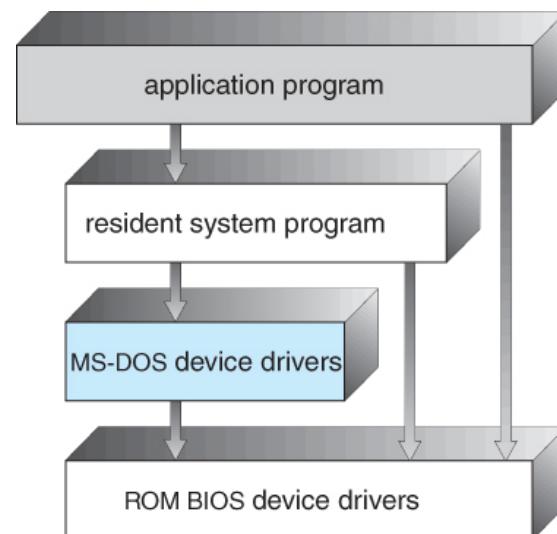
Different OS Organisational Principles

- General-purpose OS is very large program
- Various ways to structure OS
 - **Simple** Structure: Only one or two levels of code (e.g MS-DOS)
 - **Monolithic** Architectures: More complex (e.g. UNIX)
 - **Layered**: Lower levels independent of upper levels (an abstraction)
 - **Microkernel**: OS built from many user-level processes (Mach)
 - **Modular**: Core kernel with Dynamically loadable modules



Simple Structure: MS-DOS

- MS-DOS – written to provide the most functionality in the least space
 - Not divided into modules
 - Although MS-DOS has some structure, its interfaces and levels of functionality are not well separated



Monolithic Architecture

- Traditionally, systems were built around monolithic kernels
 - every OS component is contained in the kernel
 - any component can directly communicate with any other (by means of direct function calls)
 - due to this they tend to be highly efficient (high performance)

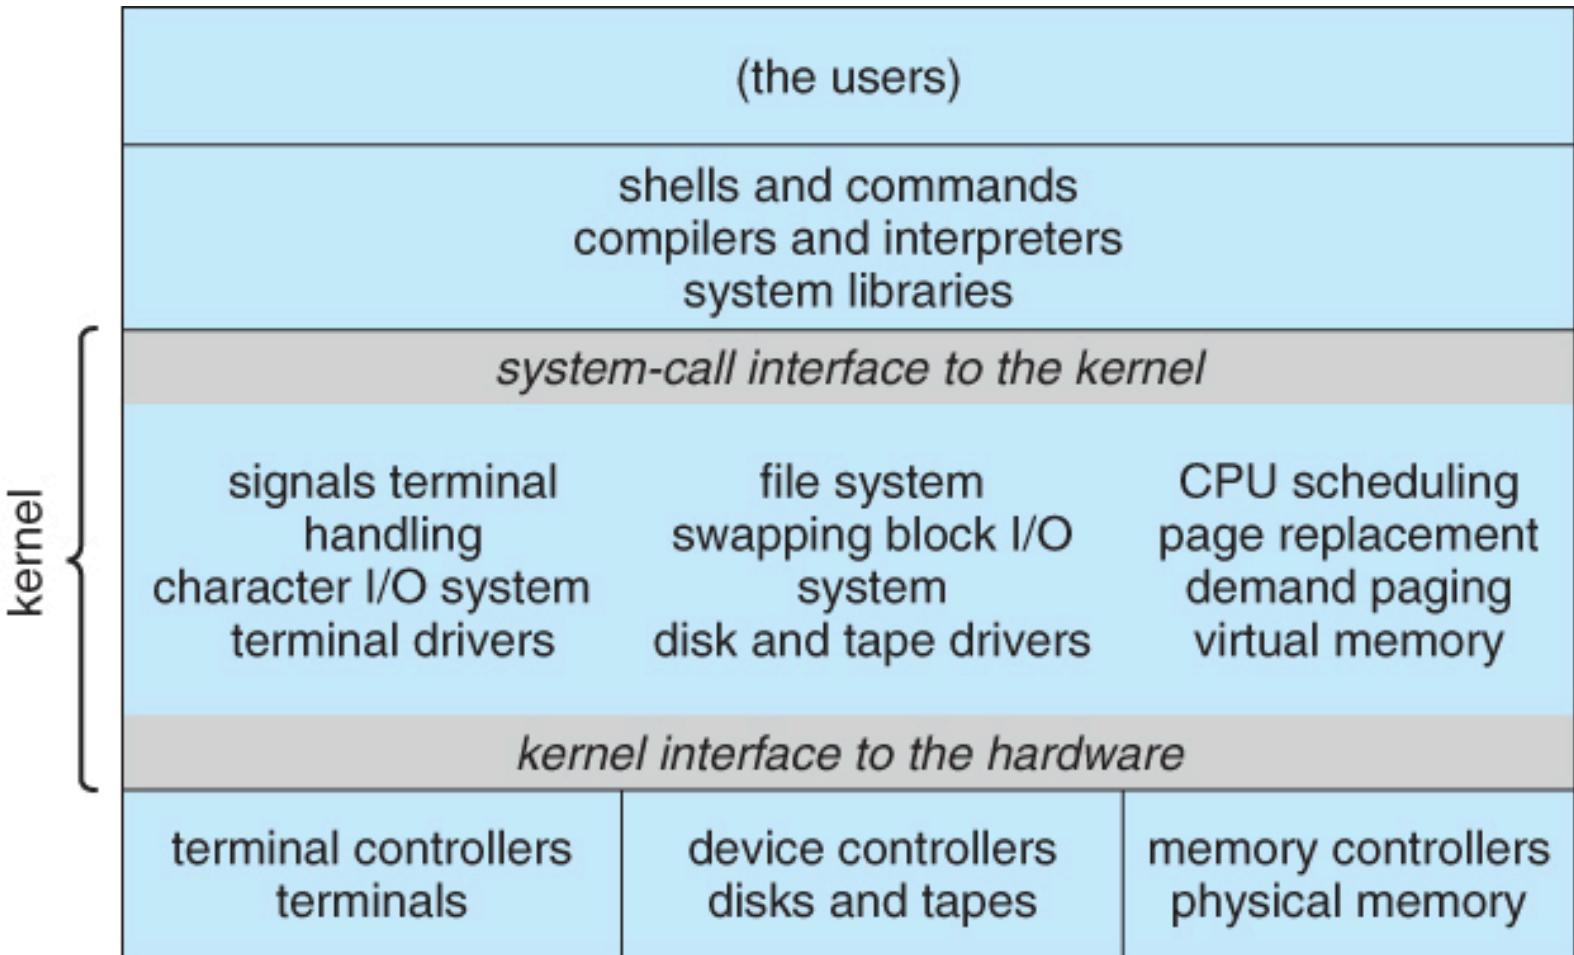


Monolithic Architecture

- Problems:
 - Unstructured → scalability makes them unwieldy
 - hard to understand, modify & maintain (e.g., debug)
 - Susceptible to damage from errant or malicious code, as all kernel code runs with unrestricted access to the system
- Examples: OS/360, VMS, Unix-like OSs (Linux, BSD)



Unix System Structure



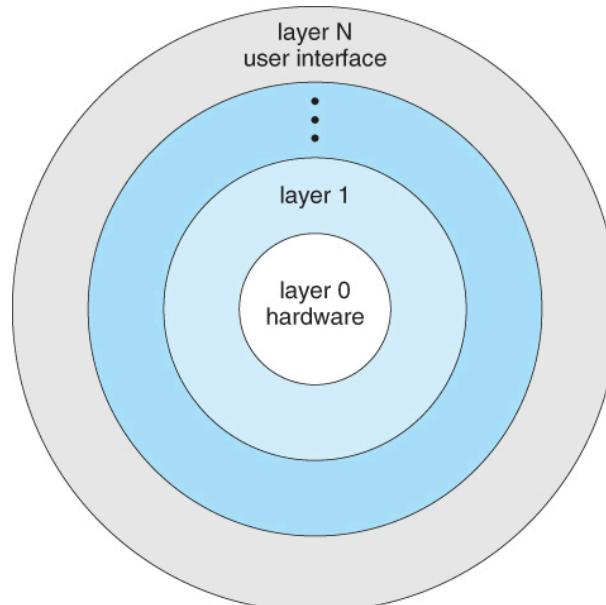
Layered Architecture

- Improves on monolithic kernel designs by adding structure
 - Components are grouped into layers that perform similar functions
 - Each built on top of lower layers
 - Bottom layer (layer 0) is hardware
 - Highest layer (layer N) is the user interface
- Advantages: modularity imposes structure and consistency (easier validation, debugging, modification, reuse)



Layered Architecture

- Each layer communicates only with layers immediately above and below it
 - each layer is a virtual machine to the layer above
 - a higher layer provides a higher-level virtual machine



First Layering-based OS

Layer	Function
5	The operator
4	User programs
3	Input/output management
2	Operator-process communication
1	Memory and drum management
0	Processor allocation and multiprogramming

Structure of the THE operating system.

- Layering was first used in Dijkstra's THE OS (1968)
layer 4 (user space) sees virtual I/O drivers
 - each layer “sees” a logical machine provided by lower layers:
 - layer 4 (user space) sees virtual I/O drivers
 - layer 3 sees virtual console
 - layer 2 sees virtual memory
 - layer 1 sees virtual processors 9/26 2nd Chapter: OS Structure
 - Based on a static set of cooperating processes
 - Each process can be tested and verified independently



Problems with Layering

- Appropriate definition of layers is difficult
 - A layer is implemented using only those operations provided by lower-level layers
 - A real system structure is often more complex than the strict hierarchy required by layering
- Performance issues
 - Processes' requests might pass through many layers before completion (layer crossing)
 - System throughput can be lower than in monolithic kernels

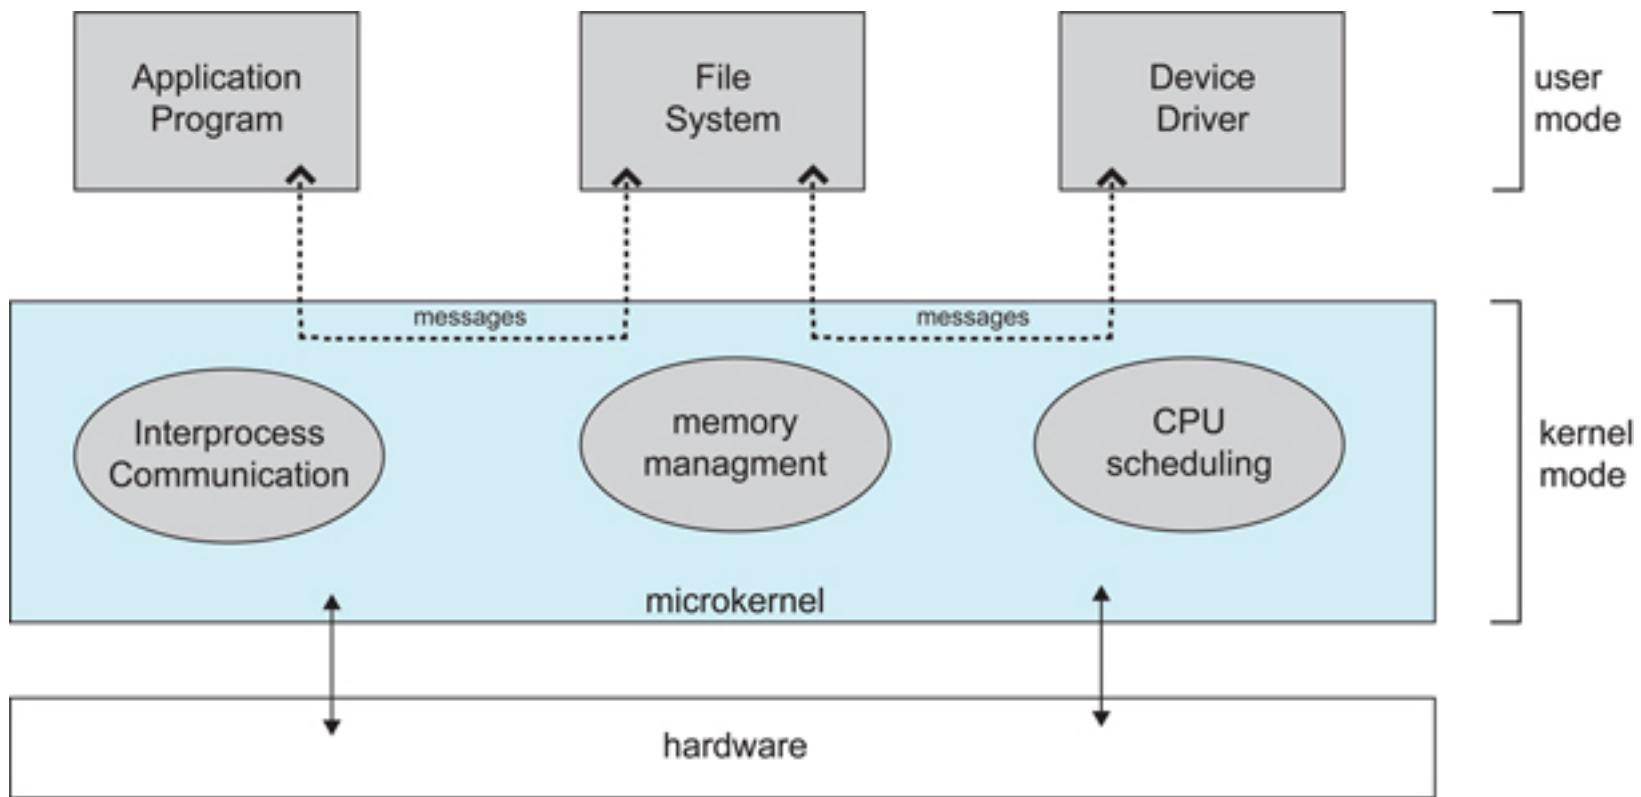


Microkernel System Architecture

- Minimise the services offered by the kernel, in an attempt to keep it small and scalable
 - Small core OS running at kernel level
 - OS Services built from many independent user-level processes
 - No consensus about minimal set of services inside micro-kernel
 - At least: minimal process and memory management capabilities, plus inter-process communications
 - Services such as networking and file system tend to run non-privileged at the user process level (i.e. out of the micro-kernel)
- Communication takes place between user modules using ***message passing***



Microkernel System Structure



Microkernel System Structure

- Benefits:
 - Easier to extend a microkernel
 - Easier to port the operating system to new architectures
 - More reliable (less code is running in kernel mode)
 - More secure
- Drawbacks:
 - Performance overhead of user space to kernel space communication
- Examples:
 - Hydra (CMU, 1970): first μ-kernel
 - Mach (CMU)
 - Mac OS X kernel (Darwin) partly based on Mach
 - Chorus/MiX (Unix-like distributed OS)
 - Windows NT (originally a layered microkernel architecture)
 - Symbian

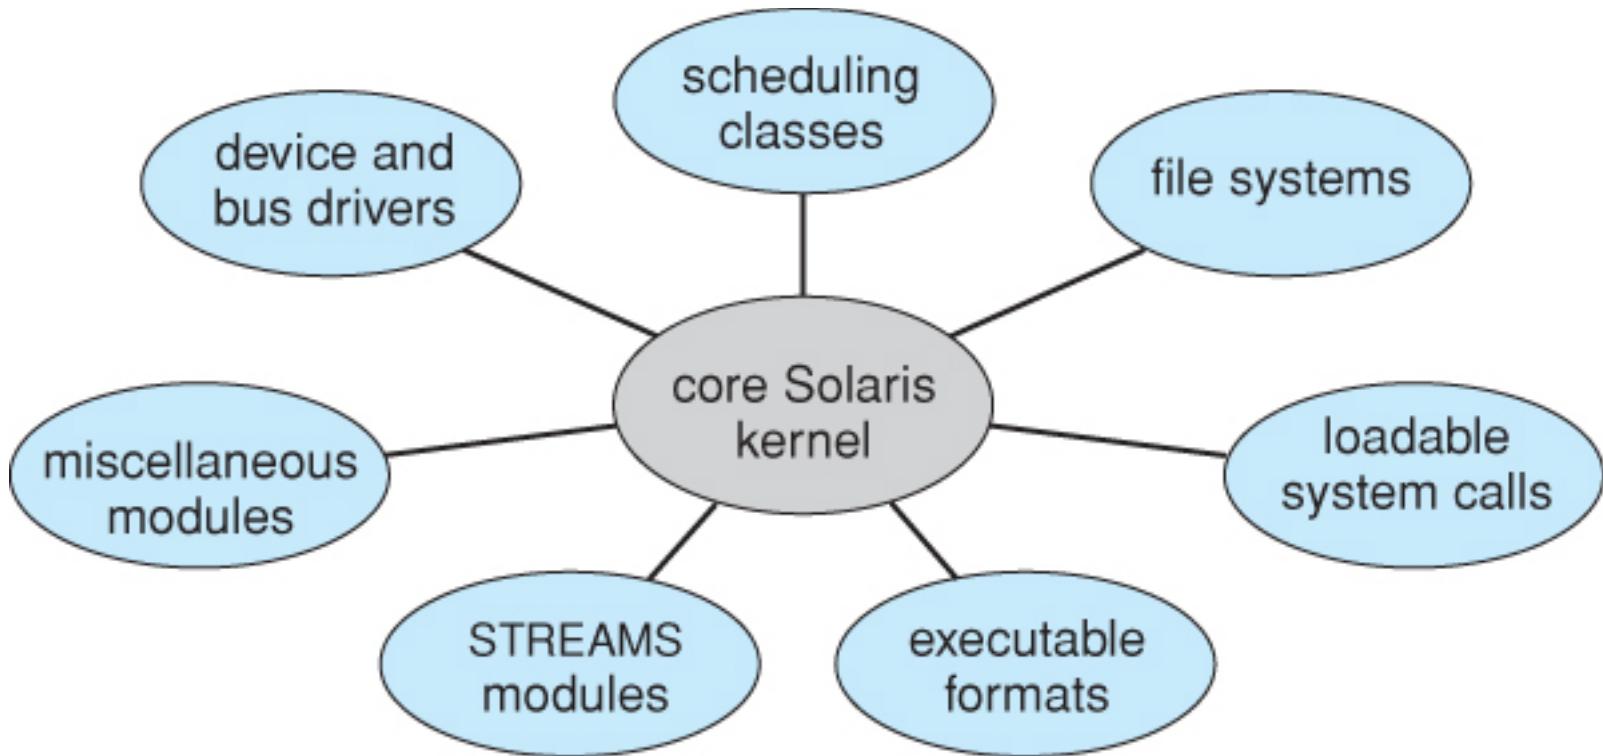


Modules-based Structure

- Many modern operating systems implement ***loadable kernel modules***
 - Uses object-oriented approach
 - Each core component is separate
 - Each talks to the others over known interfaces
 - Each is loadable as needed within the kernel
- Overall, similar to layers but with more flexibility
 - Linux, Solaris, etc.



Modules-based Structure



Conclusion

- An operating system provides a number of services/components to users, processes, and other systems (i.e. memory management, I/O, scheduling, etc.)
- Various ways to structure OS
 - **Simple** Structure: Only one or two levels of code (e.g., MS-DOS)
 - **Monolithic** Architectures: More complex (e.g., UNIX)
 - **Layered**: Lower levels independent of upper levels (an abstraction)
 - **Microkernel**: OS built from many user-level processes (Mach)
 - **Modular**: Core kernel with Dynamically loadable modules

