

Process Management III

Synchronisation



School of Computer Science,
UCD

Scoil na Ríomheolaíochta,
UCD

Outline

- Concurrent Execution
- Synchronising Concurrent Processes
- Critical Sections
- Mutual Exclusion
- Semaphores

Take home message:



Concurrent access to shared data may result in data inconsistency. Several mechanisms exist to ensure the orderly execution of cooperating processes

Basic Example of Concurrent Execution

int A;

Process 1

```
A = 1;  
if(A == 1)  
    printf("Process 1 wins");
```

Process 2

```
A = 2;  
if(A == 2)  
    printf("Process 2 wins");
```

- Variable A is shared by the two processes
- When the two processes are run concurrently on one processor, which one “wins”?
 - the outcome of the concurrent execution depends on which assignment takes place first (***race condition***)
- Note: the same situation would apply if we would consider two threads of a single process using a shared variable (instead of two processes)



Atomic Operations

- Can we possibly get some value different from either 1 or 2 in A due to concurrent execution
 - if the two instructions “A=1” and “A=2” happened to be executed concurrently?
- NO: ***References and assignments (i.e., read & write operations) are atomic in all CPUs***
 - An atomic operation cannot be interrupted, in order to avoid illogical outcomes
- This basic atomicity is provided by the hardware
 - However higher-level constructs are not atomic in general
 - Higher-level construct: any sequence of two or more CPU instructions



Higher-level Example of Concurrent Execution

int B;

Process 1

```
B = 0;  
while(B < 10) B++;  
printf( "Process 1 finished" );
```

Process 2

```
B = 0;  
while(B > -10) B--;  
printf( "Process 2 finished" );
```

- Variable B is shared
 - Post increments/decrements ("B++", "B--") are atomic
- Are there race conditions in this example?
 - Will the processes finish?
- Issue: "**while**" sections above do not behave atomically
 - ***Process synchronisation is all about getting high-level constructs to behave atomically***



Motivation: “Too Much Milk”

- Two flatmates sharing a fridge
 - Problem: they would like to have at most one bottle of milk in the fridge at any given time

Time	Person A	Person B
3:00	Look in Fridge. Out of milk	
3:05	Leave for store	
3:10	Arrive at store	Look in Fridge. Out of milk
3:15	Buy milk	Leave for store
3:20	Arrive home, put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive home, put milk away



- A and B must synchronise (cooperate) to achieve their goal
- Main issue: behaviour of A and B is not atomic

Important Definitions

- ***Race Condition*** (RC): output of a concurrent program depends on the order of operations between process/threads
- ***Synchronization***: using atomic operations to ensure cooperation between process/threads
- ***Mutual Exclusion*** (ME): ensuring that only one process/thread at a time holds or modifies a shared resource
- ***Critical Section*** (CS): piece of code that only one process/thread can execute at once
 - Critical section is the result of mutual exclusion
 - Critical section and mutual exclusion are two ways of describing the same thing



Locking and mutual Exclusion

- Achieving ME in a CS always involve some sort of ***locking mechanism***
 - ***Locking***: preventing someone else from doing something with the resource shared with them (“monopolise CS for a while”)
- Locking involves three rules:
 - Must lock before entering CS
 - Must unlock when leaving CS
 - Must wait if lock is locked when trying to enter CS



Too Much Milk: Solution 1

- Use a note to avoid buying too much milk:
 - Leave a note before buying (kind of “lock”)
 - Remove note after buying (kind of “unlock”)
 - Don’t buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```

- Result?
 - Still too much milk but only occasionally!
 - Thread can get context switched after checking milk and note but before buying milk!
 - Makes it really hard to debug



Too Much Milk: Solution 2

- As a workaround, let us change the meaning of the note:
- B buys if there is a note and A buys if there is no note

Process A	Process B
<pre>if(no_note) { if(no_milk) buy milk; put note; }</pre>	<pre>if(note) { if(no_milk) buy milk; remove note; }</pre>

- Does this really work?
 - This attempt creates true ME in the CS
 - But assume that B goes on holidays: A will starve
 - Similar issue if B is very slow
 - The relative speed of the processes is an issue
 - Processes must take turns to be in the CS (i.e. the same flatmate cannot buy milk twice in a row)



Too Much Milk: Solution 3

- In order to try to avoid the last issue, let us use **two notes** (**note A** and **note B**) and some basic courtesy protocol

```
Process A
-----
put note_A;
if(no_note_B) {
    if(no_milk)
        buy milk;
}
remove note_A;
```

```
Process B
-----
put note_B;
if(no_note_A) {
    if(no_milk)
        buy milk;
}
remove note_B;
```

- Each process can now examine each other's status, but not modify it. Does this solution work?
 - Better than before: there is ME *and* none of the processes starves if the other one goes on holidays
 - But relative speed still an issue: we need to decide who buys when both of them leave notes simultaneously. . .



Too Much Milk: Solution 4

- Here is a possible two-note solution:

Thread A

```
leave note A;
while (note B) {\\X
    do nothing;
}
if (noMilk) {
    buy milk;
}
remove note A;
```

Thread B

```
leave note B;
if (noNote A) {\\Y
    if (noMilk) {
        buy milk;
    }
}
remove note B;
```

- Does this work? Yes. Both can guarantee that:
 - It is safe to buy, or
 - Other will buy, ok to quit
- At X:
 - if no note B, safe for A to buy,
 - otherwise wait to find out what will happen
- At Y:
 - if no note A, safe for B to buy
 - Otherwise, A is either buying or waiting for B to quit



Solution 4 Discussion

- Protects a single “Critical-Section” piece of code for each thread:

```
if (no_milk) {  
    buy milk;  
}
```

- Solution 4 works. It was actually the first solution ever given to the mutual exclusion problem (Dekker), but it is **not satisfactory**:
 - Really complex – even for this simple an example
 - Hard to convince yourself that this really works
 - A’s code is different from B’s – what if lots of threads?
 - Code would have to be slightly different for each thread
 - While A is waiting, it is consuming CPU time
 - This is called “busy-waiting”
- There’s a better way
 - Have hardware provide better (higher-level) primitives than atomic load and store
 - Build even higher-level programming abstractions on this new hardware support
 - It is necessary to have standard synchronisation mechanisms in an OS, which can automatically fulfill some minimum requirements



Conditions for True Solution to CS Problem (Dijkstra)

1. **Mutual exclusion**

- One process at most inside the CS at any time

2. **Progress**

- A process in execution out of a CS cannot prevent other processes from entering it
- If several processes are attempting to enter a CS simultaneously the decision on which one goes in cannot be indefinitely postponed
- A process may not remain in its CS indefinitely (neither terminate inside it)

3. **Bounded waiting** (no starvation)

- A process attempting to enter its CS will eventually do so

Notes:

- These are **necessary and sufficient** conditions, provided that basic operations are atomic
- No assumptions are made about: number of processes, relative speed of processes, or underlying hardware



Desirable Properties of a ME mechanism

- Simple:
 - Systematic and easy to use (e.g., just bracket the CS): avoid unexpected “catches” in purported solutions
 - Easy to maintain
- Efficient:
 - Do not use up substantial amounts of resources when waiting (e.g., no busy-waiting)
 - The overhead due to entering and leaving the CS has to be small, compared to the work done inside it
 - Scalable: it should work when many agents share the CS



Mechanisms for Implementing ME in a CS

Three basic mechanisms:

1. Semaphores
 - Simple, but hard to program with (low level)
2. Monitors
 - More abstract, higher level mechanism (language support)
3. Messages
 - Very flexible and simple method of interprocess communication (IPC) & synchronisation




Semaphores

- Semaphores are a kind of generalized lock
 - Main synchronization primitive used in original UNIX
- Definition: a Semaphore has a non-negative integer value and supports the following two operations:
 - **P()**: an atomic operation that waits for semaphore to become positive, then decrements it by 1
 - Think of this as the wait() operation
 - **V()**: an atomic operation that increments the semaphore by 1, waking up a waiting P, if any
 - Think of this as the signal() operation
- Dijkstra, who first proposed semaphores (1965), was Dutch:
 - **P**roberen = to probe
 - **V**erhogen = to increment



Semaphores Like Integers Except

- Semaphores are like integers, except
 - No negative values
 - Only operations allowed are P and V – can't read or write value, except to set it initially
 - Operations must be atomic
 - Two P's together can't decrement value below zero
 - Similarly, thread going to sleep in P won't miss wakeup from V – even if they both happen at same time
- Semaphore from railway analogy
 - Here is a semaphore initialized to 2 for resource control:



Semaphores

- **Definition:** a semaphore is a protected integer variable S with an associated queue of waiting processes, upon which only two atomic operations $P()$ and $V()$ may be performed
- It is represented using an abstract data type which includes:
 - protected integer variable S (i.e., protected counter)
 - queue of waiting processes: processes in the queue are blocked
 - atomic operations:

$P(S)$

```
if( $S > 0$ )  
     $S--$ ;  
else  
    enqueue_calling_proc();
```

$V(S)$

```
if(queue_not_empty)  
    resume_enqueued_proc();  
else  
     $S++$ ;
```



Semaphores and Mutual Exclusion

- A CS may be protected by a semaphore → we may implement ME by means of semaphores; example:
 - Initialise $S = 1$
 - To enter the CS, execute **P** on its semaphore
 - When leaving the CS, execute **V** on its semaphore
 - Therefore, two or more processes sharing a CS and this semaphore achieve ME by executing



Semaphores and Mutual Exclusion

- Some remarks:
 - The initial value of S is typically non-negative ($S \geq 0$)
 - If we initialise $S > 1$ more than one process at a time can get into the CS (and thus there is no ME); this is used in a particular type of semaphores called counting semaphores
- Types of semaphores, depending on possible values of S :
 - $S \in \{0, 1\} \rightarrow$ **binary semaphore** (\approx mutex)
 - $S \in \{0, 1, 2, 3 \dots\} \rightarrow$ **general (or counting) semaphore**



Example (Milk and Semaphores)

- Use of semaphores to solve the “milk problem”

```
semaphore S(1,NULL);
```

processes A & B (in a closed loop)

```
P(S);
```

```
if(no_milk)
```

```
    buy milk;
```

```
V(S);
```

- Simple, symmetric and efficient solution
(compare with the four attempts at obtaining an *ad-hoc* solution)



Conclusion

- Important concept: Atomic Operations
 - An operation that runs to completion or not at all
 - These are the primitives on which to construct various synchronization primitives
- Showed how to protect a critical section with only atomic load and store
⇒ pretty complex!
- Showed several constructions of Locks
 - Must be very careful not to waste/tie up machine resources
 - Shouldn't disable interrupts for long
 - Shouldn't spin wait for long
 - Key idea: Separate lock variable, use hardware mechanisms to protect modifications of that variable
- Semaphores: Like integers with restricted interface
 - Two operations:
 - P(): Wait if zero; decrement when becomes non-zero
 - V(): Increment and wake a sleeping task (if exists)
 - Can initialize value to any non-negative value
 - Use separate semaphore for each constraint

