

# Process Scheduling (2)



**School of Computer Science,  
UCD**

**Scoil na Ríomheolaíochta,  
UCD**

# Outline

- Understand the operating, benefits, and drawbacks of common scheduling algorithms
  - First Come First Served (FCFS) (FIFO) Scheduling
  - Round-Robin Scheduling (RR)
  - Priority Scheduling
  - Shortest Job First (SJF)
  - Shortest Remaining Time First (SRTF)
  - Multi-Level Feedback Scheduling



# Scheduling Algorithms

- The scheduling algorithm implements the system policies in order to achieve the scheduling objectives
- Its decisions may take into account
  - Preemptibility
  - Accountancy
  - Priorities
  - . . .



# First-Come, First-Served

- First-Come, First-Served (FCFS)
  - Also “First In, First Out” (FIFO) or “Run until done”
    - In early systems, FCFS meant one program scheduled until done (inc. I/O)
    - Now, means keep CPU until thread blocks

- Example:

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- Suppose processes arrive in the order:  $P_1, P_2, P_3$   
The Gantt Chart for the schedule is:




- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
  - Average waiting time:  $(0 + 24 + 27)/3 = 17$
  - Average Completion time:  $(24 + 27 + 30)/3 = 27$

- *Convoy effect*: short process behind long process
- FCFS favours CPU-bound processes



# FCFS (cont'd)

- Example continued:
  - Suppose that processes arrive in order:  $P_2$ ,  $P_3$ ,  $P_1$   
Now, the Gantt chart for the schedule is:



- Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
  - Average waiting time:  $(6 + 0 + 3)/3 = 3$
  - Average Completion time:  $(3 + 6 + 30)/3 = 13$
- In second case:
  - average waiting time is much better (before it was 17)
  - Average completion time is better (before it was 27)
- FIFO Pros and Cons:
  - Simple (+)
  - Short jobs get stuck behind long ones (-)
    - Stuck behind full trolley of small items at supermarket





# Round Robin (RR)

- FCFS Scheme: Potentially bad for short jobs!
  - Depends on submit order
  - If you are first in line at supermarket with milk, you don't care who is behind you, on the other hand...
- Round Robin Scheme
  - Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds
  - After quantum expires, the process is preempted and added to the end of the ready queue
  - $n$  processes in ready queue and time quantum is  $q \Rightarrow$ 
    - Each process gets  $1/n$  of the CPU time
    - In chunks of at most  $q$  time units
    - No process waits more than  $(n-1)q$  time units (Guaranteed)
- Performance
  - $q$  large  $\Rightarrow$  FCFS
  - $q$  small  $\Rightarrow$  Interleaved (really small  $\Rightarrow$  hyperthreading?)
  - $q$  must be large with respect to context switch, otherwise overhead is too high (all overhead)




# Round Robin vs. FCFS



# Example of RR with Time Quantum = 20

<u>Process</u>	<u>Burst Time</u>
$P_1$	53
$P_2$	8
$P_3$	68
$P_4$	24

- The Gantt chart is:



0    20    28    48    68    88    108    112    125    145    153

- Waiting time for:
  - $P_1 = (68 - 20) + (112 - 88) = 72$
  - $P_2 = (20 - 0) = 20$
  - $P_3 = (28 - 0) + (88 - 48) + (125 - 108) = 85$
  - $P_4 = (48 - 0) + (108 - 68) = 88$
  - Average waiting time =  $(72 + 20 + 85 + 88) / 4 = 66\frac{1}{4}$
  - Average completion time =  $(125 + 28 + 153 + 112) / 4 = 104\frac{1}{2}$
- Thus, Round-Robin Pros and Cons:
  - Better for short jobs, Fair (+)
  - Context-switching time adds up for long jobs (-)





# Example with Different Time Quantum

**Best FCFS:**

$P_2$ [8]	$P_4$ [24]	$P_1$ [53]	$P_3$ [68]
0    8	32	85	153

	Quantum	$P_1$	$P_2$	$P_3$	$P_4$	Average
Wait Time	Best FCFS	32	0	85	8	$31\frac{1}{4}$
	Q = 1	84	22	85	57	62
	Q = 5	82	20	85	58	$61\frac{1}{4}$
	Q = 8	80	8	85	56	$57\frac{1}{4}$
	Q = 10	82	10	85	68	$61\frac{1}{4}$
	Q = 20	72	20	85	88	$66\frac{1}{4}$
	Worst FCFS	68	145	0	121	$83\frac{1}{2}$
Completion Time	Best FCFS	85	8	153	32	$69\frac{1}{2}$
	Q = 1	137	30	153	81	$100\frac{1}{2}$
	Q = 5	135	28	153	82	$99\frac{1}{2}$
	Q = 8	133	16	153	80	$95\frac{1}{2}$
	Q = 10	135	18	153	92	$99\frac{1}{2}$
	Q = 20	125	28	153	112	$104\frac{1}{2}$
	Worst FCFS	121	153	68	145	$121\frac{3}{4}$

# Round-Robin Discussion

- How do you choose time slice?
  - What if too big? Response time suffers
  - What if infinite ( $\infty$ )? Get back to FIFO
  - What if time slice too small? Throughput suffers!
- Actual choices of timeslice:
  - Initially, UNIX timeslice one second:
    - Worked ok when UNIX was used by one or two people.
  - In practice, need to balance short-job performance and long-job throughput:
    - Typical time slice today is between 10ms – 100ms
    - Typical context-switching overhead is 0.1ms – 1ms
    - Roughly 1% overhead due to context-switching



# Comparison between FCFS and Round Robin

- Assuming zero-cost context-switching time, is RR always better than FCFS?
- Simple example:
  - 10 jobs, each take 100s of CPU time; RR scheduler quantum of 1s;  
All jobs start at the same time
- Completion Times:

Job #	FIFO	RR
1	100	991
2	200	992
...	...	...
9	900	999
10	1000	1000

- Both RR and FCFS finish at the same time
- Average response time is much worse under RR!
  - Bad when all jobs same length



# Priority Scheduling

- RR makes the implicit assumption that all processes are equally important, which is not reasonable in general
  - e.g. sending email vs real-time video playback
- Priority scheduling
  - Priority number (integer) is associated with each process
  - Priorities *quantify* the relative importance of processes
  - CPU allocated to the process with the highest priority
- Starvation issue: low priority processes might never execute
  - Solution: increase priority as time progresses (aging)



# Priority Scheduling


- Priority scheduling can be either pre-emptive or non pre-emptive
- Adaptability:
  1. Static priorities:
    - Not responsive to environment changes, which could be exploited to increase throughput and reduce latency
    - easier to implement
  2. Dynamic priorities:
    - Responsive to change; e.g.: the OS may want to temporarily
    - Decrease the priority of a process holding a key resource needed by a higher-priority process
    - More complex to implement, overheads
- SJF can be seen as a type of priority scheduling, where priority is the predicted next CPU burst time (in this case a lower number means higher priority)



# Example: Priority Scheduling

	$P_1$	$P_2$	$P_3$	$P_4$	$P_5$
burst time	10	1	2	1	5
priority	3	1	4	5	2

- (we assume here that a lower number means higher priority)



- Average waiting time:  $(0 + 1 + 6 + 16 + 18)/5 = 8.2$   
(assuming simultaneous arrival)



# Shortest Job First (SJF) & Shortest Remaining Time First (SRTF)


- If we could predict CPU usage, could we always mirror best FCFS?
- **Shortest Job First (SJF):**
  - Run whatever job has the least amount of computation to do
  - Sometimes called “Shortest Time to Completion First” (STCF)
- **Shortest Remaining Time First (SRTF):**
  - Preemptive version of SJF: if job arrives and has a shorter time to completion than the remaining time on the current job, immediately preempt CPU
  - Sometimes called “Shortest Remaining Time to Completion First” (SRTCF)
- These can be applied either to a whole program or the current CPU burst of each program
  - Idea is to get short jobs out of the system
  - Big effect on short jobs, only small effect on long ones
  - Result is better average response time



# FIFO vs. SJF


Tasks

FIFO



Tasks

SJF



Time



# Discussion


- SJF/SRTF are the best you can do at minimizing average response time
  - Provably optimal (SJF among non-preemptive, SRTF among preemptive)
    - SJF gives the minimum average waiting time for a given set of processes (throughput is maximised)
- Comparison of SRTF with FCFS and RR
  - What if all jobs the same length?
    - SRTF becomes the same as FCFS (i.e. FCFS is best can do if all jobs the same length)
  - What if jobs have varying length?
    - SRTF (and RR): short jobs not stuck behind long ones



# Example: Non pre-emptive and Pre-emptive SJF


	$P_1$	$P_2$	$P_3$	$P_4$
arrival time	0	2	4	5
burst time	7	4	1	4

- Non pre-emptive SJF




– Average waiting time:  $(0 + 6 + 3 + 7)/4 = 4$

- Pre-emptive SJF (SRTF)



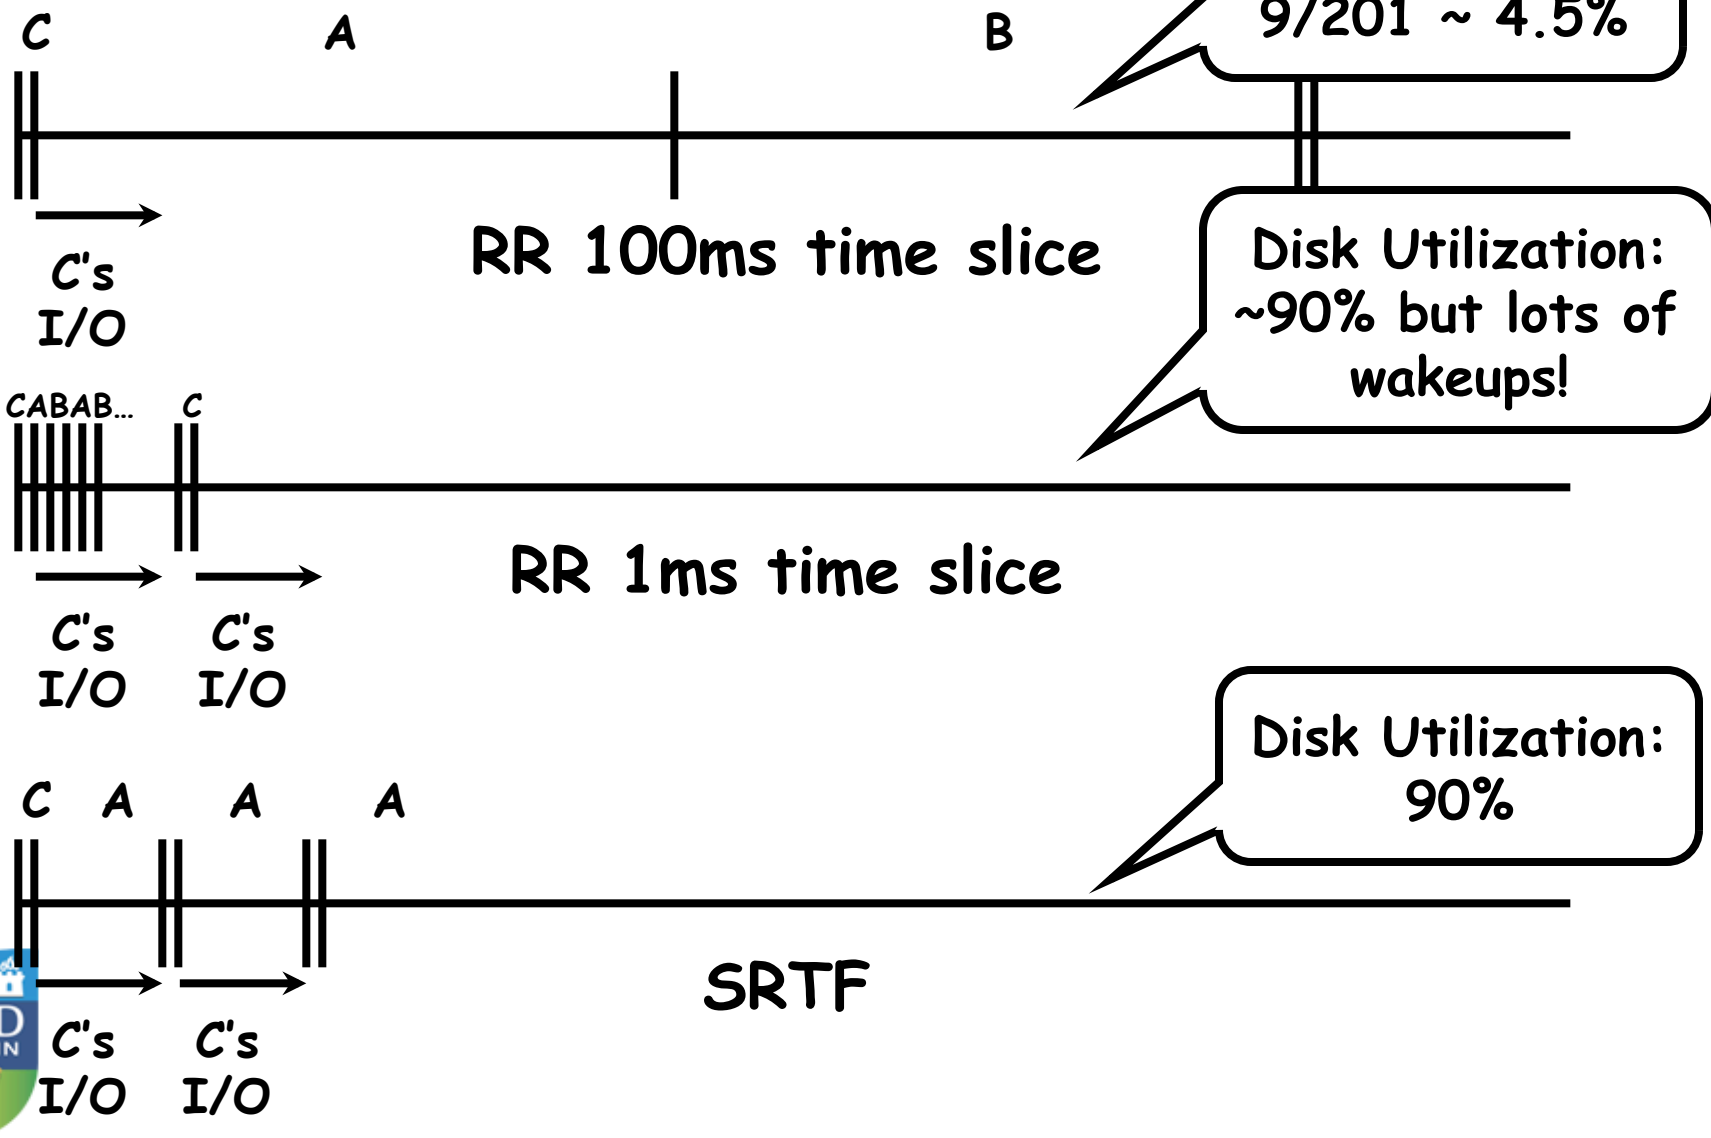
- Average waiting time:  $(9 + 1 + 0 + 2)/4 = 3$

# Example to Illustrate the Benefits of SRTF



- Three jobs:
  - A,B: both CPU bound, run for week
  - C: I/O bound, loop 1ms CPU, 9ms disk I/O
  - If only one at a time, C uses 90% of the disk, A or B could use 100% of the CPU
- With FIFO:
  - Once A or B get in, keep CPU for two weeks
- What about RR or SRTF?
  - Easier to see with a timeline

## SRTF Example (cont'd)



# SRTF Further Discussion

- Starvation
  - SRTF can lead to starvation if many small jobs!
  - Large jobs never get to run
- Somehow need to predict future
  - How can we do this?
  - Some systems ask the user
    - When you submit a job, have to say how long it will take
    - To stop cheating, system kills job if takes too long
  - But: Even non-malicious users have trouble predicting runtime of their jobs
- Bottom line, can't really know how long job will take
  - However, can use SRTF as a yardstick for measuring other policies
  - Optimal, so can't do any better
- SRTF Pros & Cons
  - Optimal (average response time) (+)
  - Hard to predict future (-)
  - Unfair (-)



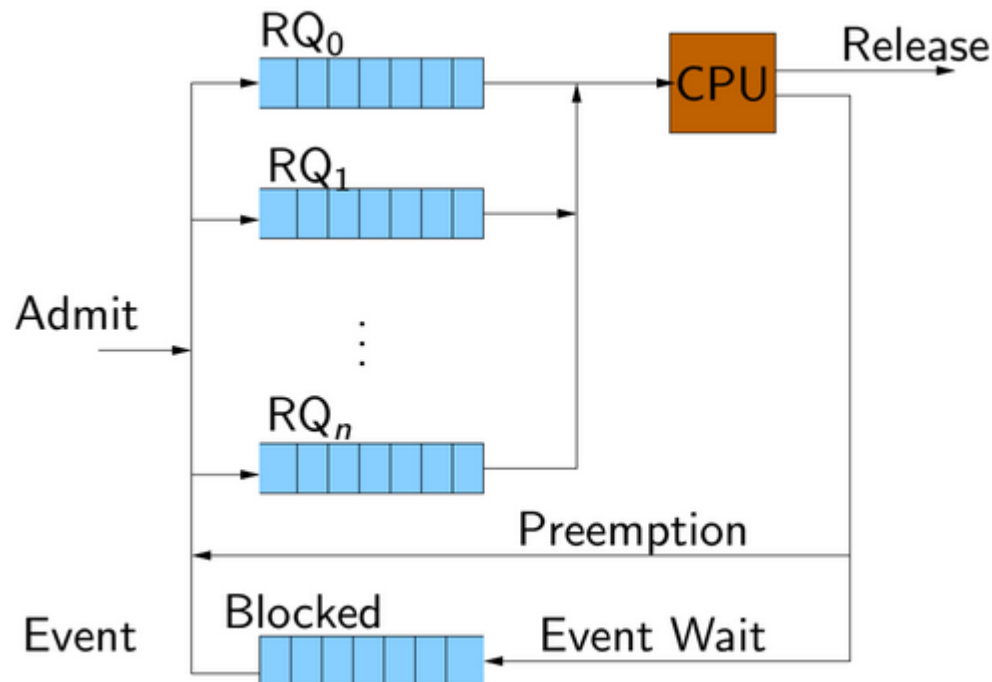
# Multilevel Queue Scheduling

- The ready queue is partitioned into separate queues with their own scheduling algorithm
  - e.g.: Foreground queue (interactive) with RR, background queue (batch) with FCFS
- In multilevel queues there must be scheduling ***among the queues too***; common strategies are:
  1. Fixed priority scheduling
    - e.g.: serve first all processes from foreground queue, then all from background queue
    - Absolute precedence of higher-priority queues: possibility of starvation
  2. Time slice
    - Each queue gets a certain amount of CPU time which it can schedule among its processes
    - e.g.: 80% to foreground in RR, 20% to background in FCFS



# Multilevel Queue

- RQ0 system processes (highest priority)
- RQ1 interactive processes
- . . .
- RQn batch processes (lowest priority)



# Multilevel Feedback Queue (MFQ)


- Goals:
  - Responsiveness
  - Low overhead
  - Starvation freedom
  - Some tasks are high/low priority
  - Fairness (among equal priority tasks)
- Not perfect at any of them!
  - Used in Linux (and probably Windows, MacOS)





# Multilevel Feedback Queue

- Allows a process to move between queues
- Example:



- Longer processes are penalised and eventually moved to lower priority queues
- It affords adaptability: when a new process acquires the processor, the system does not know its burst pattern(CPU-bound, I/O-bound)

# Traditional Unix Scheduling (SVR3, 4.3 BSD)

- Priority-based
- Multilevel queue feedback using RR within each queue
- 128 priorities in 32 queues (4 adjacent priorities); lower values, higher priority
- If a running process does not block or complete within one second, it is pre-empted
  - Kernel processes (priorities 0–49) cannot be pre-empted
  - Priorities are recomputed once per second, taking into account recent CPU usage
- A user can bias the priority of a process (towards less priority) using the nice command
  - Background processes (batch jobs) get higher nice values (i.e. less priority) automatically



# Scheduling in Real-time Systems

- ***Real-time system:*** meets the needs of processes that must produce correct output by a certain time (***deadlines*** or ***timing constraints***)
- We have seen that SJF is optimal considering the average waiting time, but it does not guarantee a fixed waiting time for any particular process
  - Therefore, it cannot be used in real-time scenarios
  - Consider CPU-bound processes verifying mission critical tasks
- If all processes can meet their deadlines regardless of their execution order, shortest deadlines first would be optimal



# Categories of Real-time Scheduling

## 1. Soft real-time scheduling

- Missing an occasional deadline is undesirable but tolerable
- e.g.: multimedia playback
- Priority scheduling required; real-time has highest priority
- Small dispatch latency required (system calls should be pre-emptible)

## 2. Hard real-time scheduling

- Absolute deadlines that always have to be met
- e.g.: air traffic control
- Special purpose software running on dedicated hardware

- Hard real-time: ***periodic*** or ***aperiodic*** (unpredictable) events

- Assume  $m$  periodic events; event  $i$  occurs with period  $T_i$  times/second and requires  $t_i$  seconds of CPU to handle
- These time constraints can only be met if
- A real-time system that meets this criterion is ***schedulable***



# Conclusion

- ***Scheduling:*** selecting a waiting process from the ready queue and allocating the CPU to it
- ***FCFS (FIFO) Scheduling:***
  - Run threads to completion in order of submission
  - Pros: Simple and minimizes overhead
  - Cons: Short jobs get stuck behind long ones
  - If jobs are variable in size can have very poor average response time
  - If jobs are equal in size, FIFO is optimal in terms of average response time.
- ***Round-Robin Scheduling:***
  - Give each thread a small amount of CPU time when it executes; cycle between all ready threads
  - Pros: Better for short jobs
  - Cons: If jobs are equal in size very poor average response time



# Conclusion

- ***Shortest Job First (SJF)/Shortest Remaining Time First (SRTF):***
  - Run whatever job has the least amount of computation to do/least remaining amount of computation to do
  - Pros: Optimal (average response time)
  - Cons: Hard to predict future, Unfair
- ***Multi-Level Feedback Scheduling:***
  - Multiple queues of different priorities
  - Automatic promotion/demotion of process priority in order to approximate SJF/SRTF
  - Can achieve a balance between responsiveness, low overhead, and fairness

