



Recommender Systems

COMP 30030 Introduction to Artificial Intelligence


Dr. Michael O'Mahony

michael.omahony@ucd.ie



Overview

- The Long Tail and Information Discovery
- Recommender Systems Overview
- Collaborative Filtering Algorithms
- Evaluation Methodology and Metrics
- Advantages and Limitations



The Long Tail and Information Discovery




The Long Tail

- A new economic model for the media and entertainment industries.
- In the past, these have been hit-driven industries:
 - The average US movie theatre will not show a film unless it can attract at least 1,500 people over a two-week run; that's the screen rental cost.
- Online stores such as Amazon, iTunes, Netflix etc. carry much greater inventory levels... Provide customers with access to niche content in a manner that is revolutionising sales.
- See: *The Long Tail*, Chris Anderson.

● ● ● | The Long Tail


- Physical vs digital stores: finite vs infinite inventory.
- Popularity no longer has the monopoly on profitability
- Products with low demand/low sales volume can collectively make up a market share that rivals the relatively few bestsellers and blockbusters, if the store/distribution channel is large enough.
- Recommender systems help to drive demand down the long-tail – **enable findability**



Blockbuster - A (Brief) History...

- 1985 - the first Blockbuster store opens in Dallas, Texas...
- At peak (circa 2004), Blockbuster employed about 60,000 people in over 9,000 stores...
- In 2000, the company declined an opportunity to purchase a new-ish company called Netflix for a modest \$50 million...
- Fast forward to 2010... due to competition from Netflix and others, Blockbuster lost significant revenue... filed for bankruptcy...





Recommender Systems – Amazon




Amazon Example...

Recommender Systems: An Introduction 1st Edition

by [Dietmar Jannach](#) (Author), [Markus Zanker](#) (Author), [Alexander Felfernig](#) (Author), [Gerhard Friedrich](#) (Author)

★★★★☆ 8 customer reviews

[Look inside](#) ↓



ISBN-13: 978-0521493369

ISBN-10: 0521493366

[Why is ISBN important?](#)

Sell yours for a Gift Card

We'll buy it for **\$24.21**

[Learn More](#)

Kindle

\$55.82

Hardcover

\$58.52 - \$63.64

Other Sellers

from \$54.95

☐ Buy used

\$58.52

☒ Buy new

\$63.64

In Stock.

Ships from and sold by Amazon.com. Gift-wrap available.

List Price: ~~\$79.99~~ Save: \$16.35 (20%)

34 New from **\$58.47**

This item ships to **Dublin, Ireland**. Want it **Friday, Jan. 29**? Order within **66 hrs 45 mins** and choose **AmazonGlobal Priority Shipping** at checkout. [Learn more](#)

Qty: 1



Add to Cart

[Turn on 1-Click ordering](#)

Ship to:

Michael O'Mahony- Dublin

More Buying Choices

34 New from **\$58.47** | 18 Used from **\$54.95**

52 used & new from **\$54.95**

[See All Buying Options](#)

amazonstudent

FREE TWO-DAY SHIPPING
FOR COLLEGE STUDENTS [Learn more](#)



Amazon Example...

Frequently Bought Together



Total price: **\$87.77**

Add both to Cart

Add both to List

- ✓ **This item:** Recommender Systems: An Introduction by Dietmar Jannach Hardcover **\$63.64**
- ✓ **Programming Collective Intelligence: Building Smart Web 2.0 Applications** by Toby Segaran Paperback **\$24.13**

Customers Who Bought This Item Also Bought

Page 1 of 15

Programming Collective Intelligence: Building Smart Web 2.0 Applications › Toby Segaran ★★★★★ 118 Paperback \$24.13 ✓Prime	Recommender Systems Handbook Francesco Ricci ★★★★★ 1 Hardcover	Advanced Analytics with Spark: Patterns for Learning from Data at... Sandy Ryza ★★★★★ 15 #1 Best Seller in Website Analytics Paperback \$23.81 ✓Prime	Algorithms of the Intelligent Web Haralambos Marmanis ★★★★★ 18 Paperback \$29.03 ✓Prime	Learning Spark: Lightning-Fast Big Data Analysis › Holden Karau ★★★★★ 40 Paperback \$29.01 ✓Prime	Building a Recommendation System with R Suresh K. Gorakala ★★★★★ 5 Paperback \$29.99 ✓Prime	Data Science for Business: What you need to know about data mining and... › Foster Provost ★★★★★ 125 Paperback \$30.39 ✓Prime

Amazon Example...

Frequently Bought Together



Total price: **\$87.77**

Add both to Cart

Add both to List

- ✓ **This item:** Recommender Systems: An Introduction by Dietmar Jannach Hardcover **\$63.64**
- ✓ **Programming Collective Intelligence: Building Smart Web 2.0 Applications** by Toby Segaran Paperback **\$24.13**


Customers Who Bought This Item Also Bought

Page 1 of 15

Programming Collective Intelligence: Building Smart Web 2.0 Applications › Toby Segaran ★★★★★ 118 Paperback \$24.13 ✓Prime	Recommender Systems Handbook Francesco Ricci ★★★★★ 1 Hardcover	Advanced Analytics with Spark: Patterns for Learning from Data at... Sandy Ryza ★★★★★ 15 #1 Best Seller in Website Analytics Paperback \$23.81 ✓Prime	Algorithms of the Intelligent Web Haralambos Marmanis ★★★★★ 18 Paperback \$29.03 ✓Prime	Learning Spark: Lightning-Fast Big Data Analysis › Holden Karau ★★★★★ 40 Paperback \$29.01 ✓Prime	Building a Recommendation System with R Suresh K. Gorakala ★★★★★ 5 Paperback \$29.99 ✓Prime	Data Science for Business: What you need to know about data mining and... › Foster Provost ★★★★★ 125 Paperback \$30.39 ✓Prime

Amazon Example...

Frequently Bought Together



- ✓ **This item:** Recommender Systems: An Introduction by Dietmar Jannach Hardcover **\$63.64**
- ✓ **Programming Collective Intelligence: Building Smart Web 2.0 Applications** by Toby Segaran Paperback **\$24.13**

Customers Who Bought This Item Also Bought

Page 1 of 15



 <p>Programming Collective Intelligence: Building Smart Web 2.0 Applications › Toby Segaran ★★★★★ 118 Paperback \$24.13 ✓Prime</p>	 <p>Recommender Systems Handbook Francesco Ricci ★★★★★ 1 Hardcover</p>	 <p>Advanced Analytics with Spark: Patterns for Learning from Data at... Sandy Ryza ★★★★★ 15 #1 Best Seller in Website Analytics Paperback \$23.81 ✓Prime</p>	 <p>Algorithms of the Intelligent Web Haralambos Marmanis ★★★★★ 18 Paperback \$29.03 ✓Prime</p>	 <p>Learning Spark: Lightning-Fast Big Data Analysis › Holden Karau ★★★★★ 40 Paperback \$29.01 ✓Prime</p>	 <p>Building a Recommendation System with R Suresh K. Gorakala ★★★★★ 5 Paperback \$29.99 ✓Prime</p>	 <p>Data Science for Business: What you need to know about data mining and... › Foster Provost ★★★★★ 125 Paperback \$30.39 ✓Prime</p>
--	---	--	---	---	---	---



Amazon Example...

Frequently Bought Together



- ✓ **This item:** Recommender Systems: An Introduction by Dietmar Jannach Hardcover \$63.64
- ✓ **Programming Collective Intelligence: Building Smart Web 2.0 Applications** by Toby Segaran Paperback \$24.13

**Non-personalised
Recommendations**

Customers Who Bought This Item Also Bought

Page 1 of 15

Book Title	Author	Format	Price
Programming Collective Intelligence: Building Smart Web 2.0 Applications	Toby Segaran	Paperback	\$24.13 ✓Prime
Recommender Systems Handbook	Francesco Ricci	Hardcover	\$63.64 ✓Prime
Advanced Analytics with Spark: Patterns for Learning from Data at Scale	Sandy Ryza	Paperback	\$23.81 ✓Prime
Algorithms of the Intelligent Web	Haralambos Marmanis	Paperback	\$29.03 ✓Prime
Learning Spark: Lightning-Fast Big Data Analysis	Holden Karau	Paperback	\$29.01 ✓Prime
Building a Recommendation System with R	Suresh K. Gorakala	Paperback	\$29.99 ✓Prime
Data Science for Business: What you need to know about data mining and...	Foster Provost	Paperback	\$30.39 ✓Prime

Amazon Example...

Your Recently Viewed Items and Featured Recommendations

Inspired by your browsing history

Page 2 of 9 | [Start over](#)



The German War: A Nation Under Arms, 1939-45
Nicholas Stargardt
★★★★★ 30
Kindle Edition
£12.34



The Silent Deep: The Royal Navy Submarine...
> James Jinks
★★★★★ 68
Kindle Edition
£12.34



Silo 49: Dark Till Dawn
> Ann Christy
★★★★★ 29
Kindle Edition
£1.99



SPQR: A history of Ancient Rome
> Mary Beard
★★★★★ 120
Kindle Edition
£8.99



The Great Swindle
Pierre Lemaitre
★★★★★ 29
Kindle Edition
£5.99



Silo 49: Deep Dark
> Ann Christy
★★★★★ 31
Kindle Edition
£1.99





Amazon Example...

Your Recently Viewed Items and Featured Recommendations

Inspired by your browsing history

Page 2 of 9 | [Start over](#)



The German War: A Nation Under Arms, 1939-45
Nicholas Stargardt
★★★★★ 30
Kindle Edition
£12.34



The Silent Deep: The Royal Navy Submarine...
> James Jinks
★★★★★ 68
Kindle Edition
£12.34



Silo 49: Dark Till Dawn
> Ann Christy
★★★★★ 29
Kindle Edition
£1.99



SPQR: A history of Ancient Rome
> Mary Beard
★★★★★ 120
Kindle Edition
£8.99



The Great Swindle
Pierre Lemaitre
★★★★★ 29
Kindle Edition
£5.99



Silo 49: Deep Dark
> Ann Christy
★★★★★ 31
Kindle Edition
£1.99







Amazon Example...


Your Recently Viewed Items and Featured Recommendations

Inspired by your browsing history

Page 2 of 9 | [Start over](#)




The German War: A Nation Under Arms, 1939-45
Nicholas Stargardt
★★★★★ 30
Kindle Edition
£12.34




The Silent Deep: The Royal Navy Submarine...
James Jinks
★★★★★ 68
Kindle Edition
£12.34




Silo 49: Dark Till Dawn
Ann Christy
★★★★★ 29
Kindle Edition
£1.99




SPQR: A history of Ancient Rome
Mary Beard
★★★★★ 120
Kindle Edition
£8.99



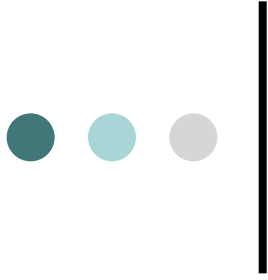
The Great Swindle
Pierre Lemaitre
★★★★★ 29
Kindle Edition
£5.99



Silo 49: Deep Dark
Ann Christy
★★★★★ 31
Kindle Edition
£1.99



Personalised Recommendations



Recommender Systems Overview



Recommender Systems

- Recommendations can be personalised or non-personalised:
 - Tailored to the particular interests of users or generic
- Approaches to recommendation:
 - Collaborative filtering (CF), content-based, demographic, hybrid, group RS, conversational RS... This lecture, focus on CF
- Sources of recommendation knowledge:
 - Implicit/explicit user preferences, product metadata, product reviews, tweets, FB posts/likes...
- Recommendation output:
 - Ranked lists of items, predicted ratings, explanations (why are these items recommended?)



Content-based Recommendation

- Items are recommended which are similar in content to previously selected items
- Recommendations are based on a description of the **content** of items as opposed to users' **opinions** about items
- Comparisons between items are calculated over the features associated with each item
 - For example, in the movie domain, features include actors, genres, director, plot summary...
 - Based on the range of feature values associated with a user's past choices, further movies with similar feature values are recommended
 - *A more like this* approach to recommendation



Collaborative Filtering (CF)

- CF – automates the “word-of-mouth” process
- Assists users to make choices based on the opinions of other users
- The underlying heuristic: people who agreed or disagreed on items in the past are likely to agree or disagree on future items
- Predictions and recommendations are made for users by combining the preferences of similar users in the system:
 - Note: content descriptions are not required – recommendations are based only on user preferences

Collaborative Filtering (CF)

- Let's look at a typical problem...
- Suppose we had the following preference data...

	The Quiet Man	Casino	Star Wars	Top Gun	Dallas: The Movie
Eamon	3		3	4	2
Sharon	4	1	4	2	4
John	5	2		2	5
Trisha	2	5	3		1
Mike	?	2	4	1	5

- Q – based on this preference data, would user *Mike* like or dislike *The Quiet Man*?

Collaborative Filtering (CF)

- Let's look at a typical problem...
- Suppose we had the following preference data...

	The Quiet Man	Casino	Star Wars	Top Gun	Dallas: The Movie
Eamon	3		3	4	2
Sharon	4	1	4	2	4
John	5	2		2	5
Trisha	2	5	3		1
Mike	?	2	4	1	5


- Q – based on this preference data, would user *Mike* like or dislike *The Quiet Man*?

Collaborative Filtering (CF)

- Let's look at a typical problem...
- Suppose we had the following preference data...

	The Quiet Man	Casino	Star Wars	Top Gun	Dallas: The Movie
Eamon	3		3	4	2
Sharon	4	1	4	2	4
John	5	2		2	5
Trisha	2	5	3		1
Mike	?	2	4	1	5

- Q – based on this preference data, would user *Mike* like or dislike *The Quiet Man*?



Collaborative Filtering (CF) Algorithms

User-based CF Algorithm



Collaborative Filtering (CF)

- Many CF algorithms proposed (user-based, item-based, matrix factorisation, other model-based approaches)...
- Focus on user-based CF – similar to the example we looked at earlier...
- Remember CF approaches operate over a set of user preferences; no content descriptions are required



User-based CF Algorithm

Prediction Algorithm: objective is to predict the rating of a *target item* for a given user (*active user*)

1. Data:

Obtain preference data from users

2. Similarity Computation:

Compute the similarity between the active user and all other users in the system which have rated the target item

3. Neighbourhood Formation:

Select a subset consisting of the most similar users to the active user which have rated the target item

4. Make a prediction for the active user by aggregating the neighbour's ratings for the target item

Different approaches used in the above steps ...



Preference Data

- Need to acquire (lots of) user preferences
- Types of ratings:
 - Scalar ratings: numerical ratings (e.g. 1 - 5 stars for movies) or ordinal ratings (e.g. strongly agree, agree, neutral, disagree, strongly disagree)
 - Binary ratings: e.g. votes up/down, like/dislike or agree/disagree
 - Unary ratings: e.g. purchased/not purchased a book, visited/not visited a web page (note that not purchasing a book does not imply dislike)
- Ratings are gathered explicitly or implicitly



Preference Data

- Explicit Ratings
 - Users supply ratings:
 - Generally in the form of scalar ratings (e.g. 5 point scale) or binary ratings
 - Good reflection of user preferences
 - Cost – users may tire of providing ratings
 - Users need to be convinced there is a benefit in order to make the effort
 - CF algorithms need many ratings to function accurately
- Implicit Ratings
 - Not obtained directly from user; instead, ratings are “inferred” – e.g. a user’s browsing patterns (time spent on web page, a user’s bookmarks, links followed)
 - Removes cost of explicitly gathering ratings
 - Every user interaction can potentially contribute
 - Can be combined with explicit ratings
 - But: implicit ratings are typically noisy
- Ratings are “stored” in a user-item matrix

[illegible]



User Similarity

- How do we measure the similarity between users?
- A number of approaches have been considered...
- In most approaches, there needs to be an overlap between two users in order to compute their similarity:
 - No overlap – assume similarity is zero
 - Small overlap – skewed results
 - Large profiles (i.e. users who have rated many items, bots) will have overlaps with many other profiles – problematic, can lead to such profiles being included in neighbourhoods “by default”



Mean Squared Difference

- The Mean Squared Difference in the ratings between users a and i is computed as:

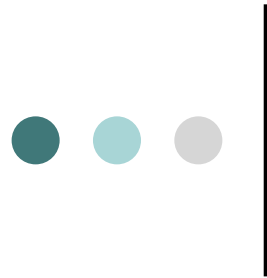
$$\text{MSD}_{a,i} = \frac{\sum_{j \in I_a \cap I_i} (r_{a,j} - r_{i,j})^2}{|\{j : j \in I_a \cap I_i\}|}$$

where:

- I_a is the set of items rated by user a
 - $r_{a,j}$ is the rating of user a for item j
- MSD computes the *difference* in ratings between two users; convert MSD into a *similarity* metric as follows:

$$w_{a,i} = 1 - \frac{\text{MSD}_{a,i}}{(r_{\max} - r_{\min})^2}$$

where r_{\min} and r_{\max} are the minimum and maximum ratings, resp.



Mean Squared Difference

- Summations over co-rated items only
- Results in a similarity value of $[0, +1]$
- MSD assumes that users rate items according to similar distributions:
 - Critical vs. 'generous' raters? Critical raters rarely assign the maximum ratings to items...
 - So while differences in magnitudes between users' ratings may exist, they may be highly correlated
 - Need a scale-invariant user similarity function to more accurately reflect rating patterns and to capture trends between users' ratings...



Pearson Correlation

- Given by:





$$w_{a,i} = \frac{\sum_{j \in I_a \cap I_i} (r_{a,j} - \bar{r}_a)(r_{i,j} - \bar{r}_i)}{\sqrt{\sum_{j \in I_a \cap I_i} (r_{a,j} - \bar{r}_a)^2 \sum_{j \in I_a \cap I_i} (r_{i,j} - \bar{r}_i)^2}}$$

where:

- I_a is the set of items rated by user a
 - $r_{a,j}$ is the rating of user a for item j
 - \bar{r}_a and \bar{r}_i are the mean ratings of users a and i (**Note: means are computed over co-rated items only**)
- Summations over co-rated items only
- Results in a value of $[-1, +1]$
 - +1 indicates total agreement on co-rated items
 - 0 indicates no similarity between users
 - 1 indicates total disagreement – useful, can exploit the preferences of these users for recommendation
- Widely used similarity metric

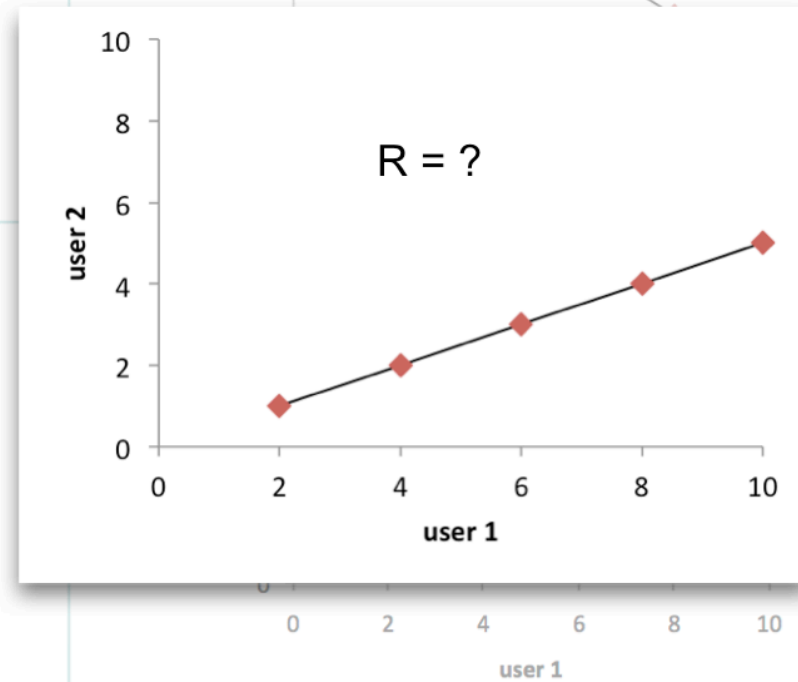




Pearson Correlation Examples






Pearson Correlation Examples





Pearson Correlation Examples





Significance Weighting

- Weights that are calculated over small numbers of co-rated items may provide an unreliable measure of the similarity between users
- Modify similarity weights based on the number of co-rated items (n) between users as follows:


$$w'_{a,i} = \begin{cases} w_{a,i} \times \frac{n}{N} & \text{if } n < N \\ w_{a,i} & \text{otherwise} \end{cases}$$

where N is a constant (typically set to 50)

- Improves reliability of similarity weights – similarity is modified as a function of the number of co-rated items

Neighbourhood Formation

- Select a subset of users to make predictions for the active user
- Different approaches – here consider **k nearest-neighbour (kNN)** approach
- Form a neighbourhood by selecting the k most similar users which have rated the target item
- Neighbourhood size – determined by experiment:
 - Large k can lead to reduced accuracy by diluting the influence of the most similar neighbours
 - Small k can result in poor accuracy for users without many (or any) close neighbours





Making Predictions

- **Resnick's** algorithm – deviation from mean approach – target item is assigned a rating that is an adjusted form of the active user's mean rating:

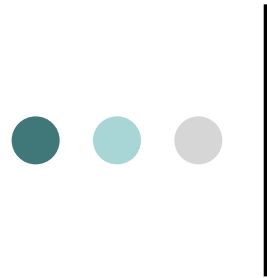
$$p_{a,j} = \bar{r}_a + \frac{\sum_{i=1}^n w_{a,i} (r_{i,j} - \bar{r}_i)}{\sum_{i=1}^n |w_{a,i}|}$$

- Where:
 - $p_{a,j}$ is the predicted rating for the active user a on item j
 - \bar{r}_a is the mean rating of the active user a (**Note: mean is calculated over all items user a has rated**)
 - \bar{r}_i is the mean rating of neighbour i (**Note: mean is calculated over all items neighbour i has rated**)
 - $r_{i,j}$ is the rating of neighbour i for item j
 - $w_{a,i}$ is the similarity between the active user a and neighbour i
 - n is the number of neighbours



Performance Evaluation

- Live-User Analysis vs Off-line Evaluation
 - Analysis of real usage and prediction/recommendation feedback. A/B testing to evaluate different algorithms/techniques. Most comprehensive approach but expensive.
 - Offline evaluations to test core recommendation algorithm components by using existing user-rating datasets.
- Off-line Evaluation Methodology – k -Fold Cross Validation:
 - Randomly partition the data set into k subsamples.
 - Of the k subsamples, a single subsample is used as test data, the remaining $k-1$ subsamples as training data.
 - Repeat k times, using each subsample in turn as the test set.
 - Compute average results (accuracy etc.) over the test set for each fold – then compute average results across all folds.



Evaluation Metric

Mean Squared Error (RMSE) – for a given fold, RMSE is calculated as follows:

$$RMSE = \sqrt{\frac{1}{N} \sum_{(a,j) \in T} (p_{a,j} - r_{a,j})^2}$$

where T is the test set, N is the number of ratings in the test set for which a prediction can be made, and $p_{a,j}$ and $r_{a,j}$ are the predicted and actual ratings for user a , item j , respectively.

RMSE penalises larger errors over smaller errors.

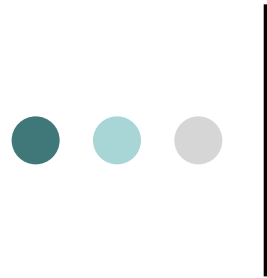
The smaller the RMSE value, the more accurate the recommender system.



Performance Evaluation

Other Considerations:

- **Coverage** (one definition) – the percentage of ratings in the test set for which predictions can be made
- **Serendipity** – recommend items the user is unaware of but would like
- **Diversity** – ensuring that top-N lists are not comprised of only “similar” items
- **Site performance metrics** – tracking system throughput (items purchased or downloaded, the number of active users etc.)
- **Robustness** – are CF systems vulnerable to manipulation? (**Yes!**)



CF – Advantages

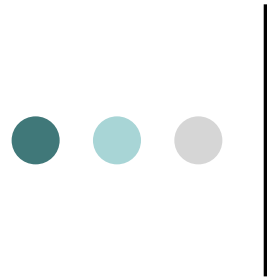
Advantages

○ **Quality & Taste:**

- Human beings have the ability to analyse information based on quality and taste
- Since CF operates on user preferences, it inherits this ability
- Other filtering techniques are less able or unable to quantify the quality that is inherent in certain items – e.g. a content-based search may return relevant documents, but some of these may be actually disliked by users...

○ **Item Features:**

- CF algorithms make predictions / recommendations based solely on user preferences
- Do not rely on any item features that need to be pre-determined and extracted prior to filtering
- Thus, CF algorithms can be readily implemented in complex domains which contain items that are difficult to analyse by automated processes

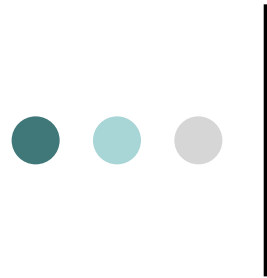


CF – Advantages

- **Serendipitous Recommendations:**
 - CF algorithms have the ability to provide serendipitous recommendations – i.e. recommend items the user likes and is unaware of (and unlikely to find unless they are recommended)

Limitations

- **Cold-Start Problem:**
 - CF algorithms rely on relationships between users and items in order to make predictions / recommendations
 - In new applications, ratings are missing / sparse
 - In some applications, new users are required to rate a subset of items...



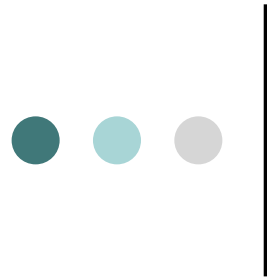
CF – Limitations

- **Early Rater Problem:**

- This issue applies to predictions that are sought for items that are new or only recently added to a system.
- For such items, predictions are problematic, since there are no (or few) ratings contained in the system on which to calculate predictions
- Similarly, newly registered users are problematic for even established systems, until they have rated at least some of the available items

- **Sparsity Problem:**

- Typically, recommender system datasets are very sparse, because users will have assigned ratings to only a relatively small number of items ($\ll 1\%$).
- This is particularly true for e-commerce systems, where very large numbers of items are frequently offered for sale (e.g. consider Amazon.com)
- Thus it may not be possible to make accurate (or any) predictions / recommendations for certain users / items



CF – Limitations

- **Scalability**

- User-based CF algorithms suffer from serious scalability problems:
 - Computations grow with both the number of users and items
 - With potentially millions of users and items, typical real-world systems running user-based CF algorithms will suffer serious scalability problems
- Solutions have been proposed to improve scalability performance:
 - Cluster the database into groups of like-minded users from which all neighbours are subsequently drawn...
 - Model-based approaches – item-based CF, matrix factorisation approaches ...



Summary

- The long tail and information discovery
- Recommender systems algorithms: collaborative filtering. content-based (many other approaches exist...)
- Focused on the one particular collaborative filtering algorithm – user-based CF
- Evaluation methodology and metrics
- Advantages and limitations