

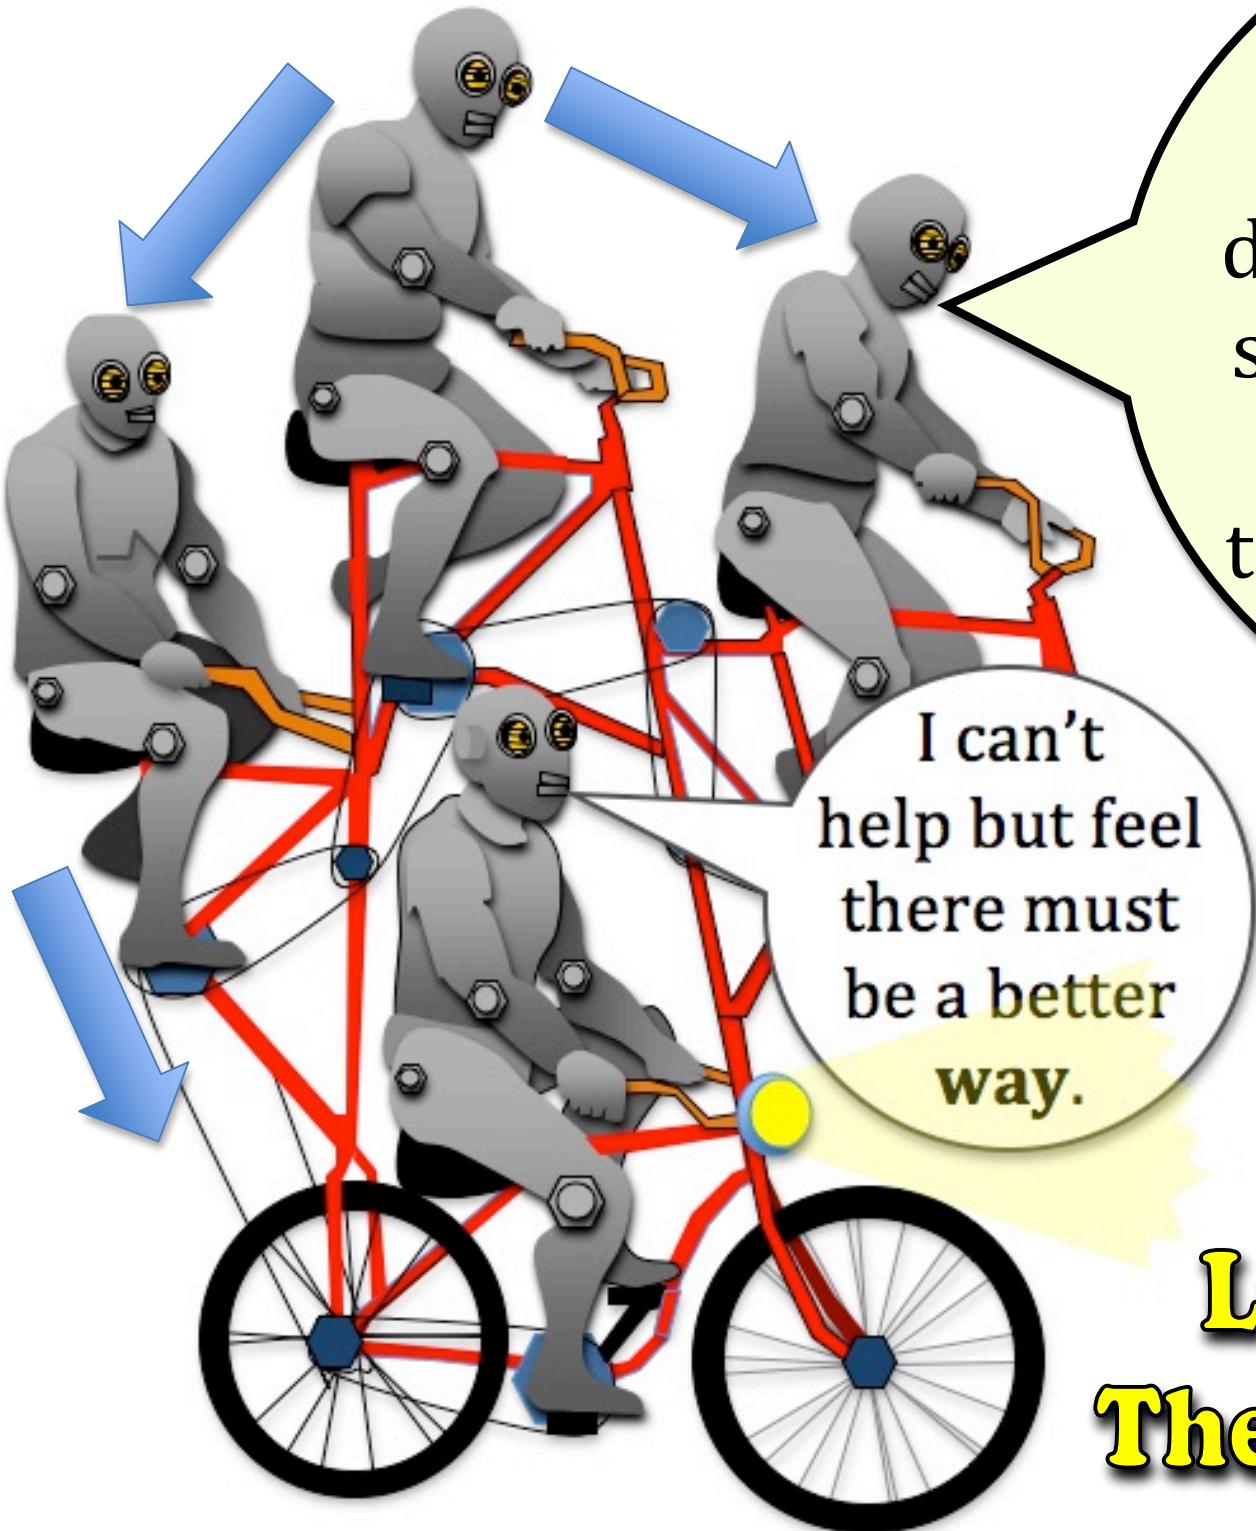
A Beginner's Guide to Relational Databases and SQL





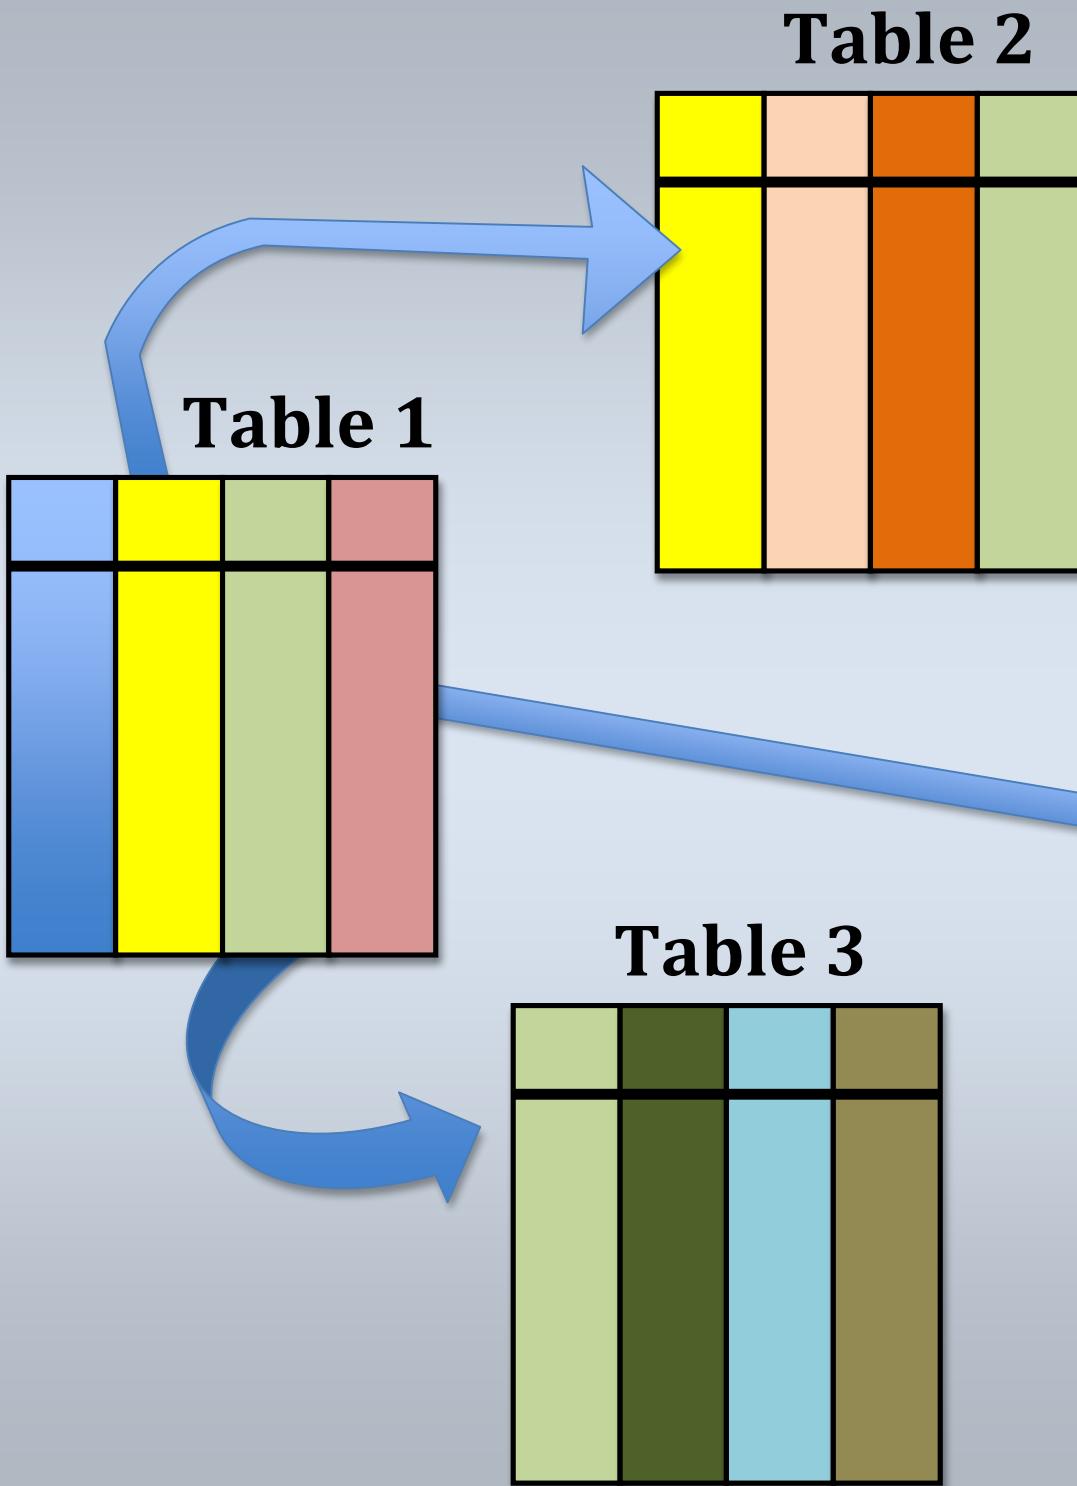
Want some
Alternative Facts?
Flat File systems?
They're for losers!
Hierarchical databases?
They're *Fake News!*
Relational Databases
are *Tremendous*,
like ME!

Seriously, though:
The Relational Model is Best!



Hierarchical Databases organize data into tree structures, such as binary trees and **B-Trees**. They conflate the *indexing* of data with its *logical form*.

**Luckily, there is:
The Relational Model**



**A table captures
a relationship
between aspects
of the data**

** Colours represent kind
of data in these tables*

The Relational Model captures the relationships in our data

This is a
"Table"

| Name | Opponent |
|-----------------|--------------------------|
| Adam Smith | <i>Karl Marx</i> |
| Abraham Lincoln | <i>John Wilkes Booth</i> |
| Hillary Clinton | <i>Donald Trump</i> |
| Al Capone | <i>Eliot Ness</i> |
| Batman | <i>The Joker</i> |

It consists
of Columns
and Rows

A Database that employs the Relational Model
is called an RDBMS (Relational DB Mgmt. System)

Each table consists of columns calls "fields" such as Name, Gender, Fictive & Opponent

| Name | Gender | Fictive | Opponent |
|-----------------|---------------|------------|---------------------|
| Adam Smith | <i>male</i> | <i>no</i> | <i>Karl Marx</i> |
| Princess Leia | <i>female</i> | <i>yes</i> | <i>Darth Vader</i> |
| Hillary Clinton | <i>female</i> | <i>no</i> | <i>Donald Trump</i> |
| Al Capone | <i>male</i> | <i>no</i> | <i>Eliot Ness</i> |
| Batman | <i>male</i> | <i>yes</i> | <i>The Joker</i> |

**Each table consists of rows calls "records"
There is one for each unique key value**

Every table can be viewed as A collection of individual records

Princess Leia

female

yes

Darth Vader

Opponent

Karl Marx

Darth Vader

Donald Trump

Eliot Ness

The Joker

Or every table
can be seen as
a set of fields

The relational value comes from
the intersection of both fields and records

It Does Not Compute!



SQL Constraints
and DATA INTEGRITY !!



It's a NULL value,
not a blank!

An RDBMS Manager can impose the following constraints on the columns/fields of a table:

| Constraint | Constraint Action |
|-------------|--|
| NOT NULL | <i>Ensures a field cannot ever have a NULL value</i> |
| DEFAULT | <i>Provides a default value for a field when none given</i> |
| UNIQUE | <i>Requires every value in a field/column to be different</i> |
| PRIMARY KEY | <i>Uniquely specifies every row/record in a table</i> |
| FOREIGN KEY | <i>Uniquely specifies every row/record in another table</i> |
| CHECK | <i>Provides an arbitrary test for the values of a field</i> |
| INDEX | <i>This field will be used to quickly retrieve data from table</i> |

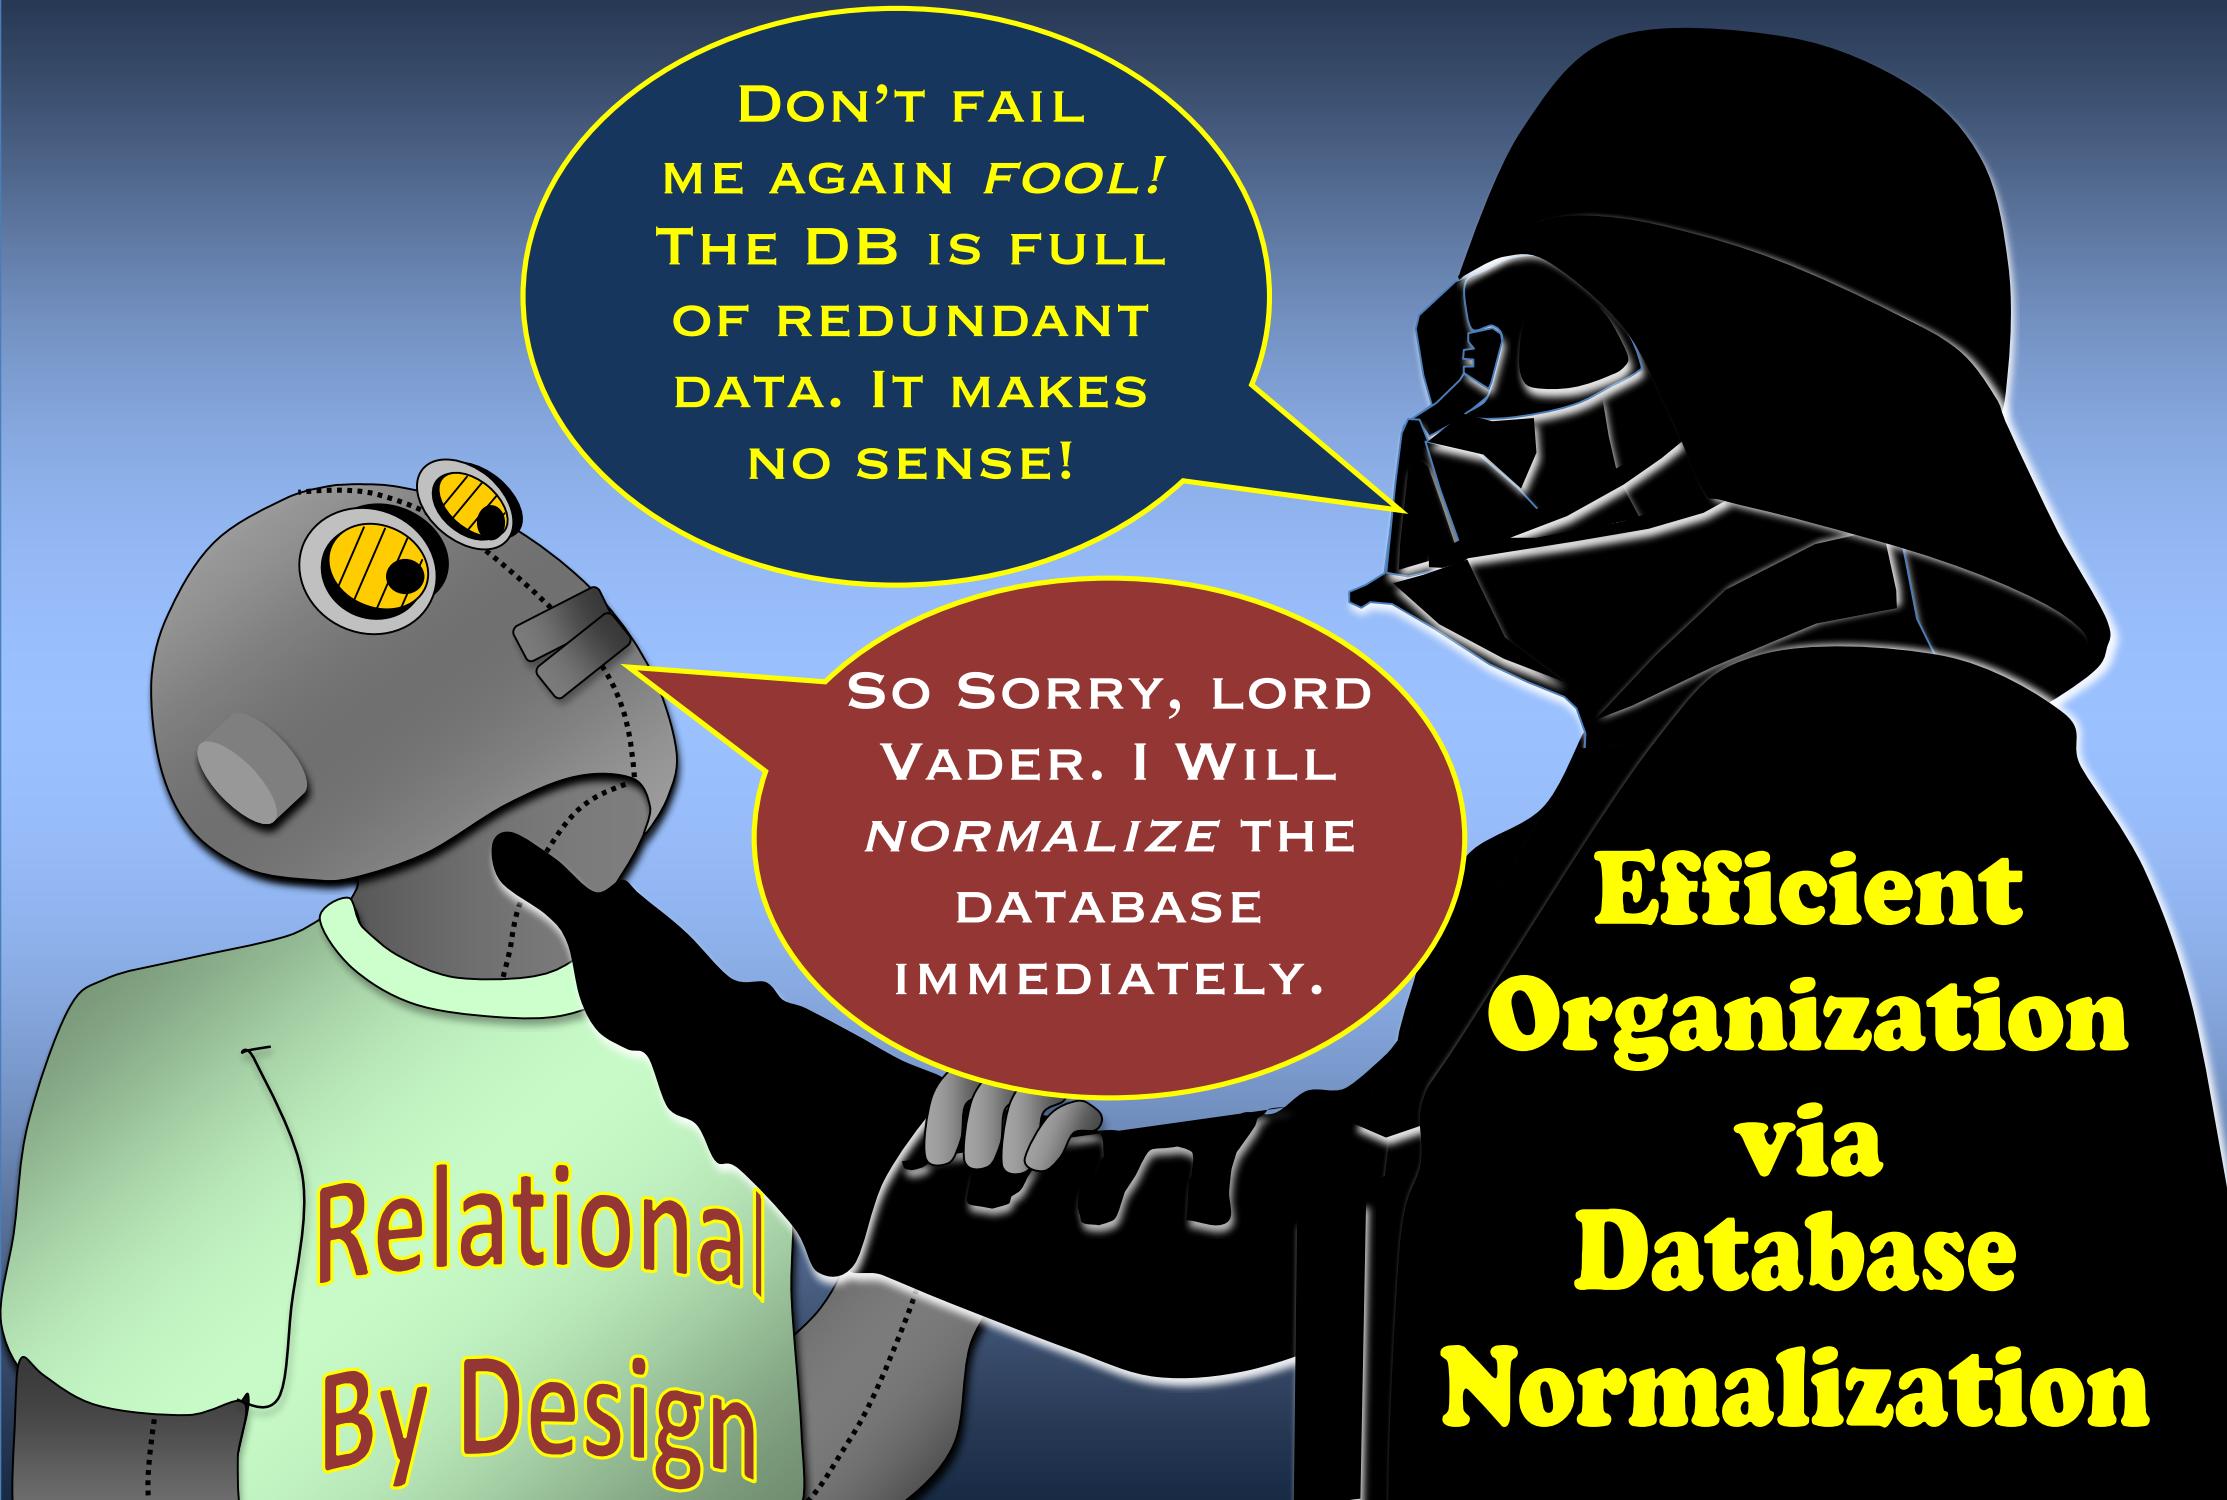
**NB: These are "column"-level constraints
We can also apply constraints to whole tables**

An RDBMS ensures "Data Integrity" by enforcing these constraints ...

| Integrity Constraint | Constraint Action |
|------------------------|---|
| ENTITY INTEGRITY | <i>There can be no duplicate records in a database table - EVER!</i> |
| DOMAIN INTEGRITY | <i>Every value of every field in every table must be of the right type, format and in the right range</i> |
| REFERENTIAL INTEGRITY | <i>A row/record cannot be deleted if it is being referenced by another row in another table</i> |
| USER-DEFINED INTEGRITY | <i>An RDBMS allows users to define own integrity constraints over tables and records</i> |

Integrity is a normal state of affairs in an RDBMS because the database is "Normalized" as such

Efficient Organization via Database Normalization



First Normal Form (**1NF**): *Ensure there are no redundant, repeating groups of data by imposing a unique **PRIMARY KEY** on every table.*

Second Normal Form (**2NF**):

2NF assumes 1NF, plus remove all **partial dependencies** of any field/column on the PRIMARY KEY.

Third Normal Form (**3NF**):

3NF assumes 2NF and 1NF. Ensures that every column/field in a table is dependent only on the PRIMARY KEY.

Normalization rasp ensures that data is stored logically and efficiently, to obviate future problems of ...

rasp ... access



First Normal Form (1NF): Ensure no repeating groups of information, as in:

C
U
S
T
O
M
E
R
S

| Customer ID | Customer NAME | DOB | ADDRESS | ORDER |
|-------------|---------------|------------|----------------------------|--|
| THX1138 | Darth Vader | 28.02.0024 | 17 Sith Plaza, Tatooine | <i>Light Sabre, model 27B/6</i> |
| THX1138 | Darth Vader | 28.02.0024 | 17 Sith Plaza, Tatooine | <i>Black helmet polish + cloth</i> |
| THX1138 | Darth Vader | 28.02.0024 | 17 Sith Plaza, Tatooine | <i>Lint remover for cape</i> |

This table lacks a PRIMARY KEY: There is no way to uniquely specify each record

1NF: Split data into separate tables

(NB: Customer ID is a Foreign Key)

| Order ID | Customer ID | Quantity | ORDER |
|----------|-------------|----------|------------------------------------|
| ORD2001 | THX1138 | 1 | <i>Light Sabre, model 27B/6</i> |
| ORD2002 | THX1138 | 1 | <i>Black helmet polish + cloth</i> |
| ORD2003 | THX1138 | 1 | <i>Lint remover for cape</i> |

| Customer ID | Customer NAME | DOB | ADDRESS |
|-------------|---------------|------------|--------------------------|
| THX1138 | Darth Vader | 28.02.0024 | 17 Sith Plaza, Tatooine |
| PLO1001 | Leia Organa | 14.02.0045 | 23 Royal Lane, Aldebaran |

ORDERS
CUSTOMERS

Each table has its own PRIMARY KEY

2NF: No Partial Dependencies among fields

S
A
L
E
S

| ORDER ID | Customer NAME | Customer ID | Order Details | Salesman |
|----------|----------------|-------------|--------------------------------------|--------------------|
| ORD2001 | Darth Vader | THX1138 | <i>Light Sabre, model 27B/6</i> | <i>Bib Fortuna</i> |
| ORD2171 | Leia Organa | PLO1001 | <i>Bunomatix Hair Curlers</i> | <i>Greedo Maxx</i> |
| ORD3092 | Luke Skywalker | LSW2003 | <i>Sand Speeder, model DSX74</i> | <i>Jabba Hutt</i> |

Notice redundancy of Name and Details
These fields are redundant or dependent on others

Remove Customer Name (in CUSTOMERS table already)



| ORDER ID | Customer NAME | Customer ID | Order Details | Salesman |
|----------|----------------|-------------|--------------------------------------|--------------------|
| ORD2001 | Darth Vader | THX1138 | <i>Light Sabre, model 27B/6</i> | <i>Bib Fortuna</i> |
| ORD2171 | Leia Organa | PLO1001 | <i>Bunomatix Hair Curlers</i> | <i>Greedo Maxx</i> |
| ORD3092 | Luke Skywalker | LSW2003 | <i>Sand Speeder, model DSX74</i> | <i>Jabba Hutt</i> |

S
A
L
E
S



Remove Order Details (in ORDERS table already)

2NF removes redundancy, reduces inconsistency

3NF: All Non-primary fields are dependent on Key

CUSTOMERS

| CUSTOMER ID | NAME | STREET | CITY | STATE | ZIP |
|-------------|-------------|---------------|------------|-----------|-------|
| THX1138 | Darth Vader | 17 Sith Plaza | Mos Eisley | Tatooine | 90210 |
| PLO1001 | Leia Organa | 23 Royal Lane | Windu City | Aldebaran | 45872 |
| OBI2007 | Ben Kenobi | 82 Rocky Way | Sandy Ave. | Tatooine | 90734 |

No Transitive Dependencies between Non-Key fields

In the above table, all fields are dependent on the unique key. But aspects of the address are also dependent on each other. Since the Street or City or State and we must change ZIP too.

3NF: Remove Transitive Dependencies

ADDRESSES

| ADDRESS ID | STREET | CITY | STATE | ZIP |
|------------|---------------|------------|-----------|-------|
| 9021017MET | 17 Sith Plaza | Mos Eisley | Tatooine | 90210 |
| 4587223WSA | 23 Royal Lane | Windu City | Aldebaran | 45872 |
| 9073482SAT | 82 Rocky Way | Sandy Ave. | Tatooine | 90734 |

CUSTOMERS

| Customer ID | Customer NAME | DOB | ADDRESS ID |
|-------------|---------------|------------|------------|
| THX1138 | Darth Vader | 28.02.0024 | 9021017MET |
| PL01001 | Leia Organa | 14.02.0045 | 4587223WSA |

We link ZIP to Street & City & State in One Table

1NF: A Primary Key is required for every Table

2NF: Non-Key Attributes are dependent on whole key

3NF: Non-Key Attributes are dependent on nothing but the whole key*



* ***The "whole key" refers to the possibility that a Primary Key may jointly comprise multiple fields in a row***

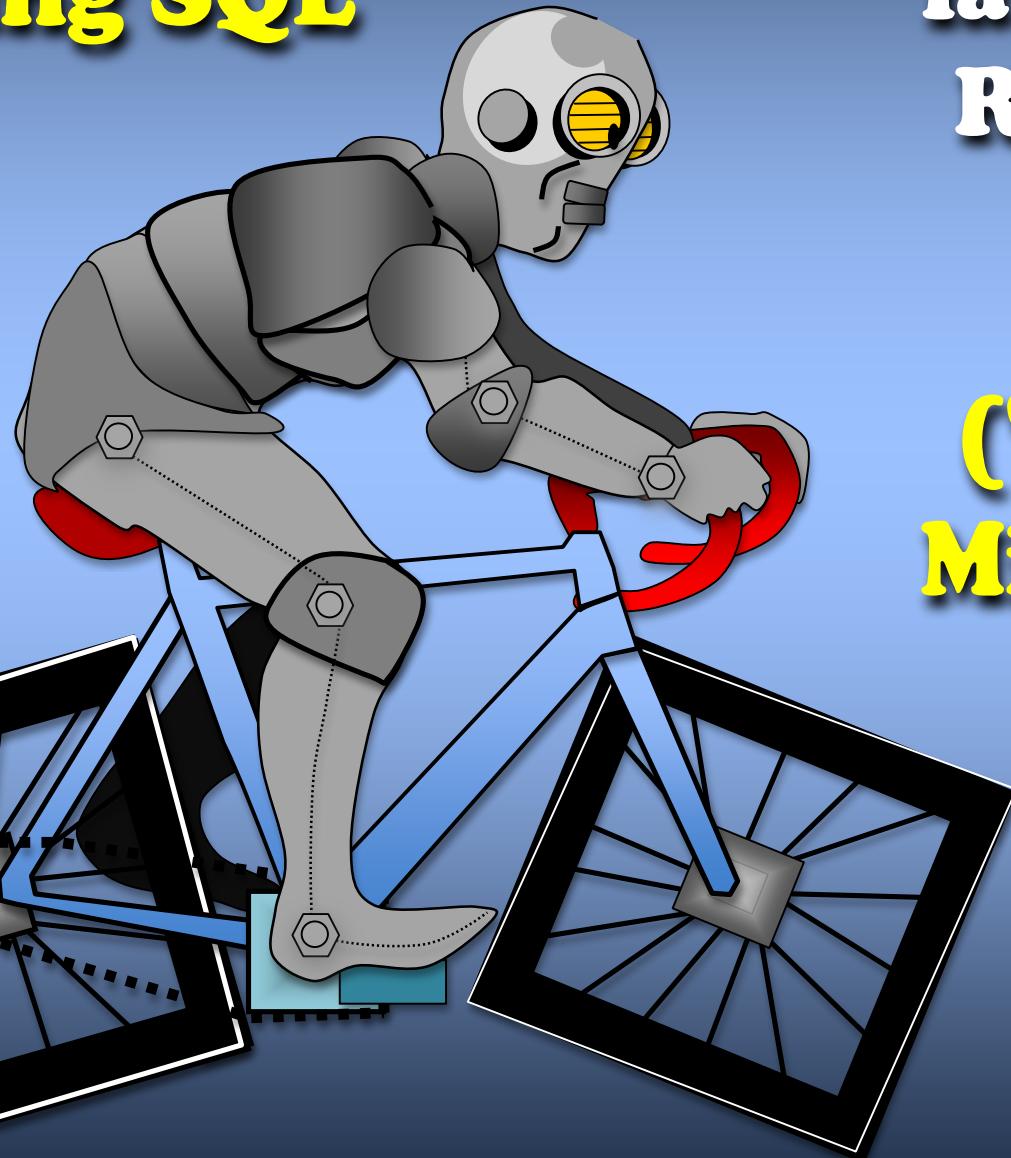


MySQL
A relational
database
employing the
SQL structured
query language

**Free & Paid
Versions**

**Efficient
and Robust**

Other RDBMS Solutions Using SQL



ORACLE
**(founded 1977,
largest commercial
RDBMS provider)**

MS SQL Server
("Sequel" Server)
Microsoft's RDBMS

Microsoft ACCESS
Entry-Level RDBMS
uses "Jet" SQL dialect

Data Definition Language (**DDL**):

Statements for creating a database
and its tables; for deleting tables too

Data Manipulation Language (**DML**):

Statements for adding data to tables
and for retrieving data from tables

Data Control Language (**DCL**):

Statements for granting and revoking
privileges to users of the database

Transaction Control Language (**TCL**):

Statements for maintaining the state
of the DB, via *commits* and *rollbacks*

*SQL provides FOUR
types of statements in
four sub-languages ...
rasp ... each of which
has a different logical
role in an RDBMS*



It Lives!



Create A New Database and its Contents

Create a new SQL database:

CREATE DATABASE database_name;

USE database_name;

Drop (Remove) an existing SQL database:

DROP DATABASE database_name;

Create a New Table in a Database:

```
CREATE TABLE table_name(  
    column1 datatype,  
    column2 datatype,  
    column3 datatype,  
    .....  
    columnN datatype,  
PRIMARY KEY(one or more columns)  
);
```

A table's **PRIMARY KEY** can combine multiple fields into a single unique key field

Create an index for your table for fast access:

```
CREATE UNIQUE INDEX index_name  
ON table_name ( column1,  
                column2,  
                ...,  
                columnN);
```

An **index** is a mapping from field values
to the table rows that contain them

When an **index** is available, any **SELECT**
query will use it for fast lookup of data

DROP a table and its contents from the database:

```
DROP TABLE table_name;
```

DROP an index for a table from the database:

```
ALTER TABLE table_name  
DROP INDEX index_name;
```

ADD a new column to a table:

```
ALTER TABLE table_name  
ADD column_name datatype
```

DROP an existing column from a table:

```
ALTER TABLE table_name  
DROP COLUMN column_name
```

RENAME a Table with a new name:

```
ALTER TABLE table_name  
RENAME TO new_table_name;
```

Change the Data Type of a Table column:

```
ALTER TABLE table_name  
MODIFY COLUMN column_name datatype
```

INSERT a row of column values into a Table:

```
INSERT INTO table_name( col1, ...., colN)  
VALUES ( value1, ..., valueN);
```

DML

Change the contents of a specific table row:

```
UPDATE table_name  
SET col1 = value1, ...., colN=valueN  
[ WHERE {CONDITION} ];
```

DML

Delete any row matching Boolean criteria:

**DELETE FROM table_name
WHERE {CONDITION};**

DML

The WHERE clause specifies Boolean criteria:

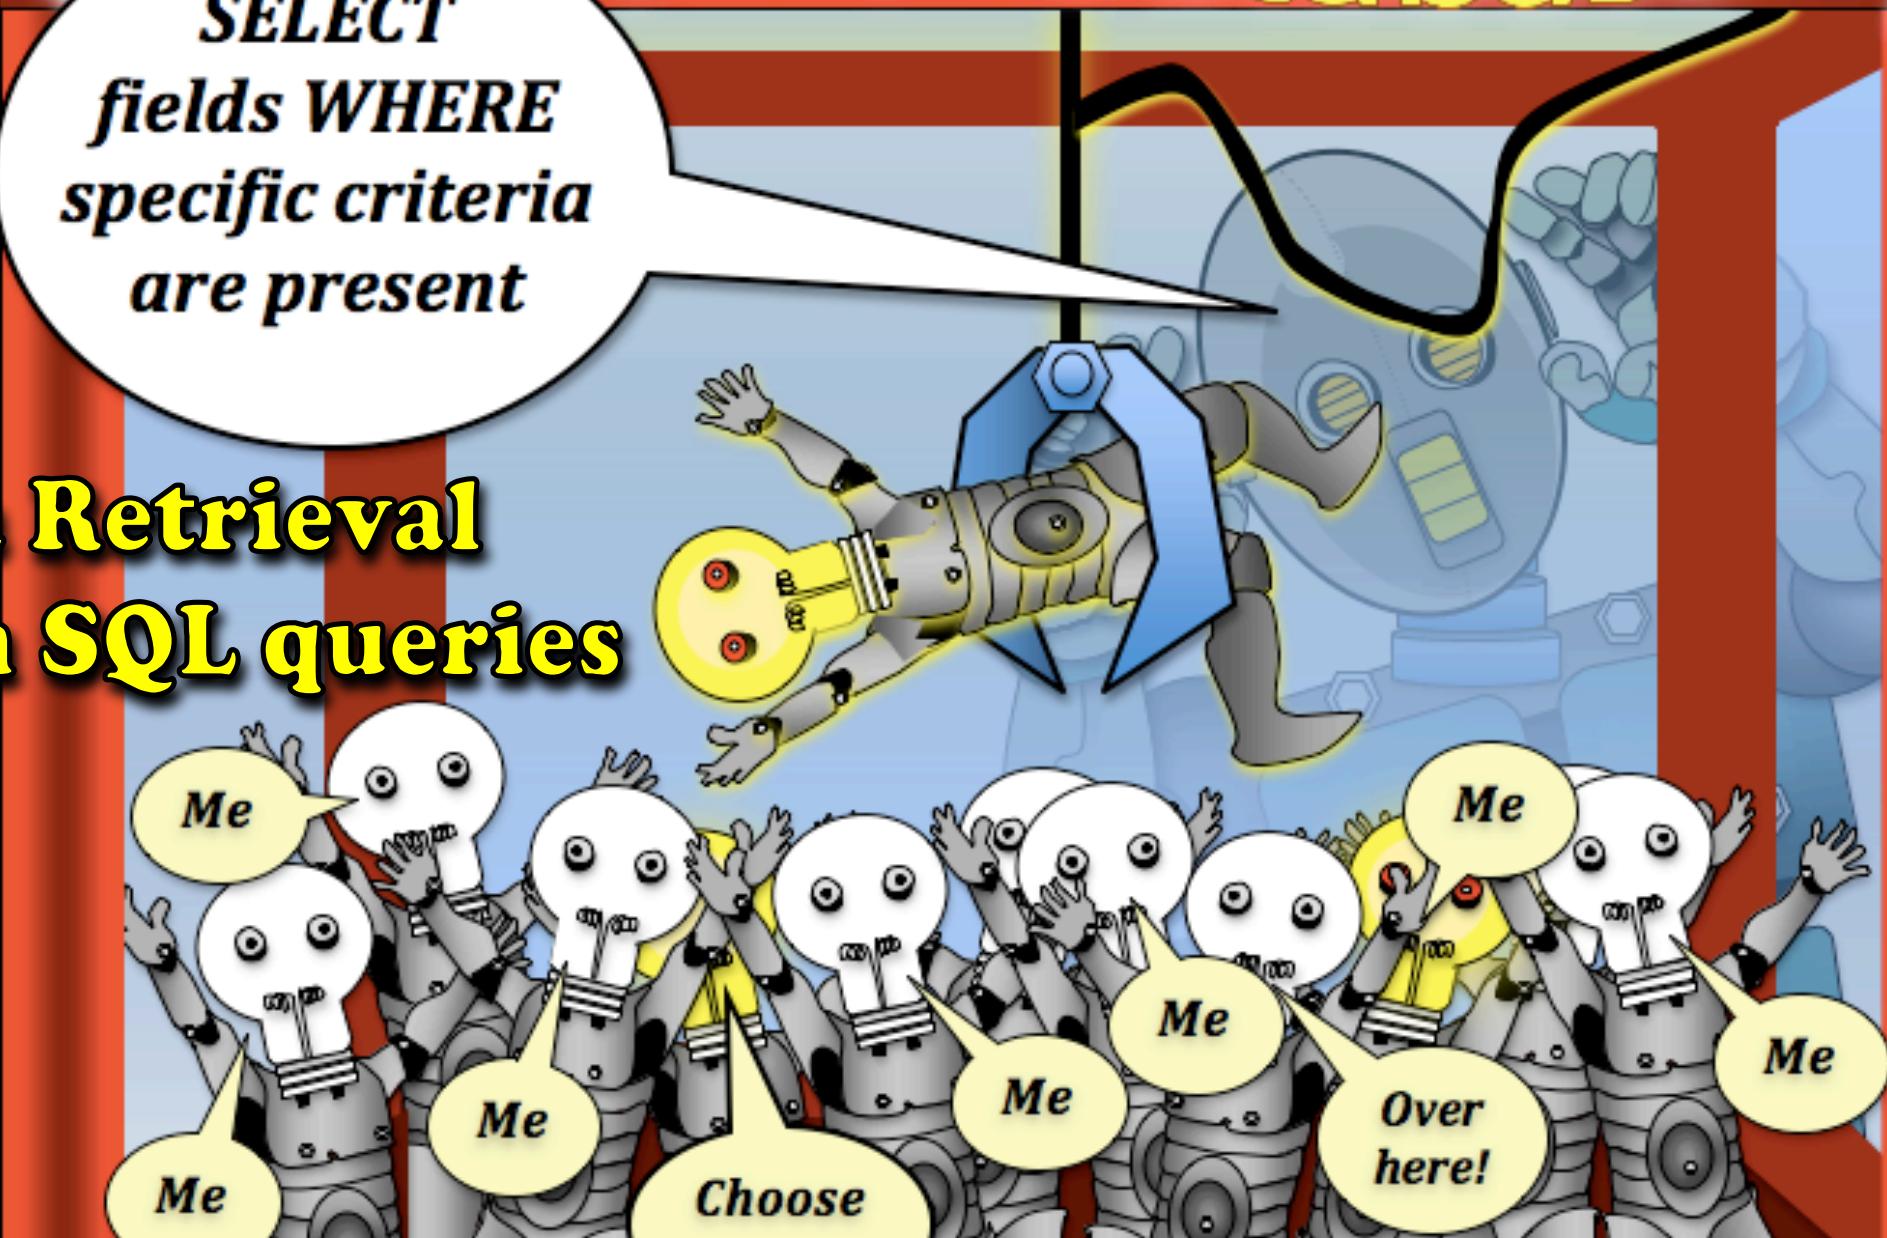
**[SELECT | DELETE FROM]
WHERE condition1
{and|or} condition2 ...
{and|or} ... conditionN;**

DML

MySQL Database

***SELECT
fields WHERE
specific criteria
are present***

Data Retrieval with SQL queries



MySQL Database

\$1

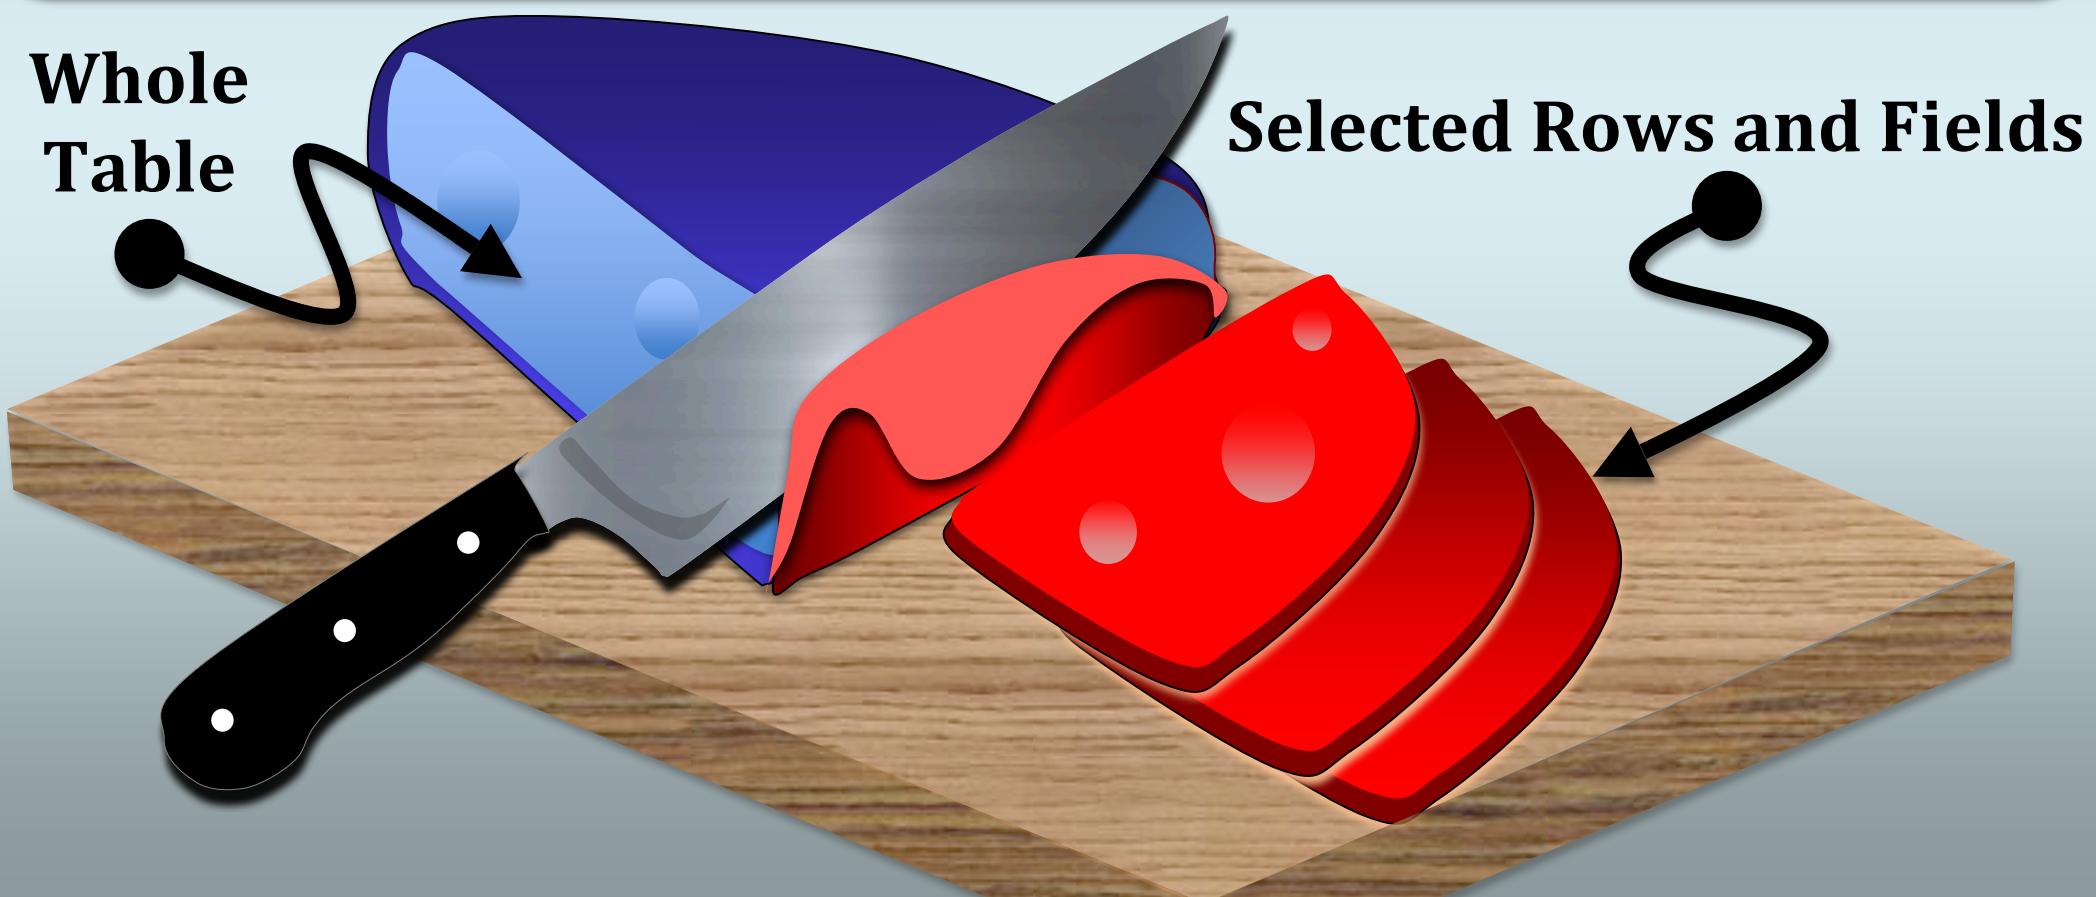
Using SQL for Queries

So what's
your query,
PIG?

SELECT *Crime*
from *NewsTable*
WHERE *perp* = *you*

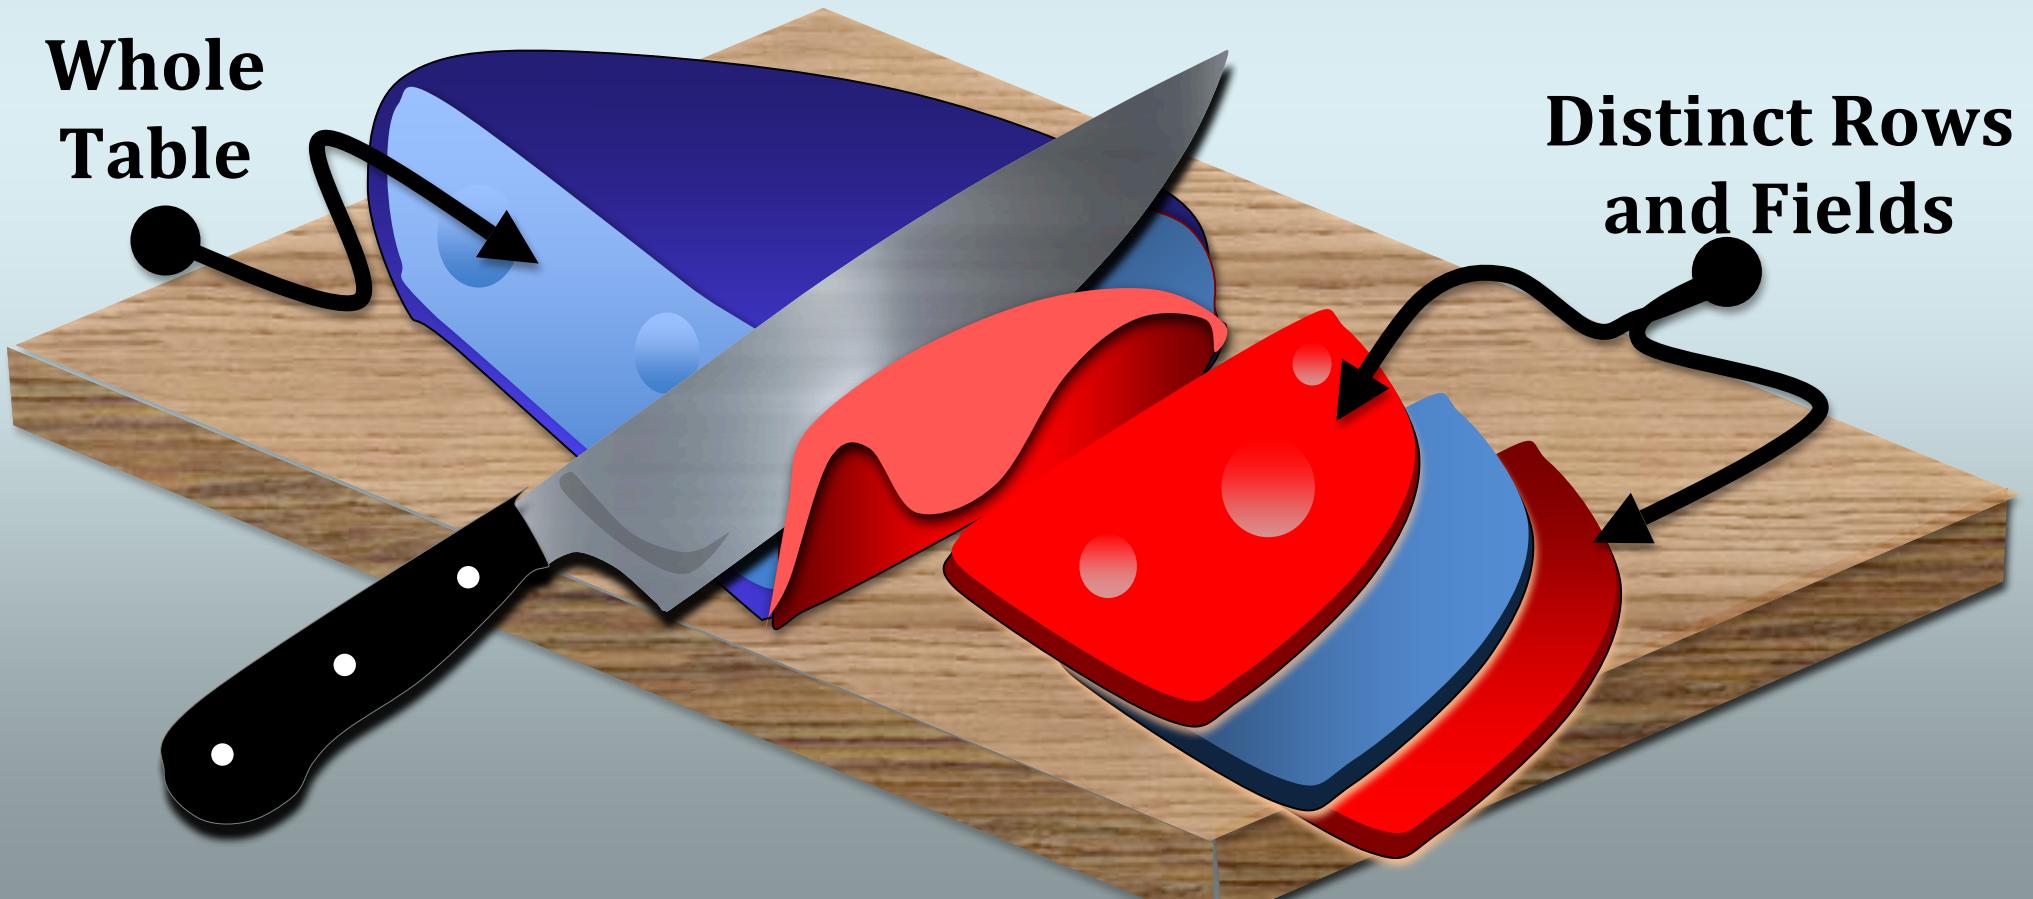
SELECT specific fields from a table:

**SELECT column1, column2...columnN
FROM table_name;**



SELECT and filter for uniqueness of results:

**SELECT DISTINCT column1...columnN
FROM table_name;**



SELECT data that meets specific criteria:

SELECT column1, ..., columnN

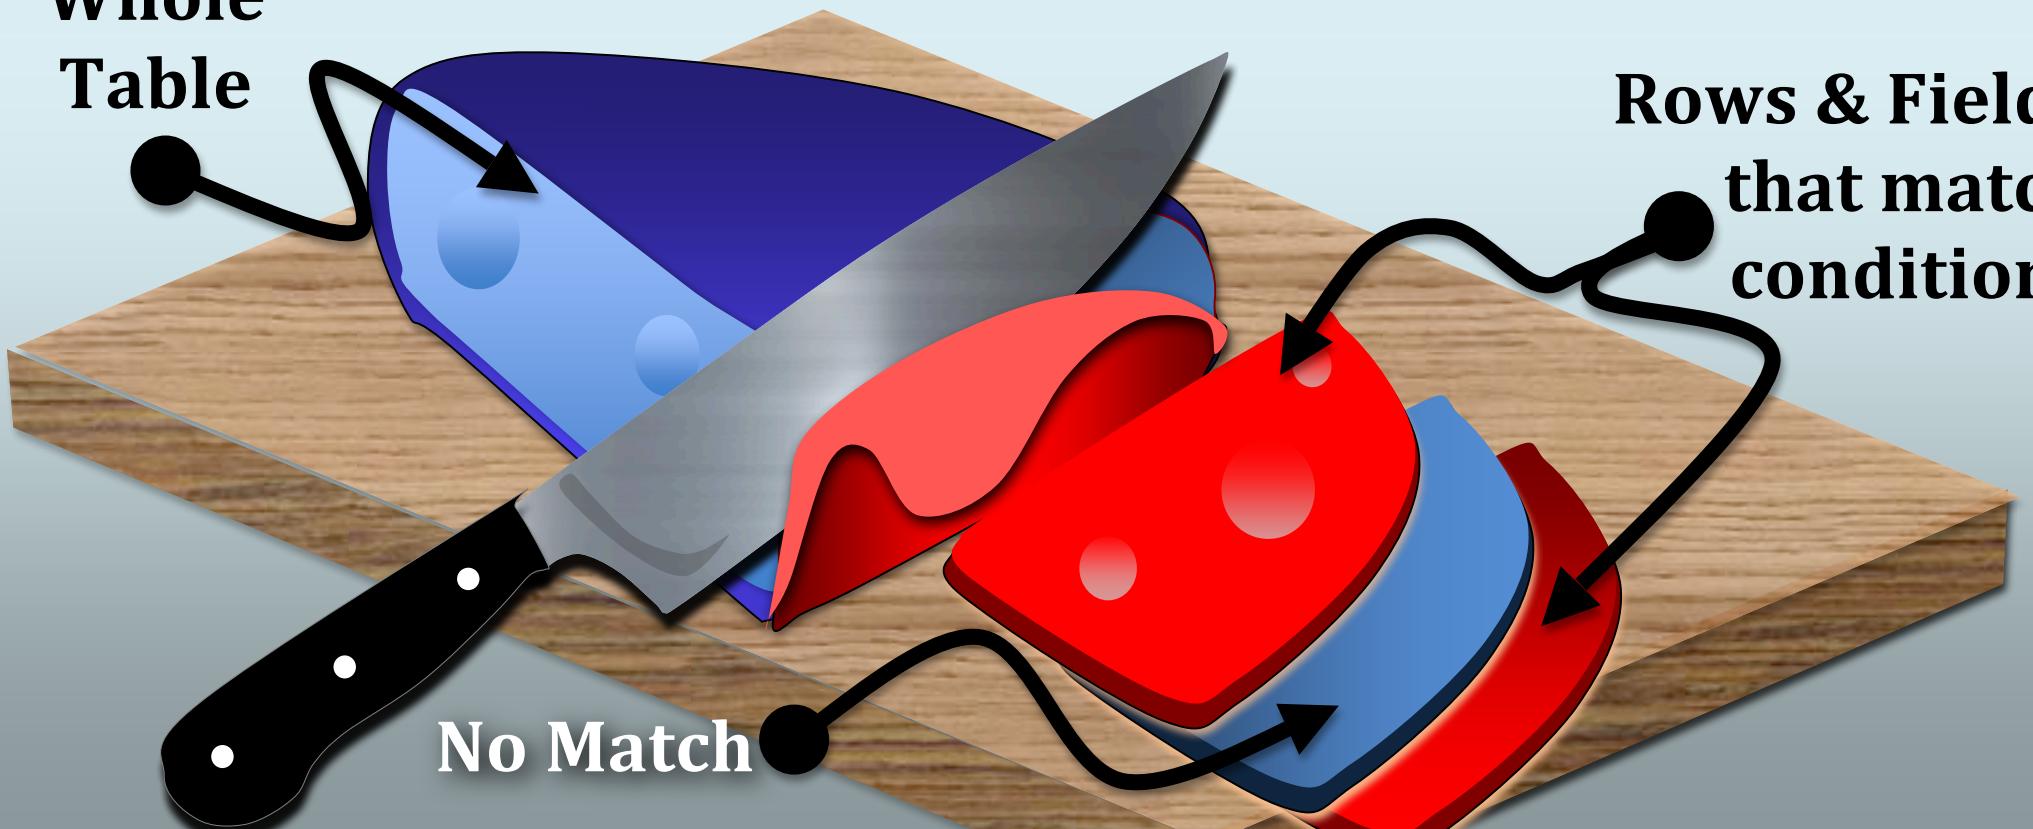
FROM table_name

WHERE condition1 {and|or} conditionN

Whole
Table

Rows & Fields
that match
conditions

No Match



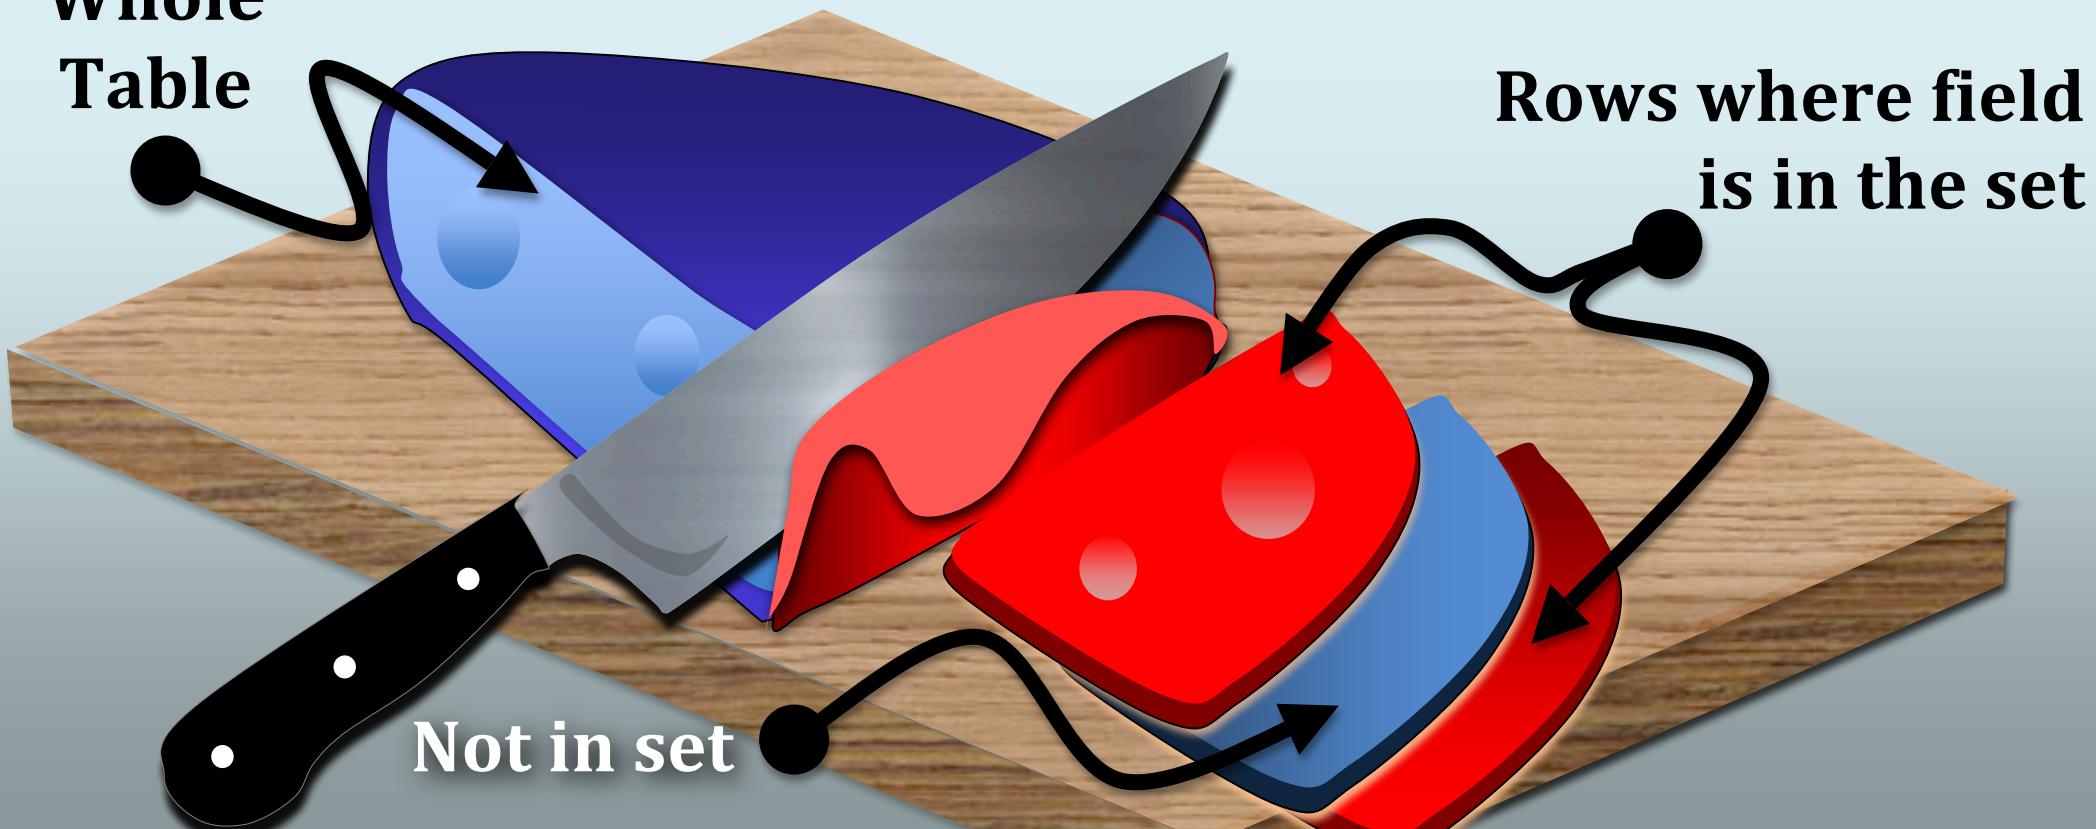
SELECT rows where field is in a set of values:

SELECT column1, column2....columnN

FROM table_name

WHERE column_name IN (val-1, ...,val-N);

Whole
Table





Transaction Control Language
Saving the state of a Database
Rolling back to an earlier state

We can commit changes and mark the new state

COMMIT;

(commit all changes to DB)

SAVEPOINT [StateName];

TCL

To undo changes, rollback back to an earlier state

ROLLBACK; *(to last commit)*

ROLLBACK to [StateName];

TCL

Exact Data Types for Columns/Fields:

| DATA TYPE | FROM | TO |
|-----------------|----------------------------|---------------------------|
| bigint | -9,223,372,036,854,775,808 | 9,223,372,036,854,775,807 |
| int | -2,147,483,648 | 2,147,483,647 |
| smallint | -32,768 | 32,767 |
| tinyint | 0 | 255 |
| bit | 0 | 1 |
| decimal | $-10^{38} + 1$ | $10^{38} - 1$ |
| numeric | $-10^{38} + 1$ | $10^{38} - 1$ |

Approximate Data Types for Fields:

| DATA TYPE | FROM | TO |
|-----------|----------------|---------------|
| float | $-1.79E + 308$ | $1.79E + 308$ |
| real | $-3.40E + 38$ | $3.40E + 38$ |

Date and Time Data Types for Fields:

| DATA TYPE | FROM | TO |
|-----------------|-------------|--------------|
| Date / datetime | Jan 1, 1753 | Dec 31, 9999 |
| smalldatetime | Jan 1, 1900 | Jun 6, 2079 |
| time | 00:00 AM | 00:00 PM |

Characters and Strings Data Types:

| DATA TYPE | DESCRIPTION |
|---------------------|--|
| char | <i>Maximum length of 8,000 characters. (Fixed length non-Unicode characters)</i> |
| varchar | <i>Maximum of 8,000 characters. (Variable-length non-Unicode data).</i> |
| varchar(max) | <i>Maximum length of 231 characters, Variable-length non-Unicode data (SQL Server).</i> |
| text | <i>Variable-length non-Unicode data with a maximum length of 2,147,483,647 characters.</i> |

Unicode Characters / Strings Data Types:

| DATA TYPE | DESCRIPTION |
|---------------|---|
| nchar | <i>Maximum length of 4,000 characters. (Fixed length Unicode)</i> |
| nvarchar | <i>Maximum length of 4,000 characters. (Variable length Unicode)</i> |
| nvarchar(max) | <i>Maximum length of 231characters (SQL Server only). (Variable length Unicode)</i> |
| ntext | <i>Maximum length of 1,073,741,823 characters. (Variable length Unicode)</i> |

Binary Data Types in SQL:

| DATA TYPE | DESCRIPTION |
|-----------------------|---|
| binary | <i>Maximum length of 8,000 bytes (Fixed-length binary data)</i> |
| varbinary | <i>Maximum length of 8,000 bytes. (Variable length binary data)</i> |
| varbinary(max) | <i>Maximum length of 231 bytes (SQL Server only). (Variable length Binary data)</i> |
| image | <i>Maximum length of 2,147,483,647 bytes. (Variable length Binary Data)</i> |

Miscellaneous Data Types in SQL:

| DATA TYPE | DESCRIPTION |
|-------------------------|---|
| sql_variant | <i>Stores values of various SQL Server-supported data types, except text, ntext, and timestamp.</i> |
| timestamp | <i>Stores a database-wide unique number that gets updated every time a row gets updated</i> |
| uniqueidentifier | <i>Store a globally unique identifier (GUID)</i> |
| xml | <i>Stores XML instances in a column or variable (SQL Server)</i> |
| cursor | <i>Reference to a cursor object.</i> |
| table | <i>Stores a result set for later processing</i> |

Joining Multiple Tables Into One Virtual Table

| A | B | C |
|---|---|---|
| a | b | c |
| d | e | f |
| g | h | i |



| X | Y | Z |
|---|---|---|
| o | p | q |
| r | s | t |
| u | v | w |
| x | y | z |

Cartesian Join



If table #1 has C1 columns and R1 rows, and table #2 has C2 columns and R2 rows, then the joined table will have C1 + C2 columns and R1 × R2 rows.

| A | B | C | X | Y | Z |
|---|---|---|---|---|---|
| a | b | c | o | p | q |
| a | b | c | r | s | t |
| a | b | c | u | v | w |
| a | b | c | x | y | z |
| d | e | f | o | p | q |
| d | e | f | r | s | t |
| d | e | f | u | v | w |
| d | e | f | x | y | z |
| g | h | i | o | p | q |
| g | h | i | r | s | t |
| g | h | i | u | v | w |
| g | h | i | x | y | z |

Remember our Customer and Order Tables?

| Order_ID | Customer_ID | Quantity | ORDER |
|----------|-------------|----------|------------------------------------|
| ORD2001 | THX1138 | 3 | <i>Light Sabre, model 27B/6</i> |
| ORD2002 | THX1138 | 1 | <i>Black helmet polish + cloth</i> |
| ORD2003 | PL01001 | 2 | <i>Bun-shaped hair curlers</i> |

| Customer_ID | Customer_NAME | DOB | ADDRESS |
|-------------|---------------|------------|--------------------------|
| THX1138 | Darth Vader | 28.02.0024 | 17 Sith Plaza, Tatooine |
| PL01001 | Leia Organa | 14.02.0045 | 23 Royal Lane, Aldebaran |

We're going to Join them to access both at once

```
SELECT Customer_NAME, ORDER, Quantity  
FROM Customers, Orders  
WHERE Customers.Customer_ID = Orders.Customer_ID  
AND Orders.Quantity > 1 ;
```

| Customer_NAME | ORDER | Quantity |
|---------------|---------------------------------|----------|
| Darth Vader | <i>Light Sabre, model 27B/6</i> | 3 |
| Leia Organa | Bun-shaped hair curlers | 2 |

The Result Set Is A Projected Subset of the Join

Full Cartesian Join:

```
SELECT Customer_NAME, ORDER, Quantity  
FROM Customers, Orders;
```

Notice the lack of join constraints above.

**This query performs projection (column filtering)
but not selection (row filtering)**

**The Result Set thus contains $R_1 \times R_2$ rows
where R_1 = Number of Rows in Table 1
and R_2 = Number of Rows in Table 2**

A Database of Poker Hands and Playing Cards

Poker_Hands

| Player_ID | Game_ID | R1 | R2 | R3 | R4 | R5 |
|-----------|---------|----|----|----|----|----|
| 12789 | 17MET | QH | QS | 3D | 3C | 3H |
| 90734 | 82SAT | 7C | 4S | 4D | 4C | 3H |

Playing_Cards

| Card_Name | Face | Type | Suit | Face_Value | Game_Value |
|-----------|------|------|------|------------|------------|
| 1H | no | A | H | 1 | 14 |
| 2H | no | 2 | H | 2 | 2 |

Quantitative Queries on Poker_Hands will require
a Join (or two) with Playing_Cards

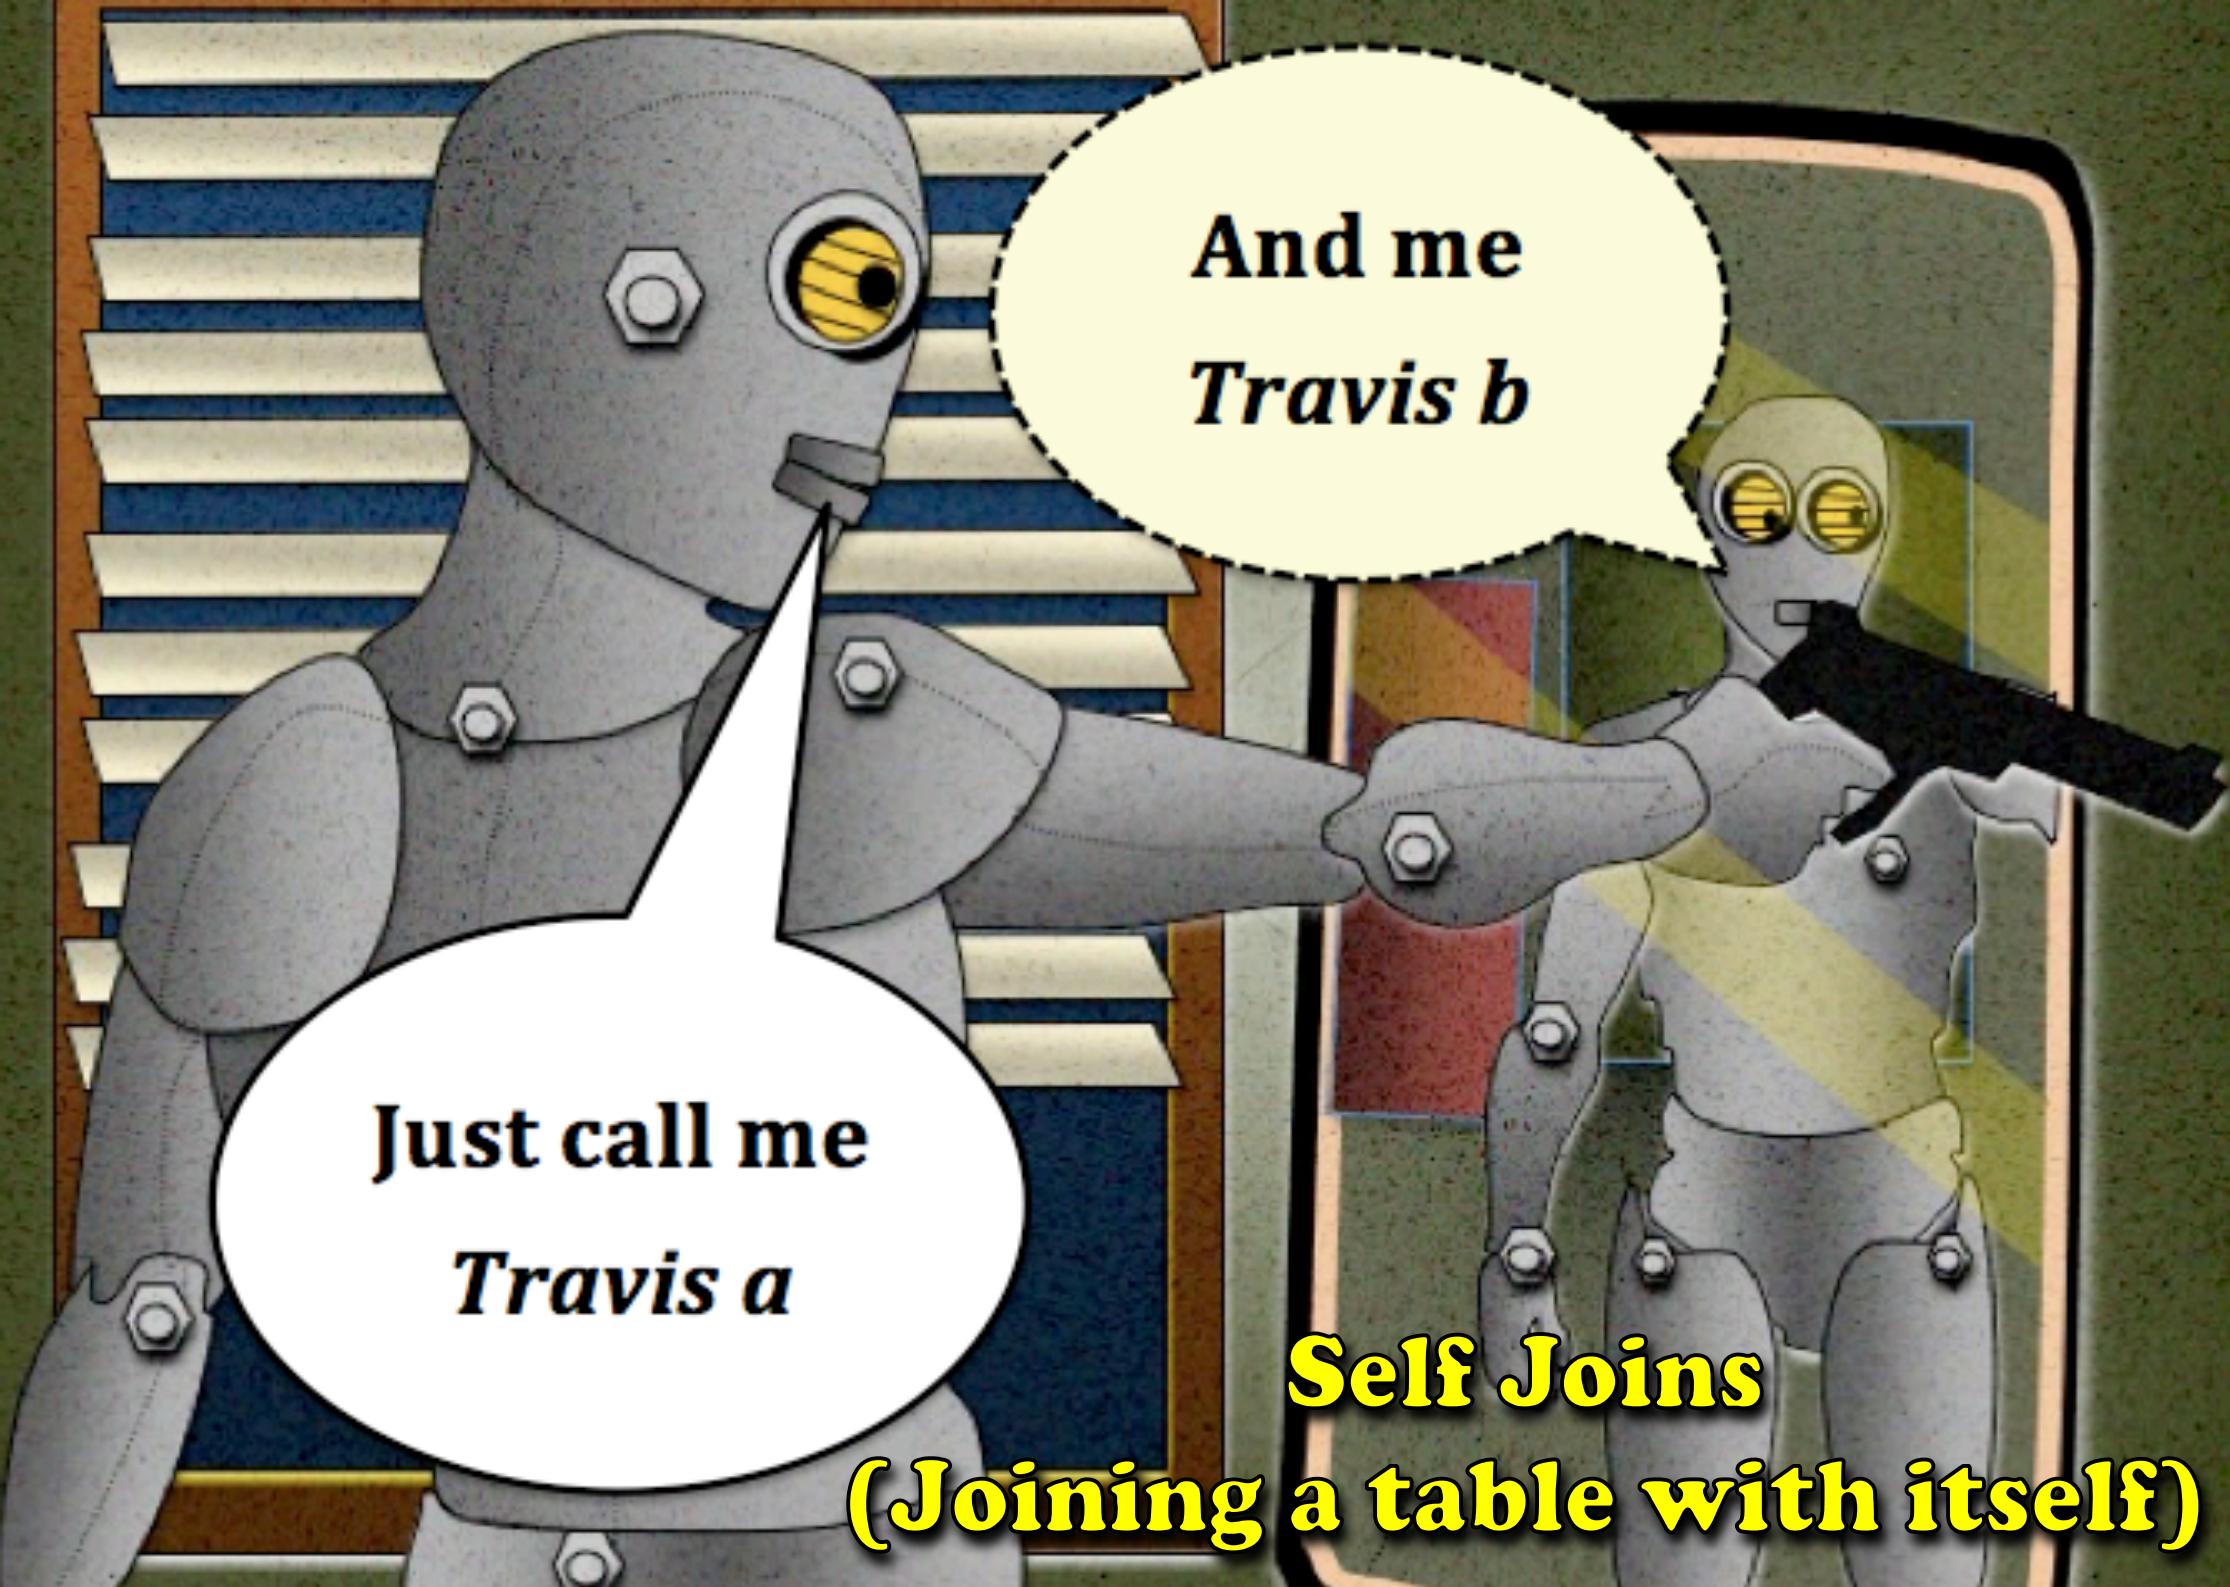
```

SELECT Player_ID, Game_ID, R1, R2, R3, R4, R5
FROM Poker_Hands, Playing_Cards
WHERE
Poker_Hands.R1 = Playing_Cards.Card_Name
AND Playing_Cards.Face = "yes";

```

Only those hands whose top card is a face card

| Player_ID | Game_ID | R1 | R2 | R3 | R4 | R5 |
|-----------|---------|----|-----|----|----|----|
| 12789 | 17MET | QH | QS | 3D | 3C | 3H |
| 12789 | 23WSA | KH | 10H | 7H | 3H | 1H |
| 56347 | 65YOB | QH | JH | 8C | 8S | 8D |



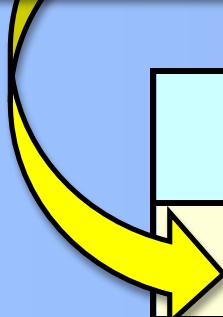
Just call me
Travis a

And me
Travis b

Self Joins
(Joining a table with itself)

```

SELECT Player_ID, Game_ID, R1, R2
      FROM
Poker_Hands, Playing_Cards a, Playing_Cards b
WHERE Poker_Hands.R1 = a.Card_Name
      AND Poker_Hands.R2 = b.Card_Name
      AND a.Face = "yes" AND b.Face = "yes";
    
```



| Player_ID | Game_ID | R1 | R2 |
|-----------|---------|----|----|
| 12789 | 17MET | QH | QS |
| 56347 | 65YOB | QH | JH |

We can join multiple versions of the same table, provided we give each version a unique handle (like a and b)

Hands whose top two cards are face cards

Regular Joins are sometimes called INNER JOINS

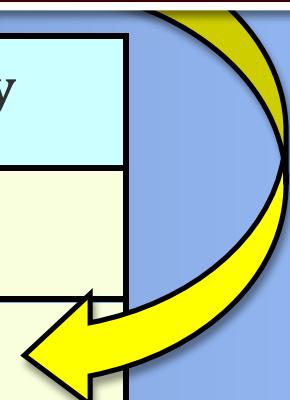
SELECT Customer_NAME, ORDER, Quantity

FROM Customers

INNER JOIN Orders

ON Customers.Customer_ID = Orders.Customer_ID;

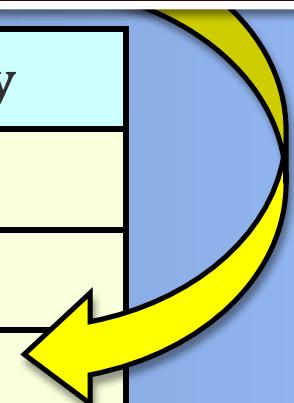
| Customer_NAME | ORDER | Quantity |
|---------------|---------------------------------|----------|
| Darth Vader | <i>Light Sabre, model 27B/6</i> | 3 |
| Leia Organa | Bun-shaped hair curlers | 2 |



Notice the alternative syntax to WHERE constraint

**SELECT Customer_NAME, ORDER, Quantity
FROM Customers
LEFT JOIN Orders
ON Customers.Customer_ID = Orders.Customer_ID;**

| Customer_NAME | ORDER | Quantity |
|-------------------|---------------------------------|----------|
| Darth Vader | <i>Light Sabre, model 27B/6</i> | 3 |
| Leia Organa | Bun-shaped hair curlers | 2 |
| Emperor Palpatine | | |
| Bib Fortuna | | |



When no matching rows in right table, return **null**

**SELECT Customer_NAME, ORDER, Quantity
FROM Customers
RIGHT JOIN Orders
ON Customers.Customer_ID = Orders.Customer_ID;**

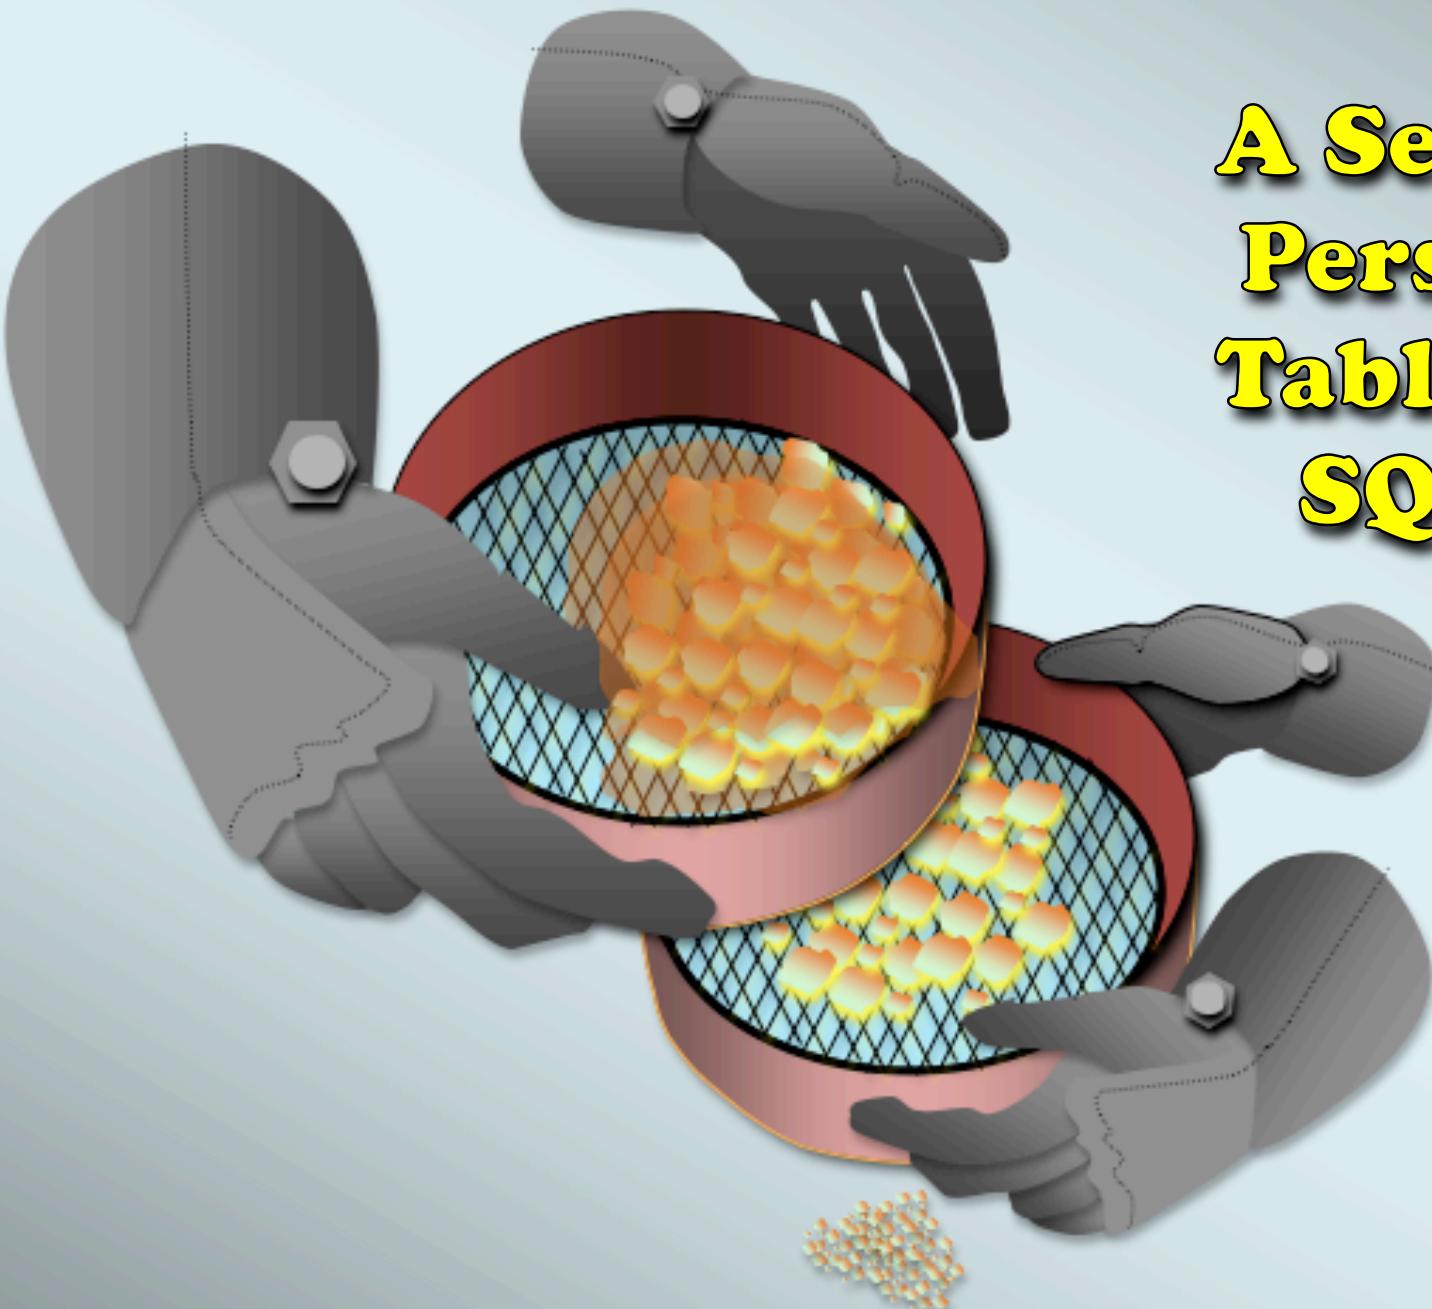
| Customer_NAME | ORDER | Quantity |
|---------------|---------------------------------|----------|
| Darth Vader | <i>Light Sabre, model 27B/6</i> | 3 |
| Leia Organa | Bun-shaped hair curlers | 2 |
| | Sand-speeder model SUX | 1 |
| | Gas-powered Harpoon gun | 2 |

When no matching rows in left table, return null

```
SELECT Customer_NAME, ORDER, Quantity  
FROM Customers  
FULL OUTER JOIN Orders  
ON Customers.Customer_ID = Orders.Customer_ID;
```

| Customer_NAME | ORDER | Quantity |
|-----------------|---------------------------------|----------|
| Darth Vader | <i>Light Sabre, model 27B/6</i> | 3 |
| Leia Organa | <i>Bun-shaped hair curlers</i> | 2 |
| | <i>Sand-speeder model SUX</i> | 1 |
| | <i>Gas-powered Harpoon gun</i> | 2 |
| Sheev Palpatine | | |
| Bib Fortuna | | |

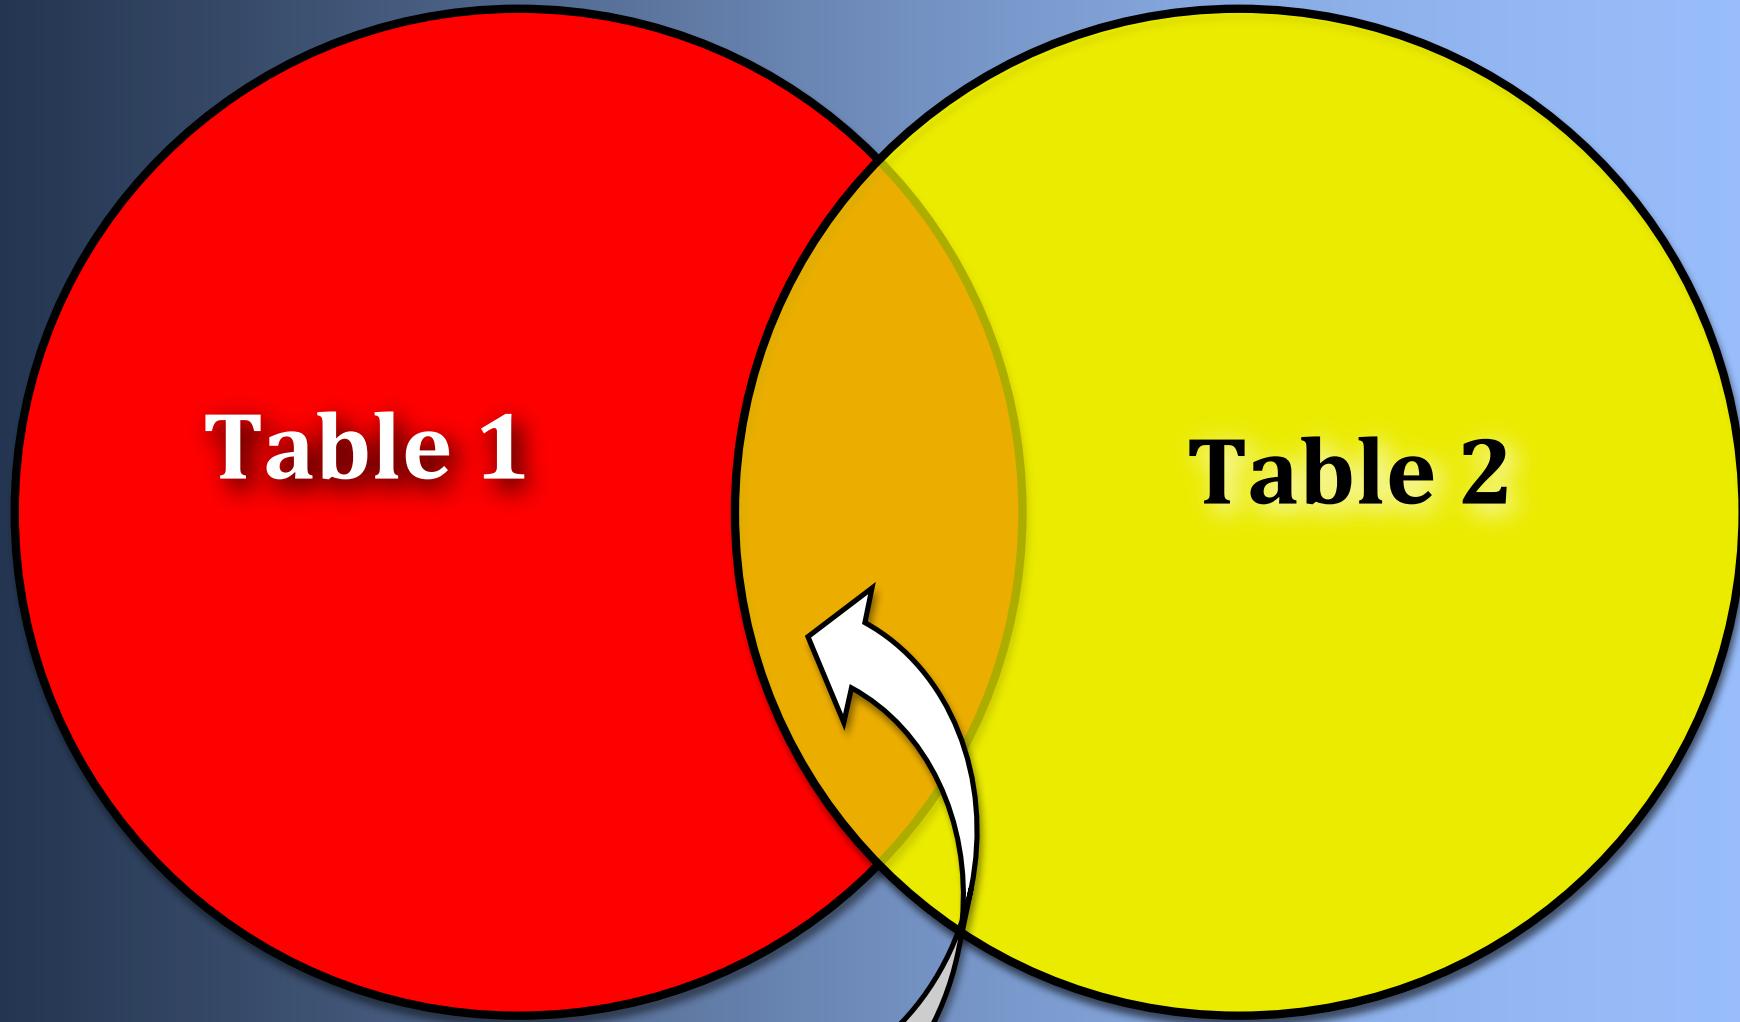
When no matching rows in either table, return **null**



A Set-Theoretic Perspective on Table Joins and SQL Queries

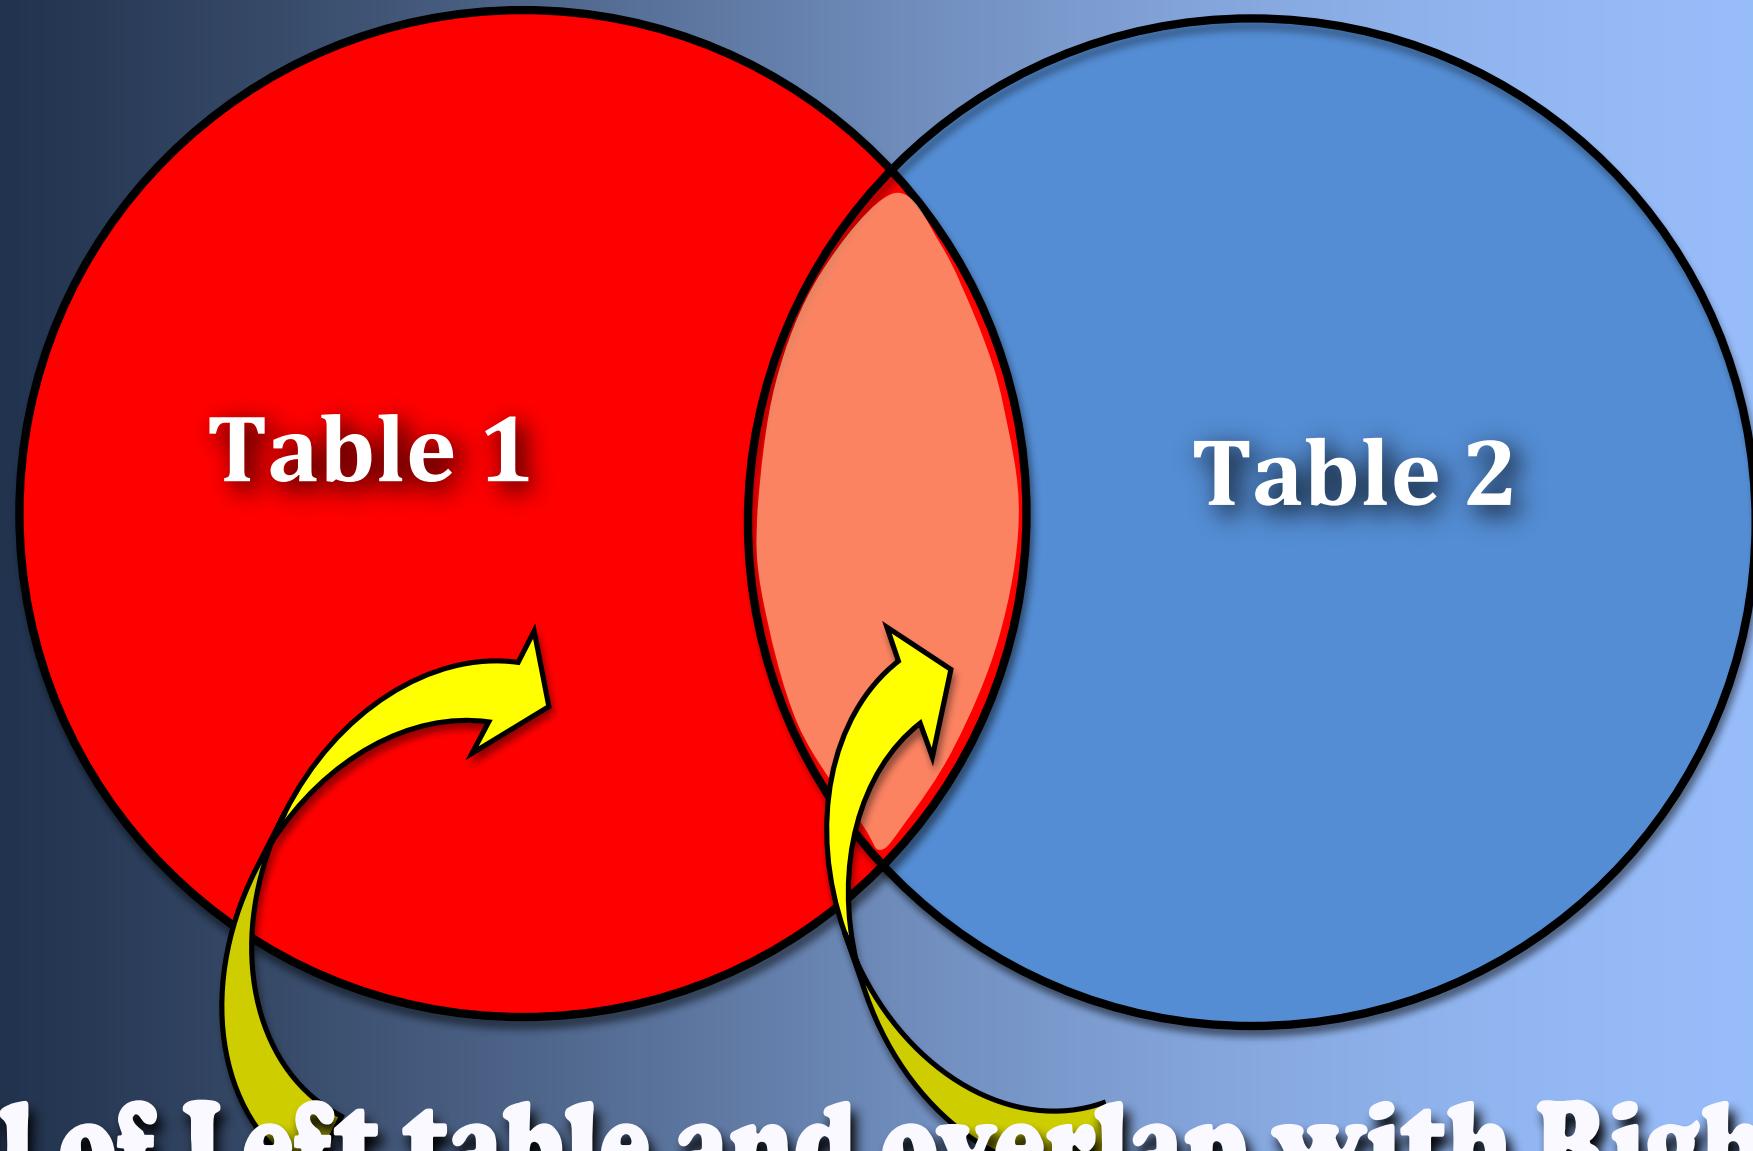
Queries sift information from data

Regular (or INNER) Joins

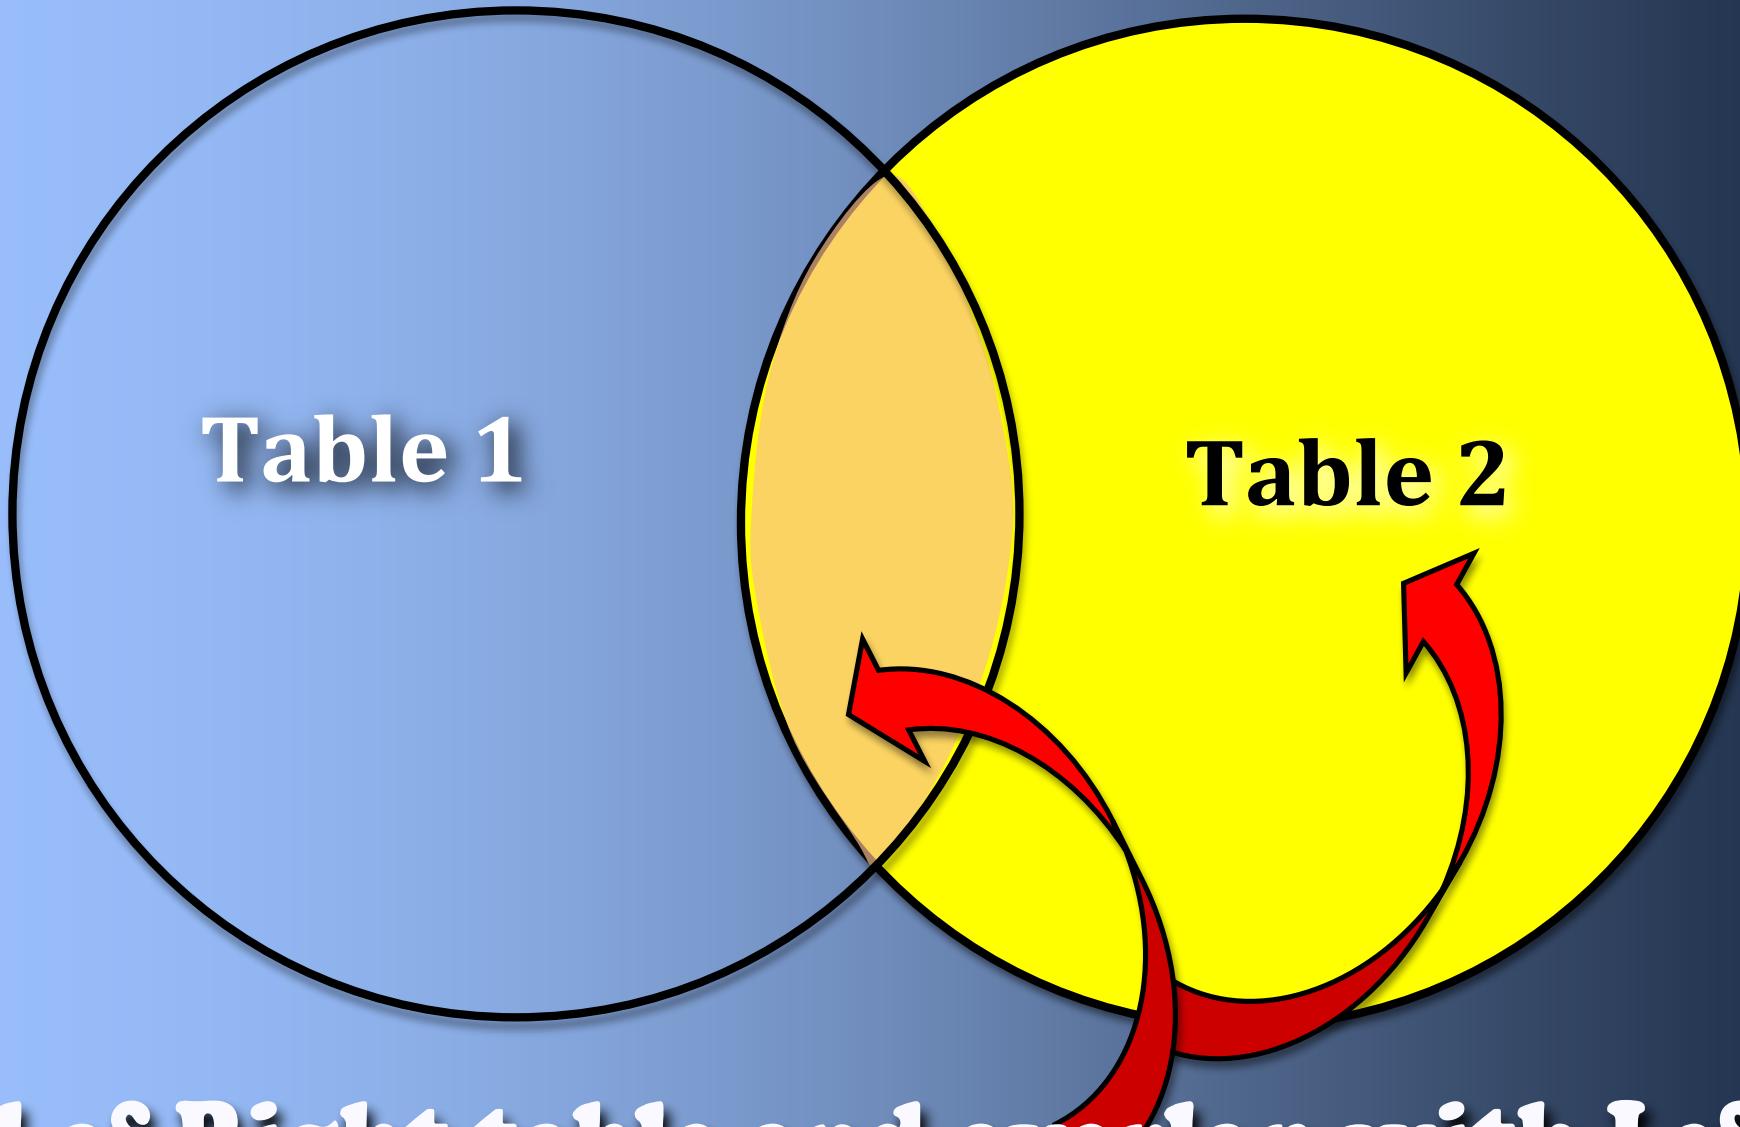


WHERE (or ON) constrained overlap

LEFT (or LEFT OUTER) Joins

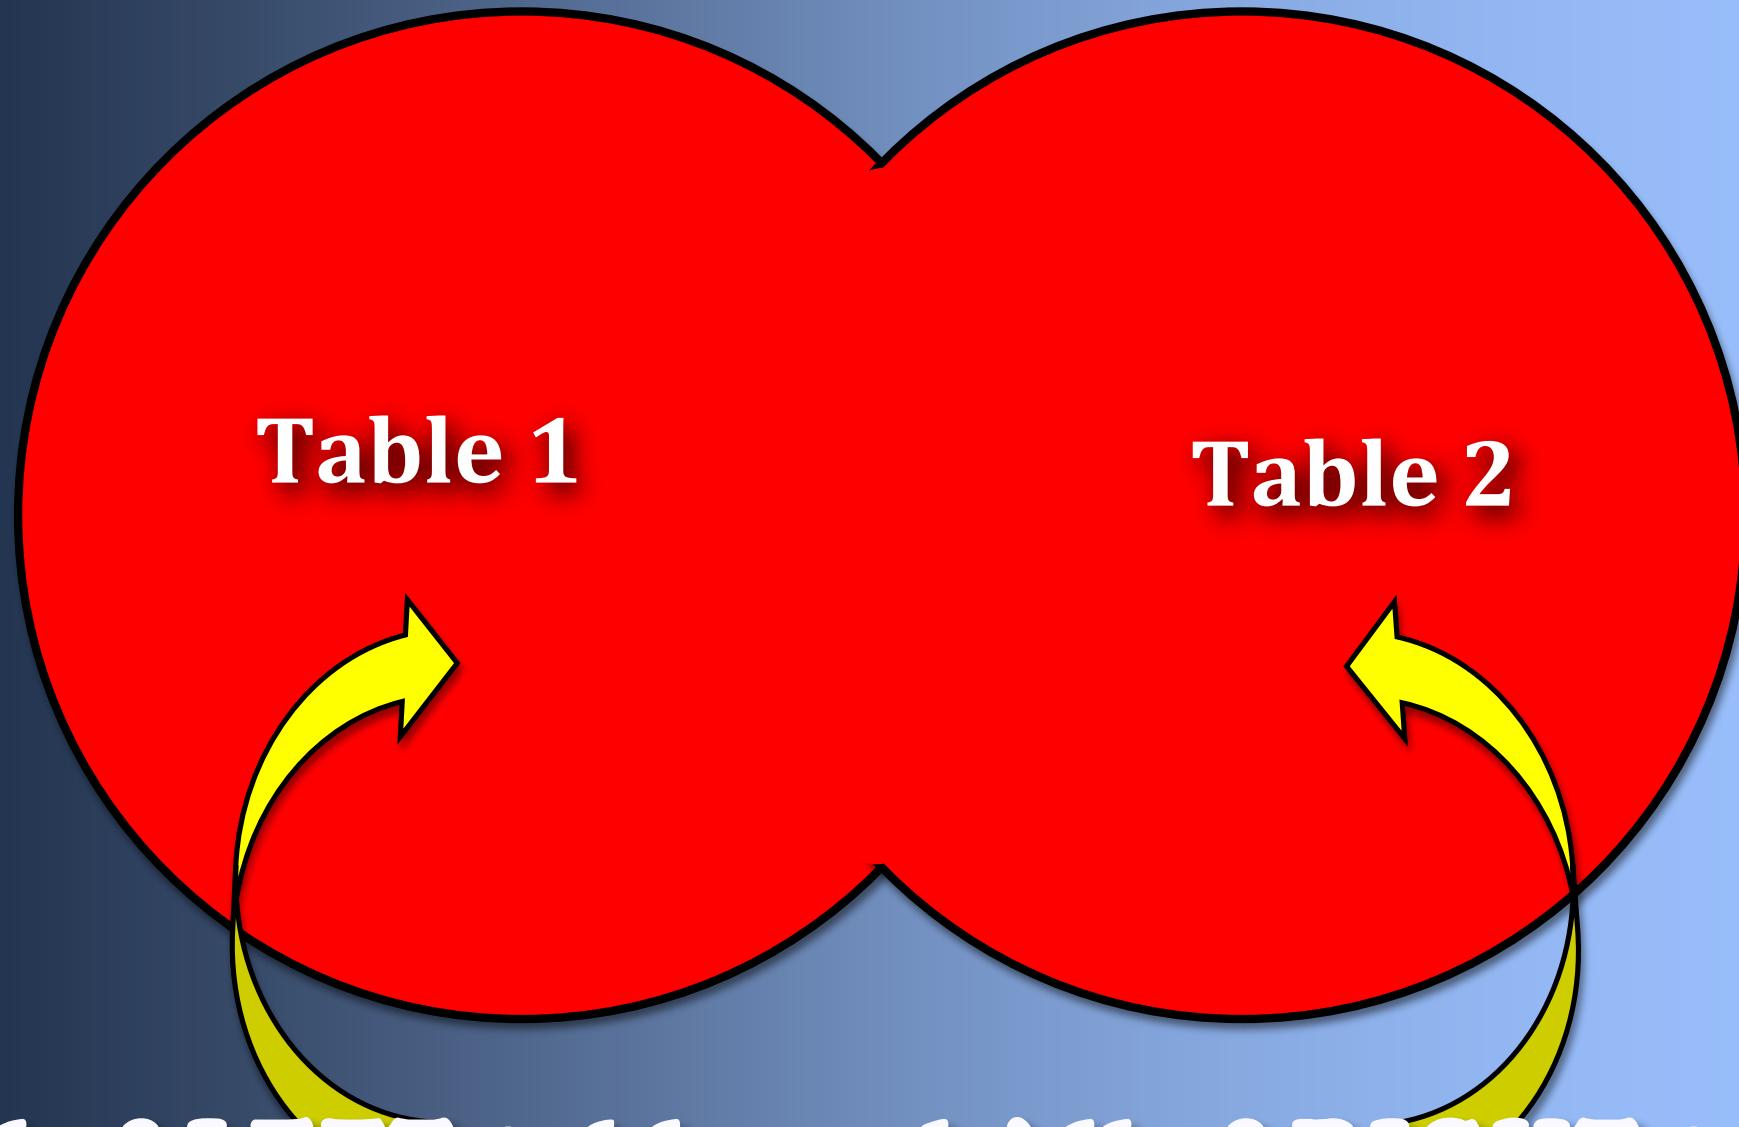


RIGHT (or RIGHT OUTER) Joins



All of Right table and overlap with Left

FULL OUTER JOINS



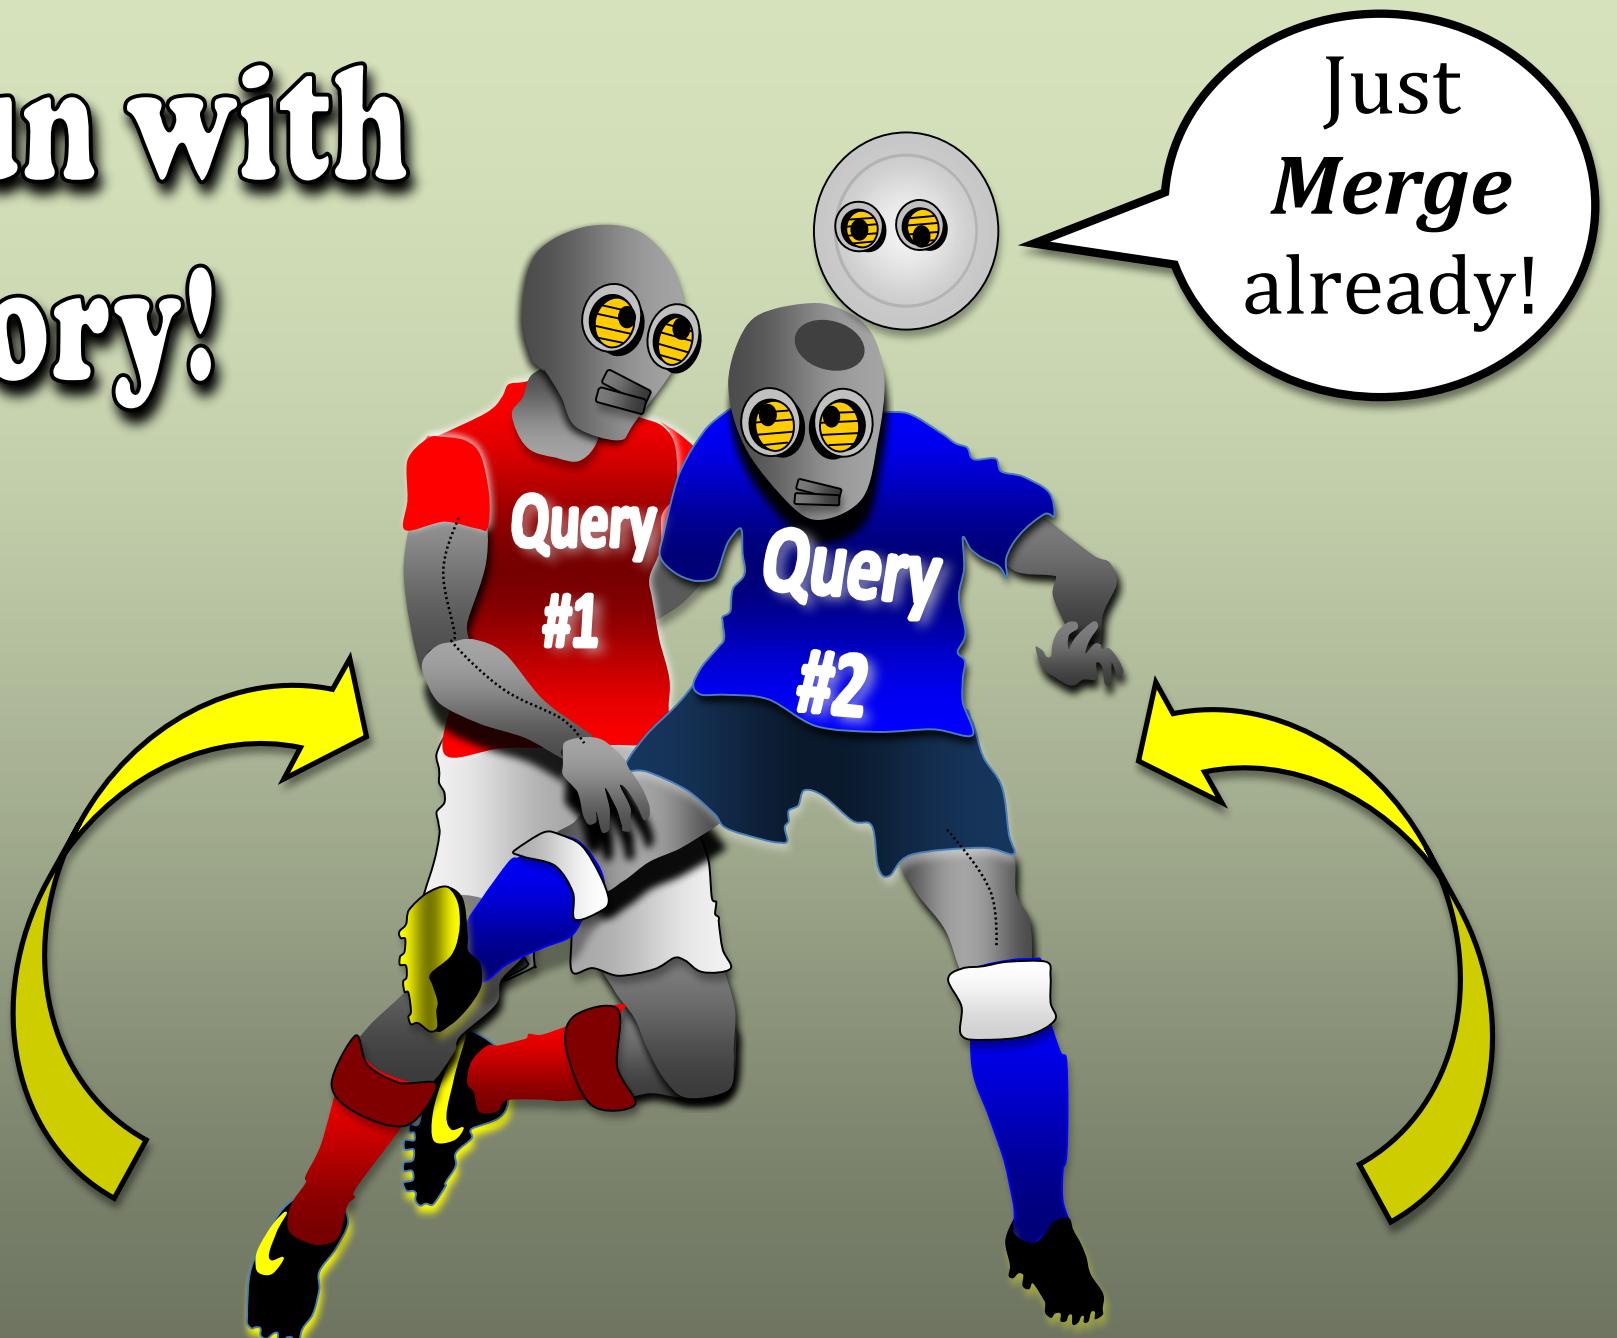
All of LEFT table and All of RIGHT table

Normalization (1NF...3NF) encourages us to factorize our data into the least redundant table structures

Table "joins" are a way of putting the pieces back together again for the purposes of querying across multiple tables and views on the data



More fun with Set theory!



Why Join tables when we can Merge queries?

Let's add another table first ...

Customers (sample rows)

| Customer_ID | Customer_NAME | DOB | ADDRESS |
|-------------|---------------|------------|--------------------------|
| THX1138 | Darth Vader | 28.02.0024 | 17 Sith Plaza, Tatooine |
| PL01001 | Leia Organa | 14.02.0045 | 23 Royal Lane, Aldebaran |

Suppliers (sample rows)

| Supplier_ID | Supplier_NAME | VAT_NUM | ADDRESS |
|-------------|----------------|----------|--------------------------|
| JTH89369 | Jabba the Hutt | 56347381 | 29 Slime Court, Tatooine |
| BF903735 | Boba Fett | 73735631 | 16 Sarlacc Pit, Tatooine |

Now let's merge queries on different tables

UNION keyword (for combining two result-sets)

*SELECT
Customer_NAME
from
Customers*

*SELECT
Supplier_NAME
from
Suppliers*

```
SELECT Customer_NAME from Customers
      UNION
SELECT Supplier_NAME from Suppliers;
```

Comparing Strings with the like constraint

```
SELECT Customer_NAME from Customers  
WHERE  
Address like "%Tatooine%";
```

It's like
<rasp>
looking into
a mirror
<rasp>

```
SELECT Customer_NAME from Customers
WHERE
Address like "%Tatooine%"
UNION
SELECT Supplier_NAME from Suppliers
WHERE
Address like "%Tatooine%";
```

Returns the unified result-set
Darth Vader, Jabba the Hutt and Boba Fett

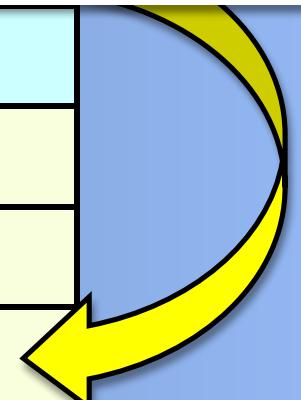
Combining Select and Insert statements

```
SELECT *
INTO Local_Customers
FROM Customers
WHERE Address like "%Tatooine%";
```

```
SELECT *
INTO Offworld_Customers
FROM Customers
WHERE Address not like "%Tatooine%";
```

**SELECT Customer_NAME, ORDER, Quantity
FROM Customers, Orders
WHERE
Customers.Customer_ID = Orders.Customer_ID
ORDER BY Quantity, Customer_NAME;**

| Customer_NAME | ORDER | Quantity |
|-------------------|---------------------------------|----------|
| Darth Vader | <i>Light Sabre, model 27B/6</i> | 3 |
| Boba Fett | <i>Gas-powered Harpoon gun</i> | 2 |
| Chewbacca | <i>Wookiee Dandruff shampoo</i> | 2 |
| Leia Organa | <i>Bun-shaped hair curlers</i> | 2 |
| Biggs Darklighter | <i>Sand-speeder model SUX</i> | 1 |



ORDER BY sorts a result-set by given column values



AGGREGATE FUNCTIONS on Table columns

MIN return smallest value

MAX return largest value

SUM return cumulative value

AVG return mean value

COUNT return num values

COUNT(*) return num rows

What is the total number of products ordered from Tatooine addresses?

```
SELECT SUM(Quantity)  
FROM Orders  
WHERE Address like “%Tatooine%”;
```

What is the average size of an order placed from offworld (outside Tatooine)?

```
SELECT AVG(Quantity)  
FROM Orders  
WHERE Address not like “%Tatooine%”;
```

Let's create a new Employees table:

| Staff_ID | Staff_NAME | Salary | Department |
|----------|-------------------|---------|-------------------------|
| THX1138 | Darth Vader | 256,000 | <i>Human Resources</i> |
| REX2910 | Sheev Palpatine | 999,666 | <i>Upper Management</i> |
| BF903735 | Boba Fett | 127,000 | <i>Contractors</i> |
| GMT5639 | Wilhuff Tarkin | 384,000 | <i>Upper Management</i> |
| OKR6758 | Orson Krennic | 372,000 | <i>Human Resources</i> |
| GOR3679 | Galen Erso | 176,000 | <i>Contractors</i> |
| MAV8368 | Maximillian Veers | 110,846 | <i>Upper Management</i> |

OK, so our “organization” is the Death Star

What is the highest salary in each department in our organization?

```
SELECT Department, MAX(Salary)  
FROM Employees  
GROUP BY Department;
```

| Department | MAX(Salary) |
|------------------|--------------------|
| Human Resources | 372,000 |
| Upper Management | 999,666 |
| Contractors | 176,000 |

We don't want the overall MAX, but the MAX by department

GROUP BY is used with aggregate functions to apply them to *groups* of column values